

# Chapter 6: CPU Scheduling

# Chapter 6: CPU Scheduling

- Basic Concepts
- Scheduling Criteria
- Scheduling Algorithms
- Real-Time CPU Scheduling
- Algorithm Evaluation

# Objectives

- To introduce *CPU scheduling*, which is the basis for multi-programmed operating systems
- To describe various CPU-scheduling algorithms
- To discuss evaluation criteria for selecting a CPU-scheduling algorithm for a particular system

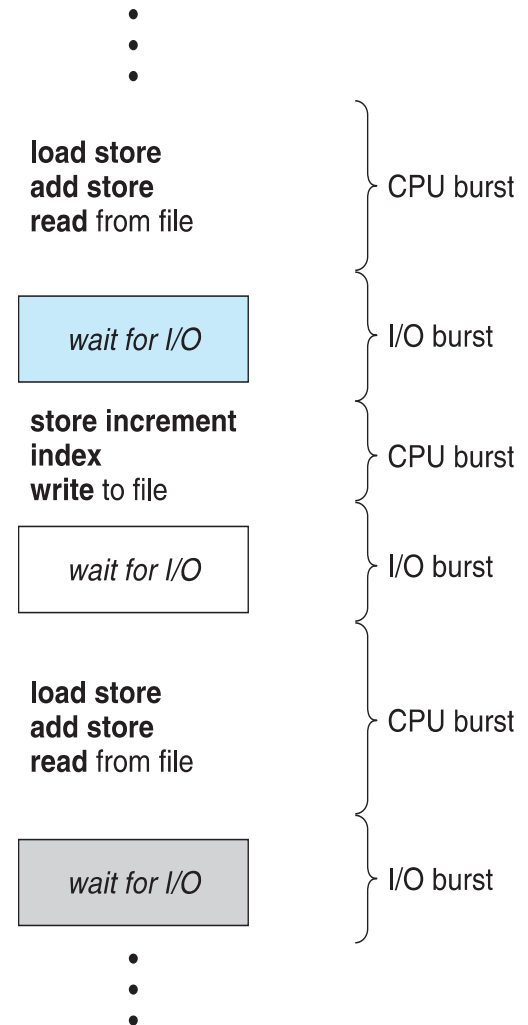
# Basic Concepts

- **Objective of multiprogramming:** Achieve maximum CPU utilization

- CPU always running a process
- No Idle CPU

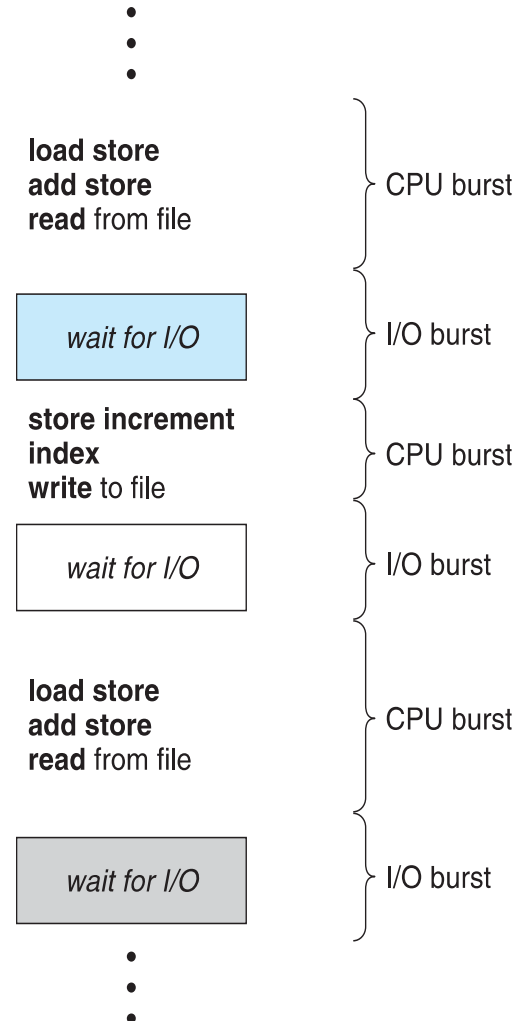

- A process runs in CPU bursts and I/O bursts

- Run instructions (CPU Burst)
- Wait for I/O (I/O Burst)




# Basic Concepts

- CPU–I/O Burst Cycle:
  - Process execution consists of a **cycle** of *CPU execution* and *I/O wait*
  - When a process is waiting, then assign the CPU to another process
  - See Process Scheduler on Chap-3
- **CPU burst** followed by **I/O burst**



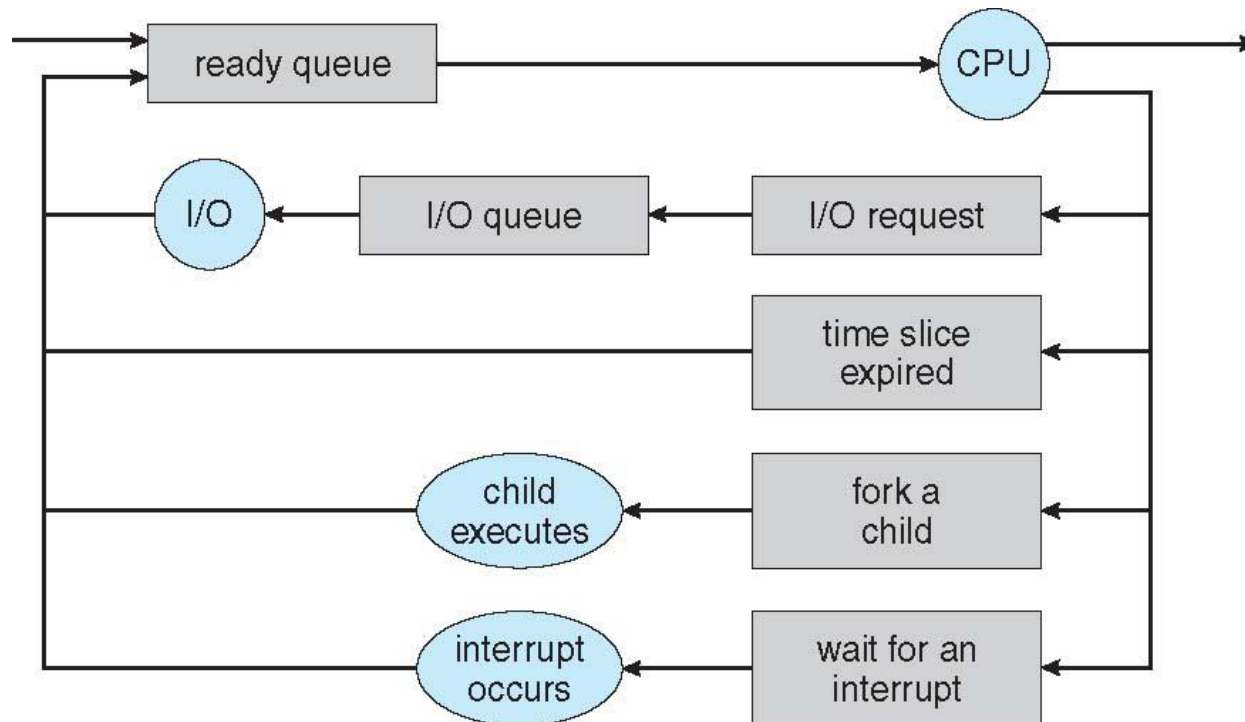
# Histogram of CPU-burst Times

- CPU burst distribution is of main concern
- CPU bursts vary from process to process, but an extensive study shows frequency patterns similar to the below diagram:
- [Scheduling](#) is aided by knowing the length of these bursts



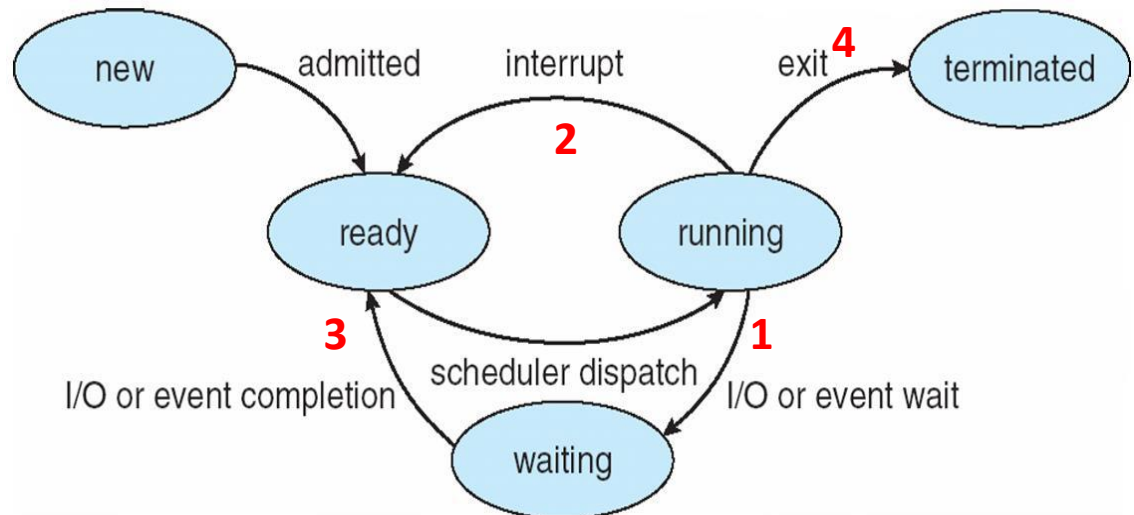
# CPU Scheduler Queues

- **Short-term scheduler** selects from among the processes in ready queue, and allocates the CPU to one of them
  - Queue may be ordered in various ways: **FIFO, LIFO, Random, Priority, ... etc**
- Types of scheduling
  - **Preemptive scheduling** - Running process may be interrupted and moved to the Ready queue
  - **Non-preemptive scheduling** - once a process is running, it continues to execute until it terminates or blocks for I/O



# Scheduling Decision Modes

- The scheduler runs when it needs to select a process to run on the processor. The **scheduling decision mode** specifies the times at which the *selection of a process to run* is made. This decision can take place at the following times:
  1. Switches from running to waiting state; *ex: as result of I/O request or wait()*
  2. Switches from running to ready state; *ex: when an interrupt occurs*
  3. Switches from waiting to ready; *at completion of I/O*. or because a new process was created and was placed in the queue.
  4. Terminates






# CPU Scheduler

- If a scheduling decision is made **only under conditions 1 and 4**, the scheduling is called **nonpreemptive scheduling**.
  - Process keeps the CPU until it either terminates or switches to waiting state
  - No choice in terms of scheduling; new process must be selected for CPU
- In circumstances 2 and 3, the OS scheduler has a choice:
  - it can either allow the current process to continue running, or it could step in and put the current process to sleep and select a different process to run. The latter operation is called "**preempting**" the current process and scheduling a new one to run.

# The Dispatcher

- The **dispatcher** is the kernel routine that performs the context switch.
- Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
  - *switching context* from the currently running process to the new one.
  - *switching to user mode* by changing the processor mode to user mode
- **Dispatch latency** – time it takes for the dispatcher to stop one process and start another running
  - Dispatcher is invoked during every process switch; hence it should be as fast as possible



# Scheduling Criteria

- **Scheduling criteria** are the objectives of scheduling algorithms.
- **CPU utilization** – keep the CPU as busy as possible
  - Ranges from 40% to 90% i.e., from light to heavy loaded
  - Scheduling algorithm optimization criteria: Max CPU utilization
- **Throughput** – # of processes that completing per time unit
  - Ranges from 10 processes/second to 1 process/hour
  - Scheduling algorithm optimization criteria: Max throughput
- **Turnaround time** – the time from the submission of a job until it **completes**;
  - Sum of times spent in job pool + ready queue + CPU execution + doing I/O
  - Scheduling algorithm optimization criteria: Min turnaround time
- **Waiting time** – amount of time a process has been waiting in the **ready queue**
  - Sum of times spent waiting in the ready queue
  - Scheduling algorithm optimization criteria: Min waiting time
- **Response time** – amount of time it takes from when a request was submitted until the first response is produced, not output
  - The time it takes to start responding to the user; in an interactive system
  - Scheduling algorithm optimization criteria: Min response time

# Scheduling Algorithm Optimization Criteria

- It is desirable that we:

- **Maximize:**

- CPU utilization
- throughput

- **Minimize:**

- turnaround time
- waiting time
- response time

# Scheduling Algorithms

- First come first serve
- Shortest Job First
- Shortest Remaining Time First
- Round Robin
- Multilevel Queue Scheduling
- Multilevel Feedback Queueing

# First- Come, First-Served (FCFS) Scheduling

- FCFS is very simple - Just a FIFO queue, like customers waiting in line at the bank or the post office or at a copying machine.

# First- Come, First-Served (FCFS) Scheduling

<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

- Suppose that the processes arrive in the order:  $P_1, P_2, P_3$ 
  - The Gantt Chart for the schedule is:[includes start and finish time of each process]



**Throughput:** 3 jobs/30 seconds = 0.1 jobs/second

- Turnaround Time:  $P_1 : 24, P_2 : 27, P_3 : 30$

- Average TT:  $(24 + 27 + 30)/3 = 27$

- Waiting time for  $P_1 = 0; P_2 = 24; P_3 = 27$

- Average waiting time:  $(0 + 24 + 27)/3 = 17$

- The simplest scheduling algorithm but usually very bad average waiting time

# FCFS Scheduling (Cont.)

Now suppose that the processes arrive in the order:

$P_2, P_3, P_1$ , (instead of  $P_1, P_2, P_3$ ) where Burst times, for  $P_2 = 3$ ,  $P_3 = 3$  and  $P_1 = 24$ .

- The Gantt chart for the schedule is: ??

- **Throughput:** ?

- **Turnaround time:** ?

- **Average TT:** ?

- **Waiting time** for  $P_1 = ?$ ;  $P_2 = ?$ ;  $P_3 = ?$

- **Average waiting time:** ?



# FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order:

$P_2, P_3, P_1$ , where Burst times, for  $P_2 = 3$ ,  $P_3 = 3$  and  $P_1 = 24$ .

- The Gantt chart for the schedule is:



- **Throughput:** 3 jobs / 30 sec = 0.1 jobs/sec
- **Turnaround time:**  $P_1 = 30$ ,  $P_2 = 3$ ,  $P_3 = 6$
- **Average TT:**  $(30 + 3 + 6)/3 = 13 \rightarrow$  much less than 27
- **Waiting time** for  $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$
- **Average waiting time:**  $(6 + 0 + 3)/3 = 3$ ; 3 instead of 17, substantial reduction in wait time
- Much better than previous case
- **Lesson:** scheduling algorithm can reduce TT
  - Minimizing waiting time can improve TT
- Can a scheduling algorithm improve throughput?
  - Yes, if jobs require both computation and I/O

# FCFS Scheduling: Convoy Effect (Cont.)

- **Convoy effect** - short processes maybe stuck behind long processes
  - CPU-bound processes hold CPU while I/O-bounds wait in ready queue for the CPU
    - I/O devices are idle until CPU released... then CPU is idle.. Then ... this repeats
  - Consider one long CPU-bound and many I/O-bound processes
    - I/O-bound processes spend most of the time waiting for CPU-bound to release CPU
    - Result in lower CPU and device utilization
- **FCFS is nonpreemptive:** process holds CPU until termination or I/O request


# Shortest-Job-First (SJF) Scheduling

- The preceding examples showed that when short processes are scheduled before long processes, the average waiting times are smaller than if they were to follow after it.
- Associate with each process the length of its next CPU burst
  - Use these lengths to schedule the process with the shortest time
  - Assign the CPU to the process that has the smallest next CPU burst
  - What if two process have the same next CPU burst length? FCFS breaks the tie
- Better term: Shortest-Next-CPU-Burst-Scheduling (SNCB) algorithm
  - But most books use SJF
- SJF is optimal – gives minimum average waiting time for a given set of processes
  - The difficulty is knowing the length of the next CPU request
  - Could ask the user to estimate their processes' time limits; for job scheduling
    - CPU scheduling can use these time limits as estimates of CPU burst lengths

# Example of SJF

<u>Process</u>	<u>Burst Time</u>
$P_1$	6
$P_2$	8
$P_3$	7
$P_4$	3

- SJF scheduling chart



- Average waiting time =  $(3 + 16 + 9 + 0) / 4 = 7$ 
  - With FCFS it would be 10.25 time units with this order P1 P2 P3 P4
- Moving a short process before a long one decreases its waiting time more than it increases the long process's waiting time. Thus, decrease of average waiting time

# Determining Length of Next CPU Burst

- Burst time can be estimated using lengths of past bursts: next = average of all past bursts
  - Then pick process with shortest predicted next CPU burst
- Can be done by using the length of previous CPU bursts, using **exponential weighted moving average**; **next = average of (past actual + past estimate)**
  - 1  $T_n$  : actual length of  $n^{th}$  CPU burst
  - 2  $S_{n+1}$  = predicted value for the next CPU burst
  3.  $\alpha, 0 \leq \alpha \leq 1$
  4. Define :
$$S_{n+1} = (1-\alpha) S_n + \alpha T_n$$
- Relative weight of recent history and past history.  $\alpha = \frac{1}{2}$  usually but anything in the range  $[0, 1]$  is acceptable
- Commonly,  $\alpha$  set to  $\frac{1}{2}$

• Note the difference between English alphabet t (pronounced TEE) and Greek letter  $\tau$  (pronounced tau)

# Prediction of the Length of the Next CPU Burst


How to find next estimated burst, if initial estimate is  $S_0 = 10$  and  $T_0 = 6$ :

$$S_{(1)} = 0.5 \times 10 + 0.5 \times 6 = 3 + 5 = 8.$$

$$S_{(2)} = 0.5 \times 8 + 0.5 \times 4 = 4 + 2 = 6$$

$$S_{(3)} = ?$$

$$S_{(4)} = ?$$



CPU burst ( $t_i$ )	6	4	6	4	13	13	13	...	
"guess" ( $\tau_i$ )	10	8	6	6	5	9	11	12	...

# Examples of Exponential Averaging

- When  $\alpha = 0$ , then the next estimated burst  $S_{n+1} = (1 - \alpha) S_n + \alpha T_n$  becomes
  - $S_{n+1} = S_n$
  - Recent history does not count
- When  $\alpha = 1$ , then the next estimated burst  $S_{n+1} = (1 - \alpha) S_n + \alpha T_n$  becomes
  - $S_{n+1} = T_n$
  - Only the actual last CPU burst counts

# Example of Shortest-Remaining-Time-First (SRTF)

- **Preemptive** version of SJF is called **shortest-remaining-time-first**
  - The next burst of new process may be shorter than that left of current process
    - Currently running process will be preempted
- Now we add the concepts of varying arrival times and preemption to the analysis

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5


Draw Gantt chart for above mentioned processes:



# Example of Shortest-remaining-time-first (SRTF)

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

- *Preemptive* SJF Gantt Chart



**Waiting Time** = Finish Time – Arrival Time – Burst Time

- **Waiting time** for  $P_1 = (10-1)$ ;  $P_2 = (1-1)$ ;  $P_3 = (17-2)$ ,  $P_4 = (5-3)$
- Average waiting time =  $[(10-1)+(1-1)+(17-2)+(5-3)]/4 = 26/4 = 6.5$  msec

- Nonpreemptive SJF will result in 7.75 time units

# SJF Limitations

- SJF doesn't always minimize average Turnaround Time (time from entering the system till completion)
  - Only minimizes waiting time
- Can lead to unfairness or starvation → newly arriving tasks have shorter bursts
- In practice, can't actually predict the future
- But can estimate CPU burst length based on past

# Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer  $\equiv$  highest priority)
  - We assume that smallest integer  $\equiv$  highest priority
    - Internally defined priority: based on criteria within OS. Ex: memory needs
    - Externally defined priority: based on criteria outside OS. Ex: paid process
  - Preemptive; a running process is preempted by a new higher priority process
  - Nonpreemptive: running process holds CPU until termination or I/O request
- SJF is priority scheduling where priority is the inverse of predicted next CPU burst time

# Priority Scheduling and Starvation

- it is possible for a process with low priority to wait indefinitely to run because a stream of higher priority processes keeps arriving. This is called **starvation**.
- Solution  $\equiv$  This problem can be overcome by using *dynamic priorities*, in which a process's priority changes over time.
- **Aging** – as time progresses increase the priority of the process waiting in the ready queue increases
  - **Example:**

```
do
    priority = priority – 1 //increase priority
every 15 minutes
```
- Conversely, processes that spend too much time on the processor can have their priorities reduced.

# Example of Priority Scheduling

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
$P_1$	10	3
$P_2$	1	1
$P_3$	2	4
$P_4$	1	5
$P_5$	5	2

- Priority scheduling Gantt Chart ( $P_2 > P_5 > P_1 > P_3 > P_4$ )



- Waiting time for  $P_1 = (16-10)$ ;  $P_2 = (1-1)$ ;  $P_3 = (18-2)$ ,  $P_4 = (19-1)$ ,  $P_5 = (6-5)$
- Average waiting time =  $(0+1+6+16+18) = 41/5 = 8.2$  msec

# Round Robin (RR)

- Each process gets a small unit of CPU time (**time quantum or time slice  $q$** ), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
  - Ready queue is treated as a circular queue
- If there are  $n$  processes in the ready queue and the time quantum is  $q$ , then each process gets  $1/n$  of the CPU time in chunks of at most  $q$  time units at once. No process **waits** more than  $(n-1)q$  time units.

# Round Robin (RR)

- Timer interrupts every quantum to schedule next process
- Performance **depends heavily on the size of the time quantum**
  - *If  $q$  very large  $\Rightarrow$  RR scheduling = FCFS scheduling*
  - *If  $q$  very small  $\Rightarrow q$  must be large with respect to context switch,*
    - **Otherwise overhead of number of context switches will be is too high**

# Example of RR with Time Quantum = 4

<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3


- The Gantt chart when time quantum is 4 seconds:
- Average Waiting Time = ?.



# Example of RR with Time Quantum = 4


<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

- The Gantt chart when time quantum is 4 seconds long:




- We assume the arrival time for each process is 0.
- **Waiting time** for  $P_1 = (10-4)$ ;  $P_2 = 4$ ;  $P_3 = 7$ ,
- Average Waiting Time =  $[(10-4)+4+7]/3 = 5.66$  milliseconds.
- **Average waiting time under RR scheduling is often long**
- Typically, higher average turnaround than SJF, but better **response**
- $q$  should be large compared to context switch time
- $q$  usually 10ms to 100ms, context switch < 10 usec

# Time Quantum and Context Switch Time



# Multilevel Queue Scheduling

highest priority



lowest priority

# Multilevel Queue Scheduling

- Ready queue is partitioned into separate queues, e.g: two queues containing
  - **foreground** (interactive) processes
    - May have externally defined priority over background processes
  - **background** (runs without active user interaction)
- Process permanently in a given queue; **no move to a different queue**  
(different than *multilevel feedback queue*, will be covered soon)

# Multilevel Queue Scheduling

- Each queue has its own scheduling algorithm:
  - foreground – RR or ...
  - background – FCFS or ...
- Scheduling must be done **between** the queues:
  - **Fixed priority scheduling;** (i.e., serve all from foreground first then from background).  
Possibility of starvation.
  - **Time slice** – each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e., 80% to foreground in RR
  - And 20% to background in FCFS

# Multilevel Feedback Queue

- A process can move between the various queues; *aging* can be implemented this way
- Demote → Move heavy foreground CPU-bound process to the background queue
- Upgrade → Move starving background process to the foreground queue


# Example of Multilevel Feedback Queue

- Three queues:

- $Q_0$  – RR with time quantum 8 milliseconds
  - Highest priority. Preempts  $Q_1$  and  $Q_2$  proc's
- $Q_1$  – RR time quantum 16 milliseconds
  - Medium priority. Preempts processes in  $Q_2$
- $Q_2$  – FCFS
  - Lowest priority

- Scheduling

- A new job enters queue  $Q_0$  which is served FCFS
  - When it gains CPU, job receives 8 milliseconds
  - If it does not finish in 8 milliseconds, job is moved(demoted) to queue  $Q_1$
- At  $Q_1$  job is again served FCFS and receives 16 additional milliseconds
  - If it still does not complete, it is preempted and moved to queue  $Q_2$



# Chapter 6: CPU Scheduling

- Basic Concepts
- Scheduling Criteria
- Scheduling Algorithms
- Real-Time CPU Scheduling
- Algorithm Evaluation




# Real-Time CPU Scheduling

- A real-time operating system is for **real-time applications** that processes data and events that have critically defined time constraints. For example patient monitoring in ICU, air traffic control, robot control in automated factory.
- In a real-time system the **correctness** of the system behavior depends not only the logical results of the computations, but also on the **time** at which these results are produced.
- **Soft real-time systems** – In these type of systems, an associated deadline is with the soft-real time task is desirable but not mandatory, i.e., missing a deadline is acceptable.
- **Hard real-time systems** – These type of OS strictly adhere to the **deadline** associated with the tasks. Missing on a deadline can have catastrophic affects.

# Priority-based Scheduling

- Real-time OS responds immediately to a real-time process when it requests CPU
  - For real-time scheduling, scheduler must support preemptive, priority-based scheduling
- For **hard real-time**, it must also provide ability to meet deadlines
- Processes have new characteristics:
  - **periodic** processes require CPU at constant intervals
    - Has processing time  **$t$** , deadline  **$d$** , period  **$p$**
    - $0 \leq t \leq d \leq p$
    - **Rate** of periodic task is  $1/p$

Admission-control: process announces its requirements, then scheduler admits the process if it can complete




# 1. Rate Monotonic Scheduling

- **Static priority** policy with preemption
  - A priority is assigned based on the inverse of its period  $p$ 
    - Shorter periods = higher priority; Longer periods = lower priority
    - Assumes processing time  $t$  is the same for each CPU burst
    - Rationale: higher priority to tasks that require the CPU more often
- Example:
  - Process  $P_1$  :  $t_1 = 20$ ,  $d_1 =$  complete CPU burst by start of next period,  $p_1 = 50$ .
  - Process  $P_2$  :  $t_2 = 35$ ,  $d_2 =$  complete CPU burst by start of next period,  $p_2 = 100$ .

# 1. Rate Monotonic Scheduling

- Example:

- Process  $P_1$  :  $t_1 = 20$ ,  $d_1 =$  complete CPU burst by start of next period,  $p_1 = 50$ .
- Process  $P_2$  :  $t_2 = 35$ ,  $d_2 =$  complete CPU burst by start of next period,  $p_2 = 100$ .
- $P_1$  is assigned a higher priority than  $P_2$ . **Why?**
- CPU utilization =  $t_i / p_i$ . Thus total CPU utilization =  $20/50 + 35/100 = 75\%$
- In this example both processes meet the deadlines successfully

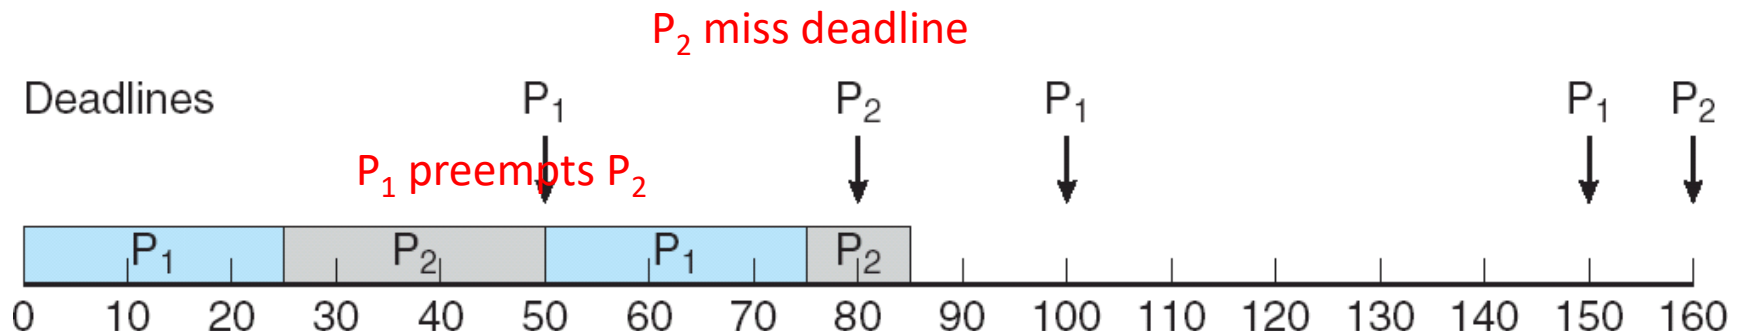


## Problem with Rate Monotonic Scheduling

- Let's change process  $P_1$  processing time from 20 to 25
- **Process  $P_1$**  :  $t_1 = 25$ ,  $d_1 =$  complete CPU burst by start of next period,  $p_1 = 50$ .
- **Process  $P_2$**  :  $t_2 = 35$ ,  $d_2 =$  complete CPU burst by start of next period,  $p_2 = 80$ .
- Total CPU utilization =  $25/50 + 35/80 = 94\%$  (in previous example  $U = 75\%$ )
- $P_1$  is assigned a higher priority than  $P_2$ . **Why?**


## Problem with Rate Monotonic Scheduling: Missed Deadlines

- Let's change process  $P_1$  processing time from 20 to 25 and period to 80
- Process  $P_1$  :**  $t_1 = 25$ ,  $d_1 =$  complete CPU burst by start of next period,  $p_1 = 50$ .
- Process  $P_2$  :**  $t_2 = 35$ ,  $d_2 =$  complete CPU burst by start of next period,  $p_2 = 80$ .
- Total CPU utilization =  $25/50 + 35/80 = 94\%$  (in previous example  $U = 75\%$ )
- $P_1$  is assigned a higher priority than  $P_2$ . **Why?**
- $P_2$  has missed the deadline for completion of its CPU burst at time 80**
- Rate Monotonic Scheduling has limitation.**



# Problem with Rate Monotonic Scheduling: Missed Deadlines

- Rate Monotonic Scheduling has limitation.
- Total CPU utilization =  $25/50 + 35/80 = 94\%$
- Reason - bounded CPU utilization; worst case for N processes:  $B = N(2^{1/N} - 1)$ 
  - Bounded at  $B = 83\%$  for  $N = 2$  processes
  - **Theory:** cannot guarantee that processes can be scheduled to meet their deadlines if actual CPU utilization > Bound B from **UB test**
    - $94\% > 83\%$  in the example above



# Class Exercise

Process	Execution Time	Period
P1	1	8
P2	2	5
P3	2	10

- What is the utilization of the system in the above scenario?
- According to UB test, what is the sufficient condition under which the system is schedulable?
- Is the system schedulable??



# Class Exercise

Process	Execution Time	Period
P1	1	8
P2	2	5
P3	2	10

- What is the utilization of the system in the above scenario?
  - 0.725
- According to UB test, what is the sufficient condition under which the system is schedulable?
  - 0.7797
- Is the system schedulable??
  - Yes:  $0.725 < 0.7797$


## 2. Earliest Deadline First Scheduling (EDF) (Solution)

- *Dynamic priority* policy with preemption
  - Priorities are dynamically assigned according to deadlines:
    - **Earlier deadline = higher priority; later deadline = lower priority**
  - New runnable process must announce its deadline requirements to scheduler
    - Scheduler will adjust current priorities accordingly
  - Process need not be periodic
  - CPU burst time need not be constant

## 2. Earliest Deadline First Scheduling (EDF) (Solution)

- *Example:*

- Process  $P_1$  :  $t_1 = 25$ ,  $d_1 = \text{complete CPU burst by start of next period}$ ,  $p_1 = 50$ .
- Process  $P_2$  :  $t_2 = 35$ ,  $d_2 = \text{complete CPU burst by start of next period}$ ,  $p_2 = 80$ .
- $P_1$  is initially assigned a higher priority than  $P_2$  since  $d_1 = 50 \leq d_2 = 80$
- But later  $P_2$  is assigned higher priority than  $P_1$ .



## 2. Earliest Deadline First Scheduling (EDF) (Solution)

- *Exercise:*


- Process  $P_1$  :  $t_1 = 25$ ,  $d_1 =$  complete CPU burst by start of next period,  $p_1 = 50$ .
- Process  $P_2$  :  $t_2 = 30$ ,  $d_2 =$  complete CPU burst by start of next period,  $p_2 = 75$ .

-

## 2. Earliest Deadline First Scheduling (EDF) (Solution)

- *Exercise:*

- Process  $P_1$  :  $t_1 = 25$ ,  $d_1 = \text{complete CPU burst by start of next period}$ ,  $p_1 = 50$ .
- Process  $P_2$  :  $t_2 = 30$ ,  $d_2 = \text{complete CPU burst by start of next period}$ ,  $p_2 = 75$ .
- $P_1$  is initially assigned a higher priority than  $P_2$  since  $d_1 = 50 \leq d_2 = 75$
- But later  $P_2$  is assigned higher priority than  $P_1$ .
- Priorities are dynamically changed.



### 3. Proportional Share Scheduling

- $T$  shares are allocated among all processes in the system
- An application receives  $N$  shares where  $N < T$
- This ensures each application will receive  $N / T$  of the total processor time
- **Example: three processes A, B and C, with  $T = 100$  shares**
  - A, B and C assigned each 50, 15 and 20 shares, respectively
  - Thus A will have 50% of total processor times, and so on with B and C
- **Scheduler must use admission-control policy:**
- **Admission-control:** process announces its requirements, then scheduler admits the process if it can complete it on time, or, reject it if it cannot be serviced
  - Admit a request only if sufficient shares are available
  - **Example: If new process D needs 30 share then scheduler will deny him CPU**
    - Current total share is  $50 + 15 + 20 = 85$
    - Only a new process with share  $< 15$  can be scheduled

# Chapter 6: CPU Scheduling

- Basic Concepts
- Scheduling Criteria
- Scheduling Algorithms
- Real-Time CPU Scheduling
- Algorithm Evaluation

# Algorithm Evaluation

- How to select CPU-scheduling algorithm for an OS?
- Determine criteria (e.g., throughput, utilization etc.), then evaluate algorithms
- **Deterministic modeling**
  - Type of **analytic evaluation**
  - Takes a particular predetermined workload and defines the performance of each algorithm for that workload


- Consider 5 processes arriving at time 0:

Process	Burst Time
$P_1$	10
$P_2$	29
$P_3$	3
$P_4$	7
$P_5$	12


- FCFS is 28ms:



- Non-preemptive SJF is 13ms:



- RR is 23ms:





## (2/4) Queueing Models

- Describes the arrival of processes, and CPU and I/O bursts probabilistically
- Computer system described as *network of servers*, each with queue of waiting processes
  - Commonly exponential, and described by mean
  - Knowing arrival rates and service rates
  - Computes utilization, average queue length, average wait time, etc

# (3/4) Simulations

- Queueing models limited
- **Simulations** more accurate
  - Programmed model of computer system
  - Clock is a variable
  - Gather statistics indicating algorithm performance
  - Data to drive simulation gathered via
    - Random number generator according to probabilities
    - Distributions defined mathematically

## (4/4) Implementation

- Even simulations have limited accuracy
- Just implement new scheduler and test in real systems
  - High cost, high risk
  - Environments vary

End of Chapter 6