

Chapter 8: Main Memory

Chapter 8: Memory Management

- **Background**

- Swapping
- Contiguous Memory Allocation
- Segmentation
- Paging
- Structure of the Page Table

Objectives


- To increase CPU utilization and response speed:
 - Several processes must be kept in main memory
 - We must share memory
 - How to manage main memory resource...?
- To provide a detailed description of various ways of organizing memory hardware
- To discuss various memory-management techniques, including *paging* and *segmentation*

Background

- Program must be brought (from disk) into memory and be placed within a process for it to be run
- Main memory and registers are the only storage CPU can access directly
- Memory unit only sees a stream of
 - addresses + read requests, or address + data and write requests

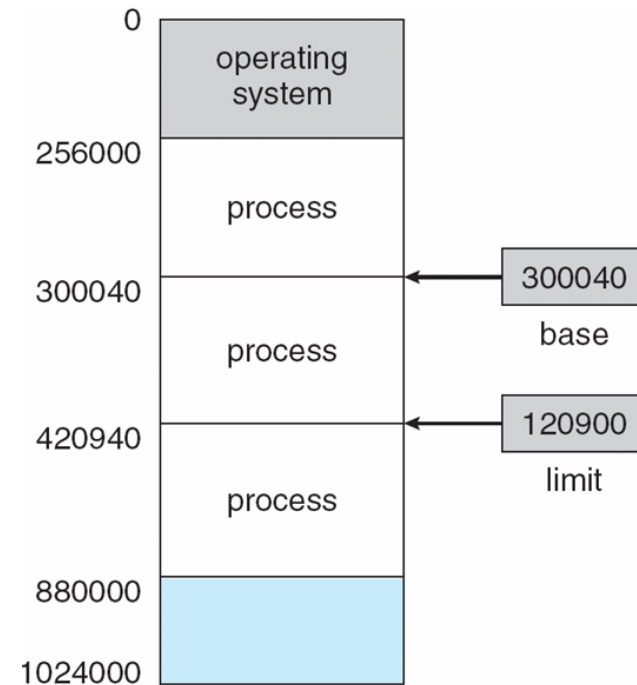
CPU, MMU and Memory

- Register access in one CPU clock (or less); **very fast memories**
- Main memory can take many cycles, causing a **stall**
- **Cache** sits between main memory and CPU registers; **solution to stall issue**
- Protection of memory required to ensure correct operation; **hardware-level**



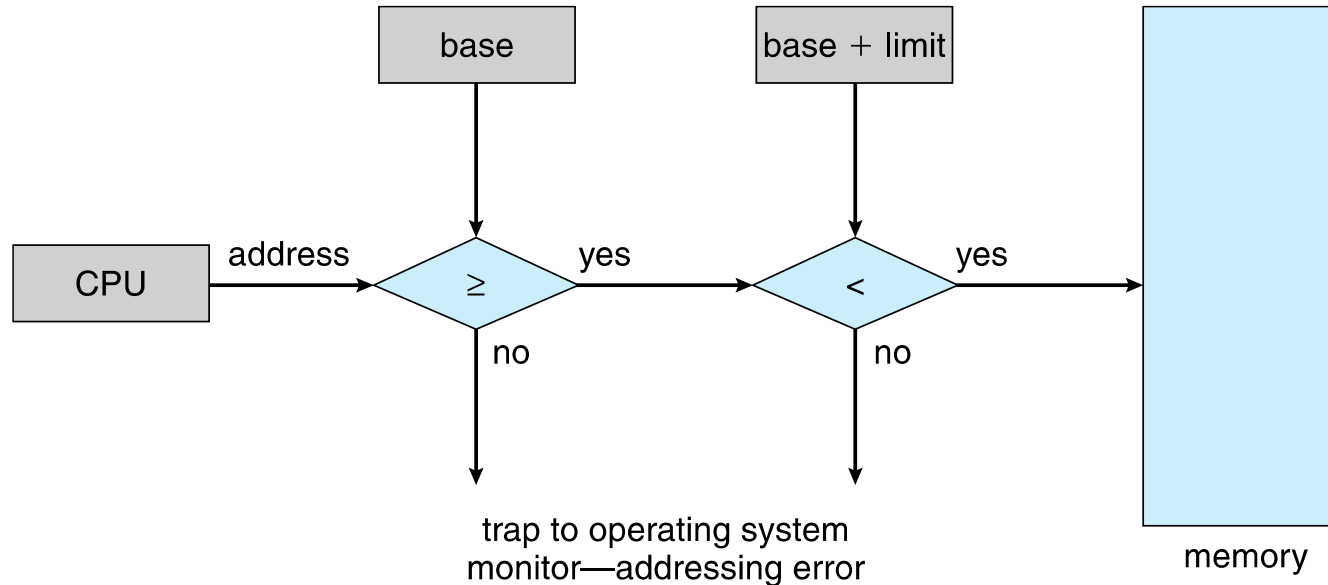
Protection of Memory: Base and Limit Registers

- Processes should not be able to *reference* memory locations in another process without permission.
- A pair of **base** and **limit registers** define the logical address space
 - To protect processes from each other ; each process has its own memory space
 - Base register** holds the smallest legal physical memory address
 - Limit register** specifies the size of the range of accessible addresses



Hardware Address Protection

- ❑ CPU must check every memory access generated in user mode to be sure it is between base and limit for that user; **to protect a process's memory space**



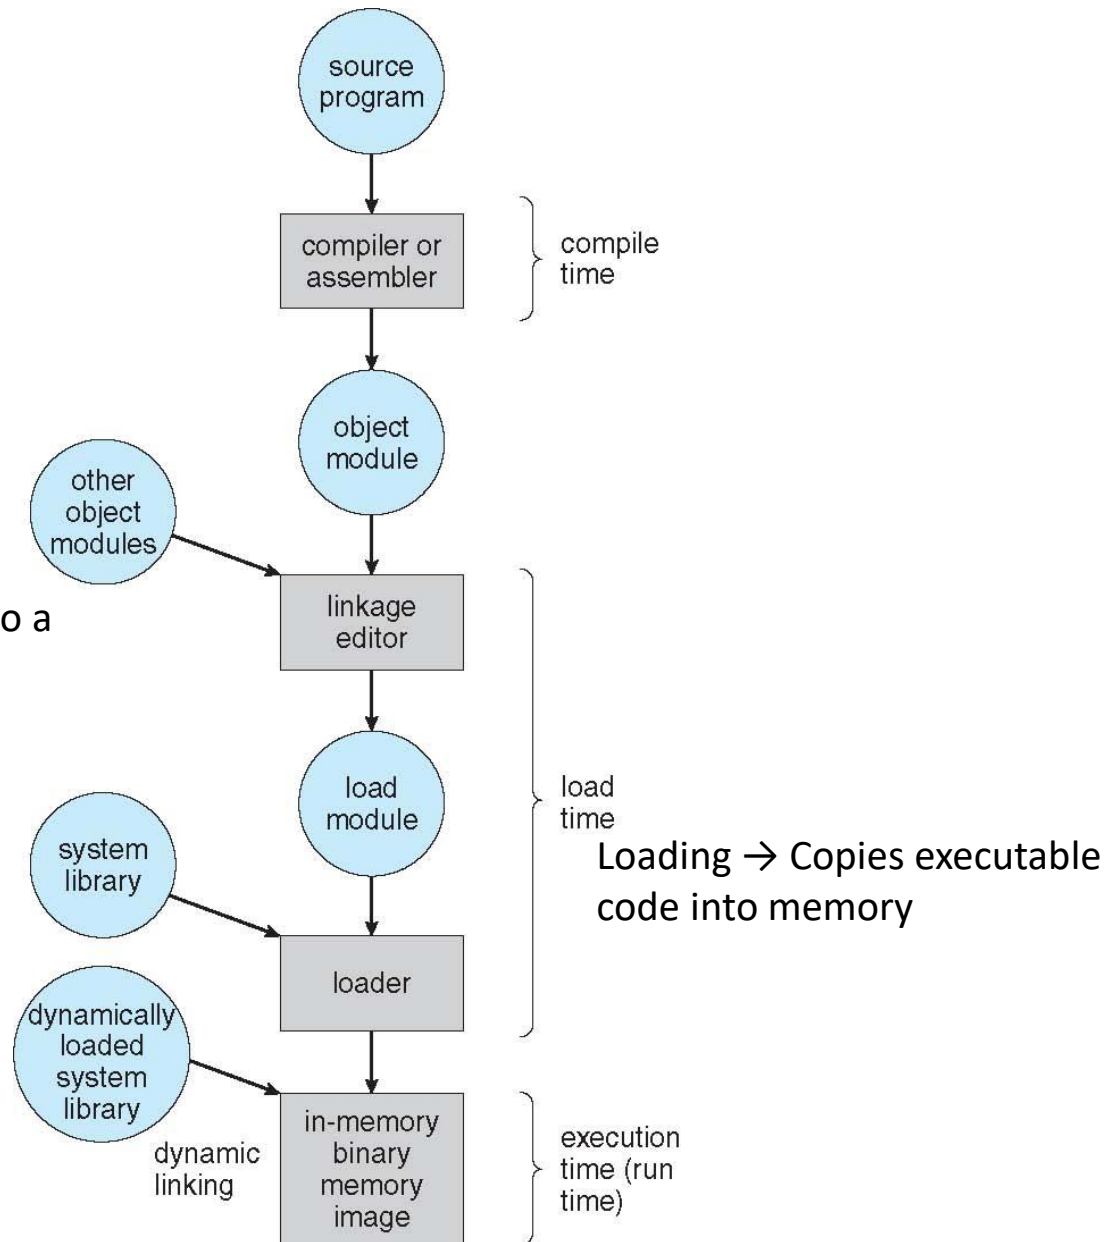
- ❑ Base and limit registers loaded only by the OS through a privileged instruction
 - To prevent users from changing these registers' contents
- ❑ OS has unrestricted access to OS memory and user memory
 - Load users' programs into users' memory, ... etc
 - Access and modify system-calls' parameter, ... etc

Multistep Processing of a User Program

Compiler → generates object code

Linker → Combines the Object code into a single self sufficient executable code

Execution → dynamic memory allocation



Address Binding

- Programs on disk, ready to be brought into memory to execute form an **input queue**
 - Without support, must be loaded into address 0000
- Inconvenient to have first user process physical address always at 0000
- Further, addresses represented in different ways at different stages of a program's life
- Source code addresses usually symbolic. They are **names of variable**, e.g., *count*
 - Compiled code addresses **bind** to ***relocatable addresses***
 - i.e. "14 bytes from beginning of this module"
 - Linker or loader will bind ***relocatable addresses*** to ***absolute addresses***
 - i.e. 74014
 - Each binding maps one address space to another

Binding of Instructions and Data to Memory

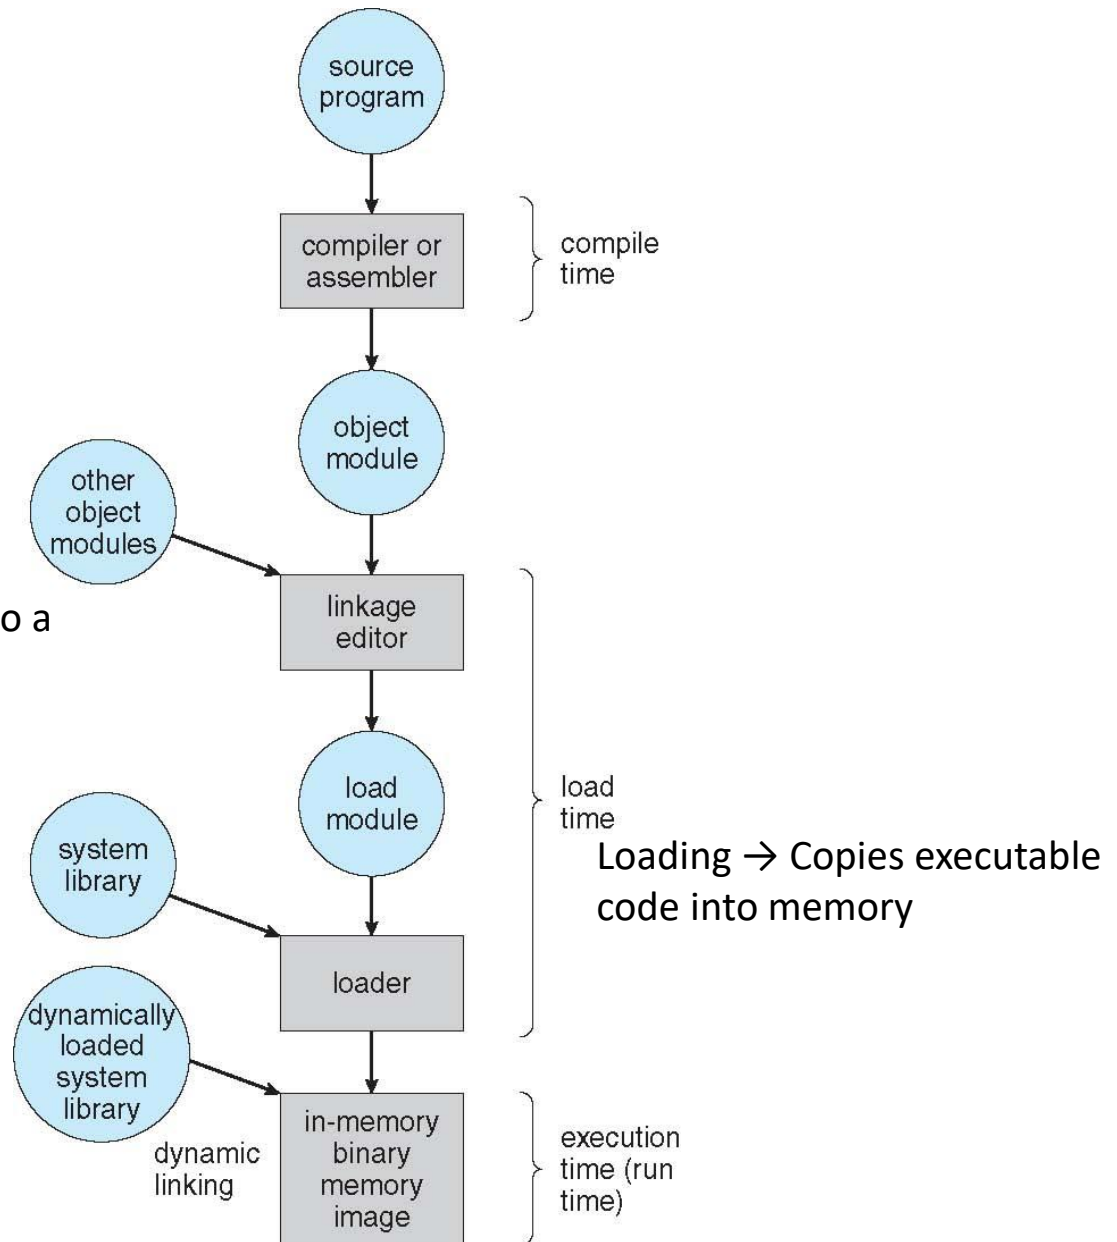
- Address binding of instructions and data to memory addresses can happen at three different stages
 - **Compile time:** If it is known at compile time where a program will reside in physical memory, then *absolute code* can be generated by the compiler, containing actual physical addresses.
 - **Load time:** If the location at which a program will be loaded is not known at compile time, then the compiler must generate *relocatable code*, which references addresses *relative* to the start of the program. If that starting address changes, then the program must be reloaded but not recompiled.
 - **Execution time:** Binding delayed until run time if the process can be moved during its execution from one memory segment to another
 - Need hardware support for address maps (e.g., base and limit registers)

Multistep Processing of a User Program

Compiler → generates object code

Linker → Combines the Object code into a single self sufficient executable code

Execution → dynamic memory allocation

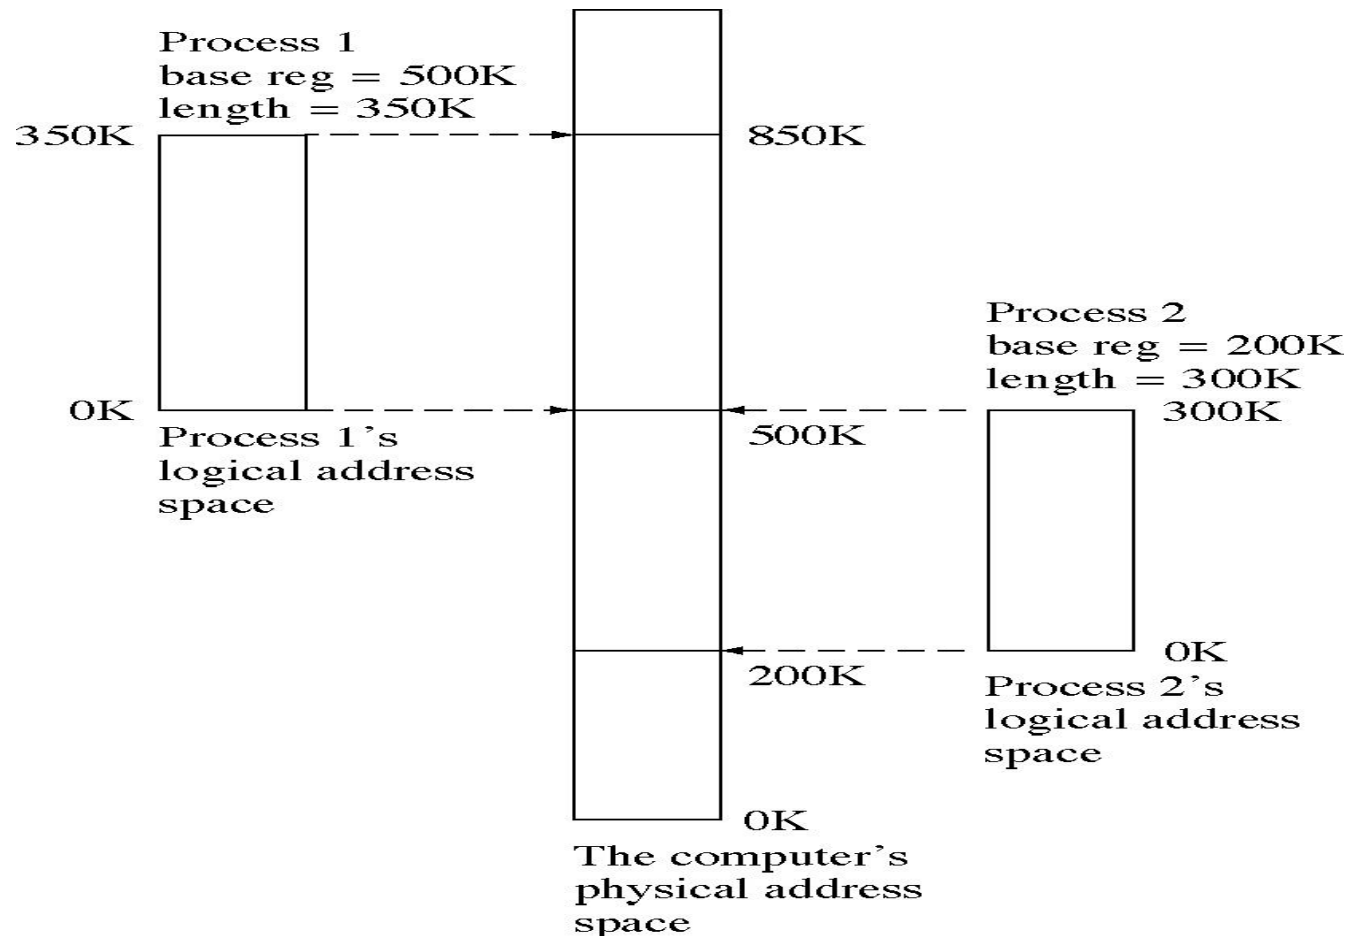


Logical vs. Physical Address Space

- The concept of a **logical address space** that is bound to a separate **physical address space** is central to proper memory management
 - **Logical address** – generated by the CPU; also referred to as **virtual address**
 - **Physical address** – address seen by the **memory-management unit** in its **memory-address register**
- The set of all logical addresses generated by a program is a **logical address space**. The set of all physical addresses corresponding to these logical addresses is a **physical address space**.

Logical and Physical Address Spaces

- Assume there are two processes, process 1 and process 2:
- Process 1's length is 350K and Process's 2 length is 300K




Memory-Management Unit (MMU)

- MMU is the hardware that does virtual-to-physical addr mappings at run-time
- Hardware device that at run-time maps virtual address to physical address
- Many methods possible, covered in the rest of this chapter
- To start, consider simple scheme where the value in the *relocation register* is added to every address generated by a user process at the time it is sent to memory
 - Base register now called *relocation register*;
- The user program deals with *logical addresses*; it never sees the *real physical addresses*
 - Logical address mapped to physical addresses before use
 - Logical addresses: in range **0 to max**
 - Physical addresses: in range **$R + 0$ to $R + \text{max}$**
 - For a base value R in the relocation register

Memory Management Unit

- The runtime mapping from virtual \rightarrow physical address




- Relocation register is added to every address generated by user process

Chapter 8: Memory Management

- Background
- **Swapping**
- Contiguous Memory Allocation
- Segmentation
- Paging
- Structure of the Page Table

Swapping


- A process can be **swapped** temporarily out of memory to a **backing store**, and then brought back into memory for continued execution
 - Total physical memory space of processes can exceed physical memory
 - Hence, swapping increases the degree of multiprogramming
- Backing store – fast disk large enough to accommodate copies of all memory images for all processes; must provide direct access to these memory images



Schematic View of Swapping

■ System maintains a ready queue of ready-to-run processes which have memory images on the backing store or in memory. **Standard swapping method:**

- Dispatcher called when the CPU scheduler selects a process P_2 from queue
 - ▶ If not enough free space in memory for P_1
 - Swap in P_2 from backing store and swap out some P_1 from memory
 - Reload registers and transfer control to P_2 ; i.e. P_2 is now in running state



Context Switch Time including Swapping


- If next processes to be put on CPU is not in memory, need to swap out a process and swap in target process
- Context switch time can then be very high
- 100MB process swapping to hard disk with transfer rate of 50MB/sec
 - Swap out time of 2000 ms = 2seconds
 - Plus swap in of same sized process
 - Total context switch swapping component time of 4000ms (4 seconds)
- **We have ignored other components of context-switch time and ignored other disk performance aspect. But, the major part of swap time is transfer time**
 - Can reduce if reduce size of memory swapped – by knowing how much memory really being used

Chapter 8: Memory Management

- Background
- Swapping
- Contiguous Memory Allocation
- Segmentation
- Paging
- Structure of the Page Table

Contiguous Memory Allocation

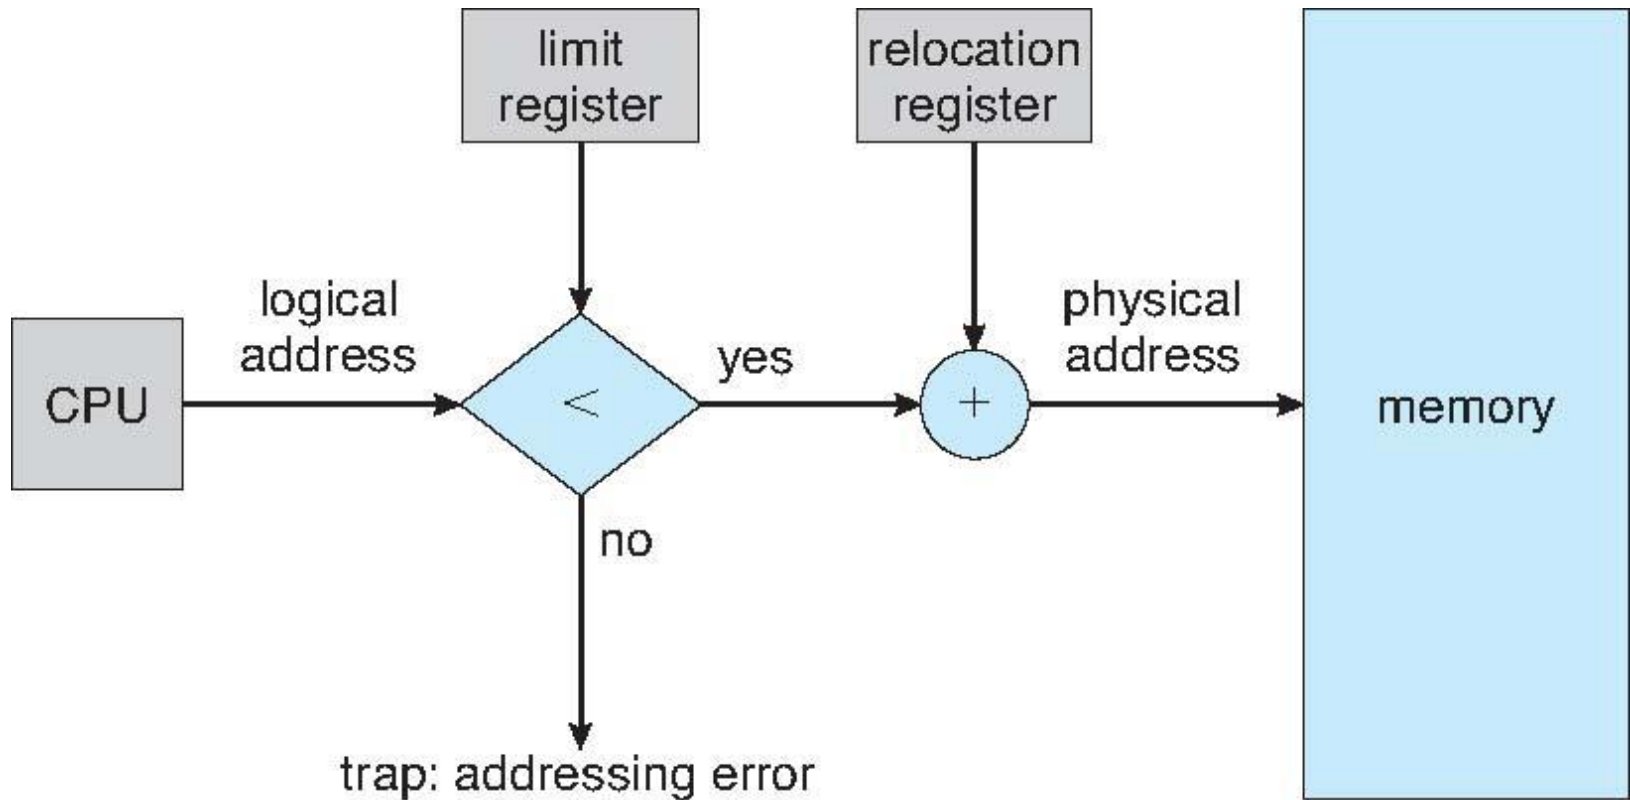
- Main memory must support both OS and user processes
- Limited resource, must allocate efficiently
- Contiguous allocation is one early method
- Main memory usually divided into two **partitions**:
 - Resident operating system usually held in low memory
 - User processes then held in high memory
 - Each process contained **in single contiguous** section of memory



Contiguous Allocation-memory protection



- We can protect a process's memory by combining ideas from previous slides
- **Relocation registers** used to protect user processes from each other, and from changing operating-system code and data
 - For each process:
 - **Relocation register** contains value of smallest physical address;
 - **Limit register** contains range of logical addresses
 - each logical address must be less than the limit register ;
 - MMU maps logical address *dynamically*
 - By adding logical address to relocation value;
 - Dispatcher always loads these two registers with their correct values
 - All logical addresses are checked against Limit register
 - Thus, protecting both each OS and each user program and data

Hardware Support for Relocation and Limit Registers



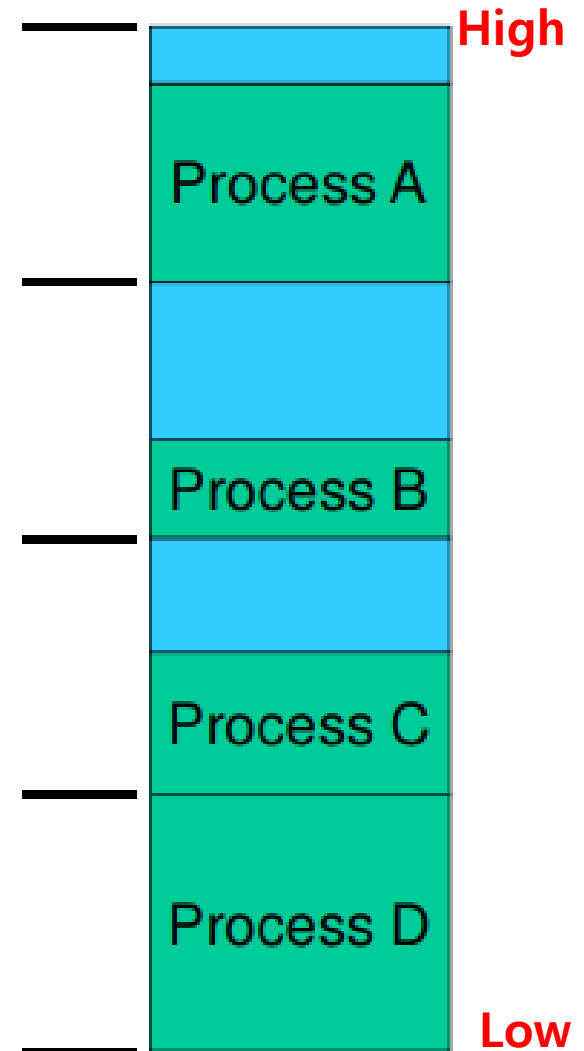
Multiple-partition allocation

- Multiple-partition allocation
 - Degree of multiprogramming limited by number of partitions
 - Example: 2 approaches
 - Fixed Partition
 - equal
 - unequal size partitions
 - Variable Partition (or Dynamic Partitioning)



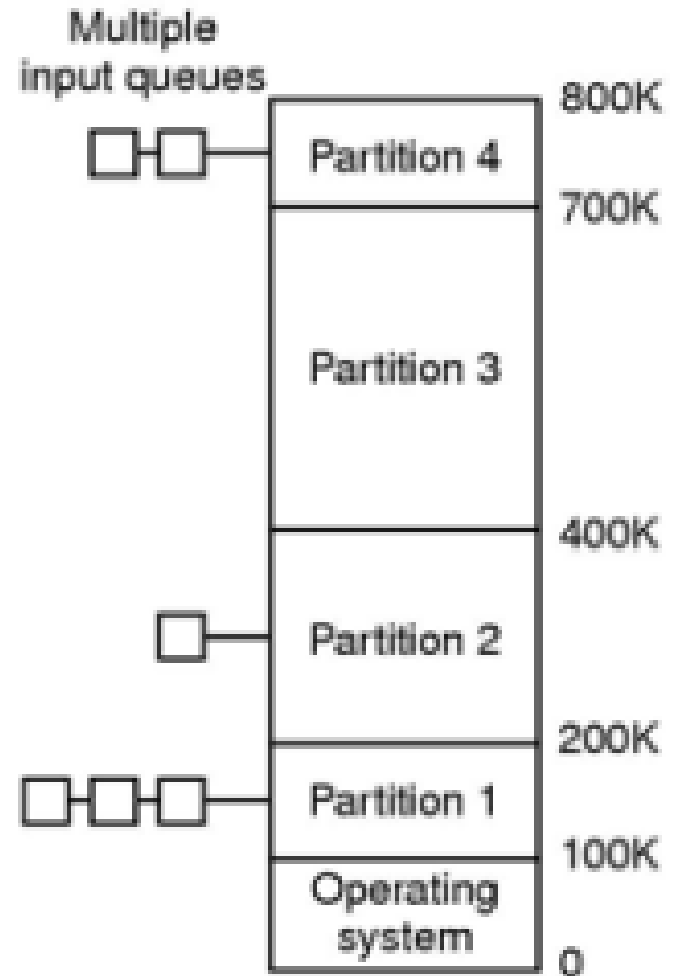
Old Technique#1: Fixed, Equal Partitions

- Physical memory is broken up into fixed, equal partitions
 - hardware requirement: **base/relocation register**, **limit register**
 - physical address** = logical address + base register
 - base register loaded by OS when it switches to a process
- All partitions are of equal size, small processes waste space, Large processes don't fit in (Solution: Fixed Unequal Partitions)



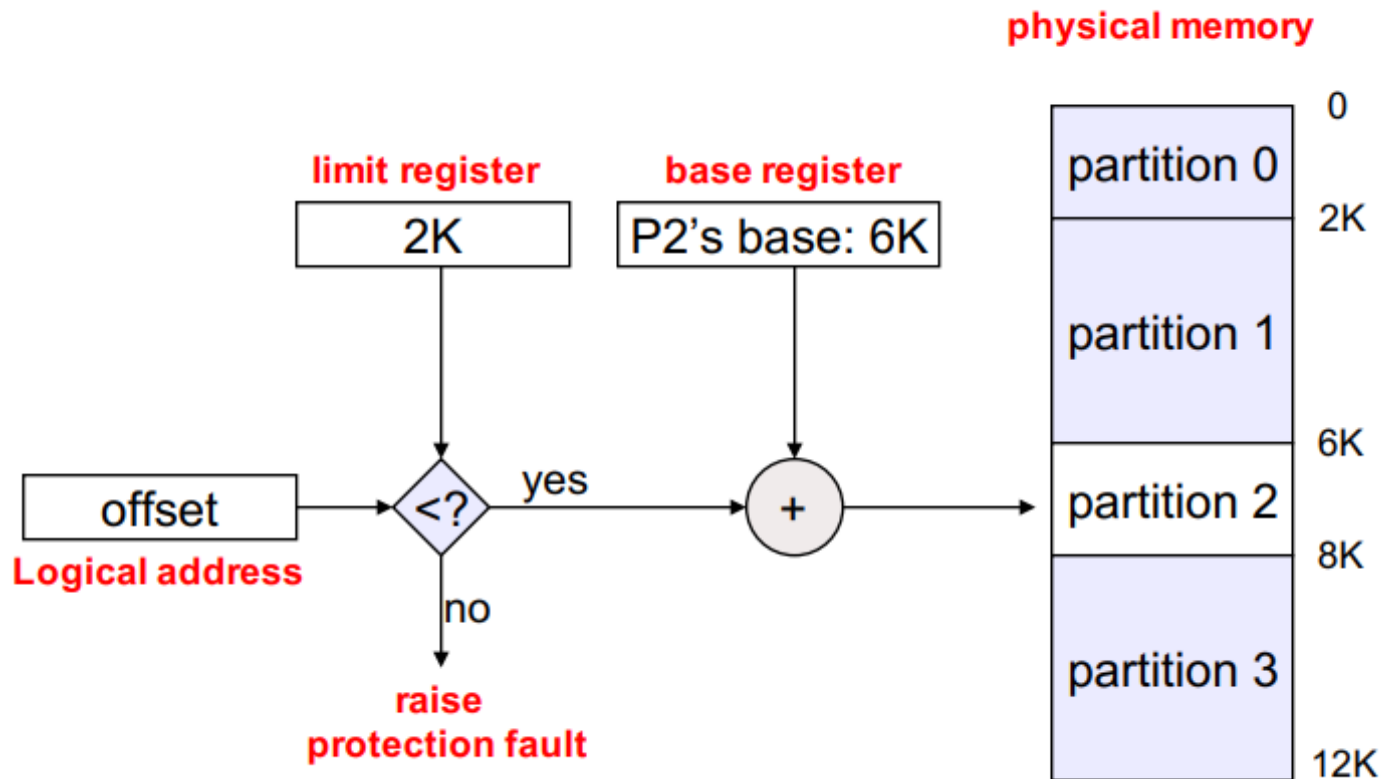
Old Technique#1: Fixed, Unequal Partitions

- Physical memory is broken up into fixed but unequal partitions
- Multiple Queues:
 - Place process in queue for smallest partition that it fits in.




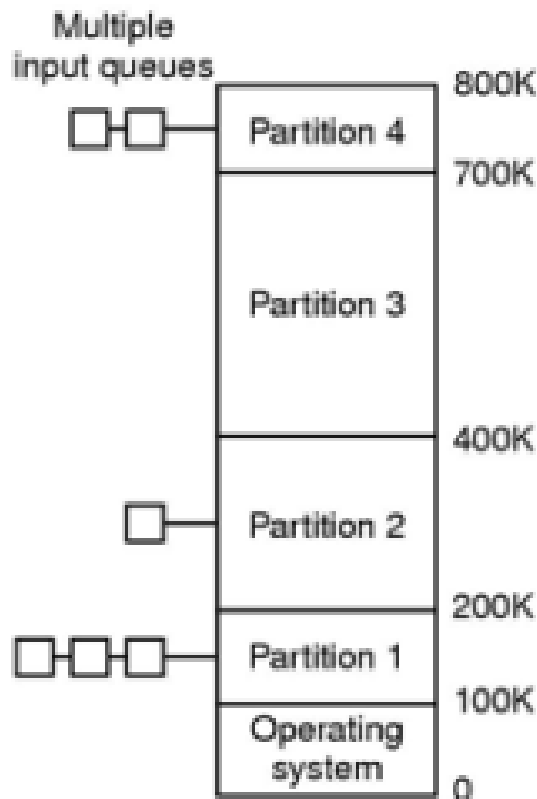
Mechanisms of Fixed Partitions (Equal, Unequal)

- Logical address is generated by the CPU.
- Suppose Process 2 is resides in partition 2.
- Limit register for process 2 is 2k and Base = 6k..
- Offset is displacement inside the user process. E.g., 1,2,3,4,..upto max 2K for the Process 2 in the below example.



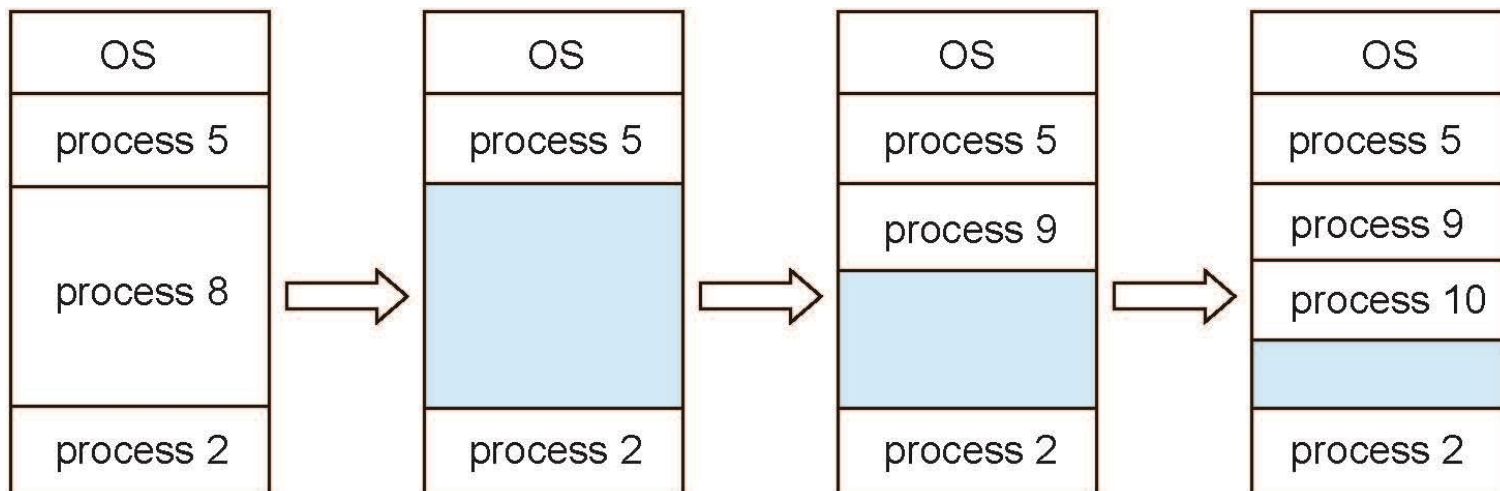
Internal Fragmentation in Old Technique#1: Fixed Partitions

- Advantages of Fixed Partitions (Equal or Unequal)
 - Simple
- Problems
 - **internal fragmentation**: the available partition is larger than what was requested. ; this size difference is memory internal to a partition, but not being used

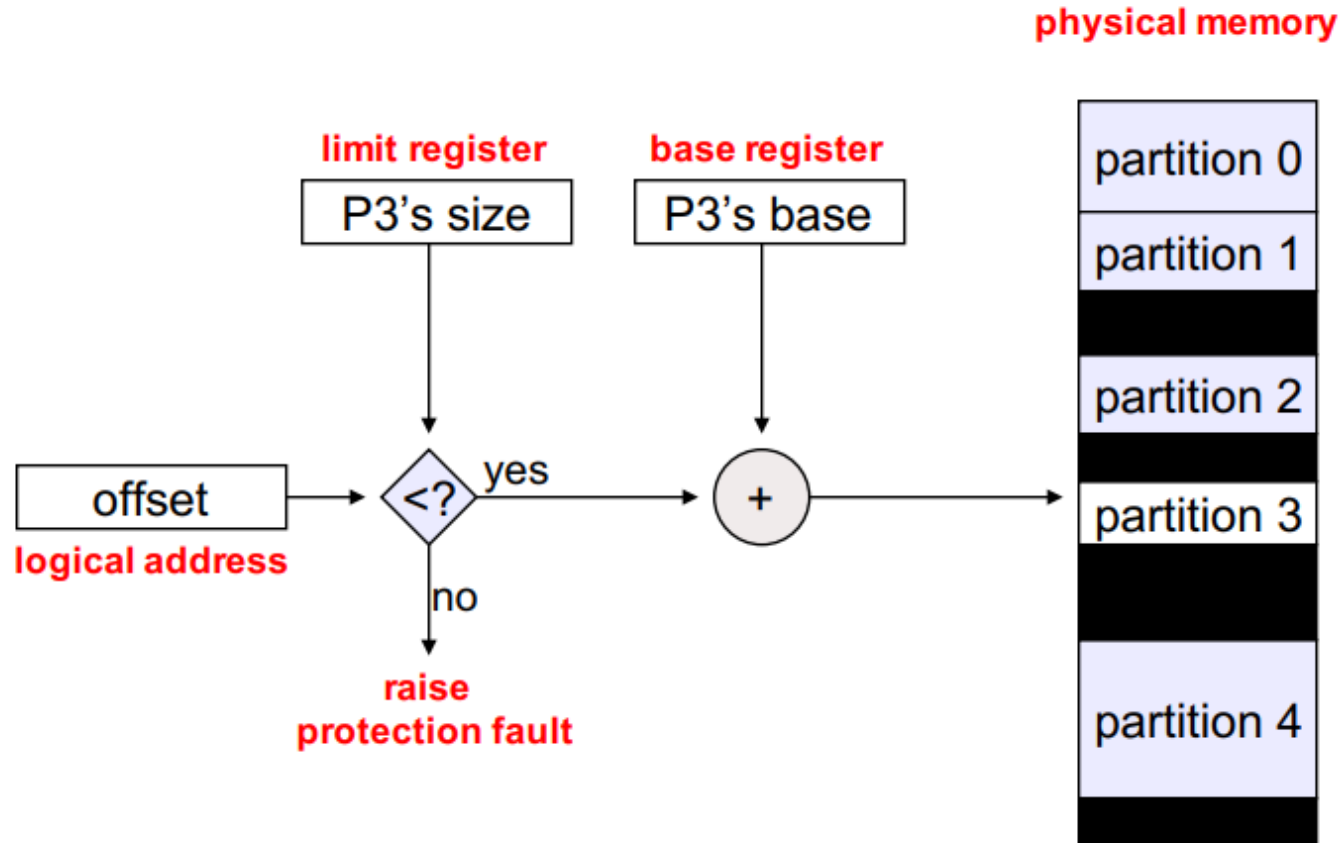


Old Technique#2: Variable Partitions (Dynamic Partitioning)

- Obvious next step: physical memory is broken up into partitions **dynamically** – partitions are tailored to processes
 - hardware requirements: **base register, limit register**
 - **physical address** = logical address + base register
- Advantages: no **internal fragmentation**
 - simply allocate partition size to be just big enough for process (assuming we know what that is!)
- Problems: **external fragmentation**
 - as we load and unload jobs, holes are left scattered throughout physical memory

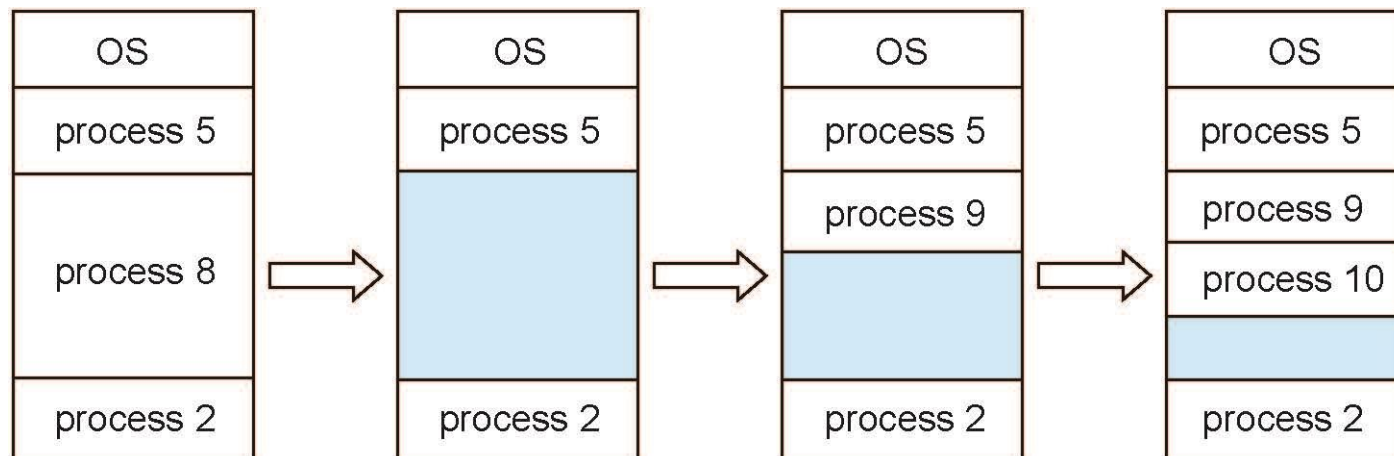


Mechanisms of variable Partitions



Multiple-partition allocation

- Multiple-partition allocation
 - **Variable-partition** sizes for efficiency (sized to a given process' needs)
 - **Hole** – block of available memory; holes of various size are scattered throughout memory
 - When a process arrives, it is allocated memory from a hole large enough to accommodate it
 - Process exiting frees its partition, adjacent free partitions combined
 - Operating system maintains information about:
 - a) allocated partitions b) free partitions (hole)**




Dynamic Storage-Allocation Problem

How to satisfy a request of size n from a list of free holes?

1. **First-fit**: Allocate the *first* hole that is big enough
2. **Next-fit**: Allocate the *next* available hole that is big enough
3. **Best-fit**: Allocate the *smallest* hole that is big enough; must search entire list, unless ordered by size
 - Produces the smallest leftover hole
4. **Worst-fit**: Allocate the *largest* hole; must also search entire list
 - Produces the largest leftover hole
 - Which may be more useful than the smaller leftover hole from a *best-fit* approach;
- First-fit and best-fit better than worst-fit in terms of speed and storage utilization
 - First-fit is generally the fastest


Dynamic Storage Allocation

- Suppose a new process requests 16 MB then which partition each would *best fit*, *next-fit* and *first-fit* algorithm pick?



Dynamic Storage Allocation

- **16 MB:** Which partition each would best fit, next-fit and first-fit algorithm pick?



Fragmentation

- **Internal Fragmentation** – allocated memory may be slightly larger than requested memory; this size difference is **memory internal to a partition**, but not being used
 - Example: a process requests 18,462 bytes and is allocated memory hole of 18,464 bytes; we are then left with a hole of 2 bytes
- **External Fragmentation** – total memory space exists to satisfy a request, but it is not contiguous
 - Free memory space is broken into pieces as processes are loaded and removed from memory
 - Statistical analysis of first-fit reveals that given N allocated blocks, $0.5N$ blocks will be lost to fragmentation

External Fragmentation

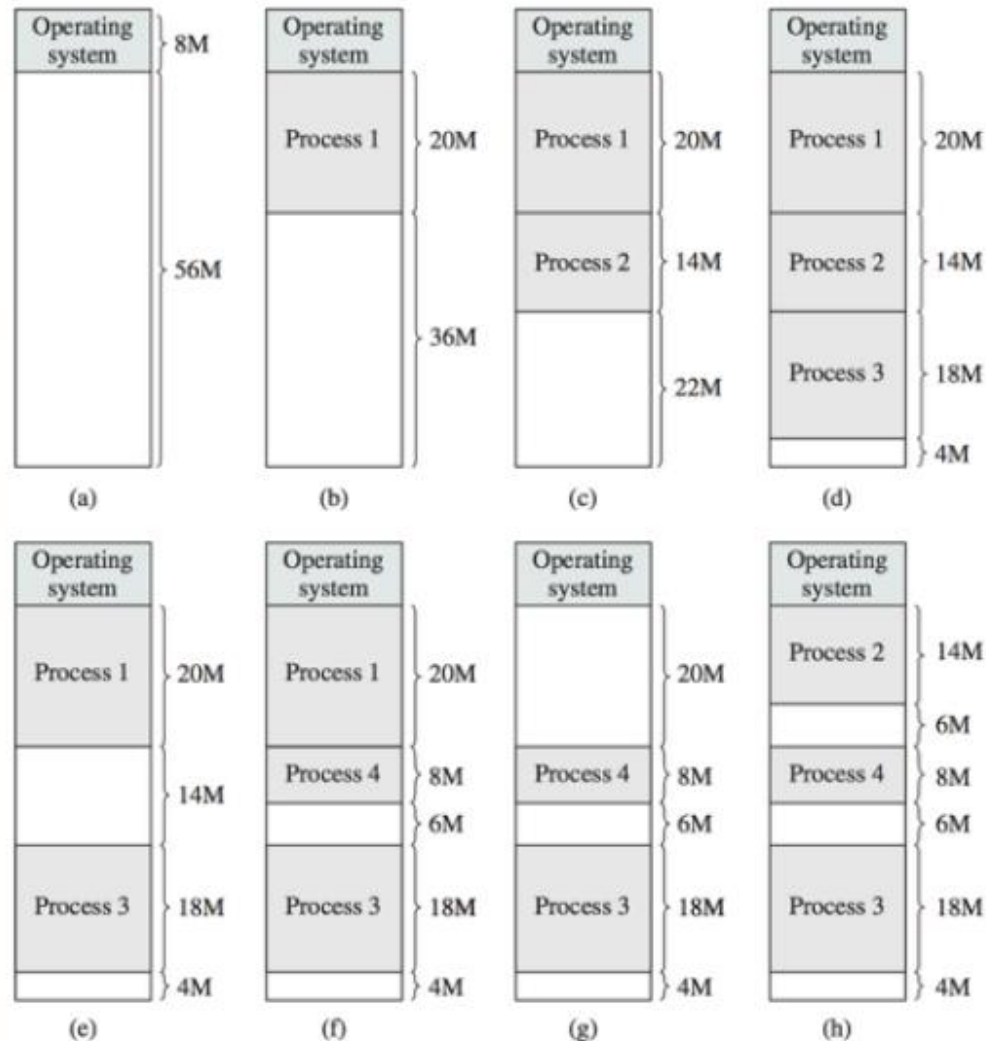
- In the end, a new process larger than 6MB and less than 16 cannot be allocated memory

Effect of Dynamic Partitioning

Process 1 needs 20M of memory, 2 needs 14, 3 needs 18, and 4 needs 8M.

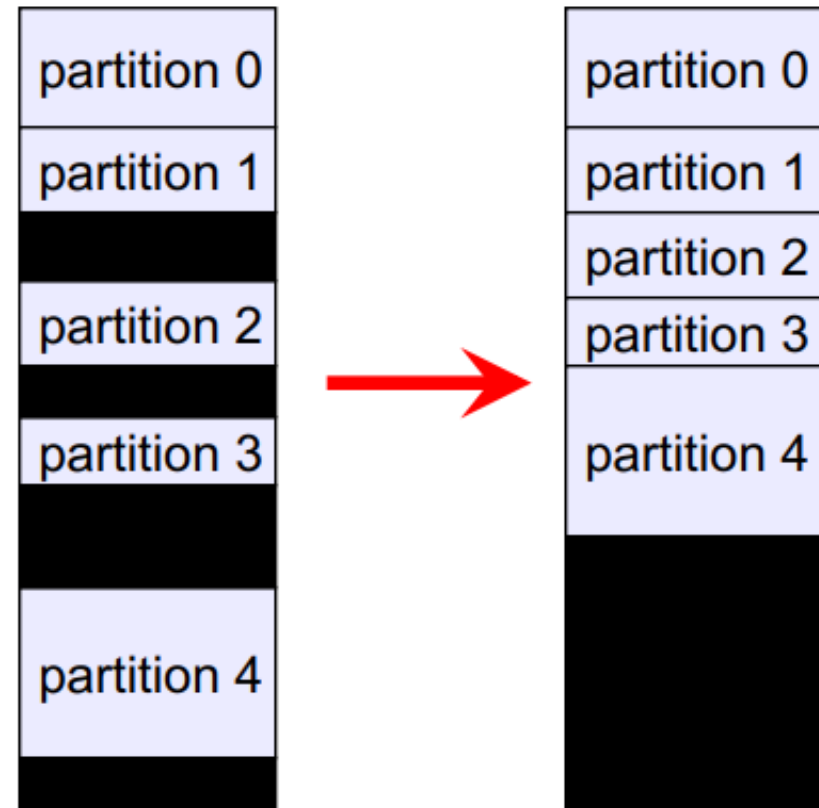
External Fragmentation:

A lot of small holes at the end!



Dealing with Fragmentation

- Reduce **external fragmentation** by **compaction**
 - Shuffle memory contents to place all free memory together in one large block
 - Compaction is possible *only* if relocation is dynamic, and is done at execution time



Dealing with Fragmentation

- Other solutions to external fragmentation problems: **Segmentation and Paging**
 - Permit the logical address space of each process to be non-contiguous
 - Process can be allocated physical memory wherever it is available


Chapter 8: Memory Management

- Background
- Swapping
- Contiguous Memory Allocation
- Segmentation
- Paging
- Structure of the Page Table

Segmentation

- Segmentation **eliminates** the need for continuous addresses
- A program is a collection of segments
 - A segment is a logical unit such as:


main program, procedure , function, method, object, local variables, global variables, common block, stack, symbol table, arrays



logical address


Segmentation: Addressing Mechanism

- Addressing consists of two parts:
 - segment number
 - an offset
- Because of the use of unequal-size segments, segmentation is similar to *dynamic partitioning*
- A natural extension of variable-sized partitions
 - *variable-sized partition* = 1 segment/process
 - *segmentation* = many segments/process




Logical View of Segmentation

- Segment 1, Segment 2... etc. are basically main program, stack, subroutine etc.



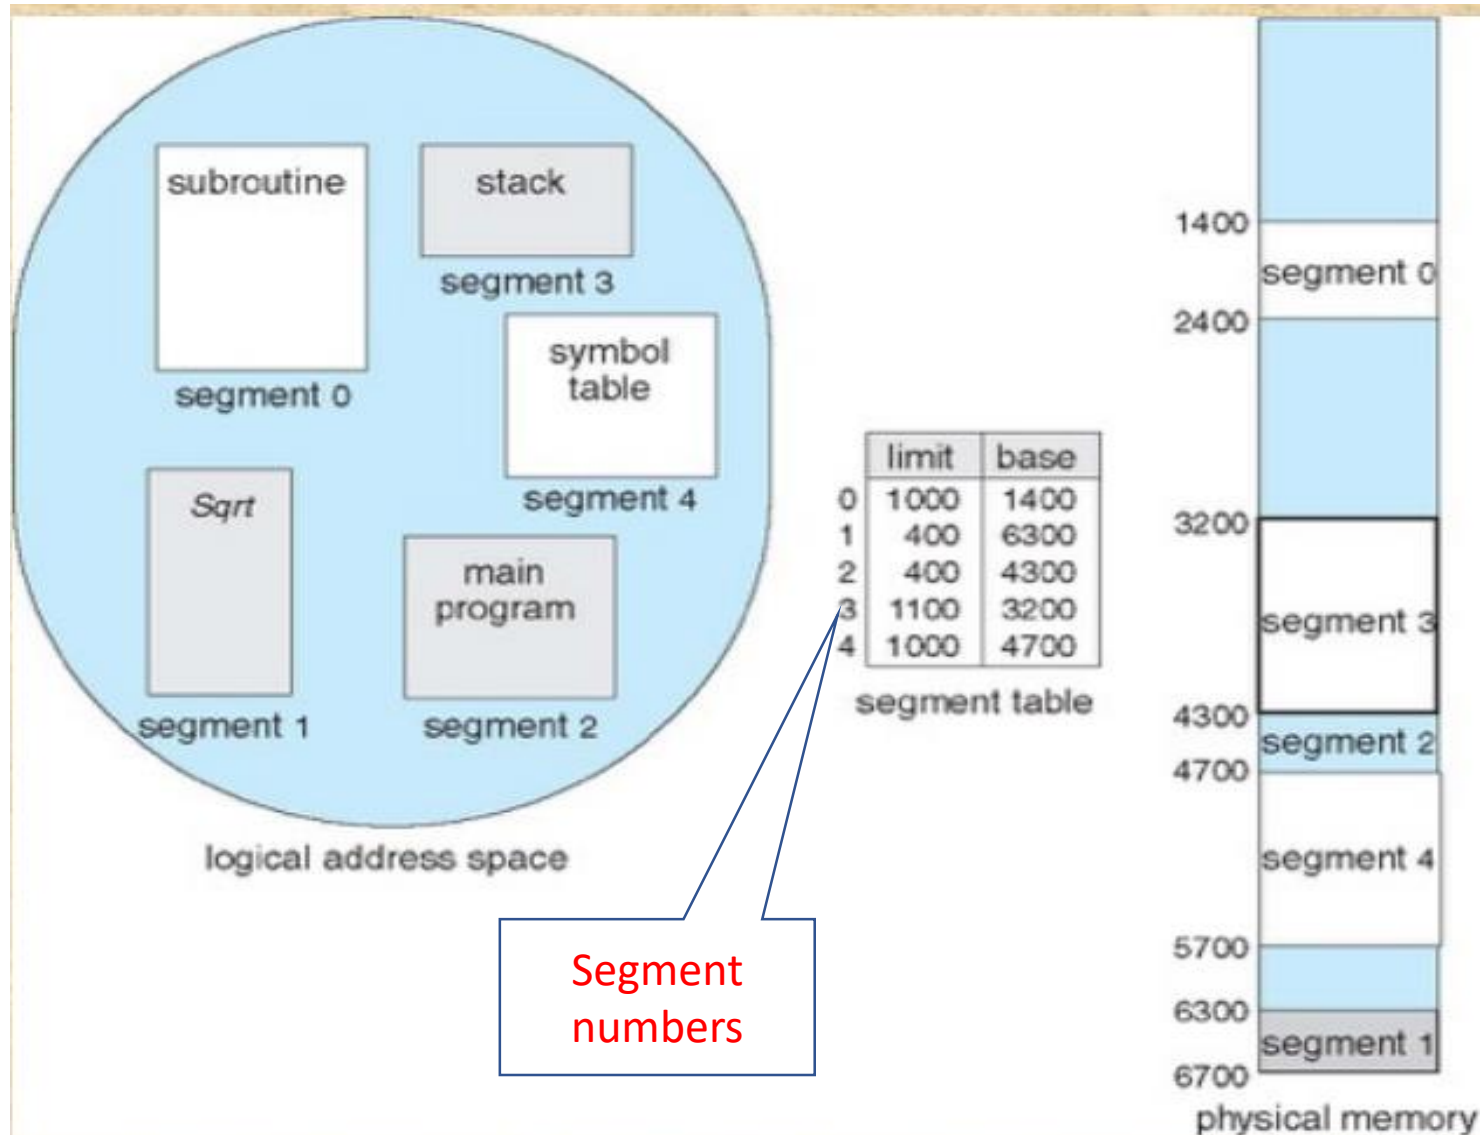
user space



physical memory space

Logical View of Segmentation

- The value in the base register represent the real address of the corresponding segment in the physical memory.

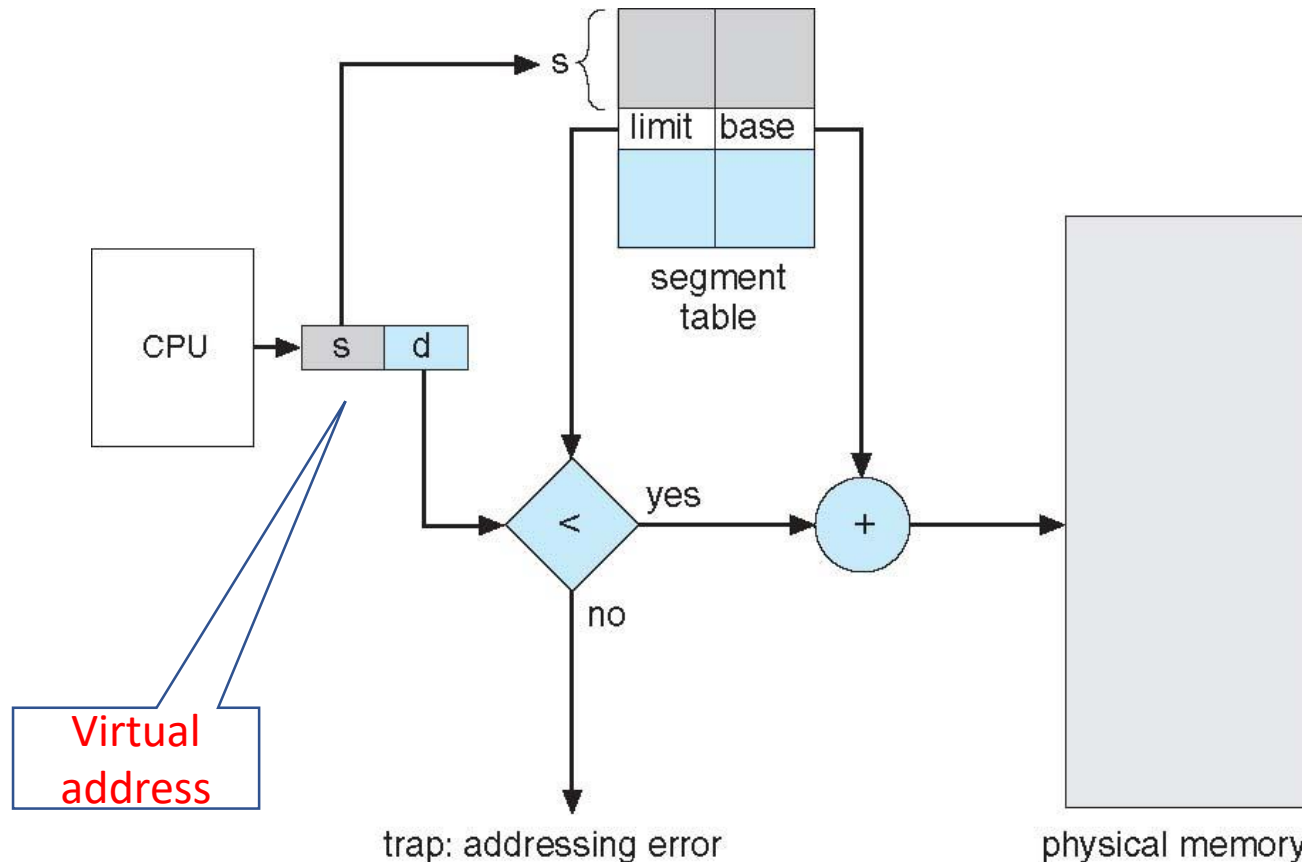


Segmentation Architecture

- Logical address consists of a two tuple:
 <**segment-number**, **offset**>,
 - **Segment table** – maps two-dimensional physical addresses; each table entry has:
 - **base** – contains the starting physical address where the segments reside in memory
 - **limit** – specifies the length of the segment
- Segment table is kept in memory:
- **Segment-table base register (STBR)** points to the segment table's location in memory
- **Segment-table length register (STLR)** indicates number of segments used by a program;
 segment number **s** is legal if **s** < **STLR**

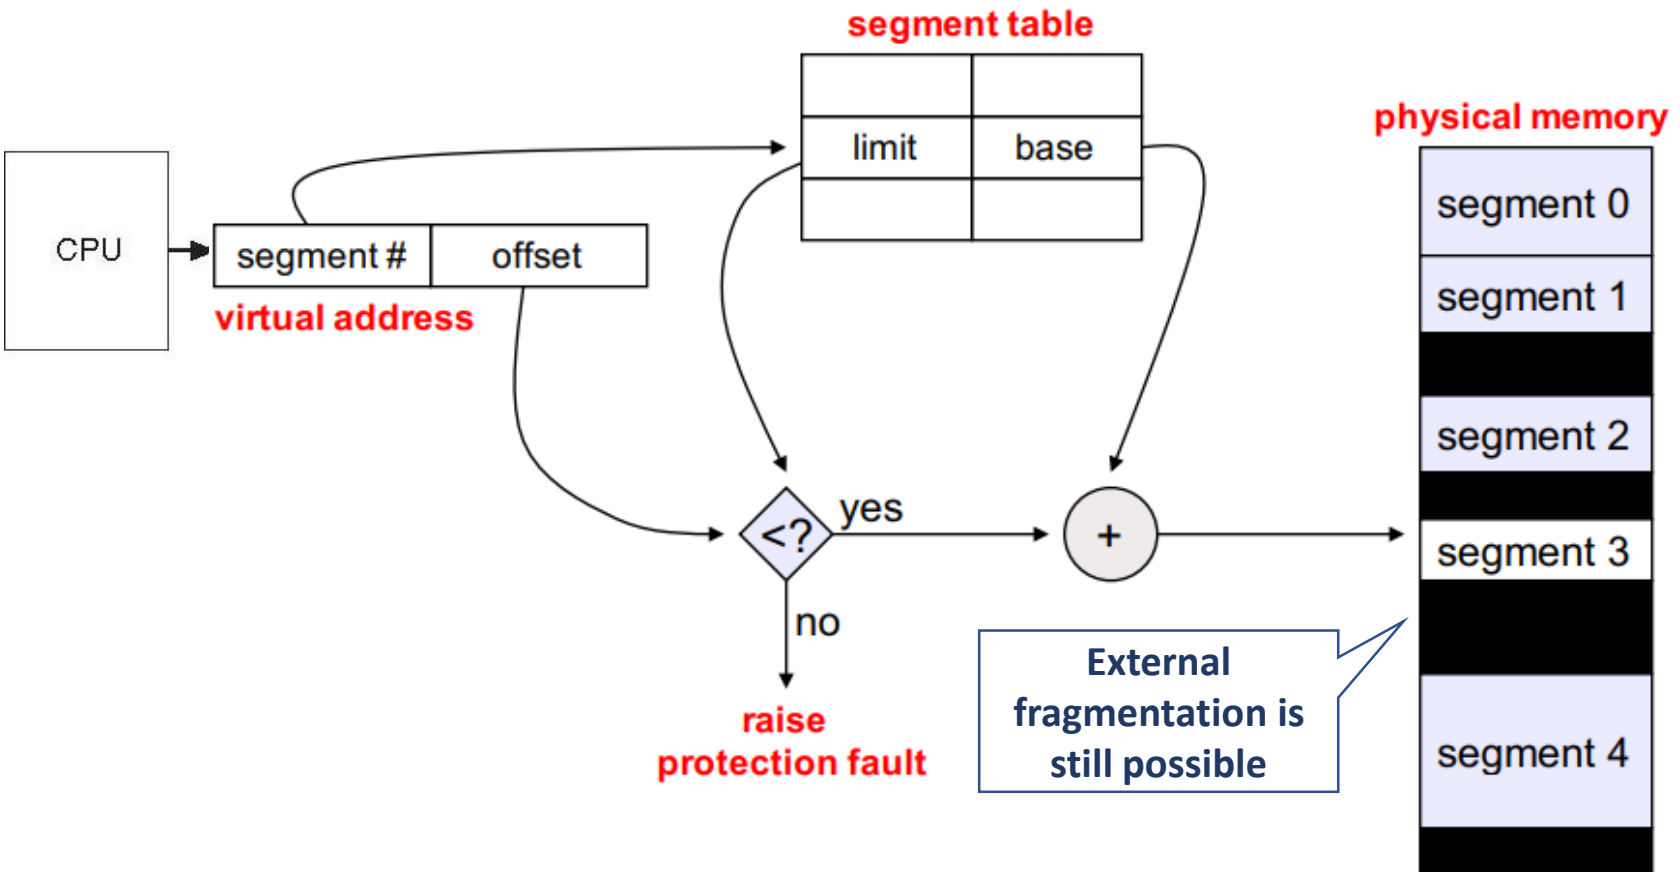
Segmentation Hardware

- Logical address consists of a two tuple:
<segment-number, offset> ,
- Each logical address <s, d> is compared with the segment table



Segmentation Hardware: Segment Lookups

- Logical address consists of a two tuple:
<segment-number, offset> ,
- Each logical address <s, d> is compared with the segment table



Chapter 8: Memory Management

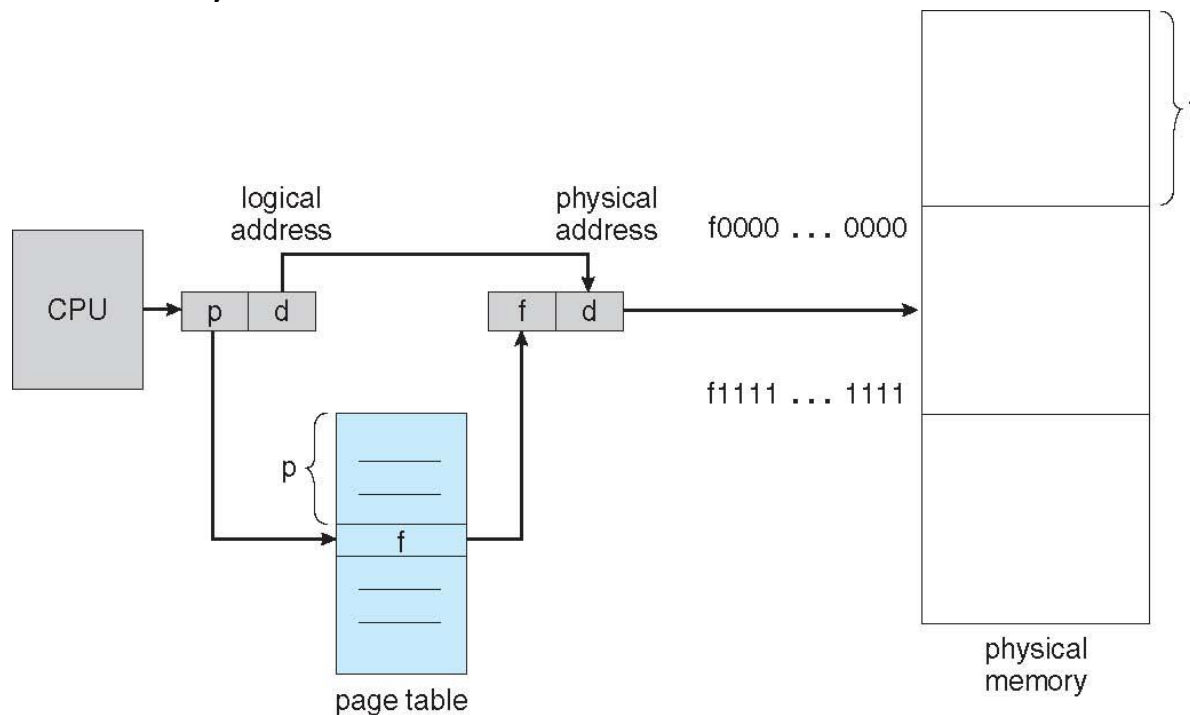
- Background
- Swapping
- Contiguous Memory Allocation
- Segmentation
- **Paging**
- Structure of the Page Table

Paging

- Physical address space of a process can be noncontiguous; process is allocated physical memory whenever the latter is available
 - Avoids **external fragmentation** and the need for compaction
 - Avoids problem of varying sized memory chunks
- **Paging**
- Divide physical memory into fixed-sized blocks called **frames**
 - Size is power of 2, between 512 bytes and 16 Mbytes
 - Divide logical address space into blocks of same size as the frames, called **pages**
 - The page size (like the frame size) is defined by the hardware; see Slide-29
 - Divide **backing store** into [clusters of] blocks of same size as the frames
- Keep track of all free frames
- To run a program of size ***N*** pages, need to find ***N*** free frames and load program
- Set up a **page table** to translate logical to physical addresses
- Limitation: Still have **Internal fragmentation**


Paging Hardware

- Logical address generated by CPU is divided into:
 - **Page number** (p) – used as an index into a **page table**
 - **Page Table** contains base address of each page in physical memory
 - **Page offset** (d) – combined with base address to define the physical memory address that is sent to the memory unit



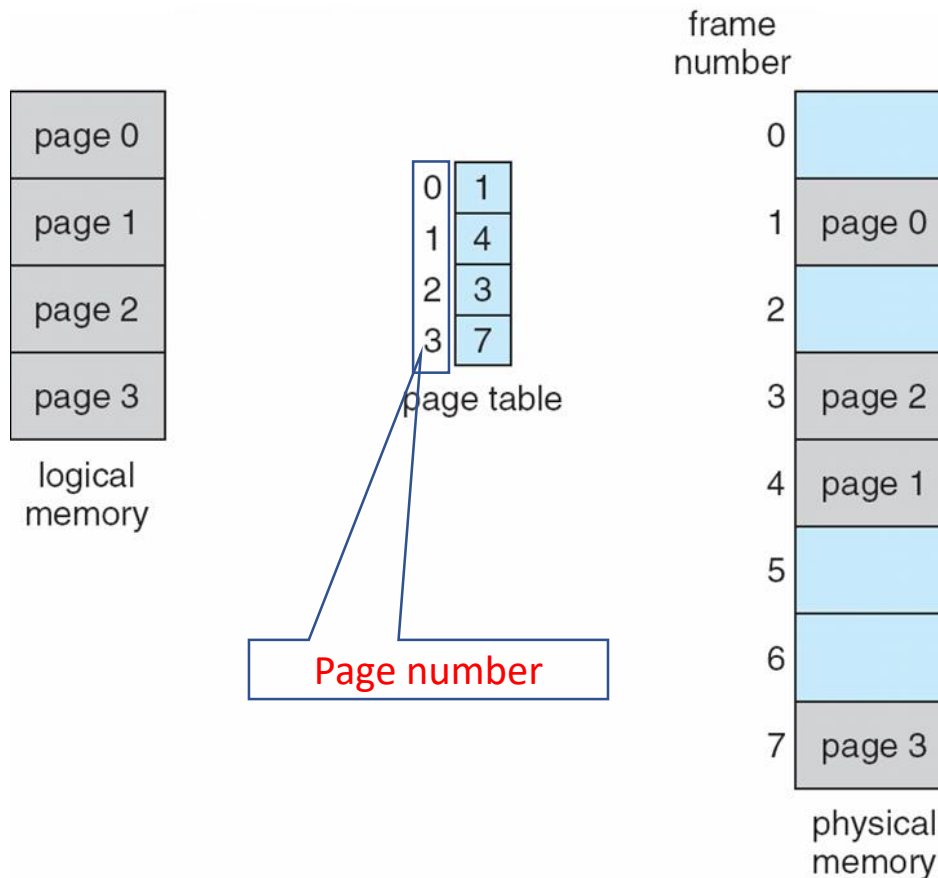
Address Translation Scheme

- Logical address generated by CPU is divided into:
 - **Page number** (p) – used as an index into a **page table**
 - **Page Table** contains base address of each page in physical memory
 - **Page offset** (d) – combined with base address to define the physical memory address that is sent to the memory unit
- If the logical address space is 2^m and the page size is 2^n
 - The binary representation of the logical address has **m bits**, such that
 - The **$m - n$** leftmost bits designate the page number **p**
 - **p** is index into the page table
 - The rightmost **n** bits designate the page offset **d**
 - **d** is displacement within the page
- For given logical address space 2^m and page size 2^n




Basic Paging Model of Logical and Physical Memory

- Page table maps virtual page number to physical frame number.
- A process's logical address space consists of 4 pages, i.e., $m = 4$.



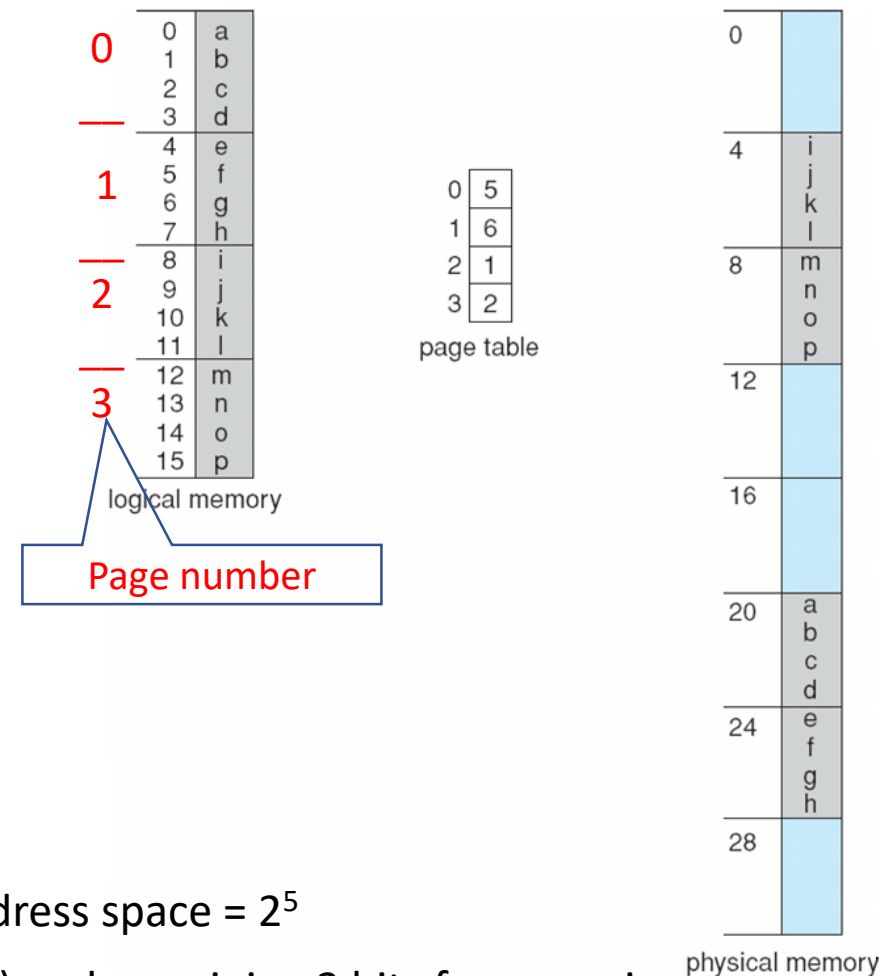
Paging Example (1)

- Example: $m = 4$ and $n = 2$
- What is logical address space of the process?
- What is Page size?
- What is the total number of pages?
- Physical memory has 5 bits.
- Now if logical address is 5 → **0101** (f)
 - What is Physical address ?
- Similarly if logical address is 13 → **1101** (n)
 - What is Physical address?



Paging Example (1)

- Example: $n = 2$ and $m = 4$
- What is logical address space?
 - $2^m = 2^4 = 16$
- What is Logical Page size?
 - $2^n = 2^2 = 4$
- What is the total number of pages?
 - Logical address space/page size
 - Or $m - n = 4 - 2 = 2$ bits $\rightarrow 2^2$ pages




- Physical memory has 5 bits \rightarrow so physical address space = 2^5
 - 3 bits for number of pages (total 8 pages) and remaining 2 bits for page size.
- Now if logical address is 5 \rightarrow 0101 (f)
 - From page table page#1 corresponds to frame#6. So Physical address is 11001
- Similarly if logical address is 13 \rightarrow 1101 (n)
 - From page table page#3 corresponds to frame#2. Physical address is 01001

Paging Example (2)

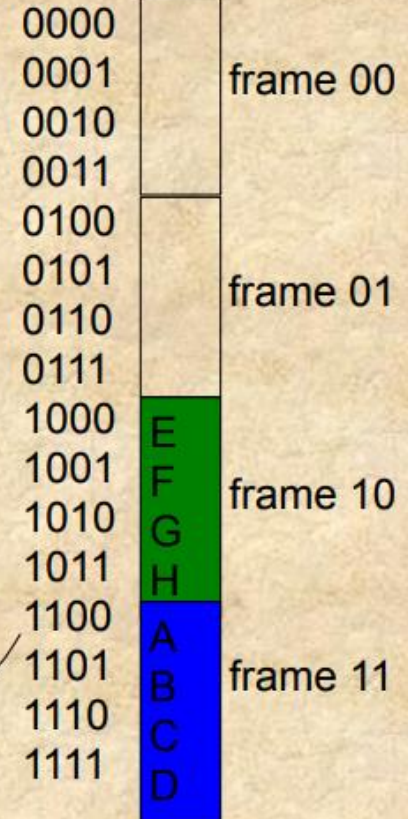

- Virtual Memory has 3 bits:
- $2^3 = 8$ logical addresses
- page size = $2^2 = 4$
- 1 bit for page # and 2 bits for offset

page 0
page 1
1 bit for page#

2 bits for offset



Logical Memory



2 bits for frame#

Physical Memory


page table

0	11
1	10

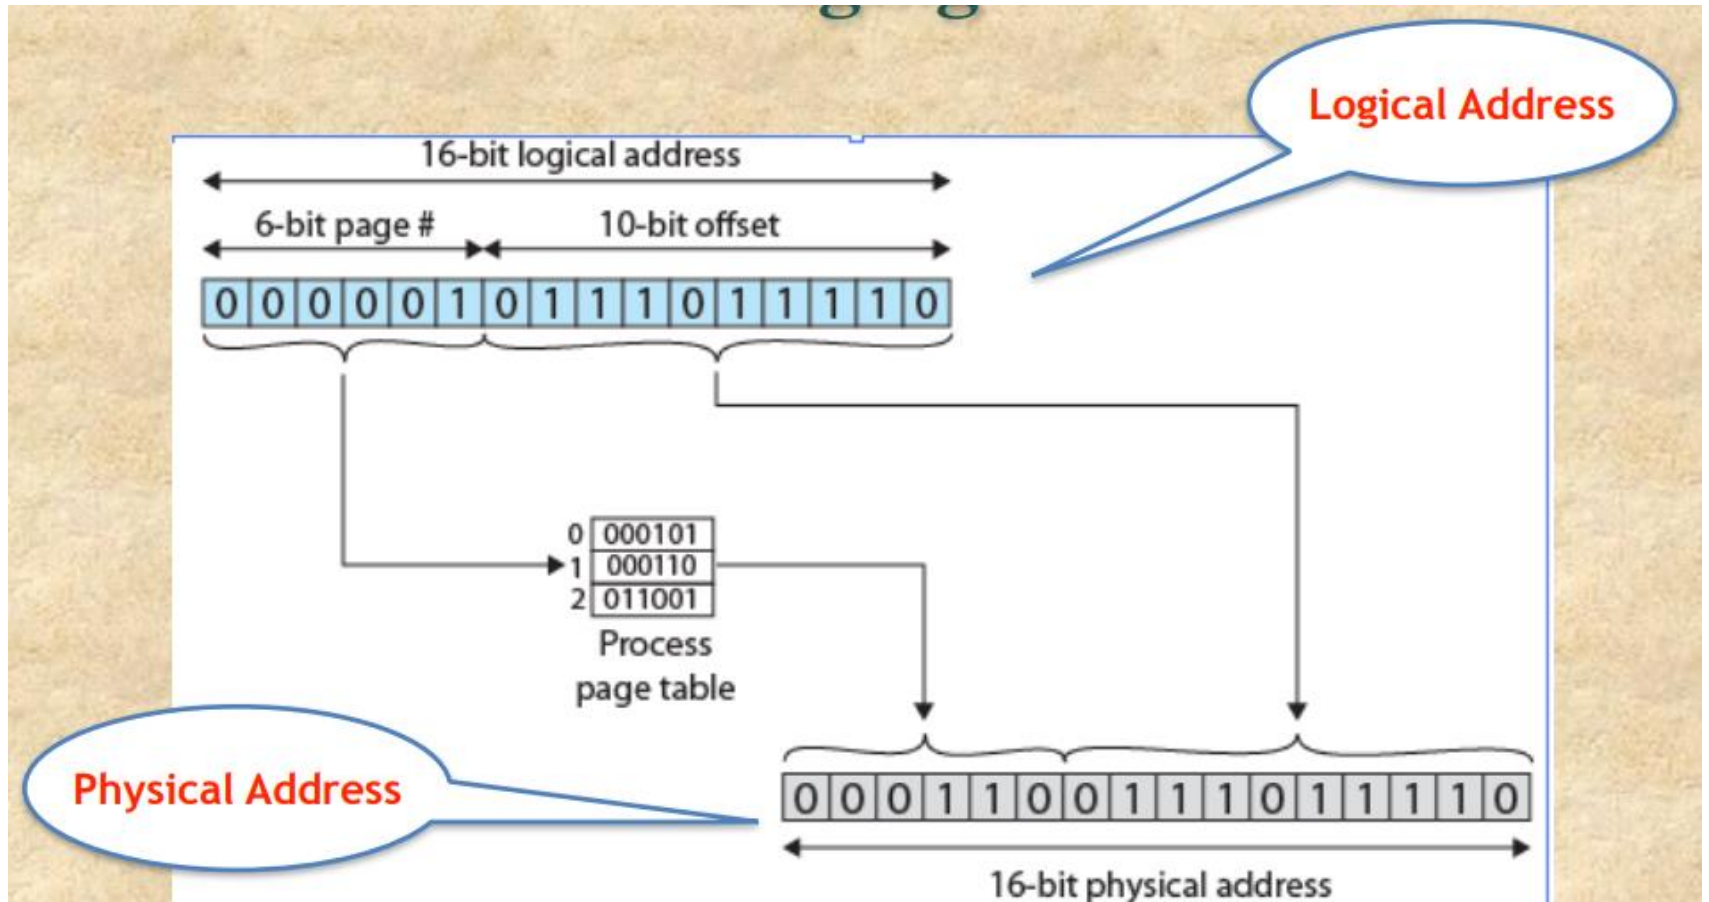
each entry is used to map
4 addresses (page size addresses)

Logical to Physical address Translation (1/2)

- For a 16-bit ($m=16$) address the logical address space is 2^{16} bytes
- Page size ($n=10$) = 1KB = 2^{10} Bytes
- Maximum 64 (2^6) pages of 1K or 1024 bytes each
- What is a Logical address of relative 1502?



Logical to Physical address Translation (2/2)




Paging (Cont.)

- There is no external fragmentation when using paging scheme
- **Internal fragmentation is possible.** Calculating internal fragmentation
 - If Page size = 2,048 bytes and
 - Process size = 72,766 bytes
 - Process has 35 pages + 1,086 bytes, **thus 36 frames required**
 - Internal fragmentation of $2,048 - 1,086 = 962$ bytes unused memory
 - Worst case fragmentation is when 1 frame contains only 1 byte used memory
 - On average fragmentation = **$1/2$ page size per process size**

Paging (Cont.)

- So **Are** small frame sizes desirable?
 - Not necessarily; each page-table entry takes memory to track (overhead)
 - Large frame sizes better when transferring data to/from disk; efficient disk I/O

Free Frames-Before and After Allocation



Before allocation

After allocation

Implementation of Page Table

- Page table is kept in main memory
 - **Page-table base register (PTBR)** points to the page table
 - **Page-table length register (PTLR)** indicates size of the page table
 - Both, PTBR and PTLR are also stored in the process's PCB
- In this scheme every data/instruction access requires two memory accesses
 - First: access the page table using PTBR value to retrieve its frame number
 - Second: access the actual memory location given the frame number
 - This is a serious time overhead that needs to be reduced
- The two memory access problem can be solved by the use of a special fast-lookup hardware cache called **associative memory** or **translation look-aside buffers (TLBs)**


Associative Memory

- Associative memory

Page #	Frame #

- Address translation (p, d)
 - If p is in associative register, get frame # out directly
 - Otherwise get frame # from page table in main memory

Paging Hardware With Translate Lookaside Buffer (TLB)



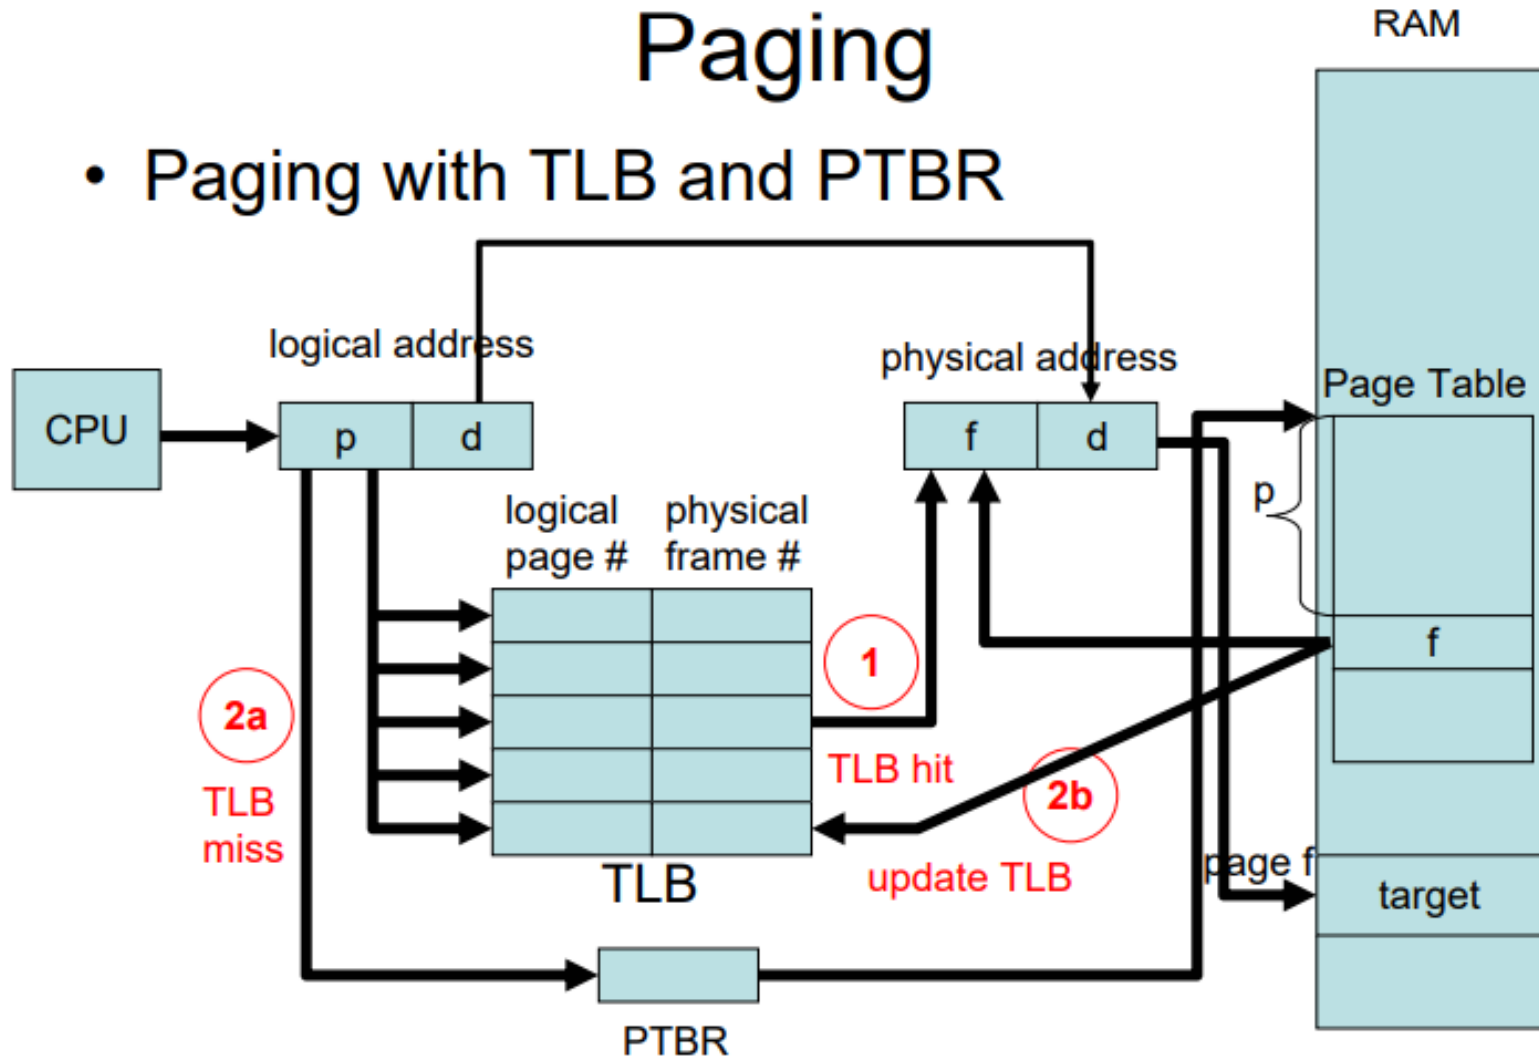
Paging hardware with TLB

- ❑ MMU in CPU first looks in TLB's to find a match for a given logical address
 - **if match found**, then quickly call main memory with physical address frame f (plus offset d)
 - this is called a **TLB hit**
 - TLB as implemented in hardware does a fast parallel match of the input page to all stored values in the cache - about 10% overhead in speed
 - **if no match found**, then
 1. go through regular two-step lookup procedure: go to main memory to find page table and index into it to retrieve frame $\#f$, then retrieve what's stored at address $\langle f, d \rangle$ in physical memory
 2. Update TLB cache with the new entry from the page table
 - if cache full, then implement a cache replacement strategy, e.g. Least Recently Used (LRU) - we'll see this later
- This is called a **TLB miss**
- Goal is to maximize TLB hits and minimize TLB misses

Paging Hardware with TLB

Paging

- Paging with TLB and PTBR



Effective Access Time (EAT)

- Associative Lookup = ε time unit
 - Can be $< 10\%$ of memory access time
- Hit ratio = α
 - Hit ratio – percentage of times that a page number is found in the associative registers; ratio related to number of associative registers

Consider $\alpha = 80\%$, and miss ratio = $1 - \alpha$, $\varepsilon = 20\text{ns}$ for TLB search, 100ns for memory access

- $\text{EAT} = 0.80 \times 20 + 0.20 \times 100 = 36\text{ns}$

Consider more realistic hit ratio $\rightarrow \alpha = 99\%$, $\varepsilon = 20\text{ns}$ for TLB search, 100ns for memory access

- **EAT = ?**

Effective Access Time (EAT)

- Associative Lookup = ε time unit
 - Can be $< 10\%$ of memory access time
- Hit ratio = α
 - Hit ratio – percentage of times that a page number is found in the associative registers; ratio related to number of associative registers

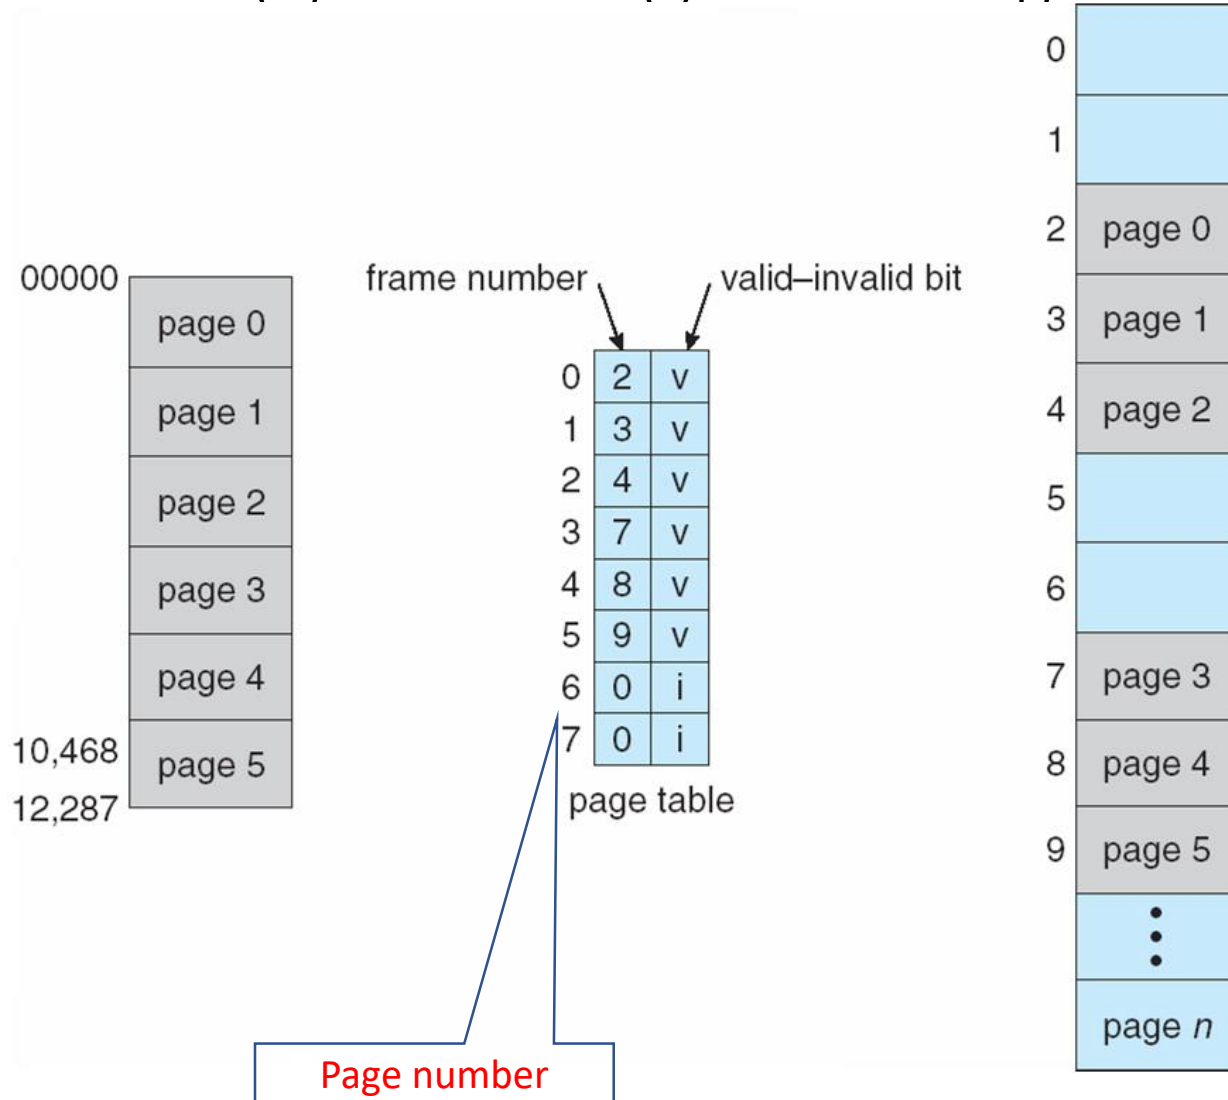
Consider $\alpha = 80\%$, and miss ratio = $1 - \alpha$, $\varepsilon = 20\text{ns}$ for TLB search, 100ns for memory access

- $\text{EAT} = 0.80 \times 20 + 0.20 \times 100 = 36\text{ns}$
- Consider more realistic hit ratio $\rightarrow \alpha = 99\%$, $\varepsilon = 20\text{ns}$ for TLB search, 100ns for memory access
 - **$\text{EAT} = 0.99 \times 20 + 0.01 \times 100 = 20.8\text{ns}$**

Memory Protection

- Memory protection implemented
 - by associating protection bit with each frame to indicate if read-only or read-write access is allowed
 - Can also add more bits to indicate page *execute-only*, and so on
- **Valid-invalid** bit attached to each entry in the page table:
 - “*valid*” indicates that the associated page is in the process’ logical address space, and is thus a legal page
 - “*invalid*” indicates that the page is not in the process’ logical address space
 - Further details in chapter#9 (virtual memory)
- Any violations result in a trap to the kernel

Valid (v) or Invalid (i) Bit In A Page Table



Shared Pages

- **Shared code**

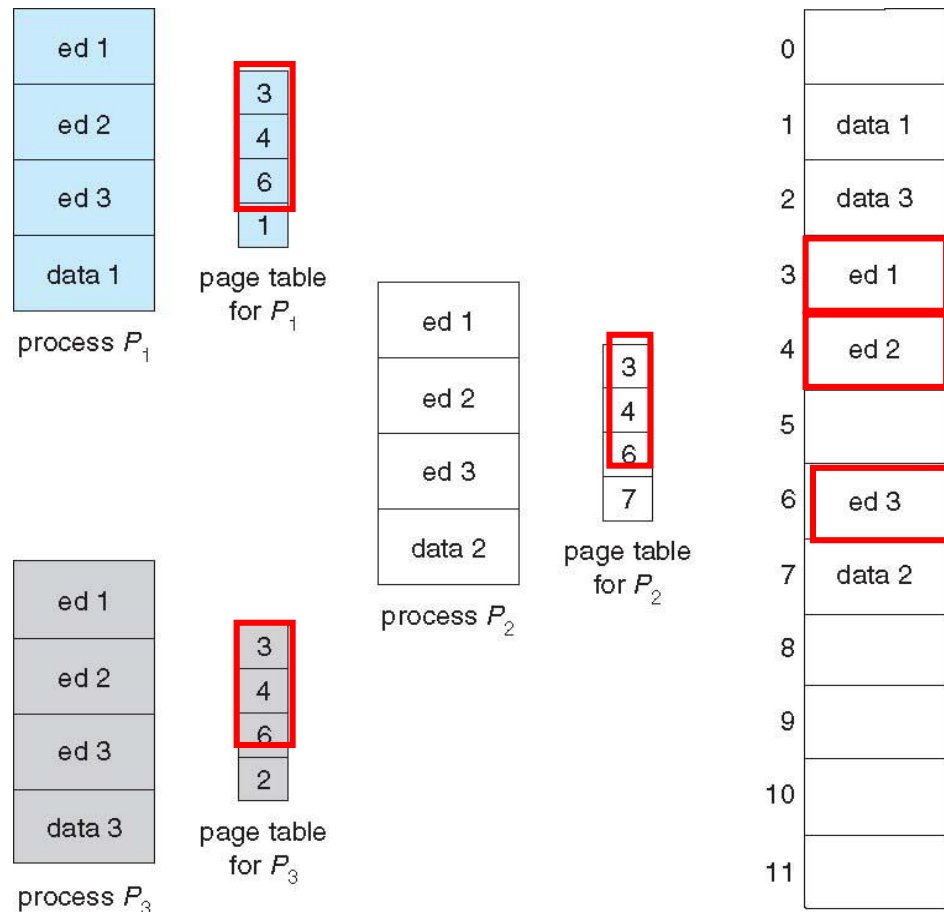
- One copy of read-only code shared among processes (i.e., text editors, compilers, window systems)
- Similar to multiple threads sharing the same process space
- Also useful for interprocess communication if sharing of read-write pages is allowed

- **Private code and data**

- Each process keeps a separate copy of the code and data
- The pages for the private code and data can appear anywhere in the logical address space

Shared Pages Example

- There are three processes. Each process has its own private data.
- However, three processes share **three read-only pages** namely ed 1, ed 2 and ed 3, loaded into frame#3, 4, and 6 respectively.



Chapter 8: Memory Management

- Background
- Swapping
- Contiguous Memory Allocation
- Segmentation
- Paging
- **Structure of the Page Table**


Structure of the Page Table

- Memory structures for paging can get huge using straight-forward methods
 - Consider a 32-bit logical address space as on modern computers
 - Page size of 4 KB (2^{12})
 - Page table would have 1 million entries ($2^{32} / 2^{12}$)
 - Assume each entry is 4 bytes → 4 MB of physical memory for page table alone
 - That amount of memory used to cost a lot if there are a lot of processes
 - Don't want to allocate that contiguously in main memory
- **Method#1: Hierarchical Paging**
- **Method#2: Hashed Page Tables**
- **Method#3: Inverted Page Tables**

Method#1: Hierarchical Page Tables

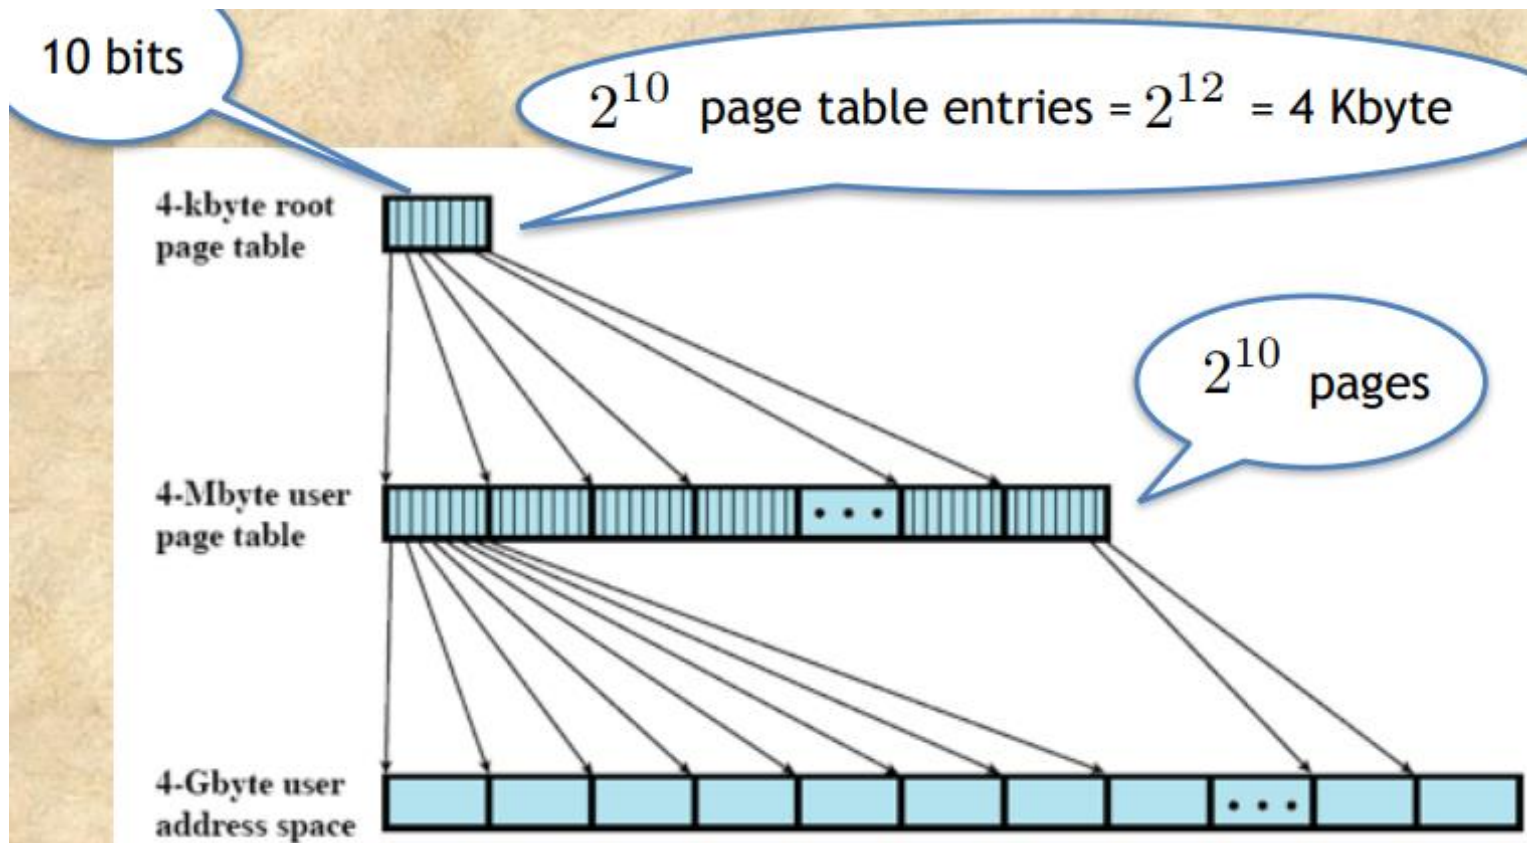
- **Paging Principle:**
 - Break up the logical address space into **multiple page tables**
 - Subdividing a process into pages helped to fit a process into memory by allowing it to be scattered in non-contiguous pieces of memory, thereby solving the external fragmentation problem
- so reapply that principle here to fit the **page table** into memory, allowing the page table to be scattered non-contiguously in memory
- A simple technique is a two-level page table
- We then **page the page table**

Method#2: Two-Level Page-Table Scheme



Two-Level Hierarchical Page Table

- Logical address space is 2^{32} and each page size is 2^{12} .
- Inner page table has 1 million entries therefore its size is 4Mbytes. Each page of the inner page table is also 4kbytes. $4\text{MB}/4\text{KB} \rightarrow 2^{10} = 1024$ pages.
- Now we need an outer page table, in which each entry points to a different page of the **inner page table**.
- In **outer page table** there are 2^{10} entries each of them is 4 bytes, therefore, total outer page table is also 4kbytes.

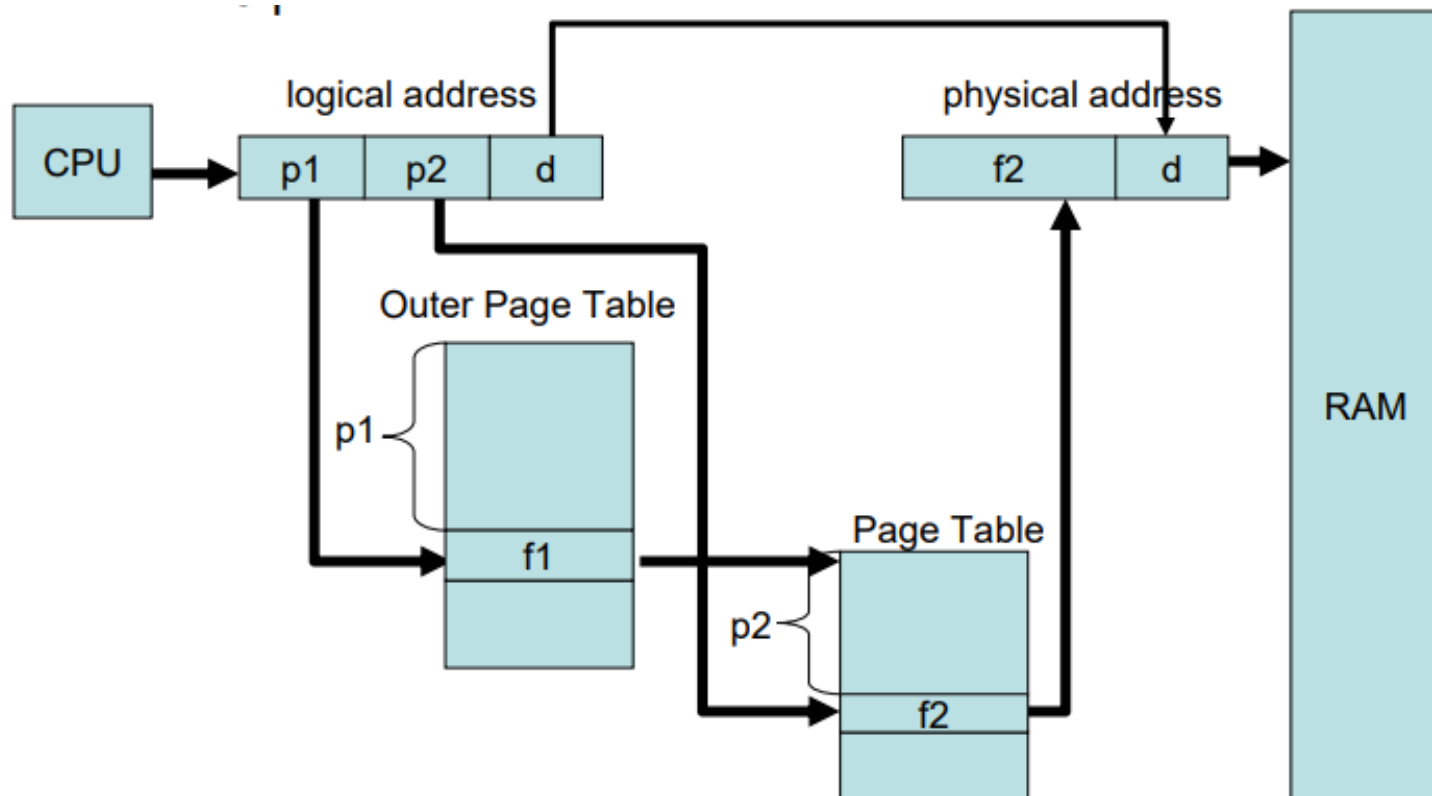


Two-Level Paging Example

- A logical address (on 32-bit machine with 4KB page size (or 2^{12})) is divided into:
 - a page number consisting of 20 bits
 - a page offset consisting of 12 bits
- Assuming each page of the **inner page table** is also 4KB, then each page holds 2^{12} bytes / 4 bytes per entry = 2^{10} entries per POPT
- Since there are 2^{20} total entries in the page table, there must be 2^{10} pages in the page table. → requires 2^{10} entries in outer page table (POPT)
- This leaves us with 10 bits left over for the **outer page table**
- Hence we have:
 - a 10-bit page number (P1)
 - a 10-bit page offset (P2)

Hierarchical Paging

- Hierarchical (2-level) paging divides the logical address into 3 parts




Method#2: Hashed Page Tables

- Hashed page tables are common when address spaces > 32 bits
- The virtual page number is hashed into a hashed page table
 - This page table contains a chain of elements hashing to the same location
- Each element contains
 - (1) the virtual page number
 - (2) the value of the mapped page frame
 - (3) a pointer to the next element
- Since there might be multiple page table indexes that hash to the same value, the hashed page table provides a linked list for each hashed entry.
- Virtual page numbers are compared in this chain searching for a match
 - If a match is found, the corresponding physical frame is extracted

Method#2: Hashed Page Table


- The benefit of this approach is that in a hash table, the lookup time should be independent of the number of entries in the table.



Method#3: Inverted Page Table

- Normally, **logical addresses** are mapped to physical addresses in a page table for each process. This requires each process to have their own page table, each with potentially millions of entries.
- The *inverted page table* maps physical frames to logical pages.
- One entry for each real frame of memory
- Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page
- **Decreases** memory needed to store each page table, but increases time needed to search the table when a page reference occurs
- Structure is called inverted because it indexes page table entries by **frame number** rather than by **virtual page number**

Inverted Page Table Architecture



End of Chapter 8