

# Chapter 1: Introduction (Part 2)

# Chapter 1: Introduction

- What Operating Systems Do
- Computer-System Organization
  - System Call
  - Interrupts
  - Interrupt handling
  - DMA
  - Storage Device Hierarchy
- Computer-System Architecture
- Operating-System Structure (Evolution of OS)
- Operating-System Operations
- Process Management
- Memory Management
- Storage Management

# Storage Definitions and Notation Review

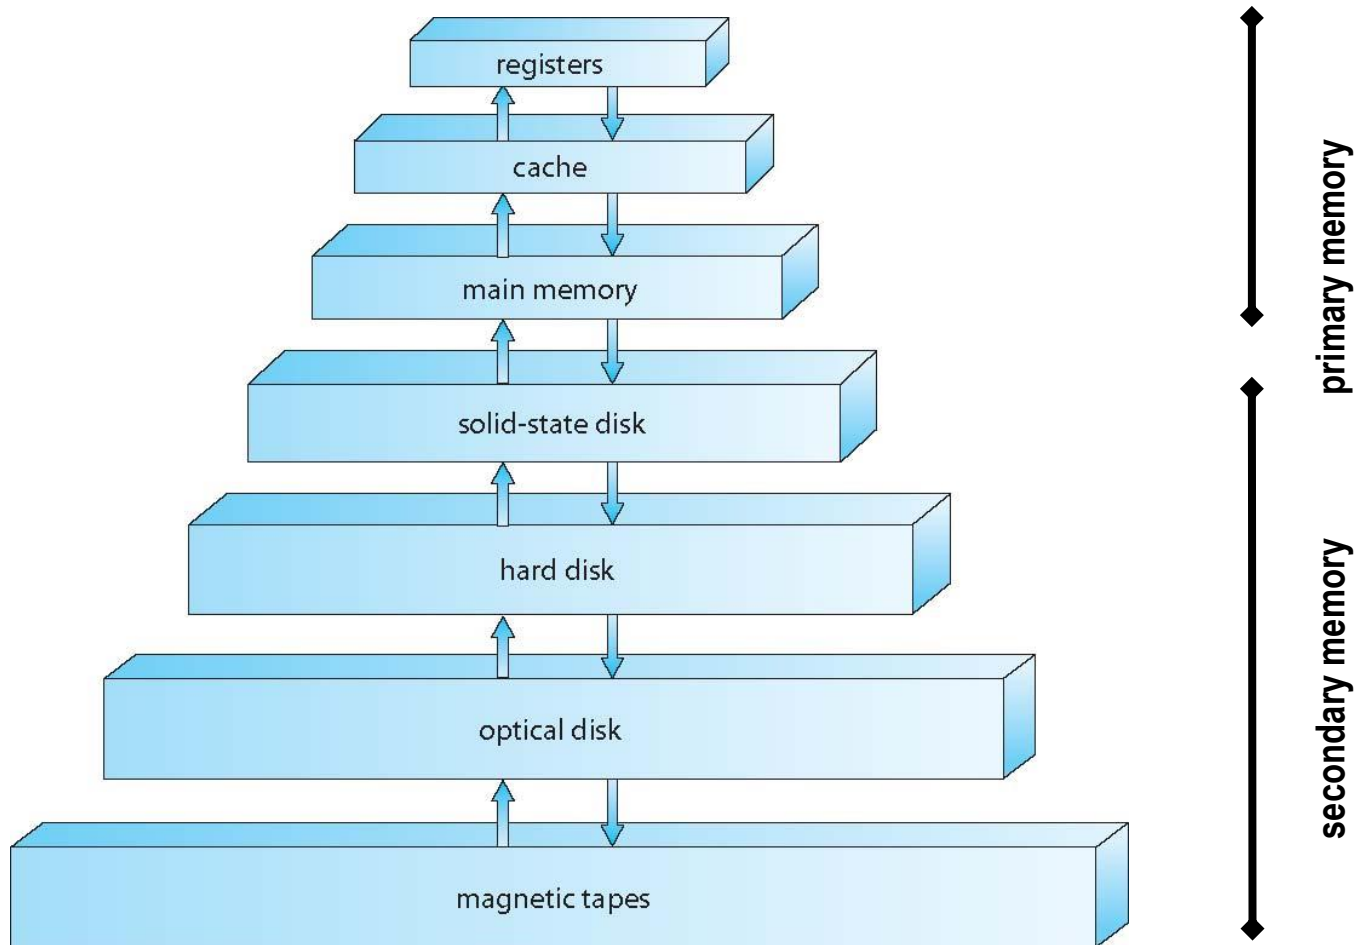
The basic unit of computer storage is the **bit**. A bit can contain one of two values, 0 and 1. All other storage in a computer is based on collections of bits. Given enough bits, it is amazing how many things a computer can represent: numbers, letters, images, movies, sounds, documents, and programs, to name a few. **A byte is 8 bits, and on most computers it is the smallest convenient chunk of storage.** For example, most computers don't have an instruction to move a bit but do have one to move a byte. A less common term is **word**, which is a given computer architecture's native unit of data. A word is made up of one or more bytes. For example, a computer that has 64-bit registers and 64-bit memory addressing typically has 64-bit (8-byte) words. A computer executes many operations in its native word size rather than a byte at a time.

Computer storage, along with most computer throughput, is generally measured and manipulated in bytes and collections of bytes.

A **kilobyte**, or **KB**, is 1,024 bytes  
a **megabyte**, or **MB**, is  $1,024^2$  bytes  
a **gigabyte**, or **GB**, is  $1,024^3$  bytes  
a **terabyte**, or **TB**, is  $1,024^4$  bytes  
a **petabyte**, or **PB**, is  $1,024^5$  bytes

Computer manufacturers often round off these numbers and say that a megabyte is **1 million bytes** and a **gigabyte is 1 billion bytes**. Networking measurements are an exception to this general rule; they are given in bits (because networks move data a bit at a time).

# Storage-Device Hierarchy



# Storage structure

- Memory is organized in a hierarchy  
**tradeoffs:** speed vs cost vs volatility
- **Primary memory**
  - fast, expensive, volatile
  - data is stored in electronic circuitry
  - Example: **Main memory** – only large storage media that the CPU can access directly
- **Secondary memory**
  - slow, cheap, permanent
  - data is stored magnetically
  - can store massive amounts of inactive data, must be copied to primary memory to be accessed
  - **Example:** Hard disks, solid state disks
  - Solid-state disks – faster than hard disks, nonvolatile
  - Various technologies
  - Becoming more popular

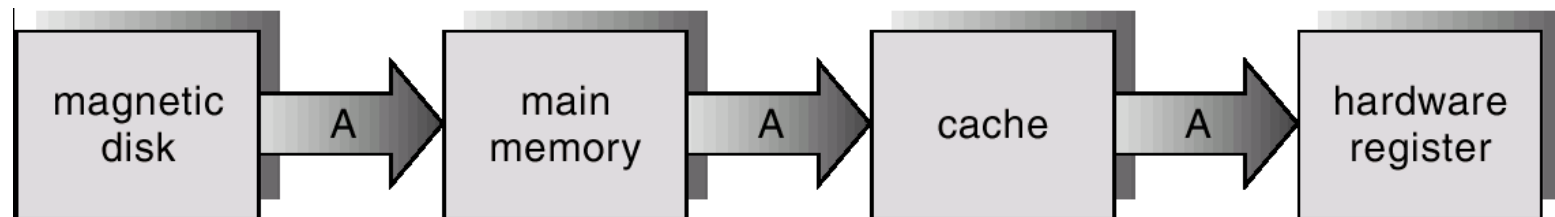
# Primary memory: RAM vs. cache

- RAM and cache both store data in (volatile) electronic circuitry
  - cache uses faster, more expensive technology
  - Level-1 cache is stored directly on the CPU chip (runs at speeds comparable to processor speed,  $\sim 10\times$  RAM access speed )
  - Level-2 cache is stored on nearby chip,  $\sim 2\times$  RAM access speed common

## common approach:

- when data from RAM is needed by CPU, first copy into cache
- CPU then accesses cache directly
- cache retains recently used (most active?) data, fast access if needed again

*note: cache is to RAM as RAM is to secondary memory*

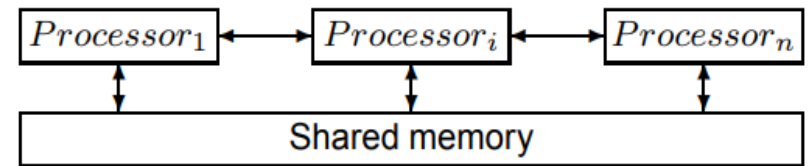


# Chapter 1: Introduction

- What Operating Systems Do
- Computer-System Organization
  - System Call
  - Interrupts
  - Interrupt handling
  - DMA
  - Storage Device Hierarchy
- **Computer-System Architecture**
- Operating-System Structure (Evolution of OS)
- Operating-System Operations
- Process Management
- Memory Management
- Storage Management

# Computer-System Architecture

- Most systems use a single general-purpose processor
  - Most systems have special-purpose processors as well
- **Multiprocessors** systems growing in use and importance
  - Also known as **parallel systems**, **tightly-coupled systems** (i.e., All processors share the same memory)

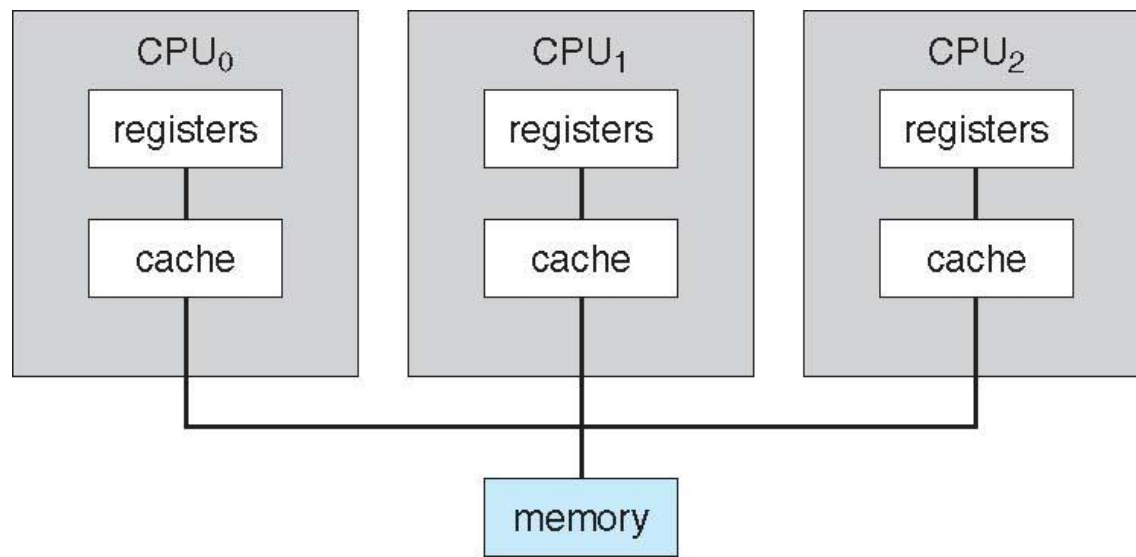


- Advantages include:
  1. **Increased throughput:** expect to get more work done in less time
  2. **Economy of scale:** share peripherals, mass storage, and power supplies
  3. **Increased reliability** – graceful degradation or fault tolerance
- Two types:
  1. **Asymmetric Multiprocessing** – Boss-Worker-Each Processor is assigned a specific task
  2. **Symmetric Multiprocessing (SMP)** – all processors are peers- each processor performs all tasks



# Symmetric Multiprocessing Architecture

- A stand-alone computer system with the following characteristics:
  - two or more similar processors of comparable capability
  - processors share the same main memory and are interconnected by a bus or other internal connection scheme
  - processors share access to I/O devices
  - all processors can perform the same functions
  - the system is controlled by an integrated operating system that provides interaction between processors and their programs at the job, task, file, and data element levels



# A Multicore Design

- Multiple computing cores (Processing Units) on a single chip (Single Physical Package)
- More efficient than multiple chips with single cores
  - On chip communication is faster than between chip communication

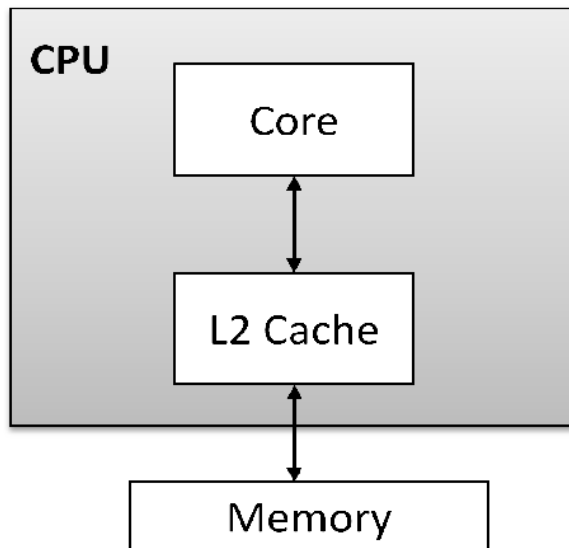


Figure 1:  
Single-core Architecture

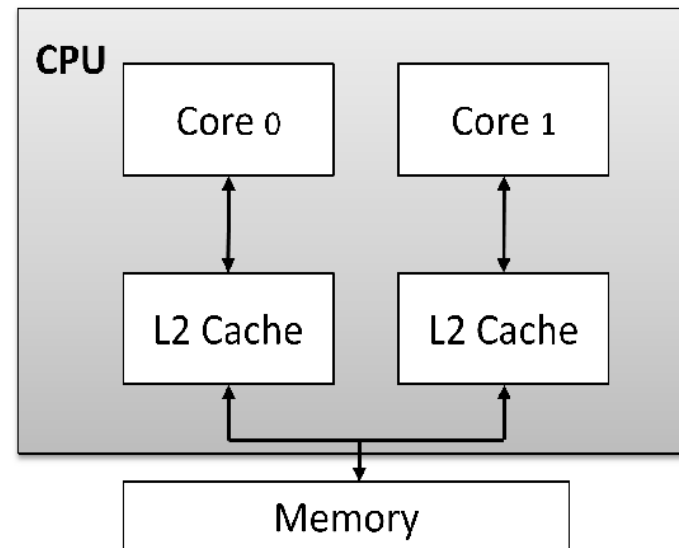
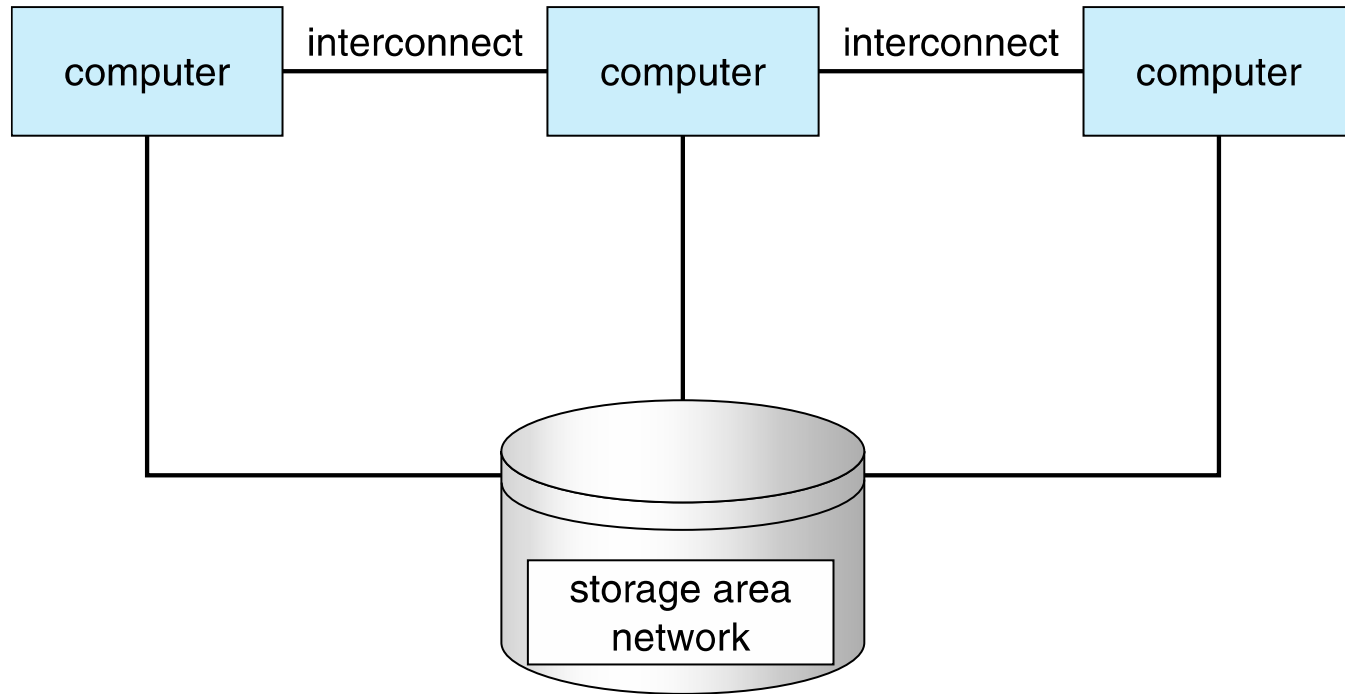


Figure 2:  
Multi-core Architecture

# Clustered Systems



# Clustered Systems

- Like multiprocessor systems, but multiple systems working together
- Usually sharing storage via a **storage-area network (SAN)**
- Provides a **high-availability** service which survives failures
  - **Asymmetric clustering** has one machine in *hot-standby mode*
    - Hot-standby host monitor the active server and takes over if the later fails
  - **Symmetric clustering** has multiple nodes running applications, monitoring each other
- Some clusters are for **high-performance computing (HPC)**
  - Applications must be written to use **parallelization**
- Hosts can perform conflicting operations which can be prevented by locking and access control mechanisms, also called **distributed lock manager (DLM)**.

# Chapter 1: Introduction

- What Operating Systems Do
- Computer-System Organization
- Computer-System Architecture
- **Operating-System Structure (Evolution of OS)**
- Operating-System Operations
- Process Management
- Memory Management
- Storage Management

# Evolution of an Operating Systems

- Must adapt to **hardware upgrades** and new **types of hardware**.

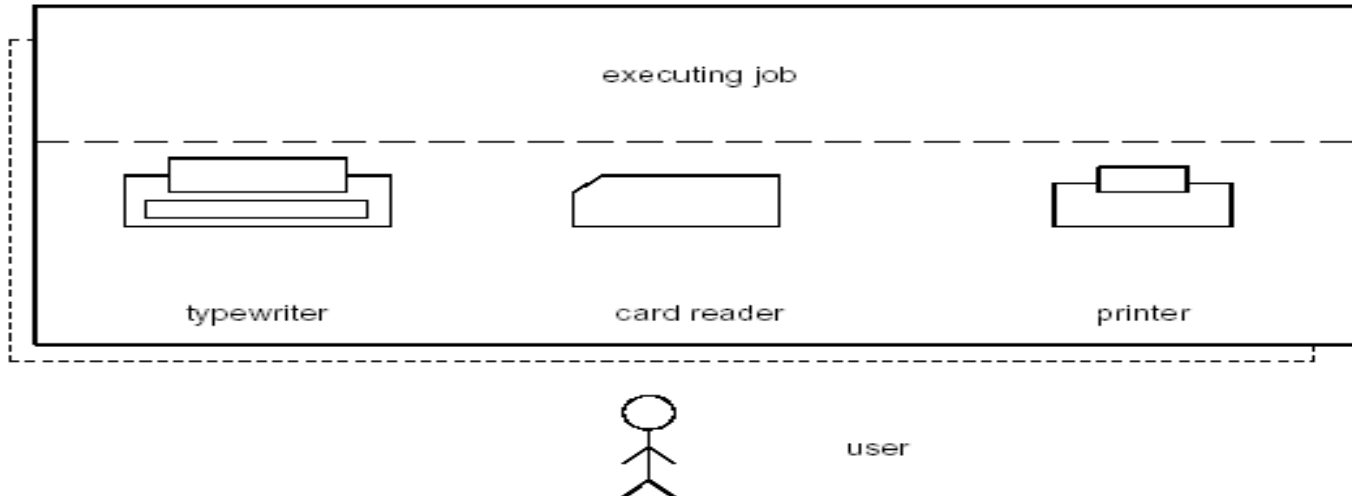
Examples:

- For example, Introduction of paging hardware
- Must offer new services, e.g., internet support.
- The need to change the OS on regular basis place requirements on it's design:
  - modular construction with clean interfaces.
  - object oriented methodology.

# Early Systems: Serial Processing ( 1940-1955)

- Structure

- Single user system.
- Programmer/User as operator
- Large machines run from console (display lights, toggle switches)
- Input received from Paper Tape or Punched cards.
- No OS



# Characteristics of Early Systems

- **Early software:** Assemblers, Libraries of common subroutines (I/O), Device Drivers, Compilers, Linkers.
- Need significant amount of setup time.
- Extremely slow I/O devices.
- Very low CPU utilization.

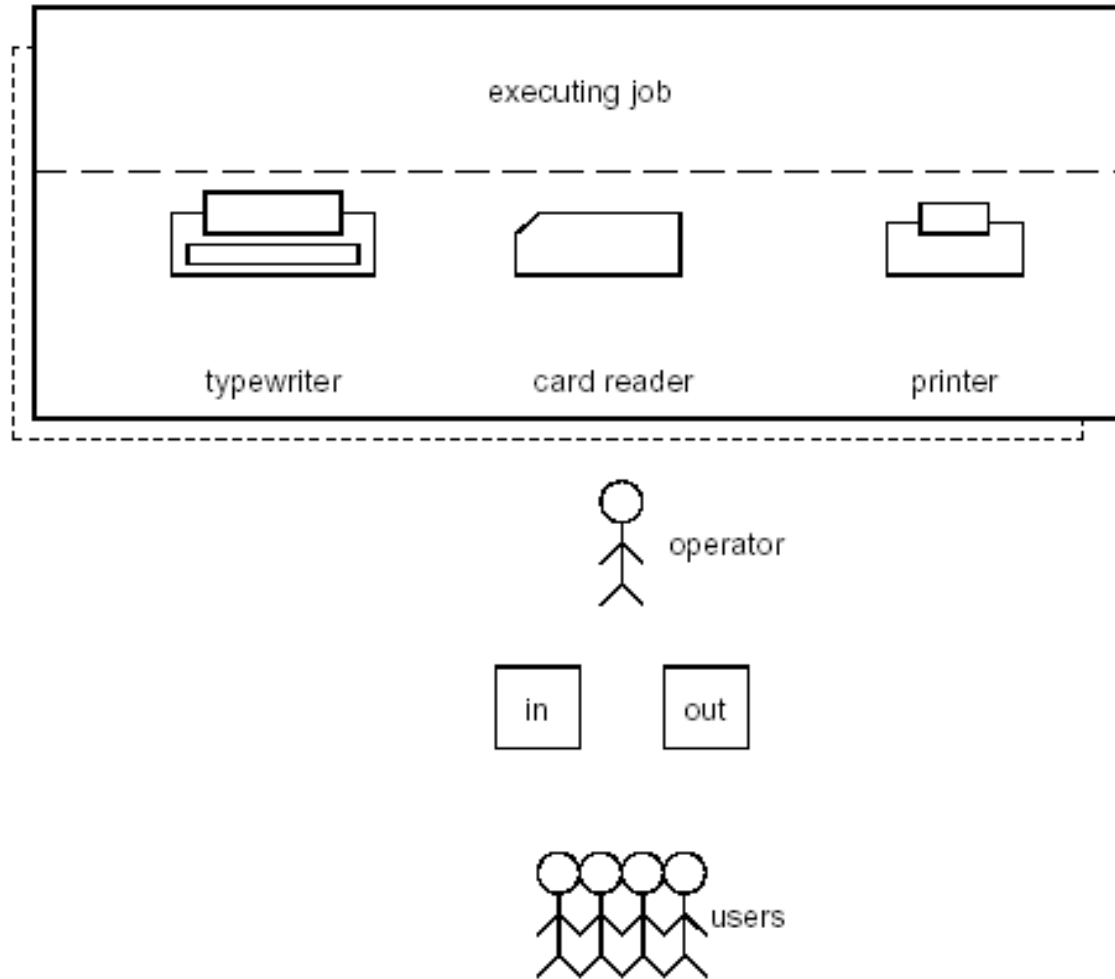


# Simple Batch Systems

- Use of high-level languages, magnetic tapes (instead of paper tapes).
- Jobs are batched together by type of languages.
- An operator was hired to perform the repetitive tasks of loading jobs, starting the computer, and collecting the output (Operator-driven Shop).
  - Overall process became a little fast
- It was not feasible for users to inspect memory or patch programs directly.



# Operator-driven Shop



# Operation of Simple Batch Systems

- The user submits a job (written on cards or tape) to a computer operator.
- The computer operator place a batch of several jobs on an input device.
- A special program, the monitor, manages the execution of each program in the batch.
- Monitor utilities are loaded when needed.
- “**Resident monitor**” is always in main memory and available for execution.
  - first OS
  - initial control is in **monitor**.
  - loads next program and transfers control to it.
  - when job completes, the control transfers back to monitor.
  - Automatically transfers control from one job to another, no idle time between programs.

# Idea of Simple Batch Systems

- Reduce setup time by **batching similar jobs**.
- Alternate execution between user program and the monitor program.
- Rely on available hardware to effectively alternate execution from various parts of memory.
- Use Automatic Job Sequencing – automatically transfer control from one job when it finishes to another one.

# Simple Batch Systems: Job Control Language (JCL)

- **JCL** is the language that provides instructions to the **monitor**:
  - what compiler to use
  - what data to use

- Example of job format: ----->>

- **'\$'** at the beginning denotes the job control instruction

- **\$FTN** loads the compiler and transfers control to it.

- **\$LOAD** loads the object code (in place of compiler).

- **\$RUN** transfers control to user program.

```
$JOB
$FTN
...
FORTRAN
program
...
$LOAD
$RUN
...
Data
...
$END
```

# Simple Batch Systems (1)

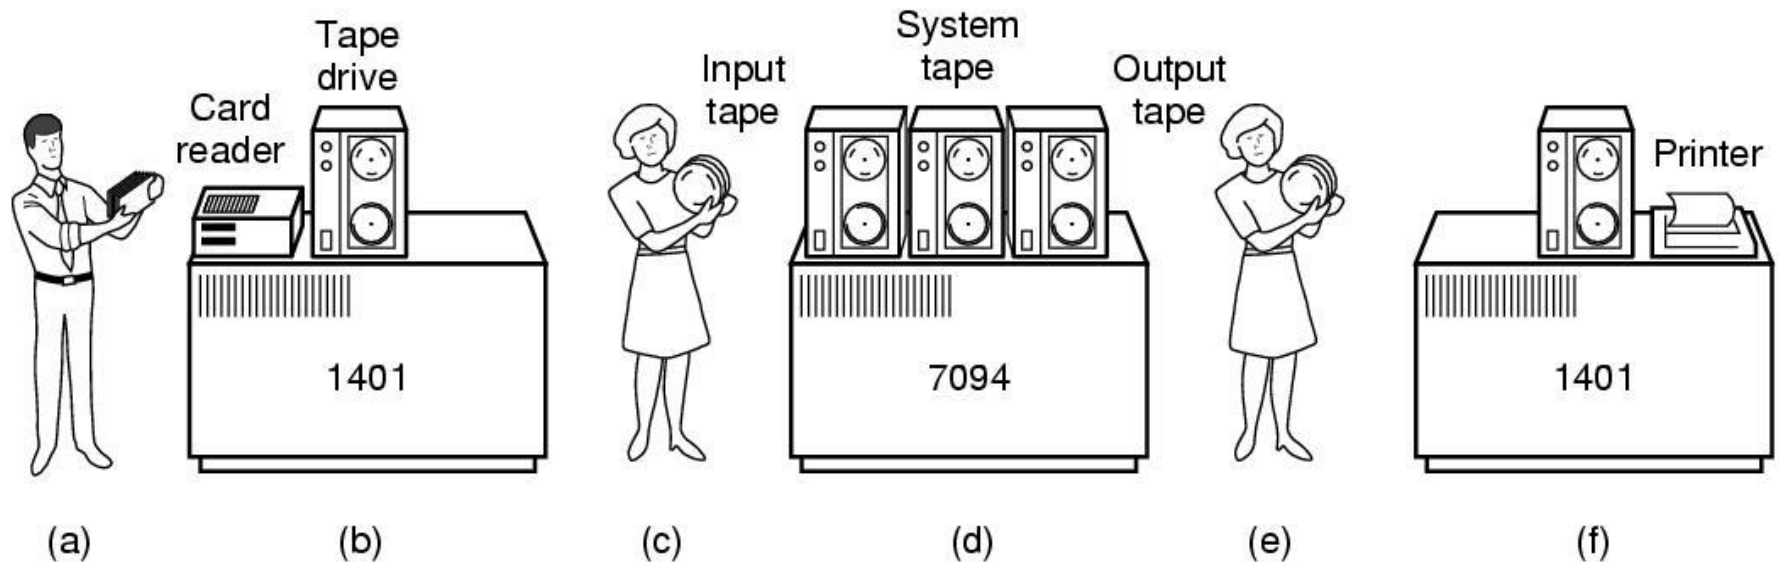


Figure 1-3. An early batch system.

(a) Programmers bring cards to 1401.

(b) 1401 reads batch of jobs onto tape.

# Simple Batch Systems (2)

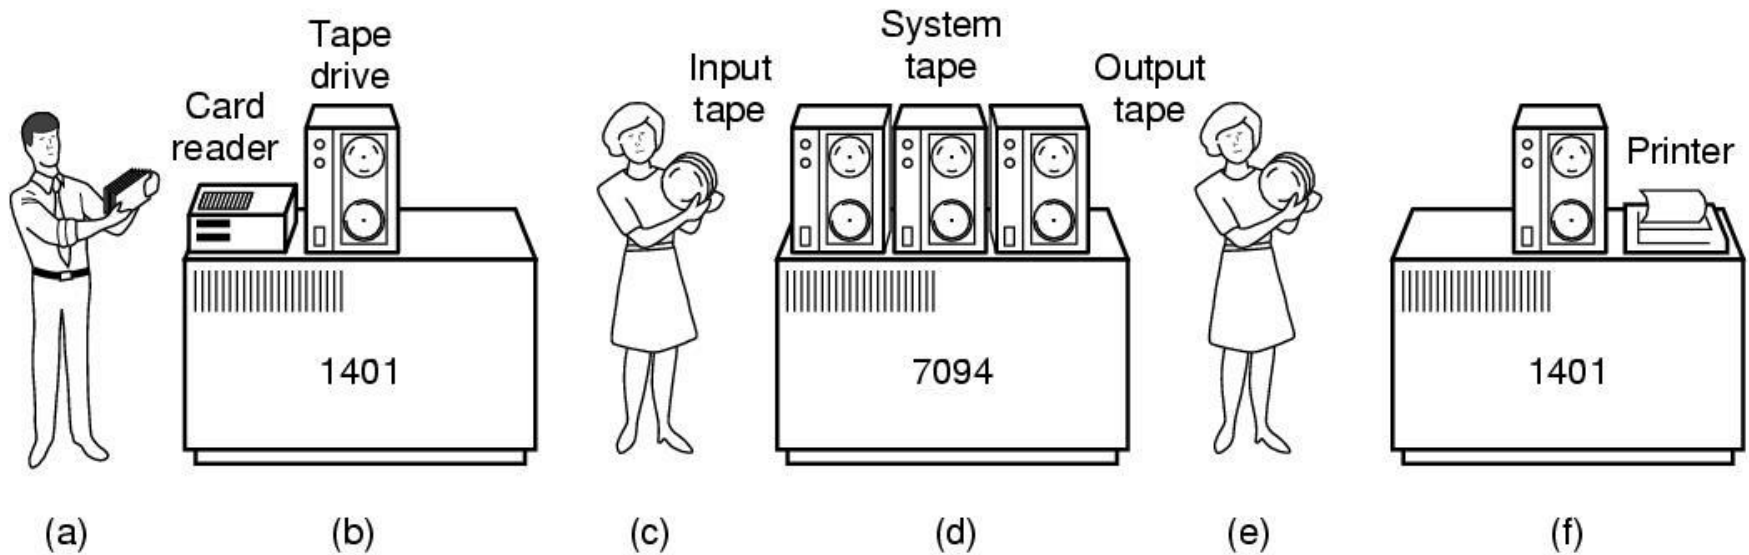
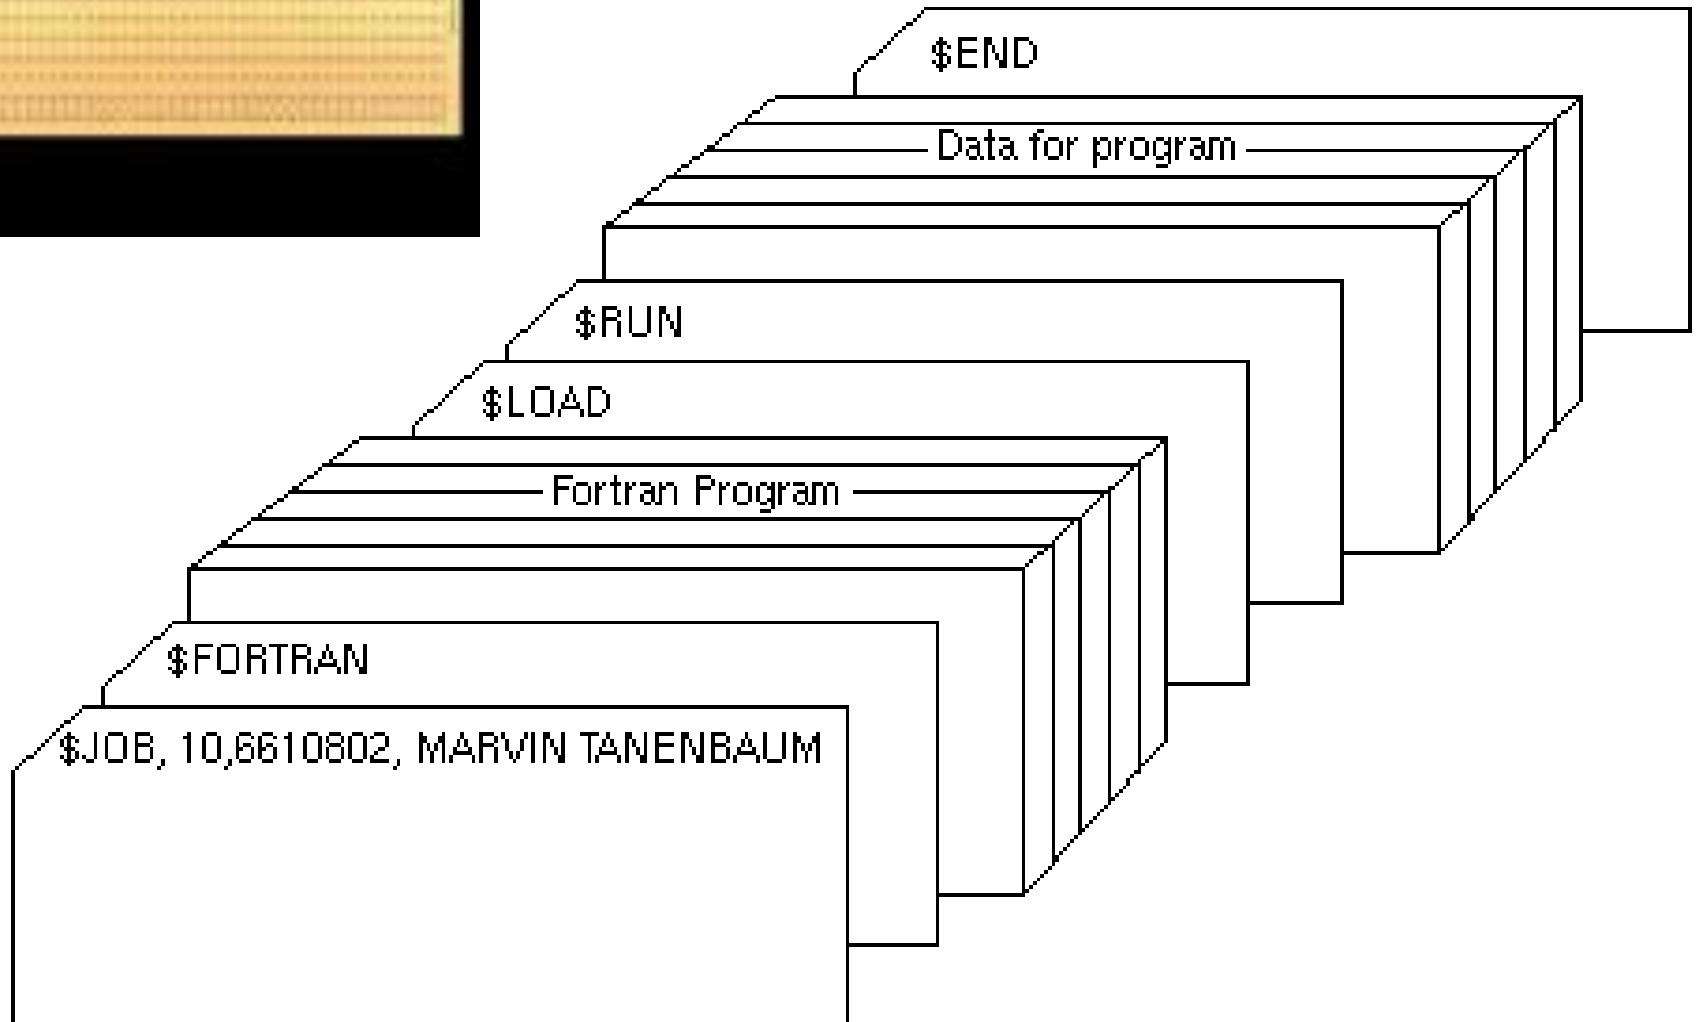


Figure 1-3. (c) Operator carries input tape to 7094. (d) 7094 does computing. (e) Operator carries output tape to 1401. (f) 1401 prints output.

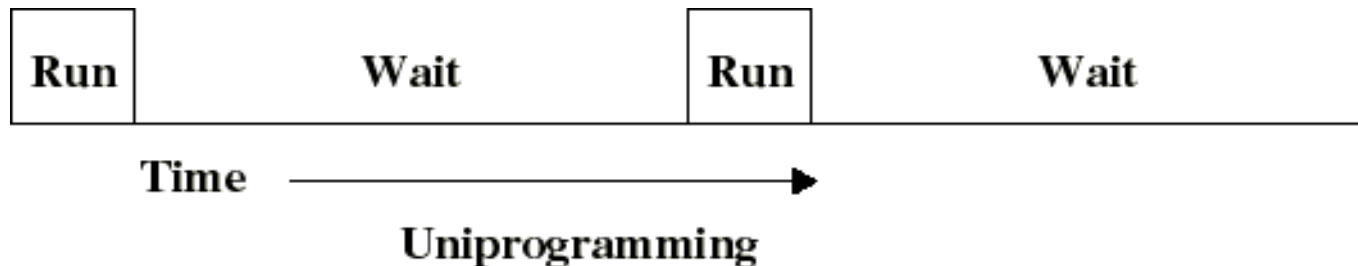
# Example card deck of a Job



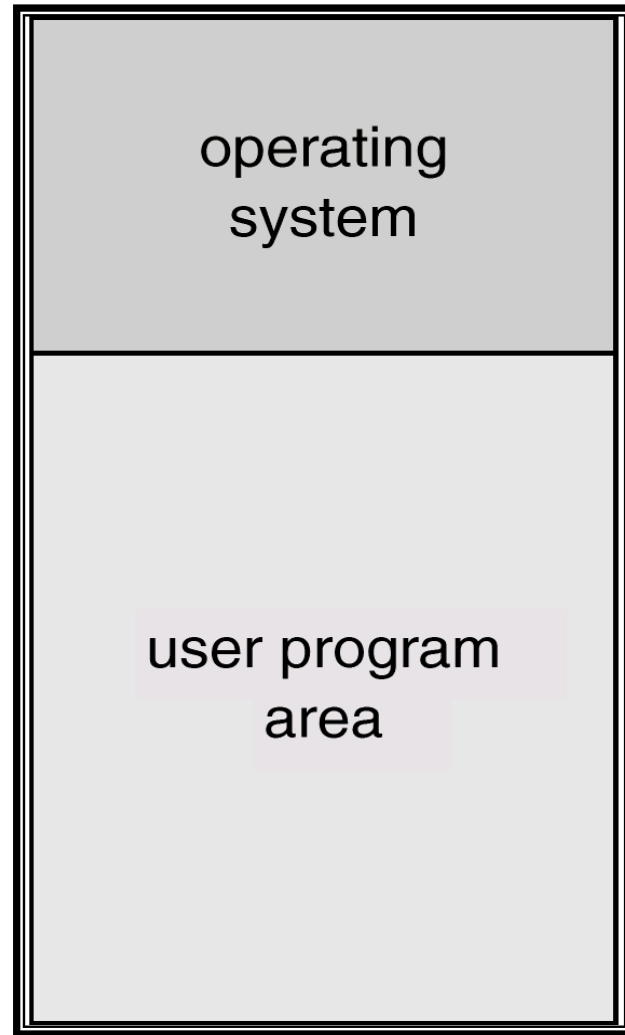


# We assumed Uniprogramming until now

- I/O operations are exceedingly slow (compared to instruction execution).
- A program containing even a very small number of I/O operations, will spend most of its time waiting for them.
- Hence: poor CPU usage when only one program is present in memory.

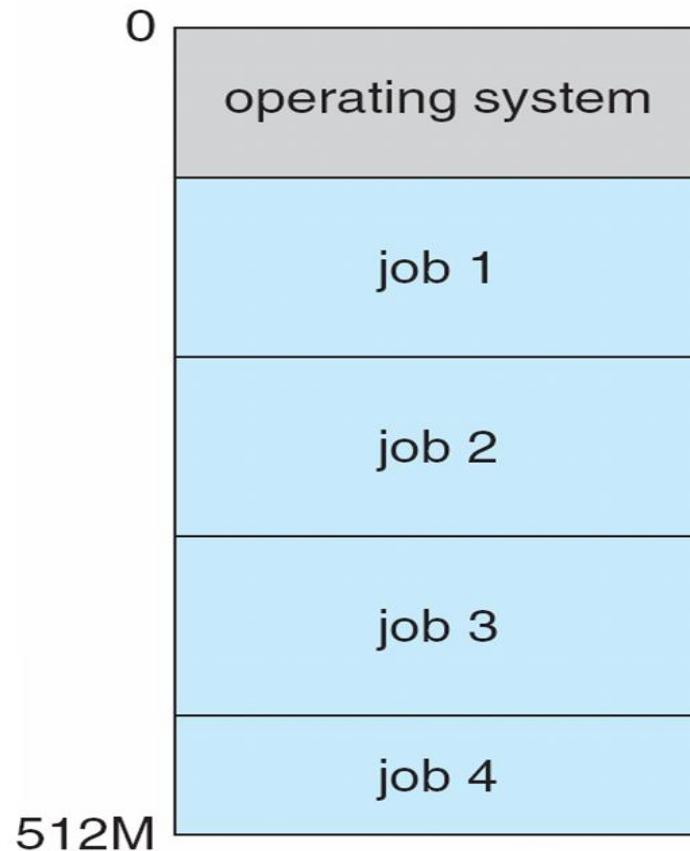


# Memory Layout for Uniprogramming

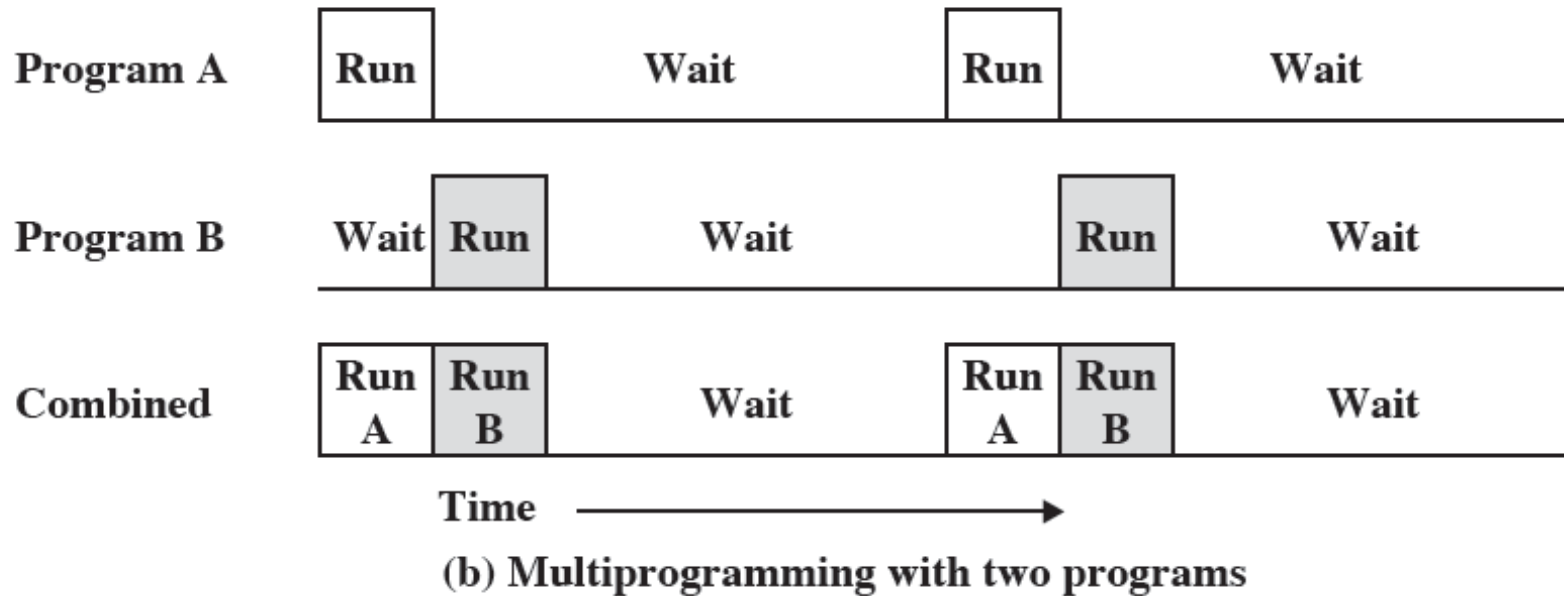
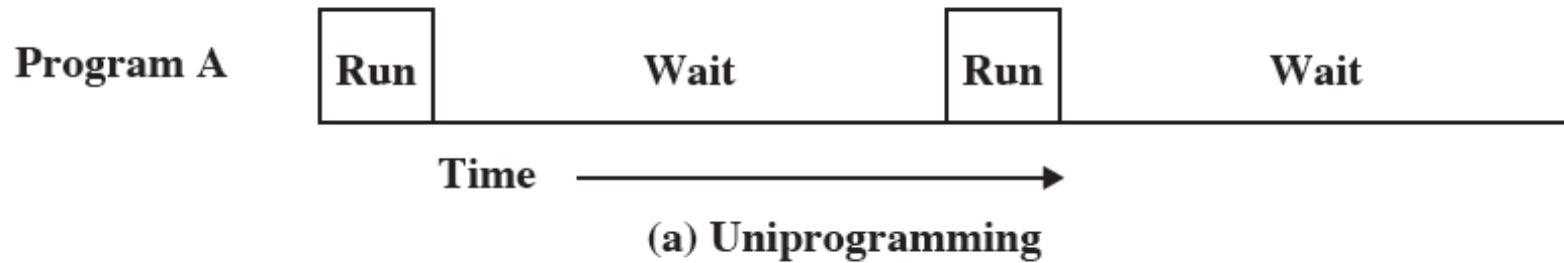


## Memory Layout for Batch Multiprogramming

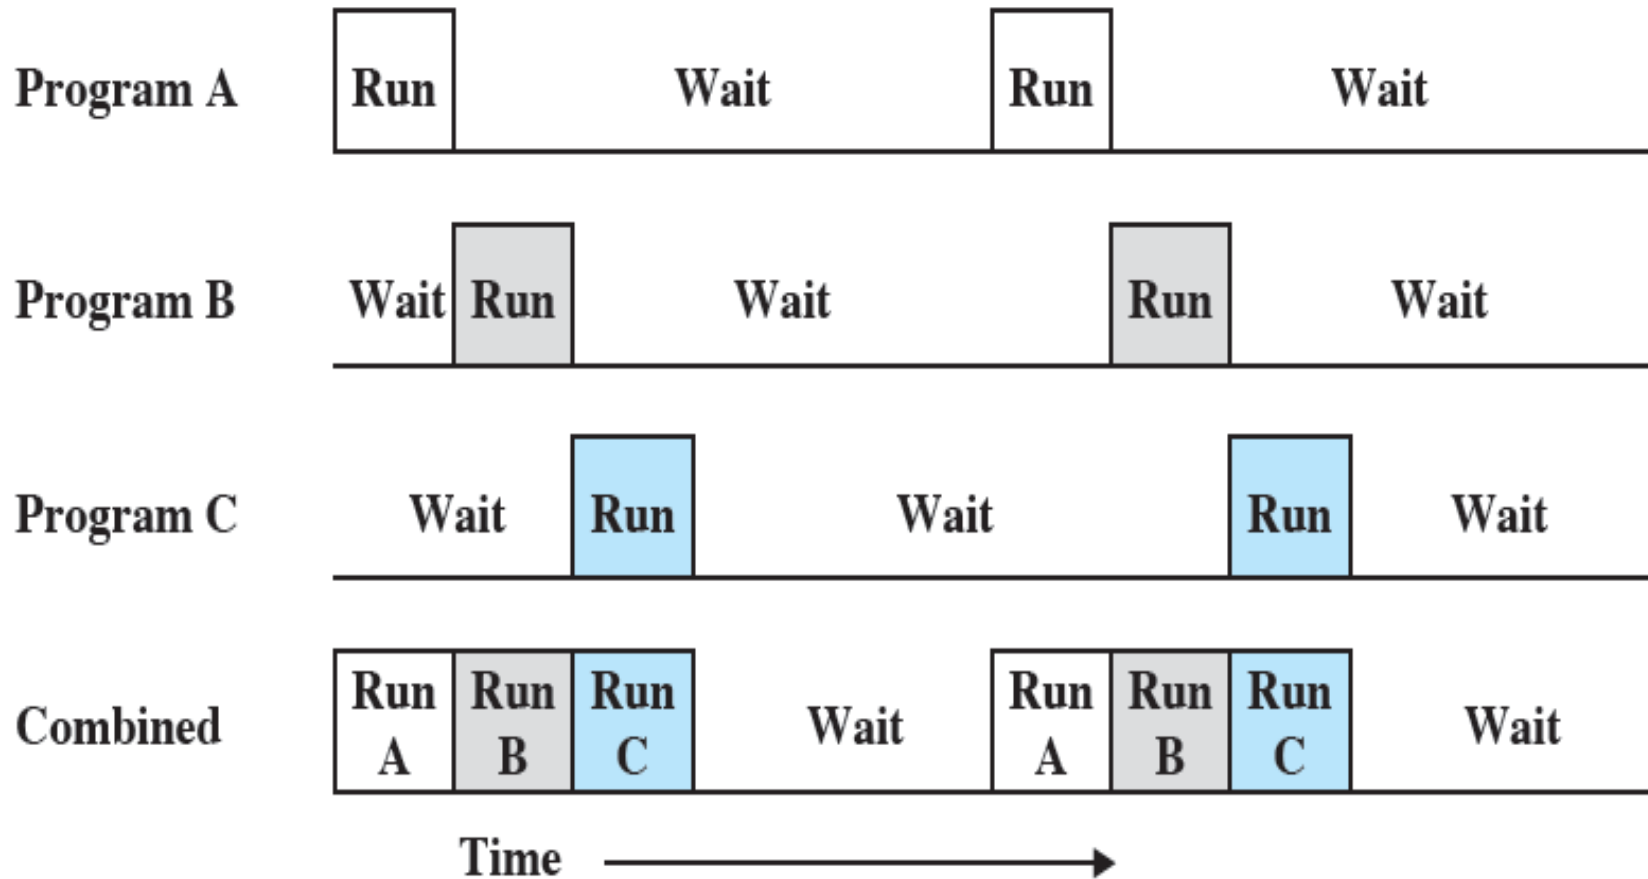
Several jobs are kept in main memory at the same time, and the CPU is multiplexed among them.



# Multiprogramming (1)



# Multiprogramming (2)



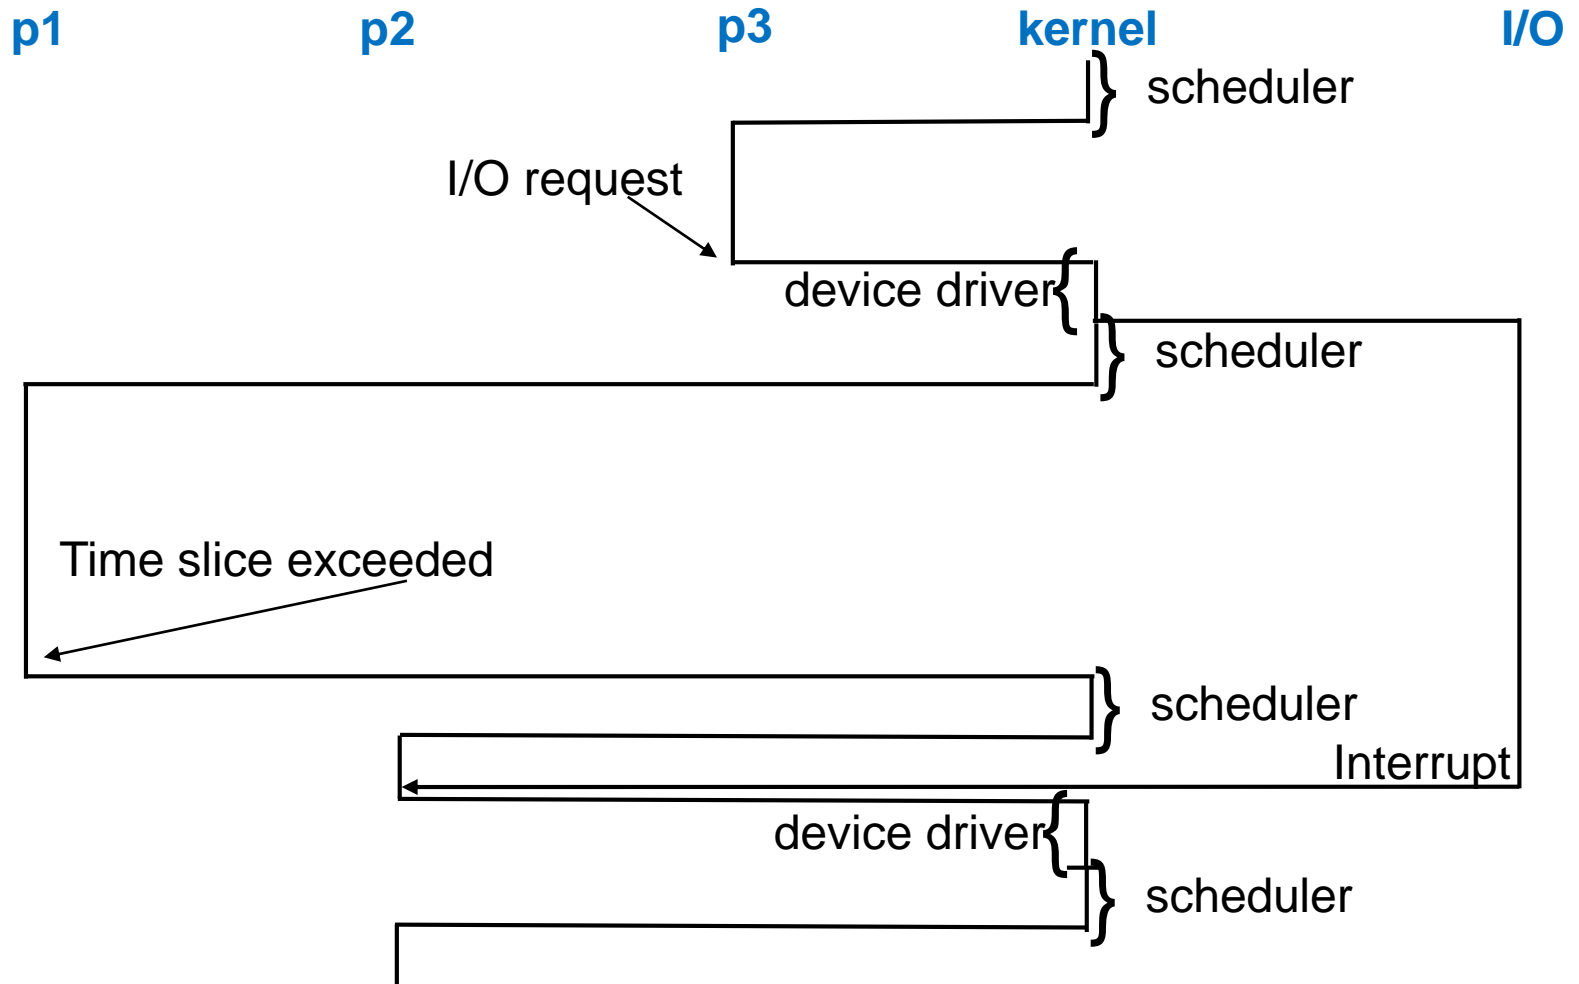
(c) Multiprogramming with three programs

# Why Multiprogramming?

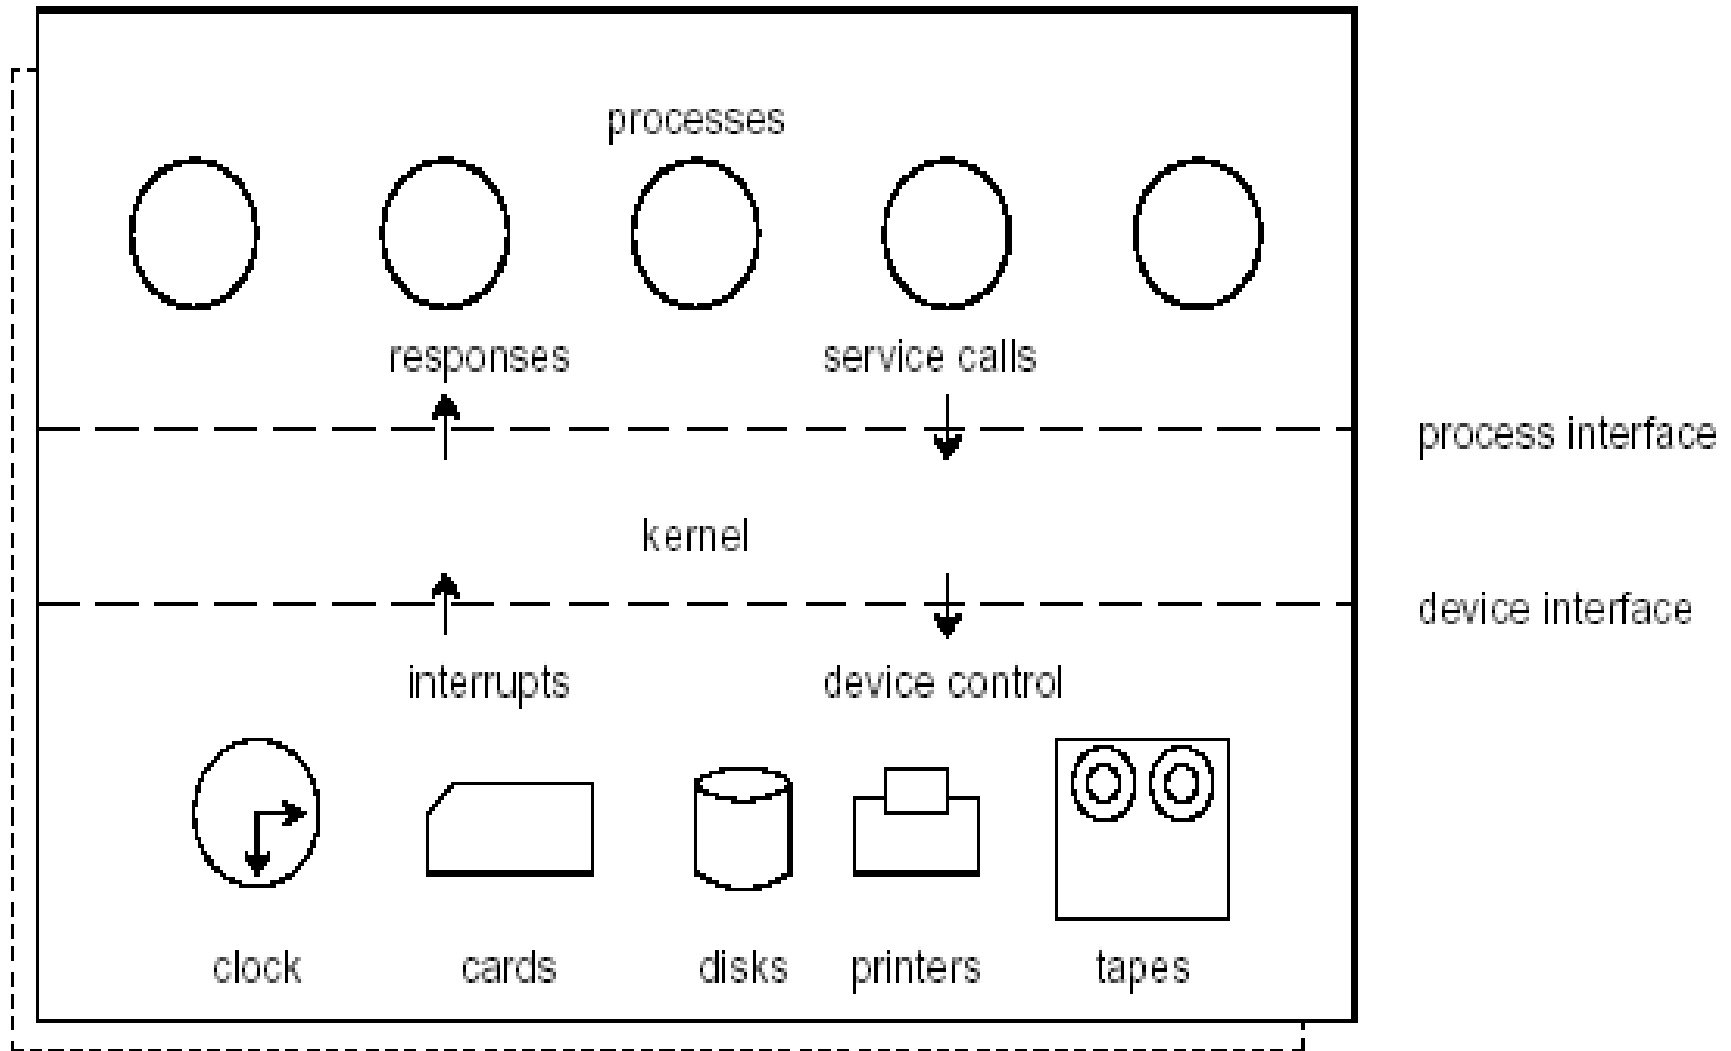
- Multiprogramming needed for efficiency:
  - Single user cannot keep CPU and I/O devices busy at all times.
  - Multiprogramming organizes jobs (code and data) so CPU always has one to execute.
  - A subset of total jobs in system is kept in memory.
  - One job selected and run via job scheduling.
  - When it has to wait (for I/O for example), OS switches to another job.

# Example of Multiprogramming

- Three processes (jobs) and OS are in memory.



# Components of Multiprogramming





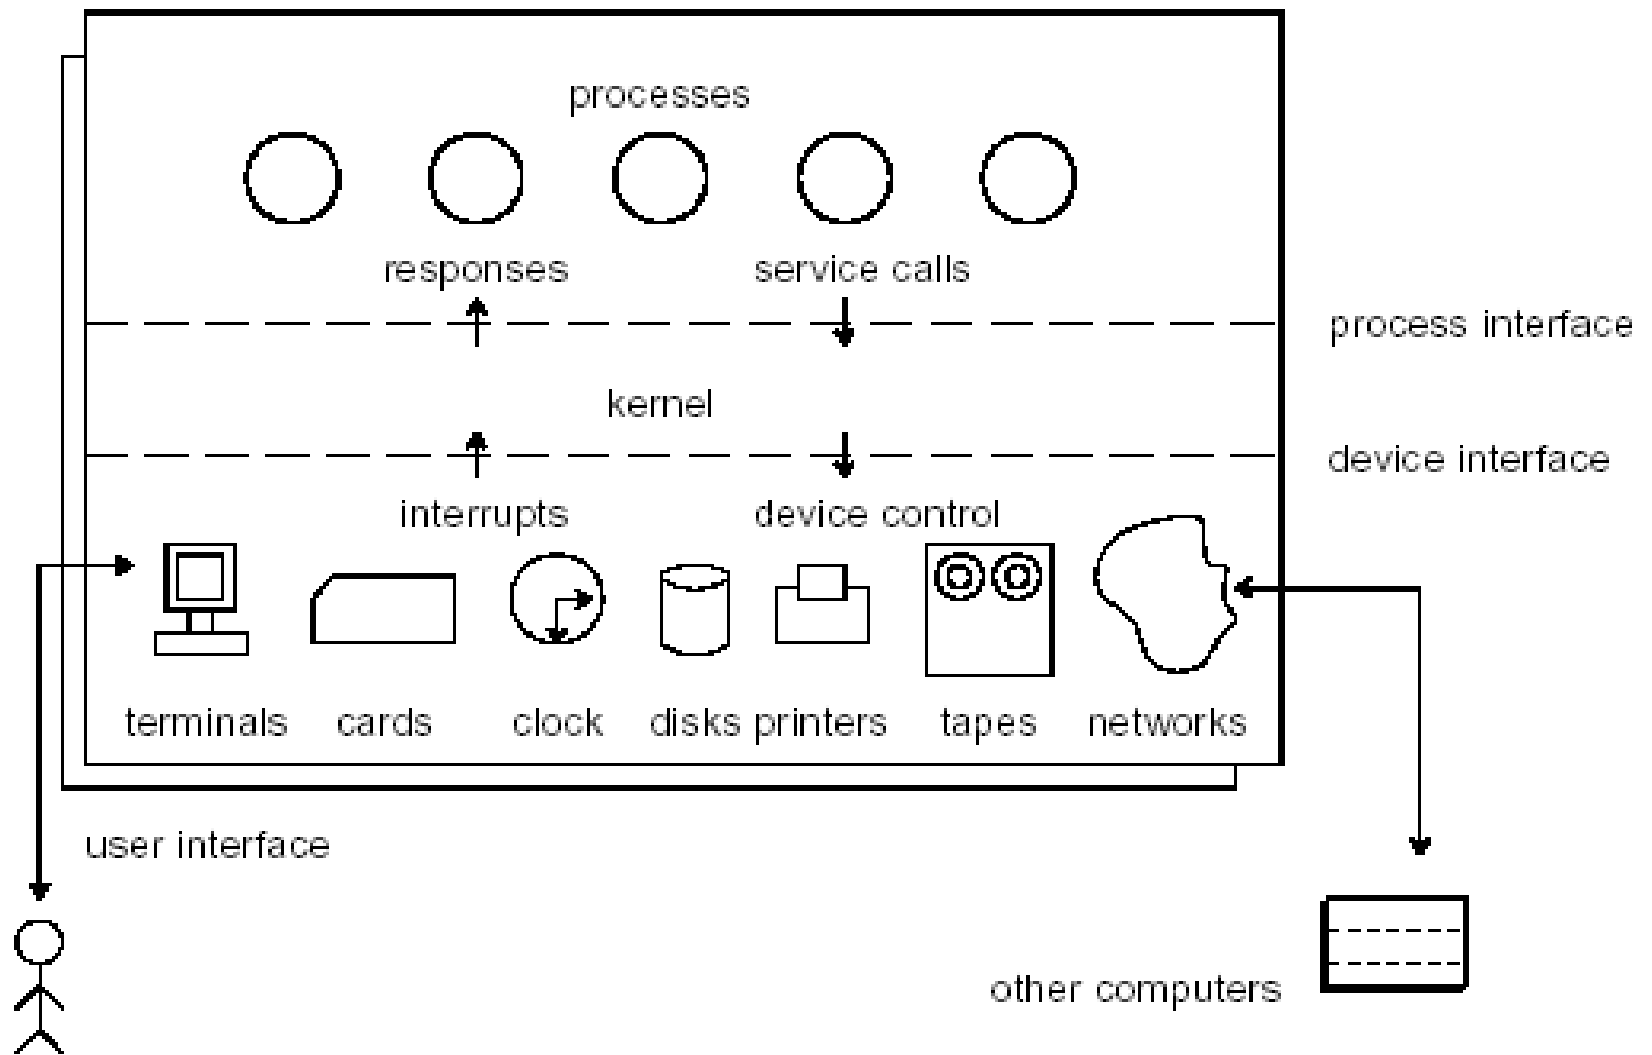
# Requirements for Multiprogramming

- **Hardware support:**
  - I/O interrupts and DMA controllers
    - in order to execute instructions while I/O device is busy.
  - Timer interrupts for CPU to gain control.
  - Memory management
    - several ready-to-run jobs must be kept in memory.
  - Memory protection (data and programs).
- **Software support from the OS:**
  - For scheduling (which program is to be run next).
  - To manage resource contention.

# Time-Sharing Systems (1970s)

- Batch multiprogramming does not **support interaction with users**.
- **Time-sharing** extends Batch Multiprogramming to handle multiple interactive jobs –  
it's Interactive Multiprogramming.
- Multiple users simultaneously access the system through commands entered at terminals.
- Processor's time is shared among multiple users.

# Time-sharing Architecture



# Why Time-sharing?

- In **Time-sharing** the CPU switches jobs so frequently that users can interact with each job while it is running, creating interactive computing:
  - System clock generate interrupt at the rate of 0.2 seconds and OS regains control and assigns processor to another user.
  - Each user has at least one program (process) executing in memory.
  - CPU scheduling supports several jobs ready to run at the same time.
  - If processes don't fit in memory, **swapping** moves them in and out to run.
  - **Virtual memory** allows execution of processes not completely in memory.
  - Both batch processing and time-sharing uses multi-programming.

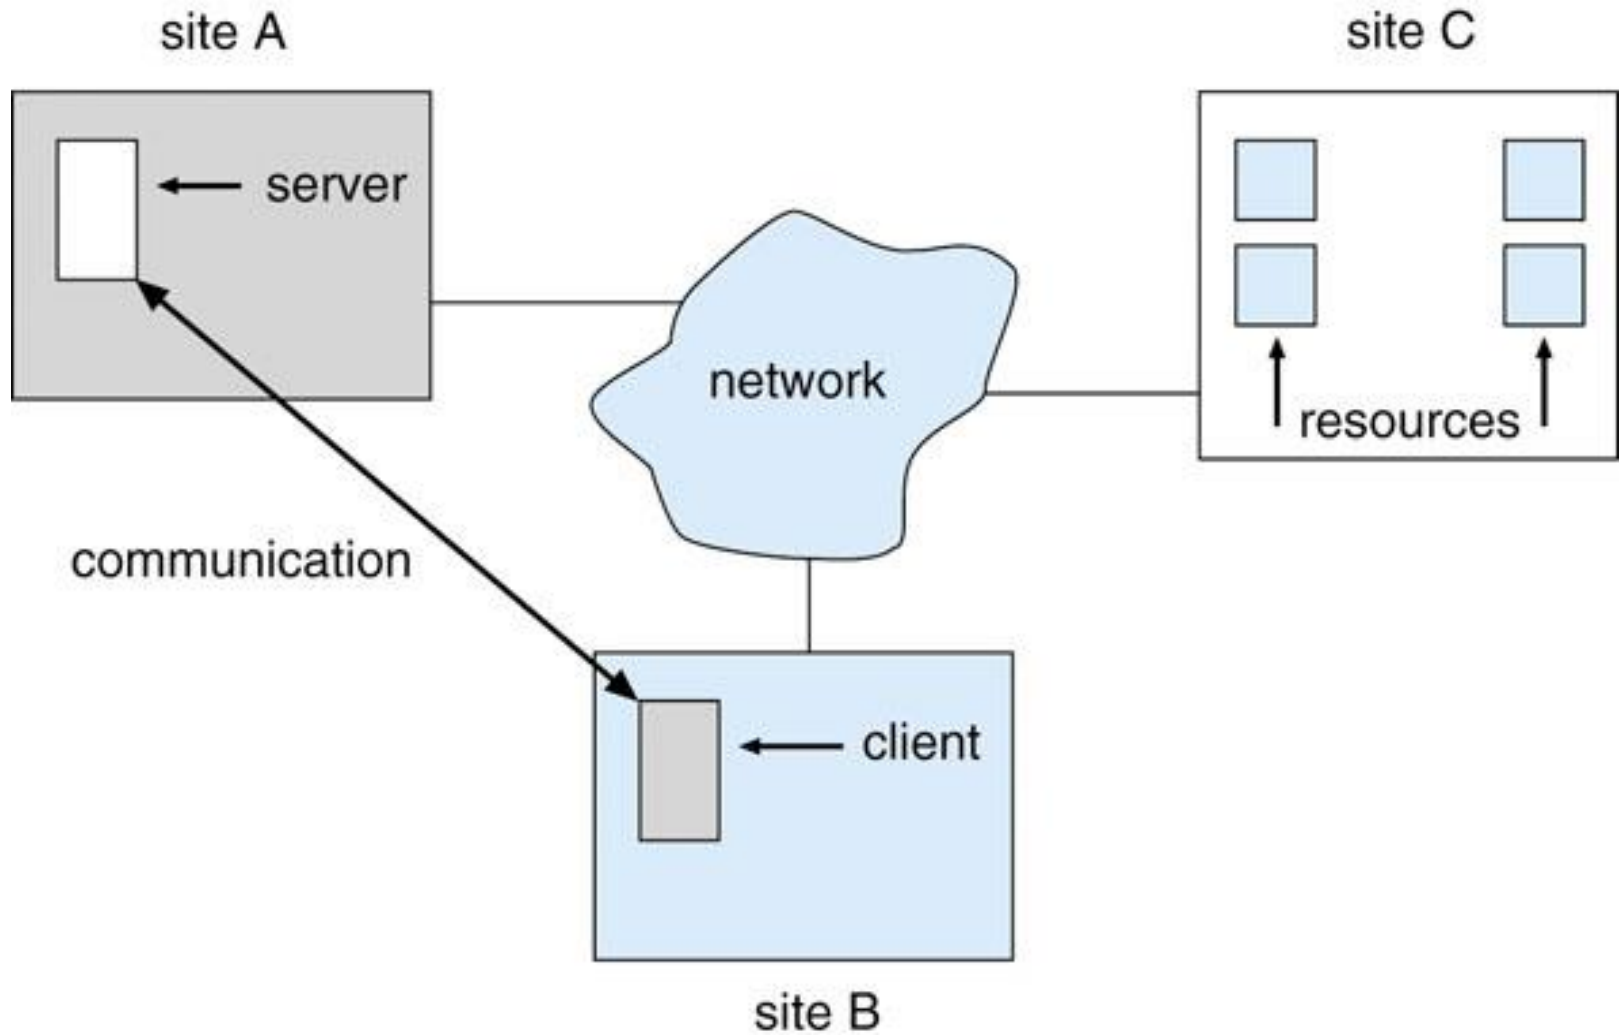
# Why did Time-Sharing work?

- Because of **slow human reaction** time, a typical user needs 2 seconds of processing time per minute.
- Then many users should be able to share the same system without noticeable delay in the computer reaction time.
- The user should get a good response time.

# Personal/Desktop Computers (1980s)

- *Personal computers* – computer system dedicated to a single user.
- **I/O devices:** keyboards, mice, display screens, small printers.
- User convenience and responsiveness.
- Can adopt technology developed for larger operating system; often individuals have sole use of computer.
- May run several different types of operating systems (Windows, MacOS, UNIX, Linux)

# Distributed Systems



# Distributed Systems

- **Distributed system** is collection of loosely coupled processors interconnected by a communications network.
- Processors variously called *nodes*, *computers*, *machines*, *hosts*.
- Reasons for **distributed systems**:
  - Resource sharing:
    - sharing files at remote sites.
    - processing information in a **distributed database**.
    - using remote specialized hardware devices.
  - Computation speedup – load sharing.
  - Reliability – detect and recover from site failure, function transfer, reintegrate failed site.
  - Communication – message passing.



# Chapter 1: Introduction

- What Operating Systems Do
- Computer-System Organization
- Computer-System Architecture
- Operating-System Structure
- **Common System Components**
  - **Process Management**
  - Memory Management
  - Storage Management

# 1. Process Management

- A process is a program in execution. It is a unit of work within the system. Program is a *passive entity*, process is an *active entity*.
- Process needs resources to accomplish its task
  - CPU, memory, I/O, files
  - Initialization data
- Process termination requires reclaim of any reusable resources

# 1. Process Management Activities

■ The operating system is responsible for the following activities in connection with **process management**:

- Creating and deleting both processes
- Suspending and resuming processes
- Providing mechanisms for process synchronization
- Providing mechanisms for process communication
- Providing mechanisms for deadlock handling

■ We will study more about **Process Management** in **Chapter#3**

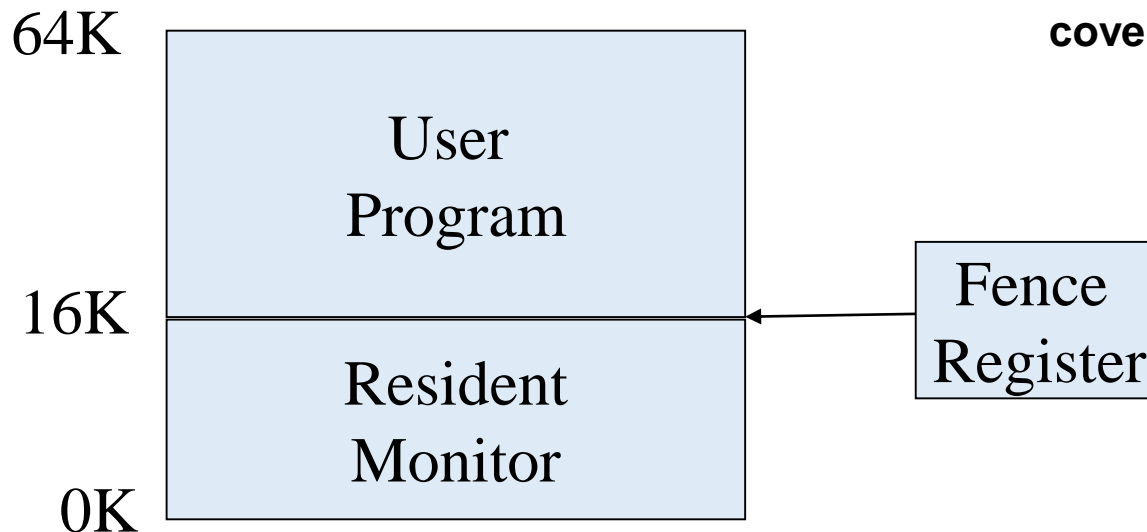
## 2. Main Memory Management

- Memory management determines what is in memory and when
  - *Optimizing* CPU utilization and computer response to users
- Memory management activities
  - Keeping track of which parts of memory are currently being used and by whom
  - Deciding which processes (or parts thereof) and data to move into and out of memory
  - Allocating and deallocating memory space as needed

# Memory Management Dynamics

- Sharing system resources requires the operating system to ensure that an incorrect program cannot cause other programs to execute incorrectly.
- **Resident Monitor** is a “Trusted Program” but how to protect it from damage by the user program?
- **Solution:** Fence Register (a dedicated register) and addressing access logic.

■ **Memory Management** will be covered in **Chapter#8 and #9**



# 3. Storage Management

- OS provides uniform, logical view of information storage
  - **Abstracts** physical properties to logical storage unit - **file**
  - Each medium is controlled by device (i.e., disk drive, tape drive)
    - Varying properties include *access speed*, *capacity*, *data-transfer rate*, *access method*
- File-System management
  - Files usually organized into directories
  - Access control on most systems to determine who can access what
  - OS activities include
    - Creating and deleting files and directories
    - **Primitives** (e.g., copy, move, delete) to manipulate files and directories
    - Backup files onto stable (non-volatile) storage media
- We will study more about **Storage Management** in **Chapter#10,11**

End of Chapter 1