

# Chapter 3: Processes

# Chapter 3: Processes

- **Process Concept**
- Process Scheduling
- Operations on Processes
- Interprocess Communication
- Examples of IPC Systems
- Communication in Client-Server Systems

# Objectives


- To introduce the notion of a process -- a program in execution, which forms the basis of all computation
- To describe the various features of processes, including *scheduling*, *creation and termination*, and *communication*
- To explore interprocess communication using shared memory and message passing
- To describe communication in client-server systems

# Process Concept (Cont.)

- Textbook uses the terms *job* and *process* almost interchangeably
- **Process** – a program in execution; process execution must progress in sequential fashion
- Program is *passive* entity stored on disk (**executable file**), process is *active*
  - Program becomes process when executable file loaded into memory
- Execution of program started via GUI mouse clicks, command line entry of its name, etc


# Process Address Space

- A list of memory locations from some min (usually 0) to some max that a process can read and write.



# Process Concept

- Multiple parts
  - The program code, also called **text section**
  - Current activity including **program counter**, processor registers
  - **Stack** containing temporary data
    - Function parameters, return addresses, local variables
  - **Data section** containing global variables
  - **Heap** containing memory dynamically allocated during run time
- One program can be several processes
  - Consider multiple users executing the same program (**Several email programs or web browsers**)
  - They are separate processes with equivalent code segment (i.e. **same text section**)
  - Program → blueprint, process → an instance



# Process Control Block (PCB)

Information associated with each process

(Process represented in OS by PCB also called **task control block**)

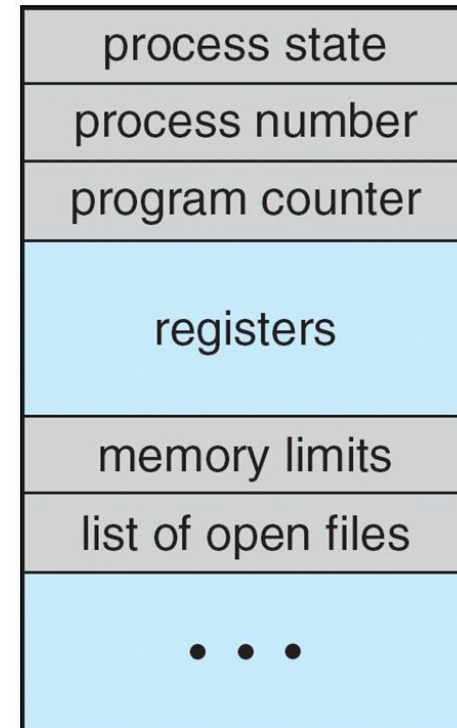
- PCB (**execution context**) is the data needed (process attributes) by OS to control process
  1. Process location information
  2. Process identification information
  3. Processor state information
  4. Process control information (mode)

# Process Control Block (PCB)

Information associated with each process

(Process represented in OS by PCB also called **task control block**)

- **Process state** – running, waiting, etc
- **Program counter** – location of instruction to next execute
- **CPU registers** – contents of all process-centric registers
- **CPU scheduling information** – priorities, scheduling queue pointers
- **Memory-management information** – memory allocated to the process
- **Accounting information** – CPU used, clock time elapsed since start, time limits
- **I/O status information** – I/O devices allocated to process, list of open files





# Role of the Process Control Block

- The most important *data structure* in an OS
  - contains all of the information about a process that is needed by the OS
  - blocks are read and/or modified by virtually every module in the OS
  - defines the state of the OS
- Protection
  - a bug in a single routine could damage process control blocks, which could destroy the system's ability to manage the affected processes
  - a design change in the structure or semantics of the process control block could affect a number of modules in the OS

# Process States (A Three-state Process Model)

- Let us start with three states:
  - 1) **Running state** –
    - the process that gets executes; its instructions are being executed.
  - 2) **Ready state** –
    - any process that is ready to be executed; the process is waiting to be assigned to a processor.
  - 3) **Waiting/Blocked state** –
    - when a process cannot execute until its I/O completes or some other event occurs.


# Process Transitions (1)

- **Ready → Running**

- When it is time, the dispatcher selects a new process to run.

- **Running → Ready**

- the running process has expired his time slot.
- the running process gets interrupted because a higher priority process is in the ready state.




# Process Transitions (2)

- Running → Waiting (blocking)

- When a process requests something for which it must wait:
  - a service that the OS is not ready to perform.
  - an access to a resource not yet available.
  - initiates I/O and must wait for the result.
  - waiting for a process to provide input.

- Waiting → Ready

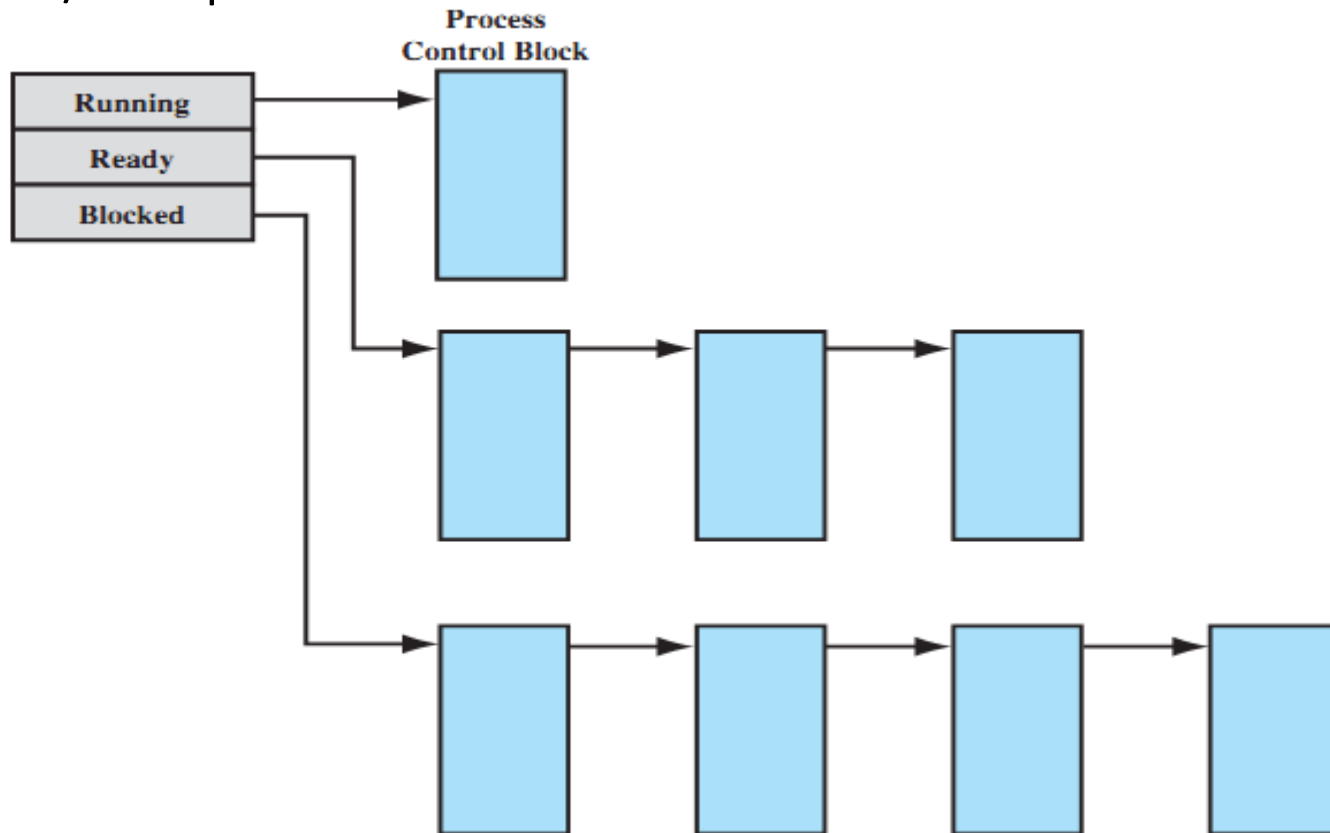
- When the event for which it was waiting occurs.



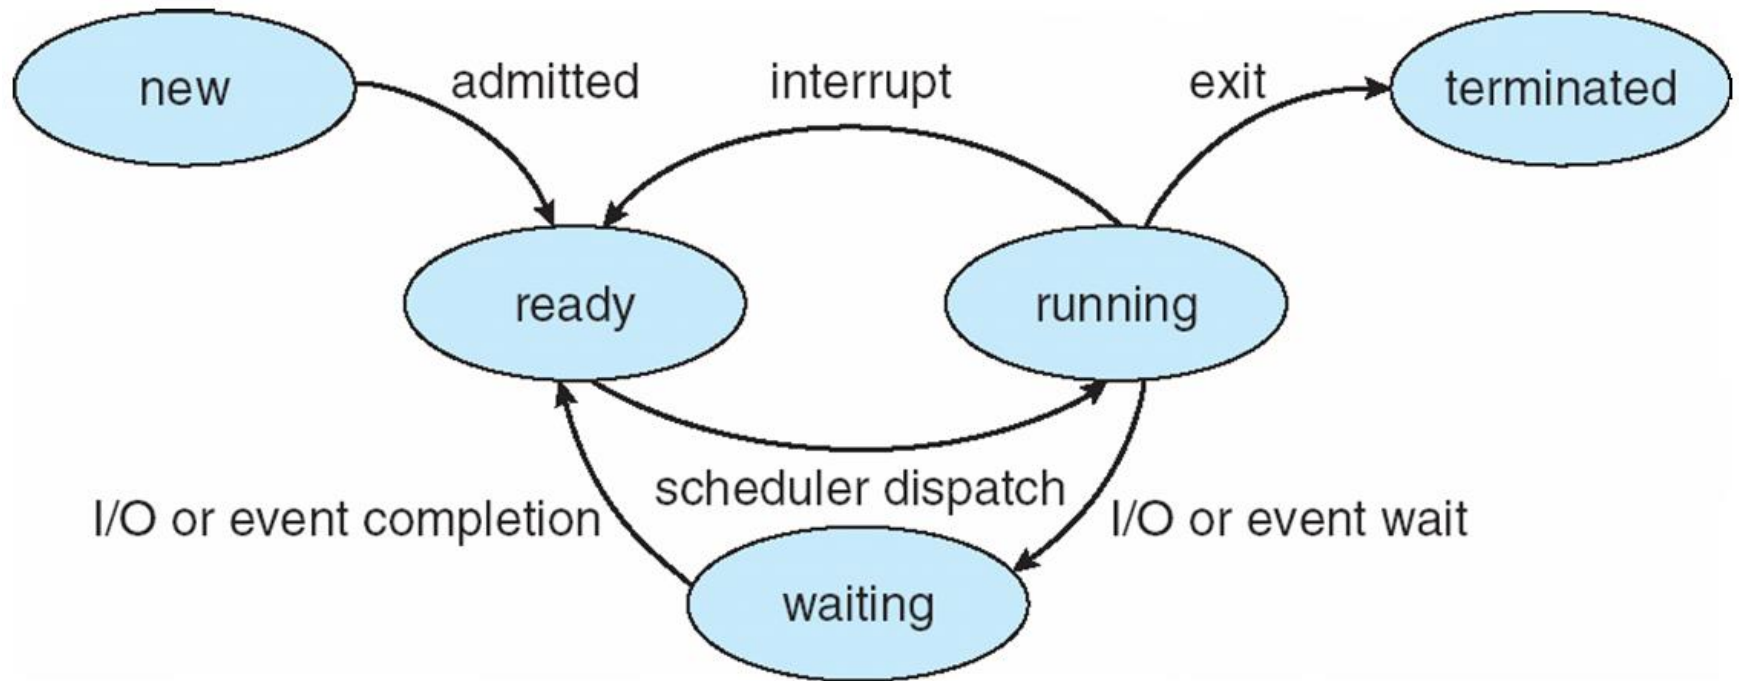
# Process List Structures

## ■ Three queues in three state process model:

- Running Queue,
- Ready Queue
- Blocked/wait queue



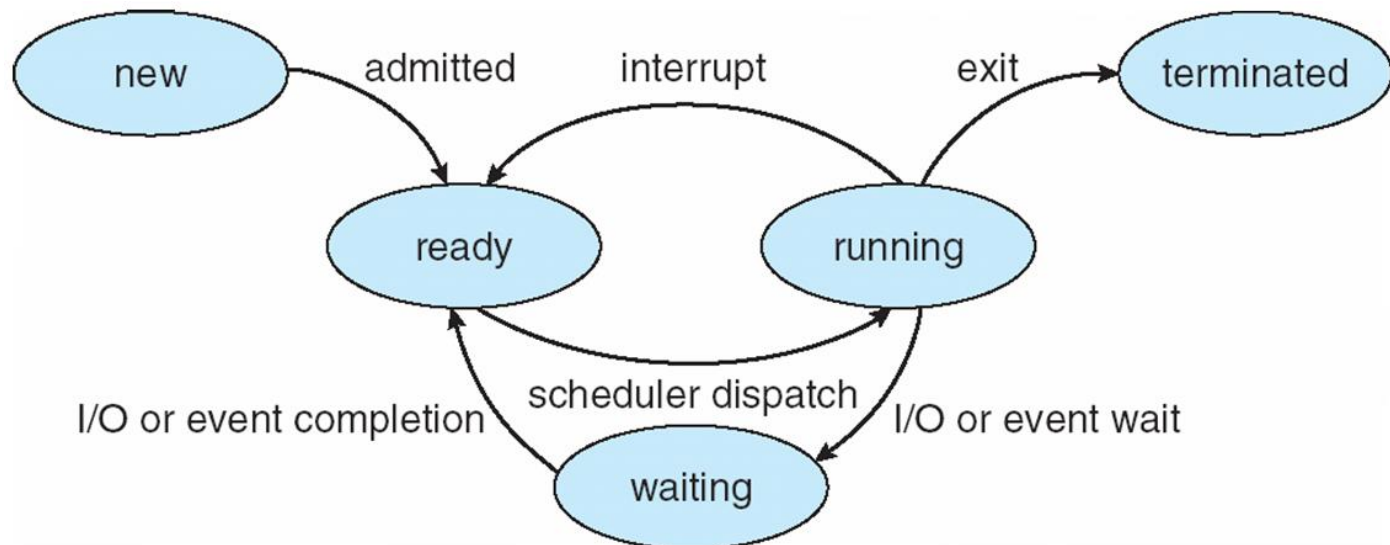
# Five State Process Model



# Five State Process Model

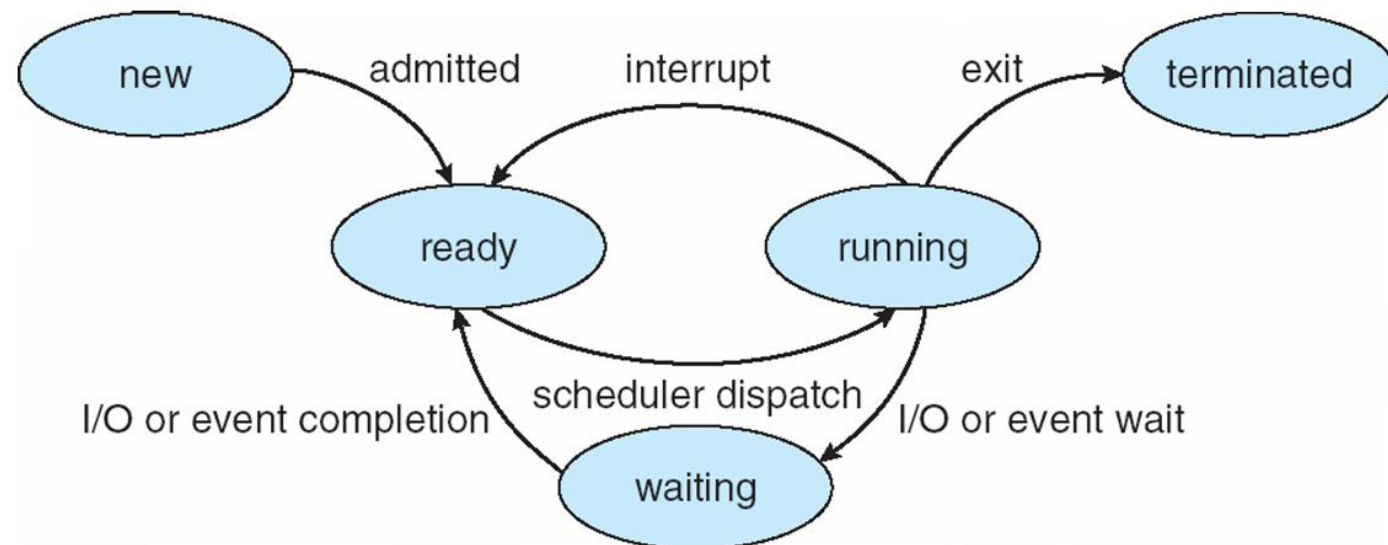
- **New state –**

- OS has performed the necessary actions to create the process:
  - has created a process identifier.
  - has created tables needed to manage the process.
- but has not yet committed to execute the process (not yet admitted):
  - because resources are limited.



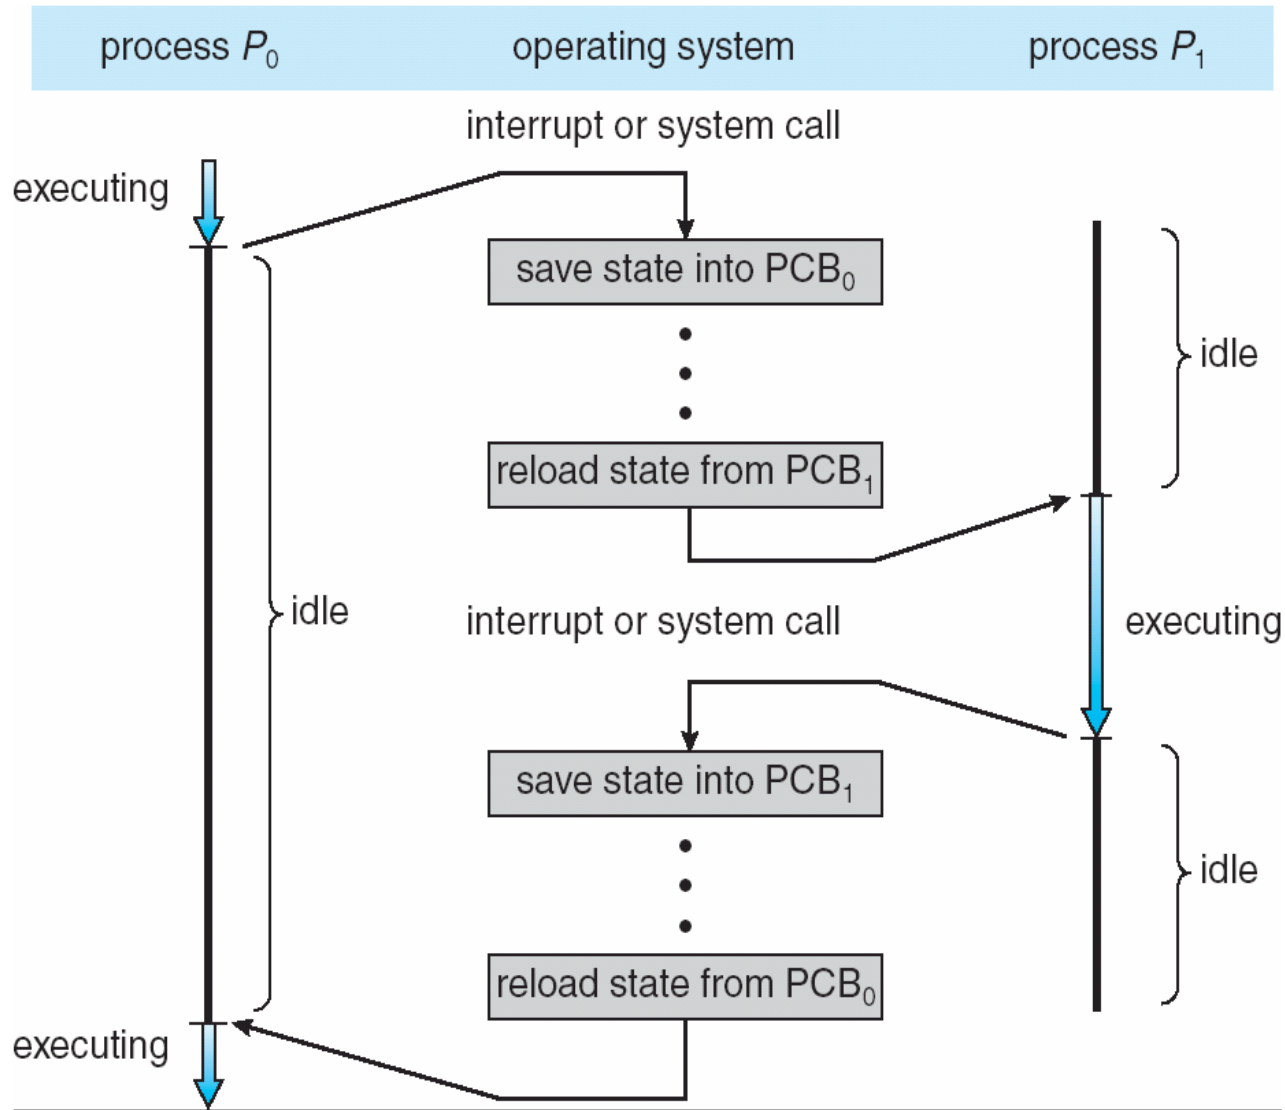
# Five State Process Model

- **Terminated state** –
  - Program termination moves the process to this state.
  - It is no longer eligible for execution.
  - Tables and other info are temporarily preserved for auxiliary program –
- **Example:** accounting program that cumulates resource usage for billing the users.
- The process (and its tables) gets deleted when the data is no more needed.






# CPU Switch From Process to Process

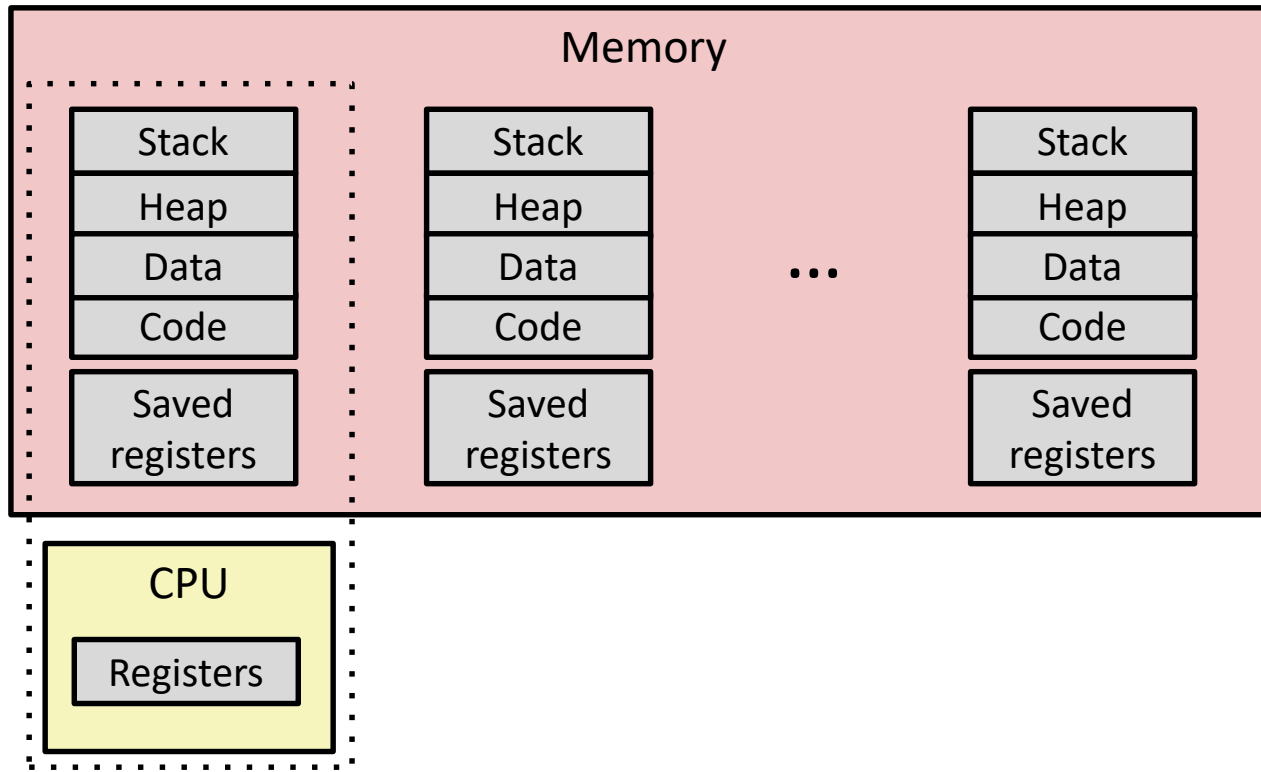


# Multiprocessing: Illusion vs Reality



- Computer runs many processes simultaneously
  - Applications for one or more users
    - Web browsers, email clients, editors, ...
    - Monitoring network & I/O devices

# Multiprocessing: Illusion vs. Reality



- Single processor executes multiple processes *concurrently*
  - Process executions interleaved, CPU runs *one at a time*
  - Address spaces managed by virtual memory system (later in course)
  - *Execution context* (register values, stack, ...) for other processes saved in memory

# What is a Context Switch?

- When CPU switches to another process, the system must **save the state** of the old process and load the **saved state** for the new process via a **context switch**
- **Context** of a process represented in the PCB
  - = process state, all register values, memory information
  - Save/restore contexts to/from PCBs when switching among processes
    - Known as **context switch**
- **Context-switch time** is overhead; the system does no useful work while switching
  - The more complex the OS and the PCB → the longer the context switch. **Typical speed is a few milliseconds**
    - **Depends on machine: memory speed, number of registers, load/save instructions**
  - Time dependent on hardware support

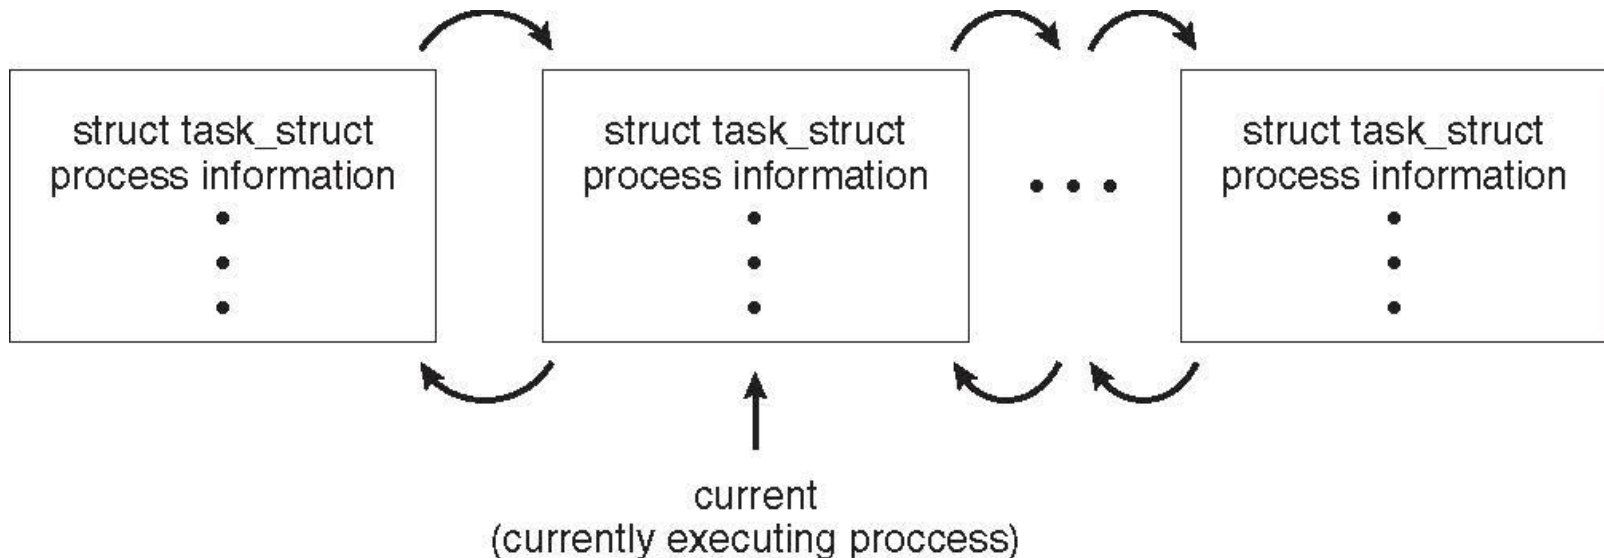
# Context switch overhead

- Assume a processor runs for a total of 10 seconds, during which it performs 4 context switches, each of them takes one second.
  - What is the overhead (in %) due to context switches?

# Process Representation in Linux

Represented by the C structure `task_struct`

```
pid_t pid; /* process identifier */  
long state; /* state of the process */  
unsigned int time_slice; /* scheduling information */  
struct task_struct *parent; /* this process's parent */  
struct list_head children; /* this process's children */  
struct files_struct *files; /* list of open files */  
struct mm_struct *mm; /* address space of this process */
```



# Chapter 3: Processes

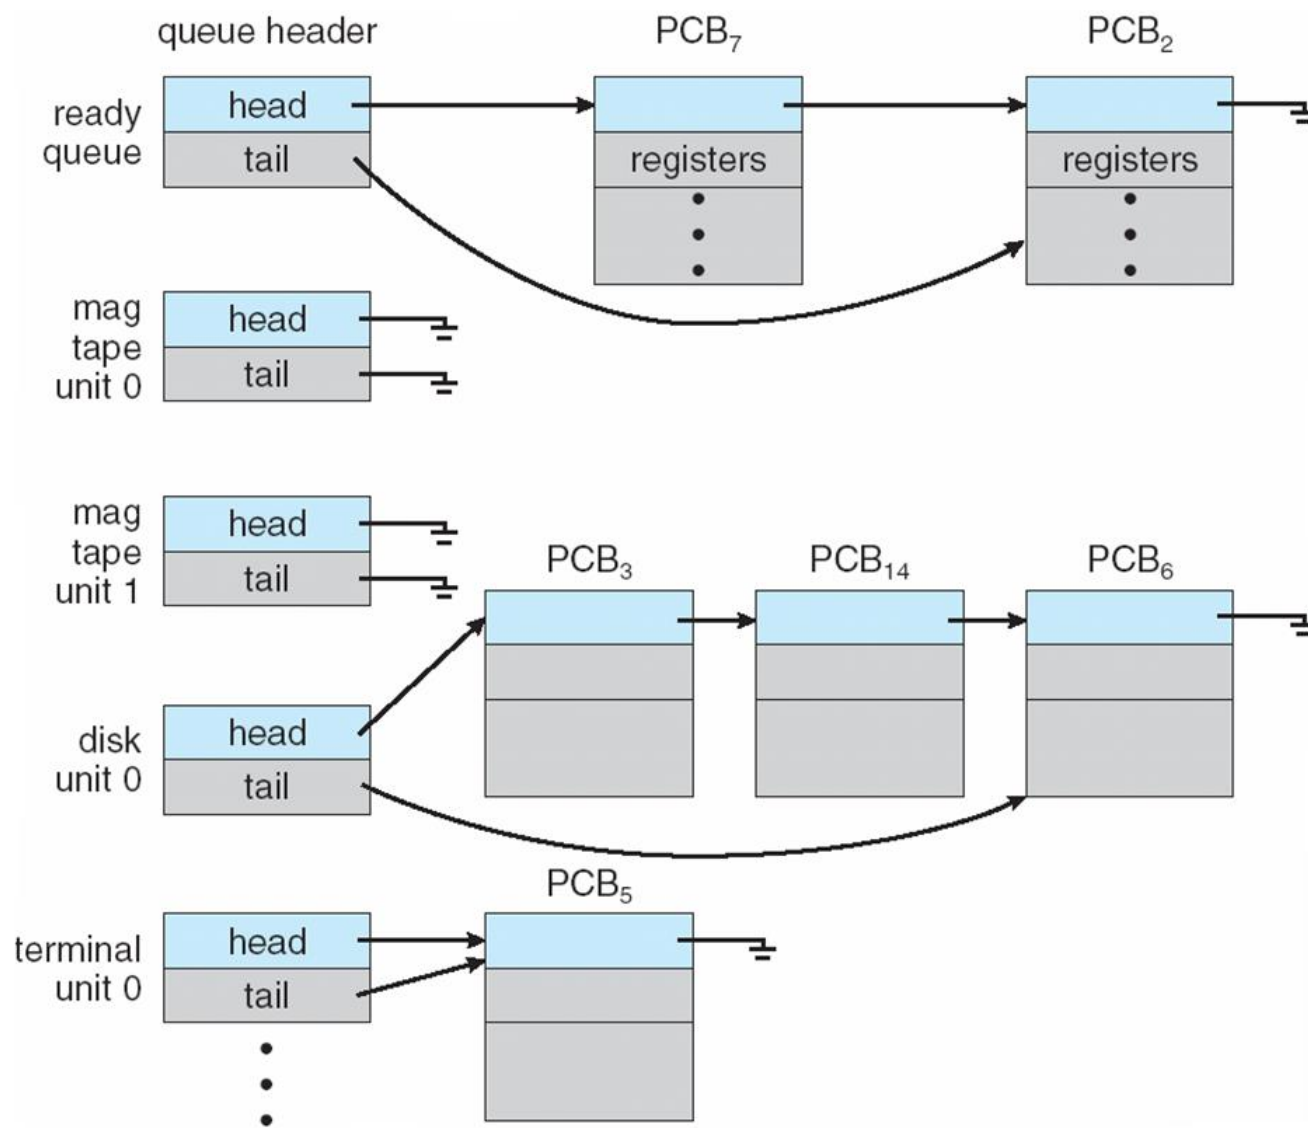
- Process Concept
- **Process Scheduling**
- Operations on Processes
- Interprocess Communication
- Examples of IPC Systems
- Communication in Client-Server Systems

# Process Scheduling

- **OS Objective:** to Maximize CPU use, quickly switch processes onto CPU for time sharing; so that users can interact with programs
- **Process scheduler** selects among available processes for next execution on CPU
  - Process scheduler = CPU scheduler + Job scheduler + other schedulers (e.g., for swapping etc. )
- Maintains **scheduling queues** of processes
  - **Job queue** – set of all processes in the system
  - **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
    - = Linked list of PCBs
  - **Device queues** – set of processes waiting for an I/O device
    - Each shared device has its associated device queue
- Processes migrate among the various queues

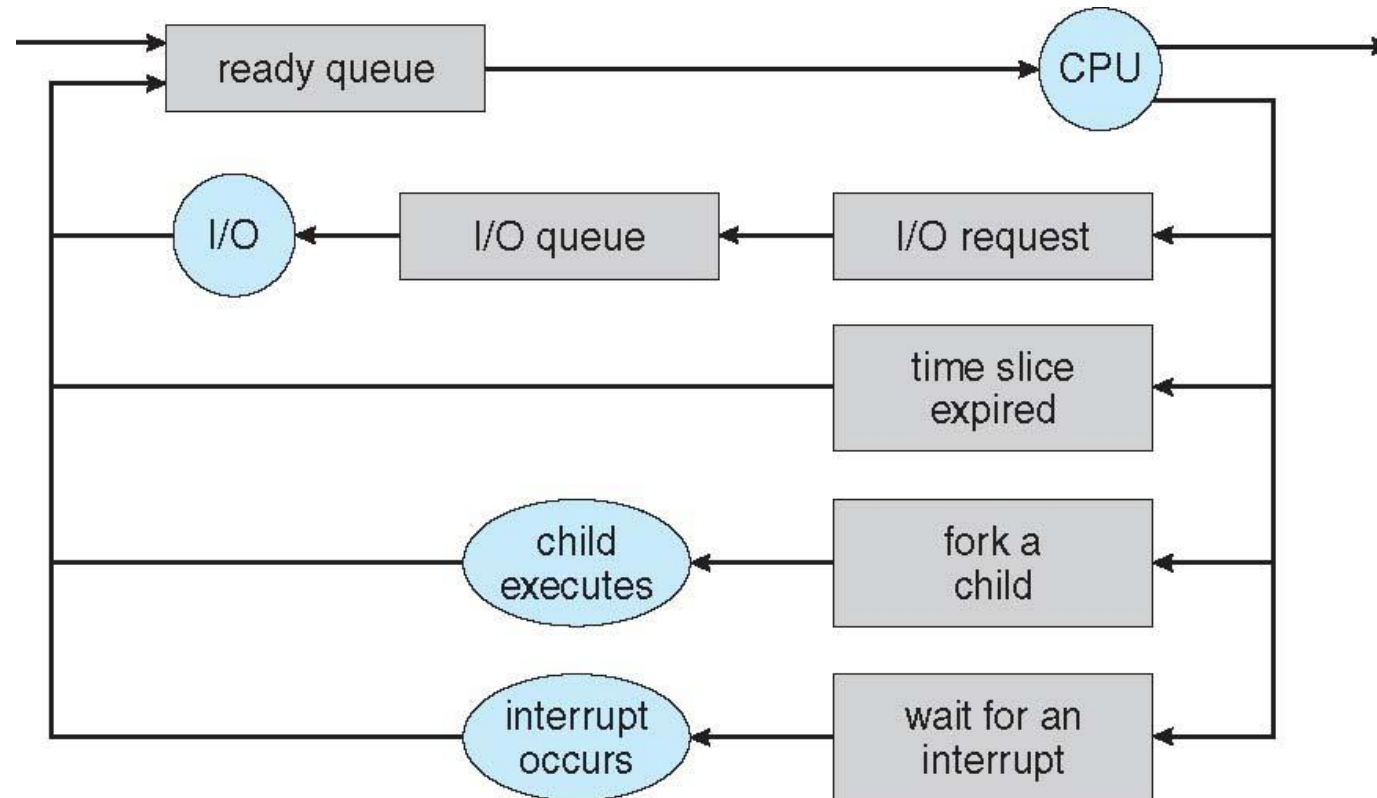


# Ready Queue And Various I/O Device Queues



# Representation of Process Scheduling

- **Queueing diagram** represents queues, resources, flows

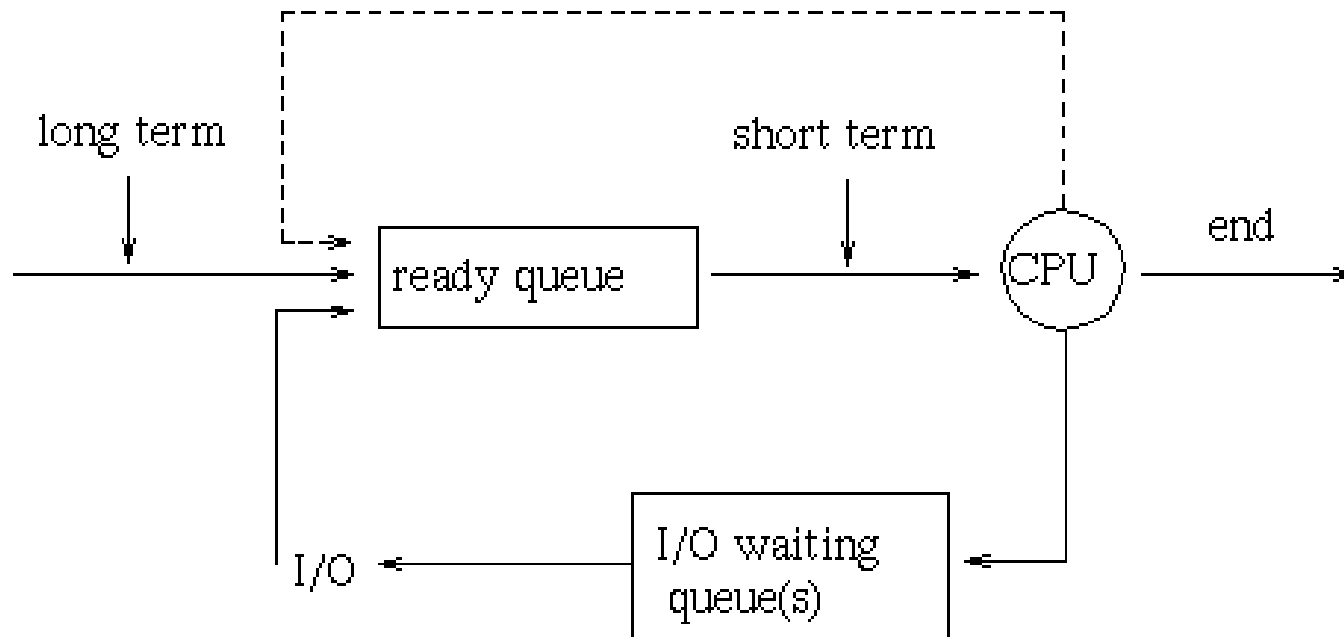


# Schedulers

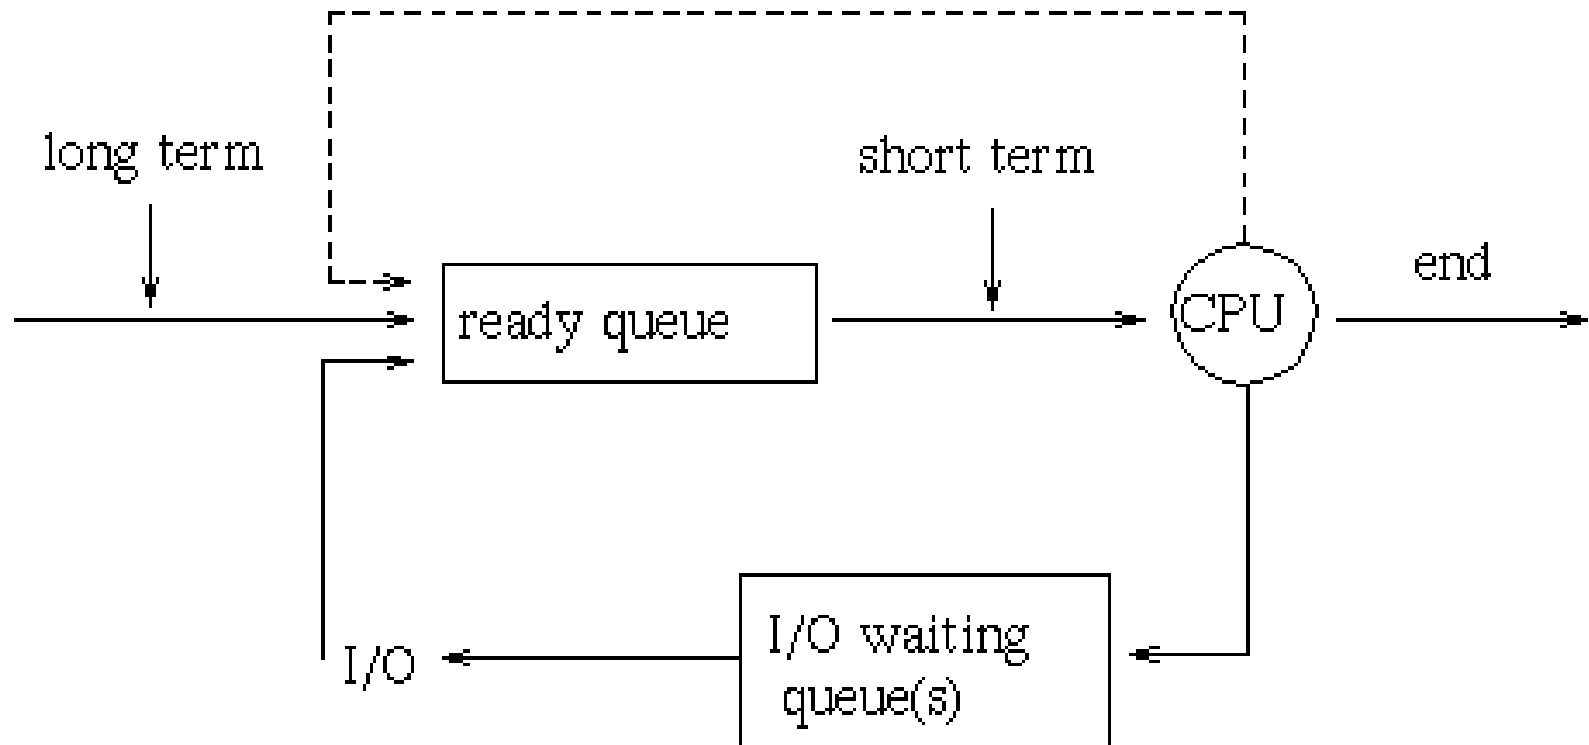
- **Short-term scheduler** (or **CPU scheduler/dispatcher**)
- **Long-term scheduler** (or **job scheduler**)
- **Medium-term scheduler**

# Schedulers

- **Short-term scheduler** (or **CPU scheduler/dispatcher**) – selects which process should be executed next and allocates CPU
  - Sometimes the only scheduler in a system.
  - Short-term scheduler is invoked frequently (milliseconds)  $\Rightarrow$  (must be fast)
  - Main objective is to increase CPU performance



# Schedulers




# Schedulers (Cont'd)

- **Long-term scheduler** (or **job scheduler**) – selects which processes should be brought into the ready queue
  - Long-term scheduler is invoked infrequently (seconds, minutes)  $\Rightarrow$  (may be slow)
  - The long-term scheduler controls the **degree of multiprogramming**
    - = Number of processes in memory (i.e., in the ready queue)
    - The **degree of multiprogramming** describes the maximum number of processes that a single-processor system can accommodate efficiently.
  - If the degree of multiprogramming is stable, then the average rate of process creation must be equal to the average departure rate of processes leaving the system, i.e., **aver no. of process creation = aver no. of process departure.**
  - Thus, invoked when a process leaves the system

# Schedulers(cont.)

- Processes can be described as either:
  - I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
    - The ready queue is almost always empty if all processes are I/O-bound
  - CPU-bound process** – spends more time doing computations; few very long CPU bursts
    - The I/O queue is almost always empty if all processes are CPU-bound
- Long-term scheduler** strives for good *process mix*
- The primary objective of the job scheduler is to provide a balanced mix of jobs, such as I/O bound and processor bound.

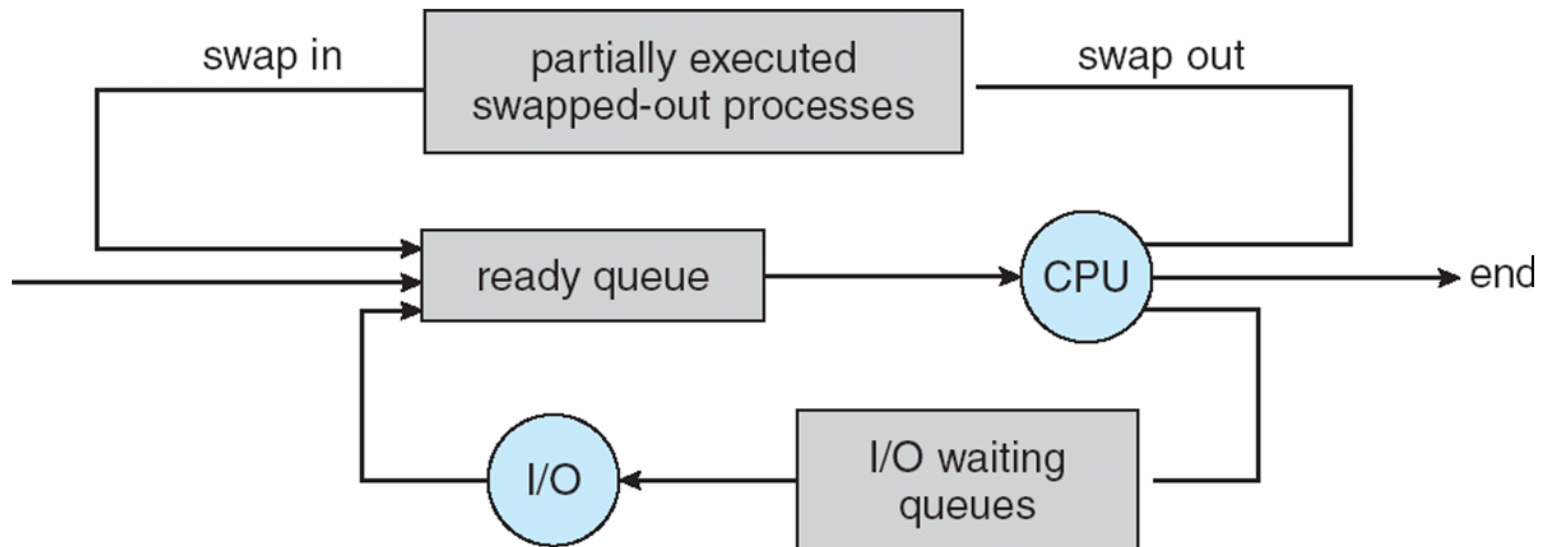


**Bursts of CPU usage alternate with periods of waiting for I/O. (a) A CPU-bound process. (b) An I/O-bound process.**

# Addition of Medium Term Scheduling

■ **Medium-term scheduler** can be added if degree of multiprogramming needs to decrease

- Remove process from memory, store on disk, bring back in from disk to continue execution: **swapping**
- **Swapping helps improve process mix**
- **Also necessary when memory needs to be freed up**





# Chapter 3: Processes

- Process Concept
- Process Scheduling
- **Operations on Processes**
- Interprocess Communication
- Examples of IPC Systems
- Communication in Client-Server Systems