

# Chapter 3: Processes (Part 2)

# Chapter 3: Processes

- Process Concept
- Process Scheduling
- **Operations on Processes**
- Interprocess Communication
- Examples of IPC Systems
- Communication in Client-Server Systems

# Operations on Processes

- There are three commonly seen operations:
  - **Process Creation:** Creates a new process. The newly created is the child of the original. Unix uses *fork()* system call to create new processes.
  - **Process Termination:** Terminate the execution of a process. Unix uses *exit()*.
  - **Process Join:** Wait for the completion of a child process. Unix uses *wait()*.
- *fork()*, *exit()* and *wait()* are system calls.


# Process Creation

- **Parent** process creates **children** processes, which, in turn create other processes, forming a **tree** of processes
- Generally, process identified and managed via a **process identifier (pid)**
- Resource (files etc. ) sharing **options**
  1. Parent and children share all resources
  2. Children share subset of parent's resources
  3. Parent and child share no resources
- Execution **options**
  1. Parent and children execute concurrently
  2. Parent waits until children terminate
- Address space options
  1. Child process is duplicate of the parent process (same program and data)
  2. Child process has a new program loaded into it.


# Process Hierarchy

- Whilst the operating system can run many processes at the same time, in fact it only ever directly starts one process called **the init** (short for initial) process. This isn't a particularly special process except that it's **PID is always 1** and it will always be running.
  - **Init** is started by the kernel during the booting process;
- All other processes can be considered children of this **initial process**. Processes have a family tree just like any other; each process has a parent and can have many siblings, which are processes created by the same parent.
- Certainly children can create more children and so on and so forth.
- The term **spawn** is often used when talking about parent processes creating children; as in "**the process spawned a child**".
- **ps tree** is a linux command that shows running processes as a tree


# Process Hierarchy




A running process



A waiting process



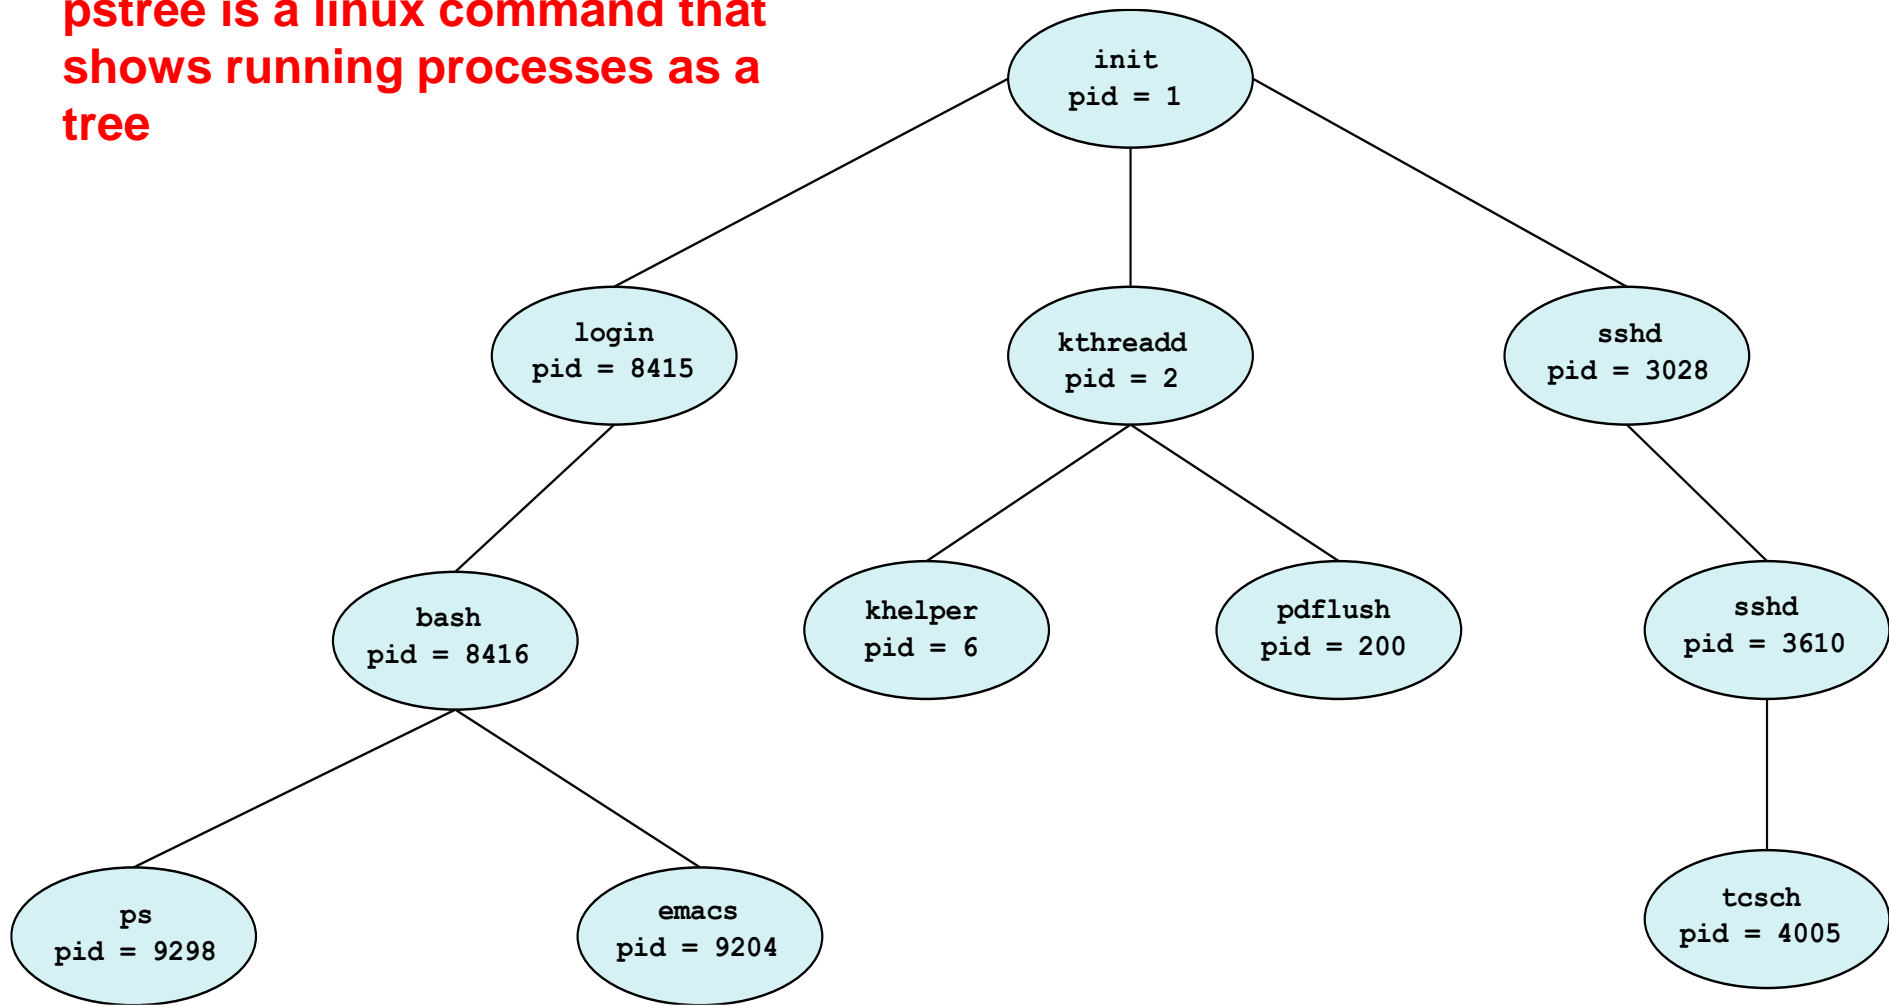
Parent waits for child  
Uniprogrammed



Parent free  
Multiprogrammed

# A Tree of Processes in Linux

**pstree is a linux command that shows running processes as a tree**



# The fork() System Call

- The purpose of **fork()** is to create a child process. The creating and created processes are the **parent** and **child**, respectively.
- Unix will make an exact copy of the parent's address space and give it to the child. Therefore, the parent and child processes have separate address spaces.
- **fork()** does not require any argument!
- If the call to **fork()** is successful, Unix creates an identical but separate address space for the child process to run.
- Both processes start running with the instruction following the **fork()** system call.




# Fork() system call

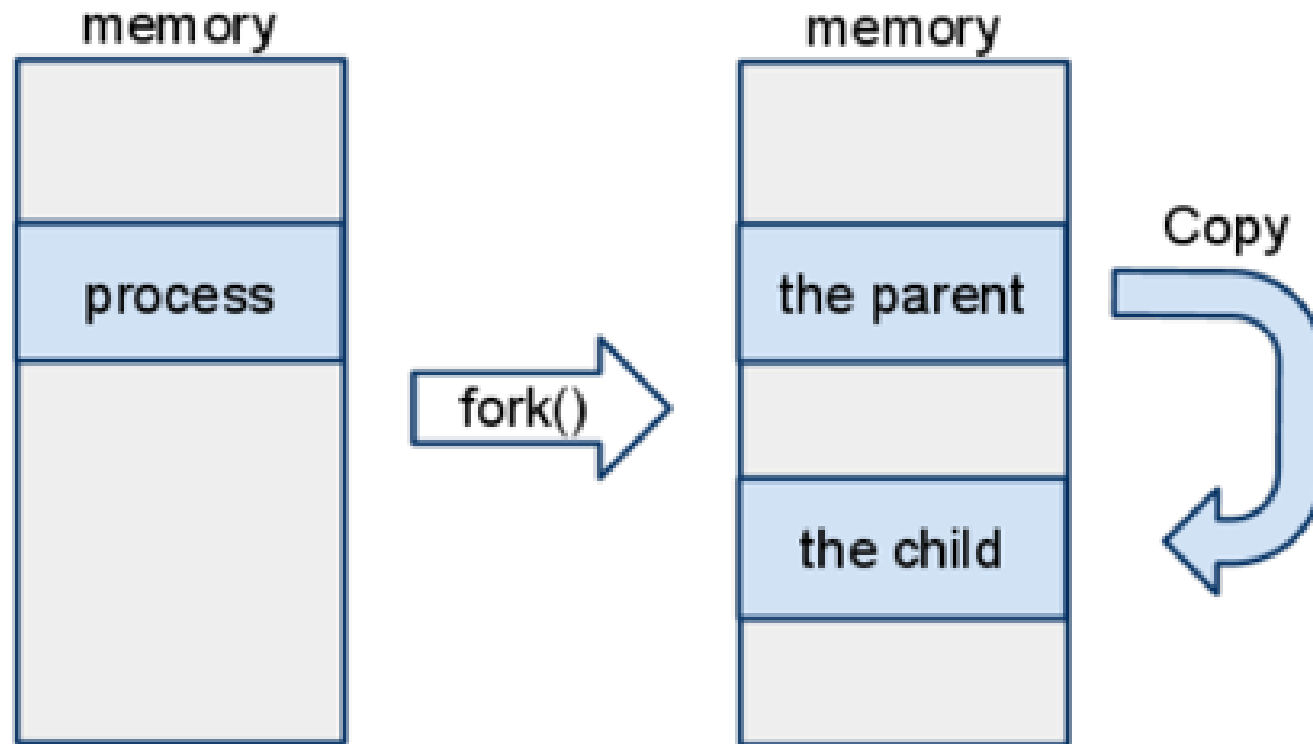
- Child process inherits
  - Stack
  - Memory
  - Open file descriptors
  - Current working directory
  - Resource limits
  - Root directory
- Child process does not inherit
  - Process ID
  - Different parent process ID.

# The fork() System Call

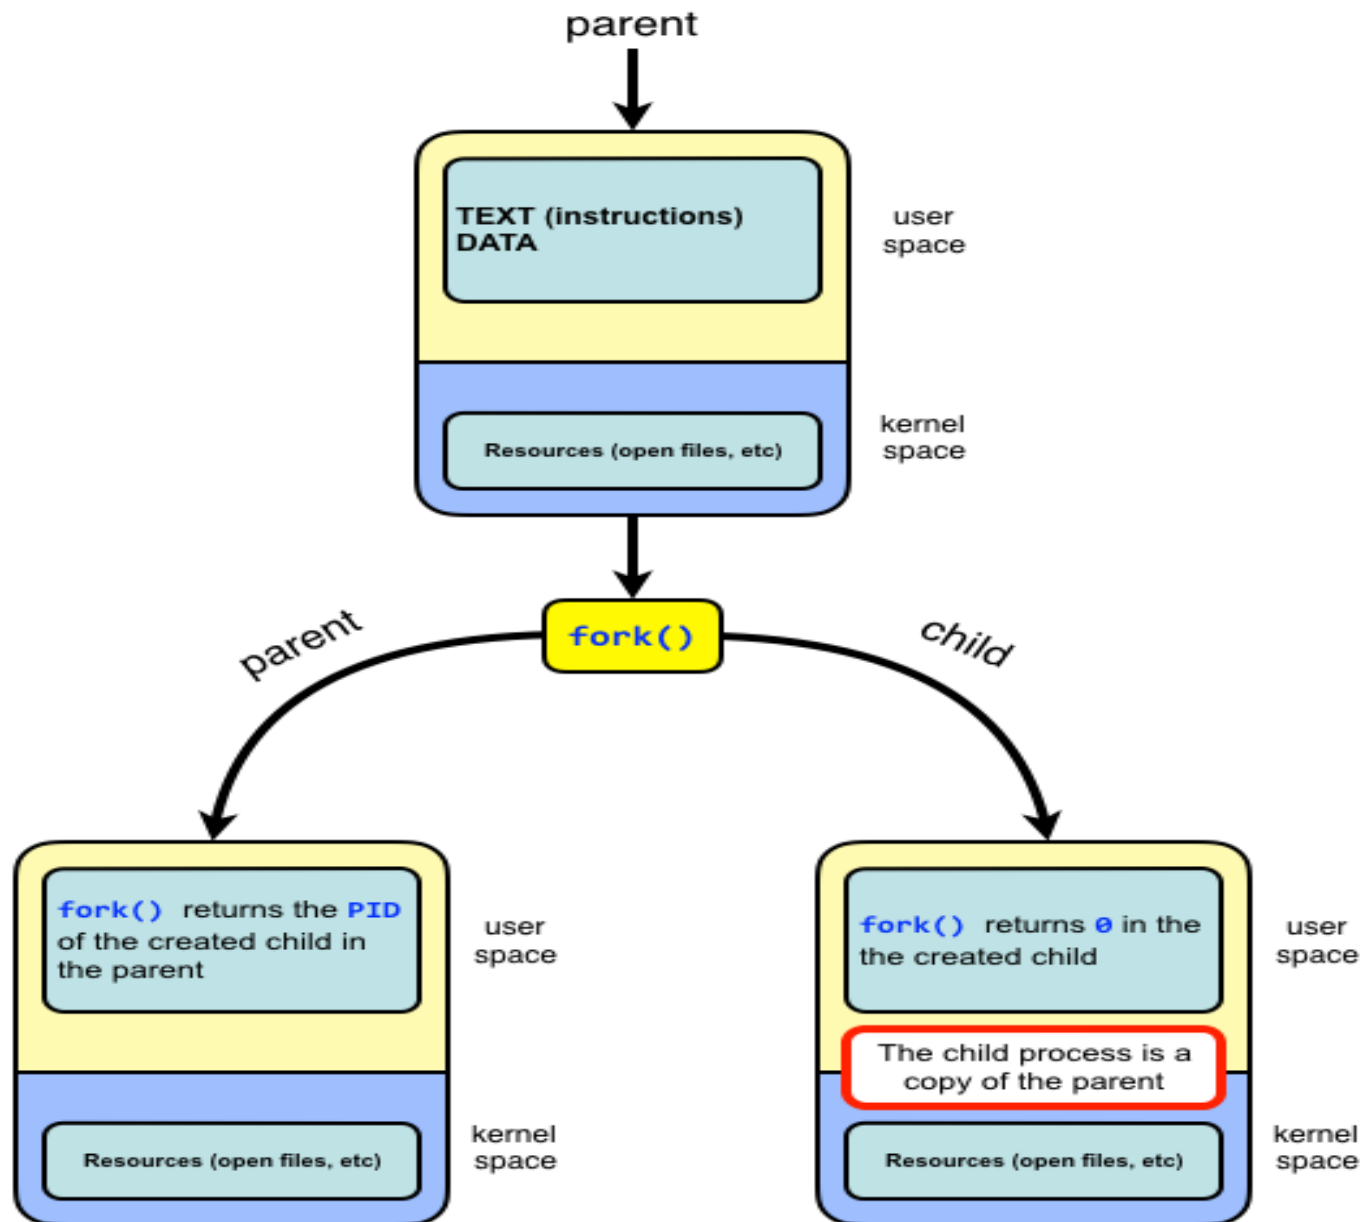
- The process invoking fork is called the **parent**. The new process created as the result of a fork is the **child** of the parent.
- Return Value:
  - **On success**, the PID of the child process is returned in the parent, and 0 is returned in the child.
  - **On failure**, -1 is returned in the parent, no child process is created.



# Before and after fork()



# The fork() System Call



# The fork() System Call

- On success fork returns twice:
  - once in the parent and once in the child. After calling fork, the program can use the **fork()** return value to tell whether executing in the parent or child.
- If the return value is 0 the program executes in the **new child process**.
- If the return value is greater than zero, the program executes in the **parent process** and the return value is the process ID (PID) of the created **child process**.
- On failure fork returns -1.
  - OS does not have enough resources to create a child process

# Example 1: 1/2

```
#include <stdio.h>
#include <sys/types.h>
int main(void) {
```

**This example does not distinguish parent and the child processes.**

```
    pid_t pid;                //variable pid of type pid_t is declared.
```

```
    //The pid_t data type is the data type used for process IDs.
```

```
    int i;
```

```
    fork(); //creates a child process. Both processes execute same copies of the below code.
```

```
    pid = getpid();           //The value of pid will be different in both child and parent
```

```
    for (i = 1; i <= 200; i++) {
```

```
        printf("This line is from pid %d, value = %d\n", pid, i); //printf is same as cout in cpp
```

```
    }
```

```
    return 0;
```

```
}
```

## Example 1: Parent and child execute same copies of code. 2/3



```
fork();  
pid = getpid();  
for (i = 1; i <= 200; i++) {  
    printf("This line is from pid %d, value = %d\n", pid, i);  
}  
return 0;  
}
```

**Parent**



```
fork();  
pid = getpid();  
for (i = 1; i <= 200; i++) {  
    printf("This line is from pid %d, value = %d\n", pid, i);  
}  
return 0;  
}
```

**child**

Example 1.2/2

```
This line is from pid 2268, value = 21
This line is from pid 2268, value = 22
This line is from pid 2268, value = 23
This line is from pid 2269, value = 3
This line is from pid 2268, value = 24
This line is from pid 2268, value = 25
This line is from pid 2269, value = 4
This line is from pid 2268, value = 26
This line is from pid 2268, value = 27
This line is from pid 2269, value = 5
This line is from pid 2268, value = 28
This line is from pid 2268, value = 29
This line is from pid 2269, value = 6
This line is from pid 2268, value = 30
This line is from pid 2269, value = 7
This line is from pid 2268, value = 31
```

Child process 2269 and  
parent process 2268, both  
run concurrently.

2268



2269

**The order of these lines are determined by the CPU scheduler.**




# fork() Return Values

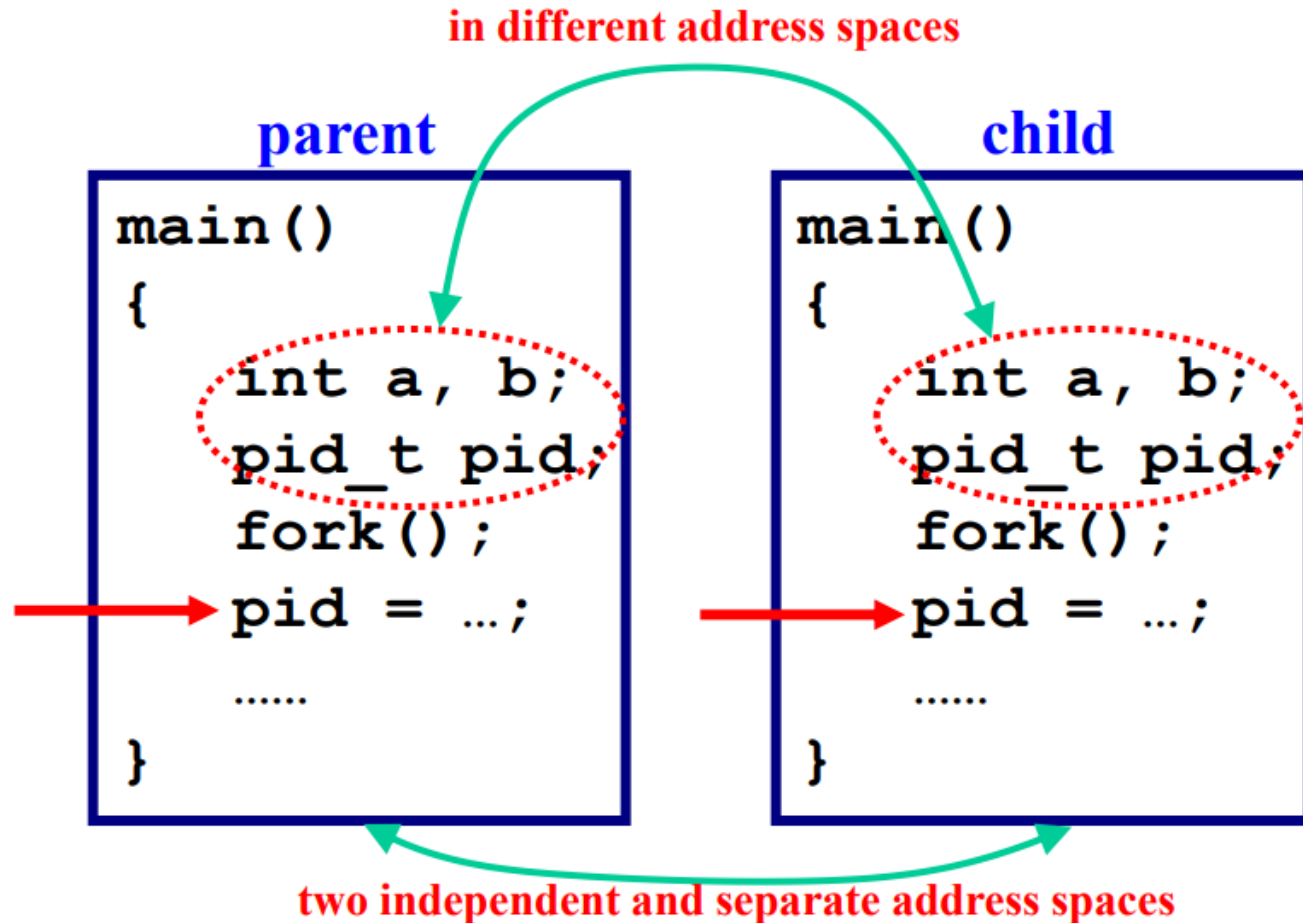
- **fork()** is unique (and often confusing) because it is called once but returns “twice”
- Child gets an identical (but separate) copy of the parent’s address space
- Child has a different PID than the parent
- Function **getpid()** returns the process ID of the caller.
- Both processes continue/start execution after fork
- Can’t predict execution order of parent and child → non deterministic

**parent**

```
main()  
{  
    int a, b;  
    pid_t pid  
    fork();  
    pid = ...;  
    .....  
}
```



# After execution of `fork()`

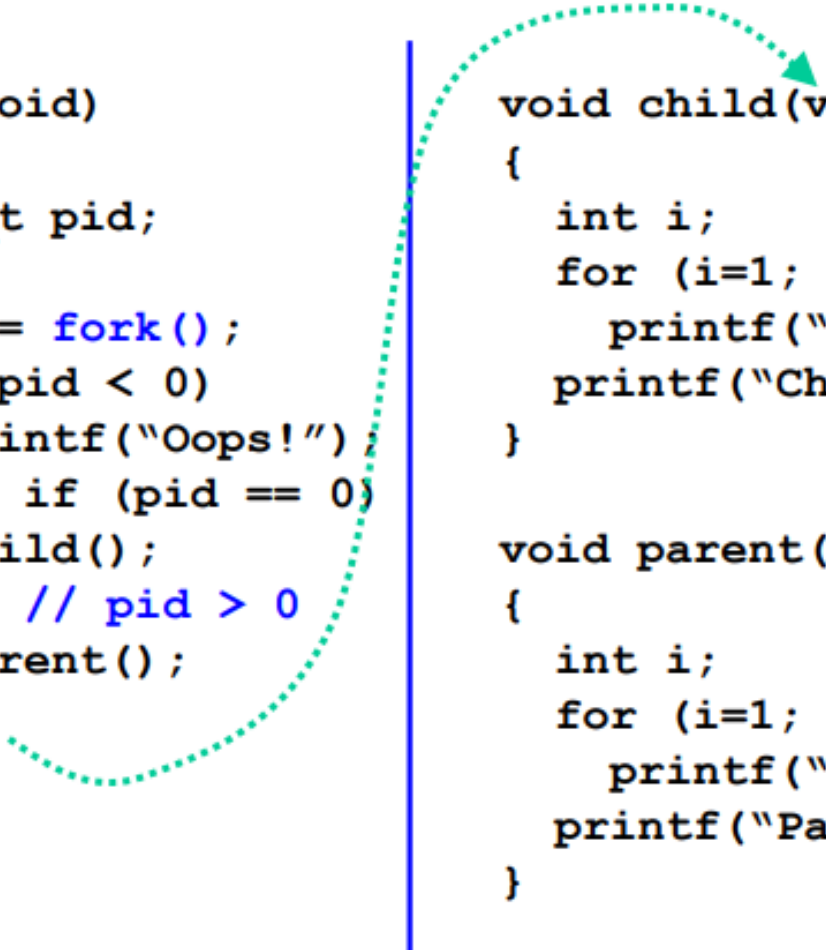


## Example 2: Separate Parent/Child Processes (1/5)

Consider a simple example, which distinguishes the parent from the child based on the value returned from `fork()`.

```
main(void)
{
    pid_t pid;

    pid = fork();
    if (pid < 0)
        printf("Oops!");
    else if (pid == 0)
        child();
    else // pid > 0
        parent();
}
```



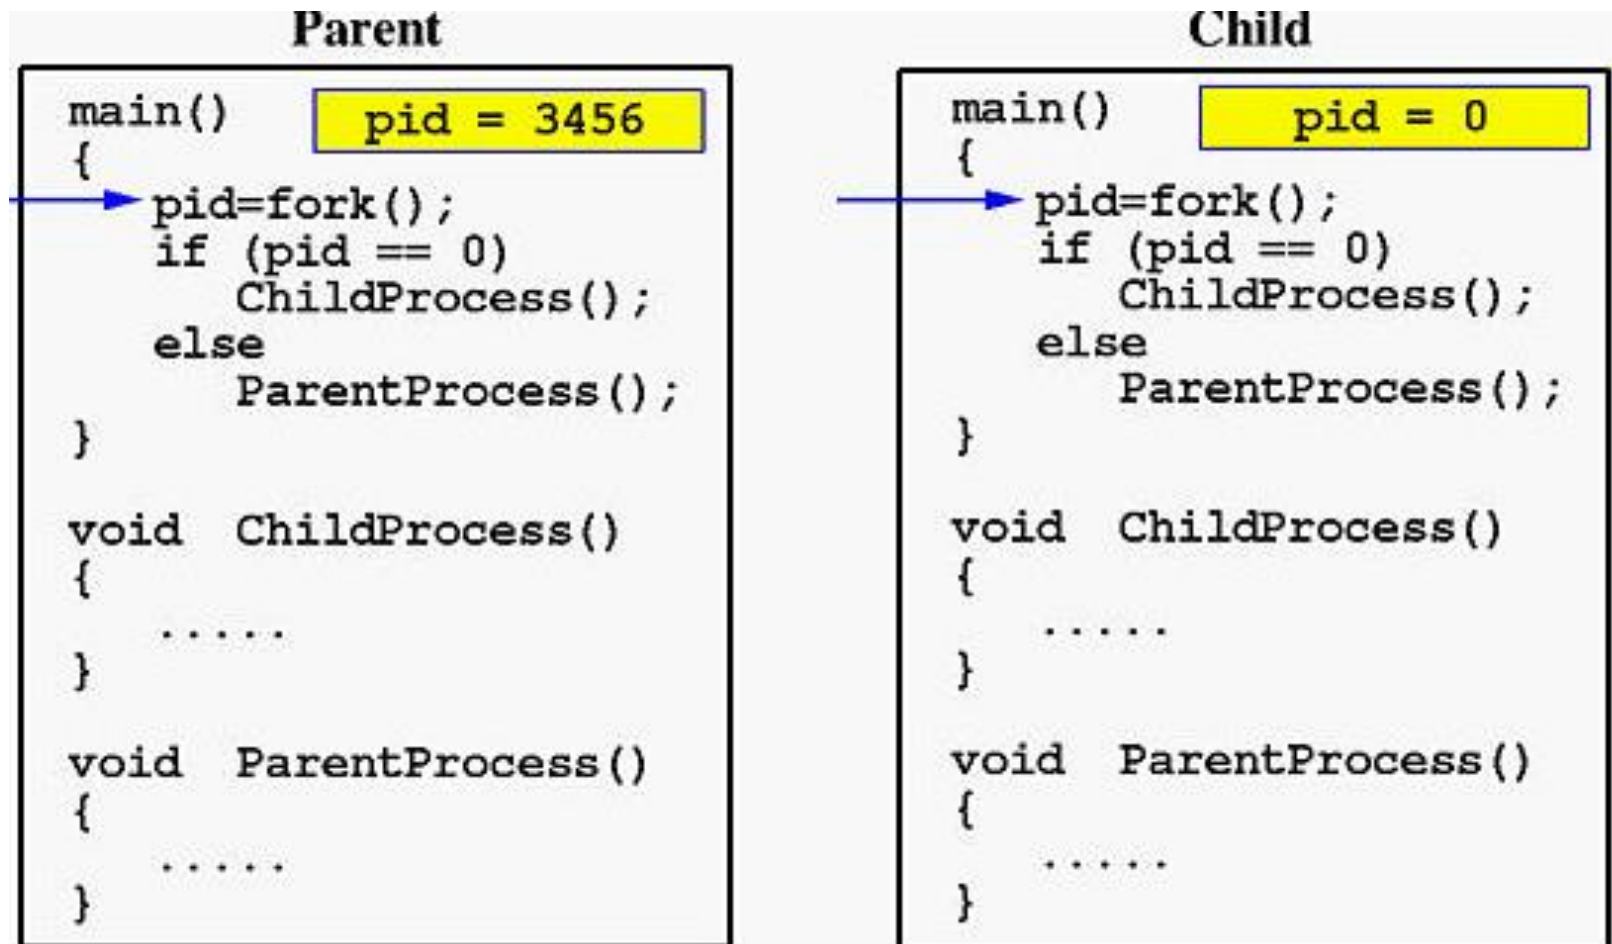
```
void child(void)
{
    int i;
    for (i=1; i<=10; i++)
        printf(" Child:%d\n", i);
    printf("Child done\n");
}

void parent(void)
{
    int i;
    for (i=1; i<=10; i++)
        printf("Parent:%d\n", i);
    printf("Parent done\n");
}
```

The diagram illustrates the execution flow. A vertical blue line separates the `main` function on the left from the `child` and `parent` functions on the right. A green dotted arrow originates from the `child()` call in the `main` function, curves upwards and to the right, and points to the `void child(void)` function definition. Another green dotted arrow originates from the `parent()` call in the `main` function, curves downwards and to the right, and points to the `void parent(void)` function definition.

## Example 2: Separate Parent/Child Processes (2/5)

**Step 1:** When the main program executes `fork()`, an identical copy of its address space, including the program and all data, is created. System call `fork()` returns the child process ID to the parent and returns 0 to the child process.



## Example 2: Separate Parent/Child Processes (4/5)

**Step 2:** In the parent, since pid is non-zero, it calls function ParentProcess(). On the other hand, the child has a pid equal to zero and calls ChildProcess():

### Parent

```
main()
{
    pid = 3456
    pid=fork();
    if (pid == 0)
        ChildProcess();
    else
        ParentProcess();
}

void ChildProcess()
{
    .....
}

void ParentProcess()
{
    .....
}
```

### Child

```
main()
{
    pid = 0
    pid=fork();
    if (pid == 0)
        ChildProcess();
    else
        ParentProcess();
}

void ChildProcess()
{
    .....
}

void ParentProcess()
{
    .....
}
```

## Example 2: Separate Parent/Child Processes (5/5)

```
#include <stdio.h>
#include <sys/types.h>
void main(void) {
    pid_t pid;
    pid = fork();
    if (pid == 0)
        ChildProcess();
    else if (pid>0)
        ParentProcess();
}
```

```
void ChildProcess(void) {
    int i;
    for (i = 1; i <= 200; i++)
        printf("This line is from child, value = %d\n", i);
    printf(" *** Child process is done ***\n"); }
```

```
void ParentProcess(void) {
    int i;
    for (i = 1; i <= 200; i++)
        printf("This line is from parent, value = %d\n", i);
    printf("*** Parent is done ***\n"); }
```

## Example 2: Separate Parent/Child Processes

```
This line is from child, value = 116
This line is from child, value = 117
This line is from child, value = 118
This line is from child, value = 119
This line is from parent, value = 153
This line is from child, value = 120
This line is from parent, value = 154
This line is from child, value = 121
This line is from parent, value = 155
This line is from child, value = 122
This line is from parent, value = 156
This line is from child, value = 123
This line is from parent, value = 157
This line is from child, value = 124
This line is from parent, value = 158
This line is from child, value = 125
This line is from parent, value = 159
This line is from child, value = 126
This line is from parent, value = 160
This line is from child, value = 127
This line is from parent, value = 161
```

**Both processes print lines that indicate (1) whether the line is printed by the child or by the parent process, and (2) the value of variable i.**

# Class exercise 1

- Write a 'C' program which uses fork().
- The main program declares two variables  $i=10$  and  $j=20$ . The program then calls fork().
  - The parent process shall print two variables and their respective values.
  - Similarly, the child shall **double** the values of the two variables  $i$ , and  $j$ , and then print them.



# Solution

```
#include <stdlib.h>
#include <unistd.h>

void main(void)
{
    pid_t pid;

    int i = 10, j = 20;
    if ((pid = fork()) == 0) { // child here
        i = 1000; j = 2000; // child changes values
        printf(" From child: i=%d, j=%d\n", i, j);

    }

    else { // parent here
        sleep(3);
        printf("From parent: i=%d, j=%d\n", i, j);

    }
}
```

# Orphan Process

- **An orphan process** is a process whose parent process has terminated, though it (child process) remains running itself.
- If parent terminated without waiting for the child process , the child process becomes **orphan**
- Or If a parent process terminates before its child terminates, the child process is automatically adopted by the **init process**

The **init** process periodically invokes **wait()** causing the release of orphan's process identifier and process-table entry.

A zombie process or defunct process is a process that has completed execution but still has an entry in the process table as its parent process didn't invoke an **wait()** system call.

## Example 3: getpid(), getppid() and Orphan Process


```
int main(void) {  
    pid_t pid;  
    pid = fork();
```

**What happens if Parent process terminates  
before child process?**

```
    if (pid==0){ //child Process  
        printf("\nFrom child %d: my Parent is %d\n", getpid(), getppid());  
        printf("From child %d: and my Parent is now %d\n", getpid(), getppid());  
    }  
    else if(pid>0){ //Parent Process  
        printf("From Parent %d: I spawned a child %d\n", getpid(), pid);  
        printf("From Parent %d: I am Done \n", getpid());  
    }  
    return 0;  
}
```

**getpid() returns process ID**

**getppid() returns parent process ID**



# Orphan Process

```
From child 3952: My parent is: 3951
```

```
From Parent 3951: I spawned a child 3952
```

```
From Parent 3951: I am Done!
```

```
From child 3952: My parent is: 1
```

While parent is still running

3951



3952

Parent terminated

1




3952


Orphan child has a new parent  
**init**, the Unix process that  
spawns all processes

# The `wait()` system call

- So far we assumed the parent process can terminate before the child process.



- The **`wait()`** system call blocks the caller (parent process) until one of its child processes exits.



# The `wait()` system call

- Synopsis of the `wait()` system call

```
pid_t wait(int *status);
```

- `wait()` returns the pid of the terminated child process.
  - If no child process is running, `wait()` returns -1.
  - In addition, `wait()` takes a pointer to an integer variable (`int status`) and returns some flags that indicate the completion `status` of the child process are passed back with the integer pointer.
- **What** If parent process has multiple children??
    - `wait` will return when any of the children terminates
  - `wait()` is used to synchronize parent/child execution

`waitpid` can be used to wait on a specific child process

# How to use `wait()` ?

```
void main(void)
{
    pid_t pid, pid_child;
    int    status;

    if ((pid = fork()) == 0)    // child here
        child();
    else {                      // parent here
        parent();
        pid_child = wait(&status);
    }
}
```

# How to use `wait()` ?

- Wait for an unspecified child process:

`wait(&status);`

- Wait for a number, say `n`, of unspecified child processes:

`for (i = 0; i < n; i++)`

`wait(&status);`

- Wait for a specific child process whose ID is known:

`while (pid != wait(&status)) ;`

```
void main(void)
{
    pid_t pid, pid_child;
    int    status;

    if ((pid = fork()) == 0)    // child here
        child();
    else {                      // parent here
        parent();
        pid_child = wait(&status);
    }
}
```



# exit() system call

- void exit(int status)
- After the program finishes execution, it calls exit()
- This system call:
  - takes the “result” of the program as an argument
  - closes all open files, connections, etc.
  - deallocates memory
  - deallocates most of the OS structures supporting the process
  - checks if parent is alive:
    - If so, it holds the result value until parent requests it, process does not really die, but it enters the zombie/defunct state
    - If not, it deallocates all data structures, the process is dead
  - Exit status is available to parent via `wait()`
    - Exit Status could be normal termination, erroneous termination, seg fault, divide by zero etc.

# Wait: Synchronization with children

```
void main() {  
    int child_status;  
  
    if (fork() == 0) {  
        printf("HC: hello from child\n");  
        exit(0);  
    } else {  
        printf("HP: hello from parent\n");  
        wait(&child_status);  
        printf("CT: child has terminated\n");  
    }  
    printf("Bye\n");  
}
```

What are Possible outputs?

# Wait: Synchronization with children

```
void main() {  
    int child_status;  
  
    if (fork() == 0) {  
        printf("HC: hello from child\n");  
        exit(0);  
    } else {  
        printf("HP: hello from parent\n");  
        wait(&child_status);  
        printf("CT: child has terminated\n");  
    }  
    printf("Bye\n");  
}
```

Two possible outputs:

HC	HP
HP	HC
CT	CT
Bye	Bye

# Process Termination (exit()) system call

- Process executes last statement and then asks the operating system to delete it using the `exit()` system call.
  - closes all open files, connections, etc.
  - deallocates memory
  - deallocates most of the OS structures supporting the process
  - Returns status from child to parent (via `wait()`)
  - Exit Status could be normal termination, erroneous termination, seg fault, divide by zero etc.

# Process Termination (cont'd)

- Under normal termination, `exit()` may be called either directly or indirectly (by a return statement ).
- The status of the child can be obtained with a call to the `WEXITSTATUS` macro:  
`WEXITSTATUS(status)`.

```
void child(void) {  
    printf("Child Process being executed\n");  
    /*      child process is about to end      */  
    exit(0);  
}
```

```
void parent(int pid) {  
    wait(&status);  
    if (WIFEXITED(status)) /* process exited normally */  
        printf("child process exited with value %d\n", WEXITSTATUS(status)); }  
}
```

# Process Termination (Cont'd)

- Some operating systems do not allow child to exist if its parent has terminated.

If a process terminates, then all its children must also be terminated.

- **cascading termination.** All children, grandchildren, etc. are terminated.

- The termination is initiated by the operating system.

- If parent terminated without invoking `wait()` , **running child** process is an **orphan**

- If a parent process terminates before its child terminates, the child process is automatically adopted by the **init process**
- The **init** process periodically invokes `wait()` causing the release of orphan's process identifier and process-table entry.

# Process Termination (Cont'd)

- **Zombie Process:**

- Dictionary meaning of “**zombie**”: One who seems more dead than alive. In Unix, a zombie process is one which has terminated before its parent had a chance to wait for it.

- If no parent waiting (did not invoke **wait()**), **the terminated** process is a **zombie**

- Need to store exit status in **process table**
- OS can't be fully free → although zombie does not use memory, but its pid cannot be assigned to another process

- **Reaping** is performed by parent on terminated child

- Parent is given exit status information and kernel then deletes zombie child process

# Example: Zombie

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

void main() {
    if (fork() == 0) {
        /* Child */
        printf("Terminating Child, PID = %d\n",
               getpid());
        exit(0);
    } else {
        printf("Running Parent, PID = %d\n",
               getpid());


        //No wait()
        while (1); /* Infinite loop */
    }
}
```

- Use `ps` command on terminal to see all the running processes in the system
- `ps` shows child process as **“defunct”**
- Killing parent allows child to be reaped by `init`




# exec ( ) system call

- After a *fork()* system call, one of the two processes typically uses the *exec()* system call to replace the process's memory space with a new program.
- The *exec()* system call loads a binary file into memory
  - They are responsible for calling external programs.
  - destroys the memory image of the program containing the *exec()* system call

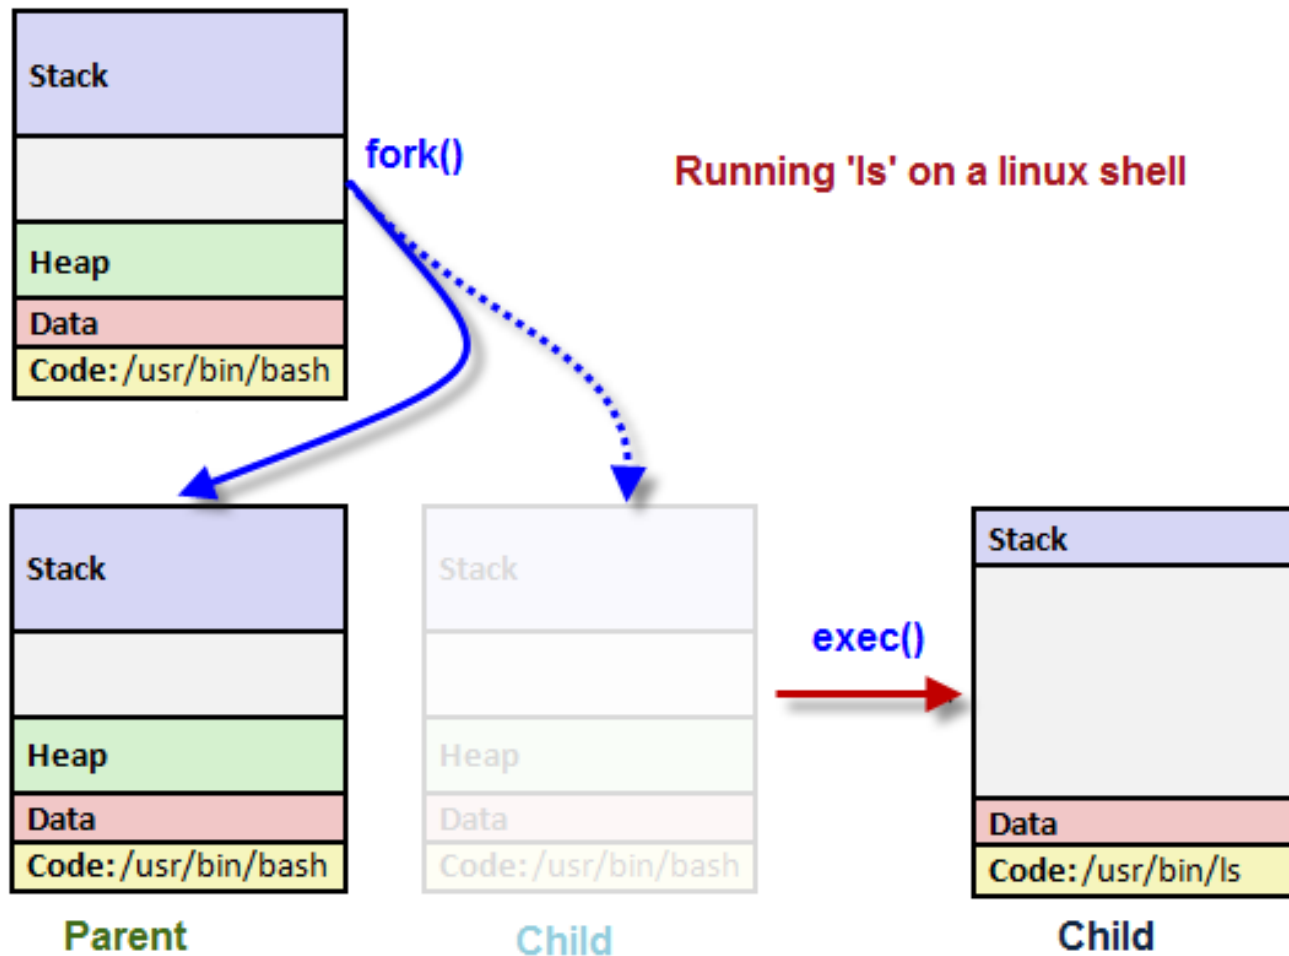


# fork-exec model:

- A newly created child process may run a different program rather than that of the parent by using one of the system call of the `exec()` family.
- The `exec()` system call will replace the currently executing program with a new executable.



# Example: Use of Exec() System Call



## exec() system call Family

- There are the family members of the exec() family.

1. int `execl`(char \* pathname, char \* arg0, ... , (char \*)0);
2. int `execv`(char \* pathname, char \* argv[]);
3. int `execle`(char \* pathname, char \* arg0, ... , (char \*)0, char envp[]);
4. int `execve`(char \* pathname, char \* argv[], char envp[]);
5. int `execlp`(char \* filename, char \* arg0, ... , (char \*)0);
6. int `execvp`(char \* filename, char \* argv[]);

# Simple `exec()` system call

- `exec()` family of system calls has several variants.
- `execlp()` is one of the members of the `exec()` family of system calls.
- 'ls' is a command (program) that we use to list all the files/directories in our system.
- We use `execlp()` to execute `/bin/ls` program.

```
#include <stdio.h>
#include <unistd.h>

int main(int argc, char* argv[])
{
    printf("Using *execlp* to exec ls -l...\n");
    execlp("/bin/ls", "ls", "-l", NULL); //takes executable path and argument list as parameters
    printf("Program Terminated\n"); // the line Program Terminated is not printed.
    exit(1); //This will be return exit code 1 if execlp is failed.
}
```

# C Program Forking Separate Process

- The first argument is the path of the executable program, cp.
- 2<sup>nd</sup> argument is the command, and 3<sup>rd</sup> and 4<sup>th</sup> argument are the name of files.
- Note that `execl()` has a variable number of arguments; the NULL pointer indicates the end of the list.

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

void main(void) {
    int status;
    pid_t pid;
    if ((pid = fork()) < 0) {
        printf("fork() failed\n");
        exit(1);
    }
    else if (pid == 0) //child process
    if (execl("/bin/cp", "cp", "src.txt", "dst.txt", NULL) < 0) {
        printf("execl() failed\n");
        exit(1);
    }
    else //parent process
    wait(&status);
}
```

# Chapter 3: Processes

- Process Concept
- Process Scheduling
- Operations on Processes
- **Interprocess Communication (Part 3)**