

Chapter 5: Process Synchronization (Part 3)

Chapter 5: Process Synchronization

- Background
- The Critical-Section Problem
- Peterson's Solution
- Synchronization Hardware
- Mutex Locks
- **Semaphores**
- Classic Problems of Synchronization
- Monitors
- Condition Variables

Semaphore

- Synchronization tool that provides more sophisticated ways (than **Mutex locks**) for processes to synchronize their activities.
- originally, semaphores were **flags or lights** for signaling between ships
- In computer science, a **semaphore** can be used to solve several different types of synchronization problems.
- **Semaphore S** – integer variable used for signaling between processes
- A semaphore S is an integer variable that, *apart from initialization*, can only be modified through **two** atomic and mutually exclusive operations:
 - **Wait (S)**
 - **Signal (S)**
 - Originally called **P ()** and **V ()**
- If a process is waiting for a signal, it is suspended until that signal is sent

Definition of Wait() and Signal()

- Definition of the **wait()** operation

```
wait(S) {  
    while (S <= 0)  
        ;  
    S--;  
}
```

- Definition of the **signal()** operation

```
signal(S) {  
    S++;  
}
```

Semaphore Usage

- **Counting semaphore** – integer value can range over an unrestricted domain
 - Can be used to control access to resources (semaphore S initialized to number of resources)
 - a resource such as a shared data structure is protected by a semaphore. You must **acquire** the semaphore before using the resource and **release** the semaphore when you are done with the shared resource.
 - Process that wishes to use resource will execute `wait(S)`
 - Process that releases a resource will execute `signal(S)`
 - If $S = 0$ then all resources are being used, and, Processes block until $S > 0$;
- **Binary semaphore** – integer value can range only between 0 and 1
 - Same as a **mutex lock**
 - Can solve various synchronization problems (See next slide to know how)

Semaphore Usage (Cont'd)

- **Process Synchronization with Semaphores** - Consider concurrent processes P_1 and P_2 that require statement S_1 to happen before S_2

Create a semaphore “**synch**” shared by P_1 and P_2 and initialized to 0

Process P1:

```
s1; //execute statement S1  
signal(synch); //increment semaphore synch
```

Process P2:

```
wait(synch); //decrement Semaphore synch  
s2; //execute statement S2
```

- Can implement a counting semaphore S as a binary semaphore if S is initialized to 1

Semaphore Implementation (wait and signal must be atomic)

- Must guarantee that **no two processes** can execute the **wait()** and **signal()** on the same semaphore at the same time
- Thus, the semaphore implementation becomes the critical section problem where the **wait** and **signal** code are placed in the critical section
 - Could now have **busy waiting** in critical section implementation

Semaphore Implementation with no Busy loop

- For no busy waiting we associate a **waiting queue** with each semaphore
- Each entry in a waiting queue has two data items:
 - value (of type integer)
 - pointer to next record in the list
- Two operations:
 - **block** – place the process invoking the operation on the appropriate waiting queue
 - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue
- **typedef struct{**

int value;

struct process *list; //list of processes waiting on the semaphore S

} semaphore;

Implementation with no Busy waiting (Cont.)

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        block();  
    }  
} //wait() operation adds a process to the S queue and suspends it  
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
} //signal() operation removes one process from S queue and awakens it
```

Example: Mutual Exclusion with semaphore

- Suppose there are three processes, **A**, **B**, and **C**.
- Every process wants to execute Critical Section.
- But CS is protected by a shared binary Semaphore called Mutex.
- Any process trying to execute CS will first call **wait(mutex)** and if not blocked will execute CS.
- Similarly, any process that executes CS will call **signal(mutex)** and will unblock the next blocked process waiting in the Mutex List.

Mutual Exclusion with Semaphores (1/34)

Mutual Exclusion Achieved Through Semaphores

Mutex Value	Mutex List	wait (mutex) { value = value - 1; if (value < 0) { add process to list; block process; } }	signal (mutex) { value = value + 1; if (value <= 0) { remove process from list; wake up process; } }
1			

Process A	Process B	Process C
start wait (mutex); critical section signal (mutex); do remainder	start wait (mutex); critical section signal (mutex); do remainder	start wait (mutex); critical section signal (mutex); do remainder

Process A begins execution.

Mutual Exclusion with Semaphores (2/34)

Mutual Exclusion Achieved Through Semaphores

Mutex Value	Mutex List	wait (mutex)	signal (mutex)
1		<pre>wait (mutex) { value = value - 1; if (value < 0) { add process to list; block process; } }</pre>	<pre>signal (mutex) { value = value + 1; if (value <= 0) { remove process from list; wake up process; } }</pre>
		Process A start wait (mutex); critical section signal (mutex); do remainder	Process B start wait (mutex); critical section signal (mutex); do remainder

Process A executes the wait (mutex) instruction.

Mutual Exclusion with Semaphores (3/34)

Mutual Exclusion Achieved Through Semaphores						
Mutex Value	Mutex List	wait (mutex)	signal (mutex)			
0	<table border="1"><tr><td></td><td></td><td></td></tr></table>				<pre>wait (mutex) { value = value - 1; if (value < 0) { add process to list; block process; } }</pre>	<pre>signal (mutex) { value = value + 1; if (value <= 0) { remove process from list; wake up process; } }</pre>
Process A	Process B	Process C				
<pre>start wait (mutex); critical section signal (mutex); do remainder</pre>	<pre>start wait (mutex); critical section signal (mutex); do remainder</pre>	<pre>start wait (mutex); critical section signal (mutex); do remainder</pre>				
The Mutex Value is decremented by 1						

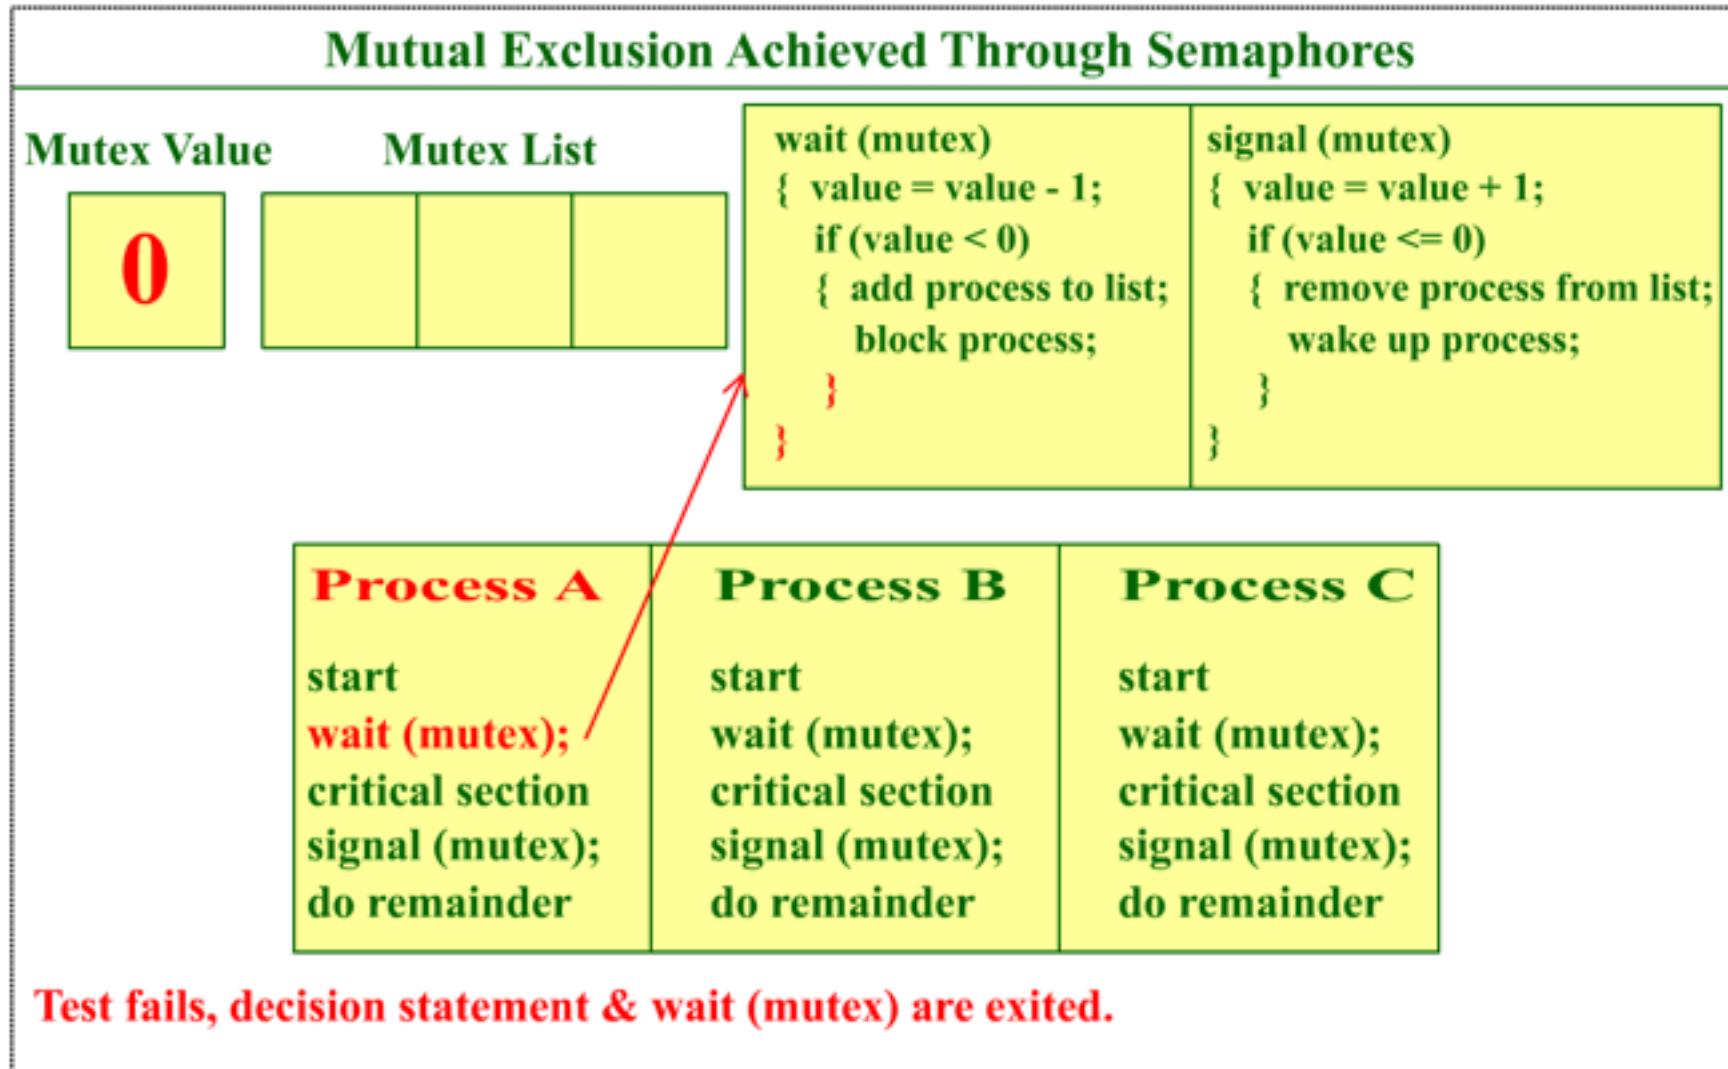
Mutual Exclusion with Semaphores (4/34)

Mutual Exclusion Achieved Through Semaphores

Mutex Value	Mutex List	wait (mutex) { value = value - 1; if (value < 0) { add process to list; block process; } }	signal (mutex) { value = value + 1; if (value <= 0) { remove process from list; wake up process; } }
0			
Process A start wait (mutex); critical section signal (mutex); do remainder		start wait (mutex); critical section signal (mutex); do remainder	start wait (mutex); critical section signal (mutex); do remainder

Comparison is made to determine if a process is already in its critical section.

Mutual Exclusion with Semaphores (5/34)



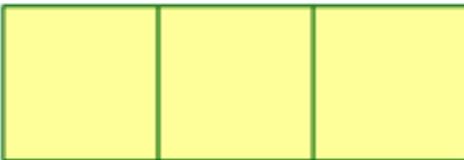
Mutual Exclusion with Semaphores (6/34)

Mutual Exclusion Achieved Through Semaphores

Mutex Value

0

Mutex List



```
wait (mutex)
{ value = value - 1;
  if (value < 0)
  { add process to list;
    block process;
  }
}
```

```
signal (mutex)
{ value = value + 1;
  if (value <= 0)
  { remove process from list;
    wake up process;
  }
}
```

Process A

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process B

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process C

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process A is allowed to enter its critical section.

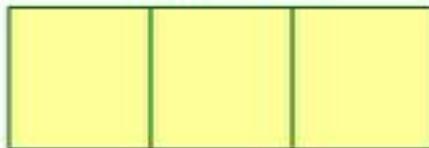
Mutual Exclusion with Semaphores (7/34)

Mutual Exclusion Achieved Through Semaphores

Mutex Value

0

Mutex List



```
wait (mutex)
{ value = value - 1;
if (value < 0)
{ add process to list;
block process;
}
}
```

```
signal (mutex)
{ value = value + 1;
if (value <= 0)
{ remove process from list;
wake up process;
}
}
```

Process A

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process B

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process C

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Context Switch by CPU Scheduler, Process B begins execution.

Mutual Exclusion with Semaphores (8/34)

Mutual Exclusion Achieved Through Semaphores

Mutex Value	Mutex List	wait (mutex) { value = value - 1; if (value < 0) { add process to list; block process; } } signal (mutex) { value = value + 1; if (value <= 0) { remove process from list; wake up process; } }
0		

Process A	Process B	Process C
start wait (mutex); critical section signal (mutex); do remainder	start wait (mutex); critical section signal (mutex); do remainder	start wait (mutex); critical section signal (mutex); do remainder

Context Switch by CPU Scheduler, Process B begins execution.

Mutual Exclusion with Semaphores (9/34)

Mutual Exclusion Achieved Through Semaphores

Mutex Value	Mutex List	wait (mutex) { value = value - 1; if (value < 0) { add process to list; block process; } } }	signal (mutex) { value = value + 1; if (value <= 0) { remove process from list; wake up process; } } }
-1			
		Process A start wait (mutex); critical section signal (mutex); do remainder	Process B start wait (mutex); critical section signal (mutex); do remainder

The Mutex Value is decremented by 1.

Mutual Exclusion with Semaphores (10/34)

Mutual Exclusion Achieved Through Semaphores					
Mutex Value	Mutex List	wait (mutex)		signal (mutex)	
-1		{ value = value - 1; if (value < 0) { add process to list; block process; } }		{ value = value + 1; if (value <= 0) { remove process from list; wake up process; } }	
Process A	Process B	start wait (mutex); critical section signal (mutex); do remainder	start wait (mutex); critical section signal (mutex); do remainder	start wait (mutex); critical section signal (mutex); do remainder	start wait (mutex); critical section signal (mutex); do remainder
Comparison is made to determine if a process is already in its critical section.					

Mutual Exclusion with Semaphores (11/34)

Mutual Exclusion Achieved Through Semaphores					
Mutex Value	Mutex List			wait (mutex)	signal (mutex)
-1	B			{ value = value - 1; if (value < 0) { add process to list; block process; } }	{ value = value + 1; if (value <= 0) { remove process from list; wake up process; } }
Process A	Process B	Process C			
start wait (mutex); critical section signal (mutex); do remainder	start wait (mutex); critical section signal (mutex); do remainder	start wait (mutex); critical section signal (mutex); do remainder			

Test passes, Process B is added to the Mutex List.

Mutual Exclusion with Semaphores (12/34)

Mutual Exclusion Achieved Through Semaphores

Mutex Value

-1

Mutex List

B | | |

```
wait (mutex)
{ value = value - 1;
  if (value < 0)
  { add process to list;
    block process;
  }
}
```

```
signal (mutex)
{ value = value + 1;
  if (value <= 0)
  { remove process from list;
    wake up process;
  }
}
```

Process A

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process B

```
start
wait (mutex);
BLOCKED
critical section
signal (mutex);
do remainder
```

Process C

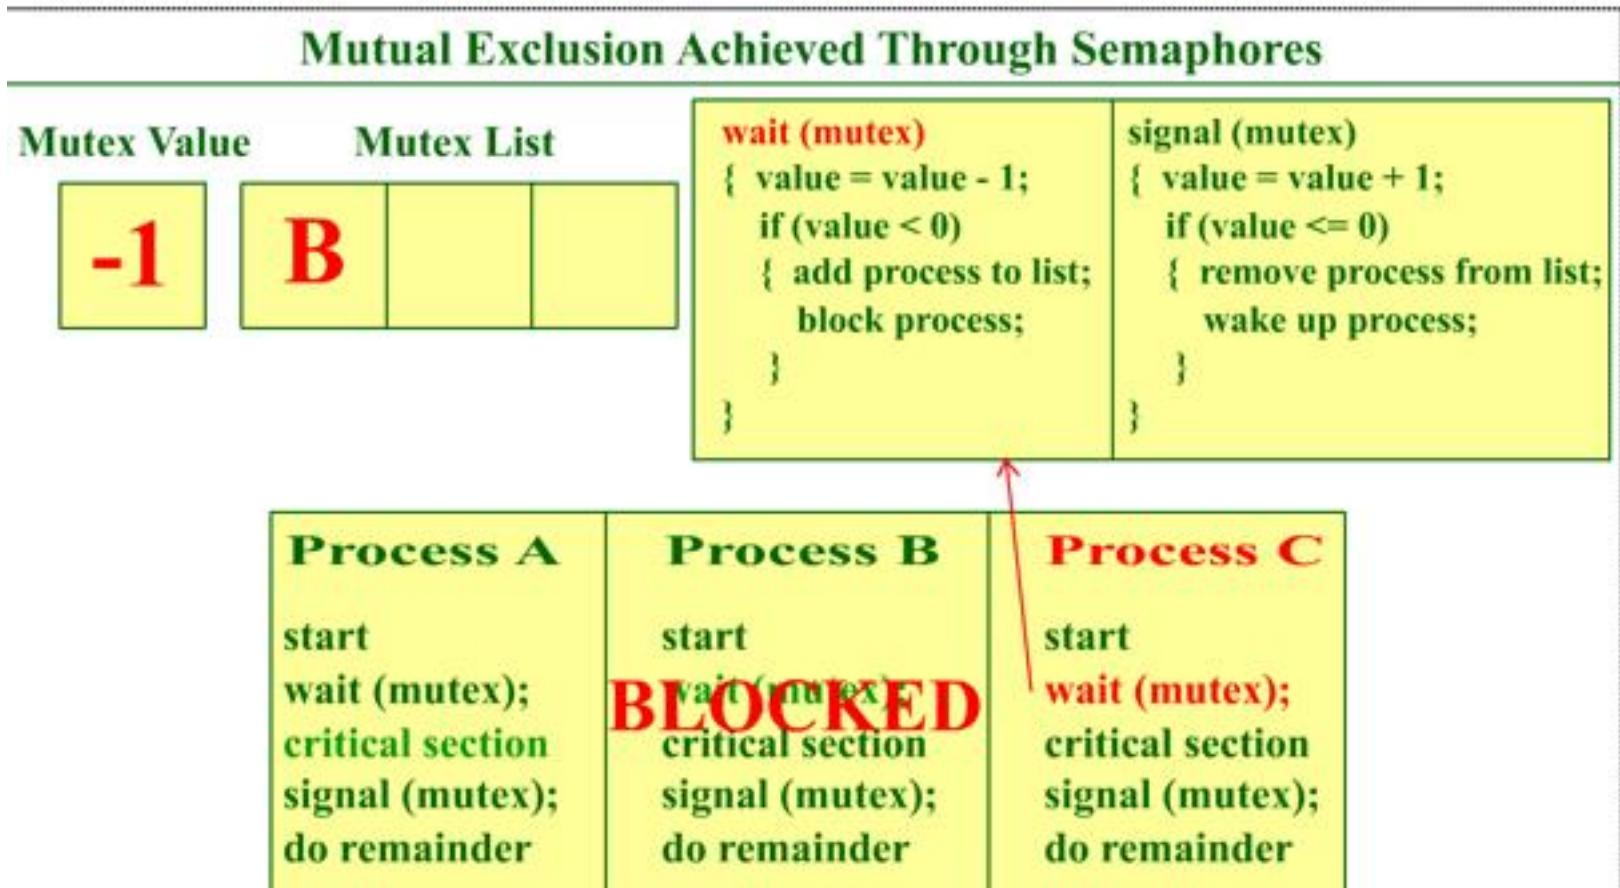
```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process B is blocked.

Mutual Exclusion with Semaphores (13/34)

Mutual Exclusion Achieved Through Semaphores					
Mutex Value	Mutex List	wait (mutex)		signal (mutex)	
-1	B	{ value = value - 1; if (value < 0) { add process to list; block process; } }		{ value = value + 1; if (value <= 0) { remove process from list; wake up process; } }	
Process A	Process B	start wait (mutex); critical section signal (mutex); do remainder		start wait (mutex); critical section signal (mutex); do remainder	
		BLOCKED		start wait (mutex); critical section signal (mutex); do remainder	
Context Switch, Process C starts execution.					

Mutual Exclusion with Semaphores (14/34)



Process C executes the wait (mutex) instruction.

Mutual Exclusion with Semaphores (15/34)

Mutual Exclusion Achieved Through Semaphores

Mutex Value

-2

Mutex List

B

```
wait (mutex)
{ value = value - 1;
  if (value < 0)
  { add process to list;
    block process;
  }
}
```

```
signal (mutex)
{ value = value + 1;
  if (value <= 0)
  { remove process from list;
    wake up process;
  }
}
```

Process A

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process B

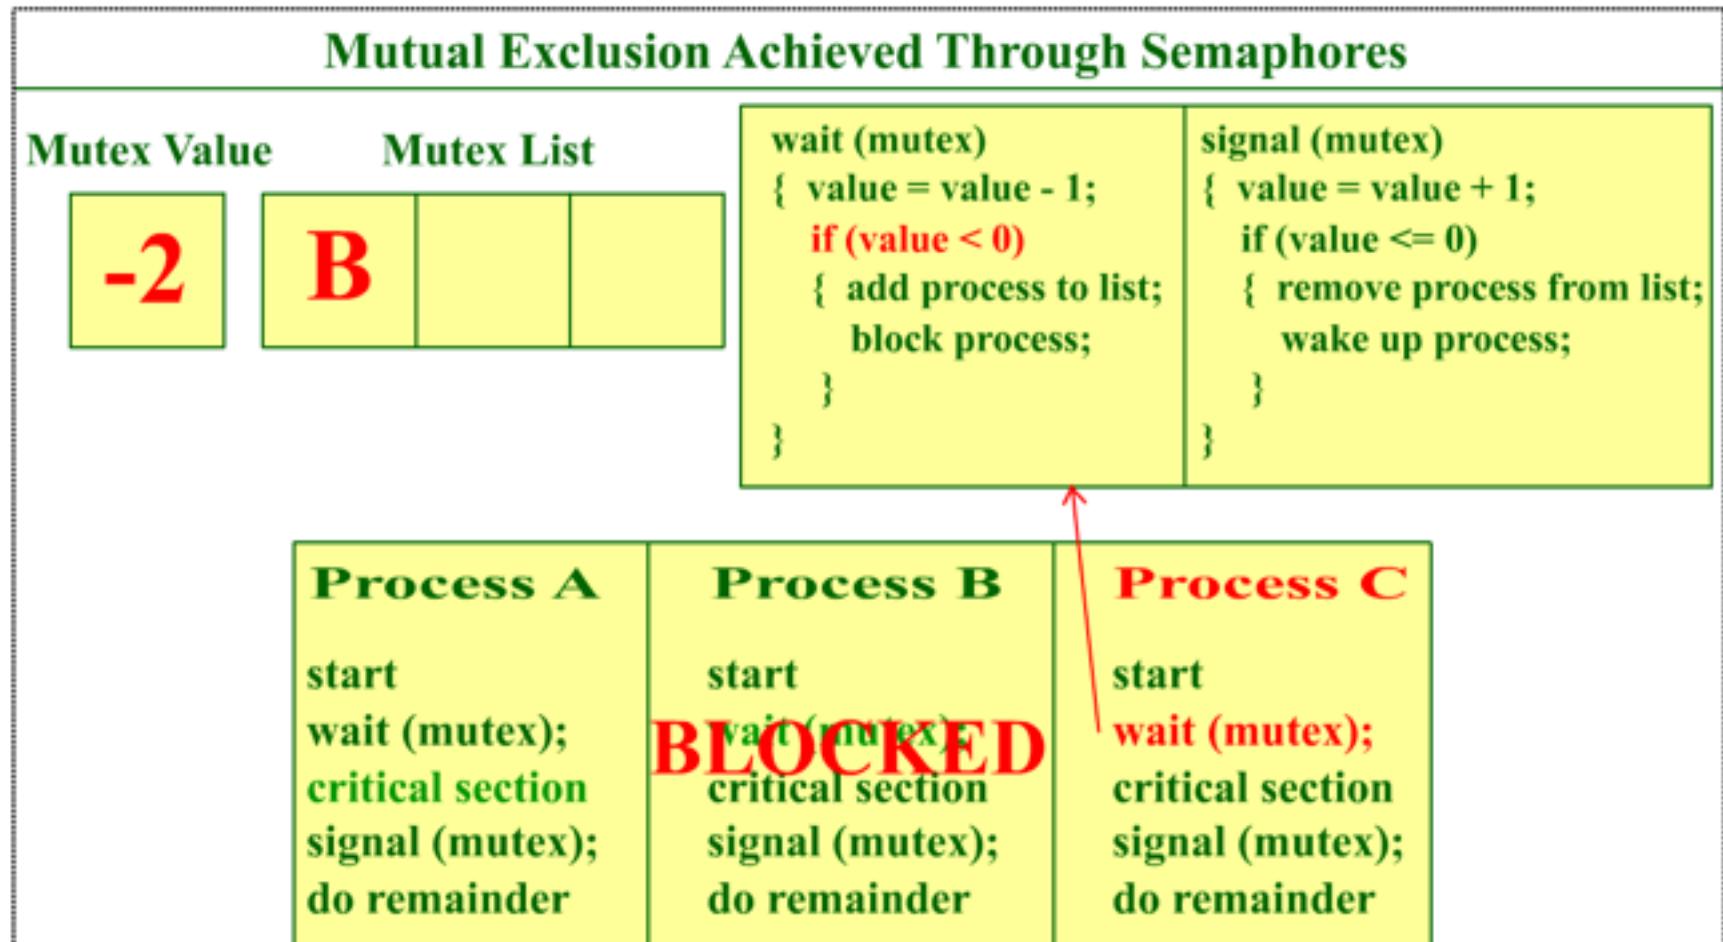
```
start
BLOCKED
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process C

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

The Mutex Value is decremented by 1.

Mutual Exclusion with Semaphores (16/34)



Comparison is made to determine if a process is already in its critical section.

Mutual Exclusion with Semaphores (17/34)

Mutual Exclusion Achieved Through Semaphores		
Mutex Value	Mutex List	
-2	B C	<pre>wait (mutex) { value = value - 1; if (value < 0) { add process to list; block process; } }</pre> <pre>signal (mutex) { value = value + 1; if (value <= 0) { remove process from list; wake up process; } }</pre>
Process A	Process B	Process C
start wait (mutex); critical section signal (mutex); do remainder	start wait (mutex) BLOCKED critical section signal (mutex); do remainder	start wait (mutex); critical section signal (mutex); do remainder

Test passes, Process C is added to the Mutex List.

Mutual Exclusion with Semaphores (18/34)

Mutual Exclusion Achieved Through Semaphores

Mutex Value

Mutex List

-2

B

C

```
wait (mutex)
{ value = value - 1;
  if (value < 0)
    { add process to list;
      block process;
    }
}
```

```
signal (mutex)
{ value = value + 1;
  if (value <= 0)
    { remove process from list;
      wake up process;
    }
}
```

Process A

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process B

```
start
wait (mutex)
critical section
signal (mutex);
do remainder
```

Process C

```
start
wait (mutex)
critical section
signal (mutex);
do remainder
```

Process C is blocked.

Mutual Exclusion with Semaphores (19/34)

Mutual Exclusion Achieved Through Semaphores

Mutex Value

-2

Mutex List



```
wait (mutex)
{ value = value - 1;
  if (value < 0)
  { add process to list;
    block process;
  }
}
```

```
signal (mutex)
{ value = value + 1;
  if (value <= 0)
  { remove process from list;
    wake up process;
  }
}
```

Process A

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process B

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process C

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Context Switch, Process A resumes execution.

Mutual Exclusion with Semaphores (20/34)

Mutual Exclusion Achieved Through Semaphores					
Mutex Value	Mutex List			wait (mutex)	signal (mutex)
-2	B	C		{ value = value - 1; if (value < 0) { add process to list; block process; } }	{ value = value + 1; if (value <= 0) { remove process from list; wake up process; } }
Process A	Process B	Process C			
start wait (mutex); critical section signal (mutex); do remainder	start wait (mutex); BLOCKED critical section signal (mutex); do remainder	start wait (mutex); BLOCKED critical section signal (mutex); do remainder			
Process A leaves its critical section, and executes the signal (mutex) instruction.					

Mutual Exclusion with Semaphores (21/34)

Mutual Exclusion Achieved Through Semaphores

Mutex Value

-1

Mutex List

B C

```
wait (mutex)
{ value = value - 1;
  if (value < 0)
  { add process to list;
    block process;
  }
}
```

```
signal (mutex)
{ value = value + 1;
  if (value <= 0)
  { remove process from list;
    wake up process;
  }
}
```

Process A

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process B

```
start
wait (mutex)
critical section
signal (mutex);
do remainder
```

Process C

```
start
wait (mutex)
critical section
signal (mutex);
do remainder
```

The Mutex Value is incremented by 1.

Mutual Exclusion with Semaphores (22/34)

Mutual Exclusion Achieved Through Semaphores		
Mutex Value	Mutex List	
-1	B C	
		<pre>wait (mutex) { value = value - 1; if (value < 0) { add process to list; block process; } }</pre>
		<pre>signal (mutex) { value = value + 1; if (value <= 0) { remove process from list; wake up process; } }</pre>
Process A	Process B	Process C
start wait (mutex); critical section signal (mutex); do remainder	start wait (mutex); critical section signal (mutex); do remainder	start wait (mutex); critical section signal (mutex); do remainder
	BLOCKED	BLOCKED
A comparison is made to determine if any processes are waiting in the Mutex List.		

Mutual Exclusion with Semaphores (23/34)

Mutual Exclusion Achieved Through Semaphores

Mutex Value **Mutex List**



```
wait (mutex)
{ value = value - 1;
  if (value < 0)
  { add process to list;
    block process;
  }
}
```

```
signal (mutex)
{ value = value + 1;
  if (value <= 0)
  { remove process from list;
    wake up process;
  }
}
```

Process A

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process B

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process C

```
start
BLOCKED
wait (mutex);
critical section
signal (mutex);
do remainder
```

Test passes, the first process in the Mutex List is removed and is woken up.

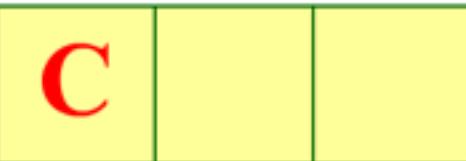
Mutual Exclusion through Semaphores (24/34)

Mutual Exclusion Achieved Through Semaphores

Mutex Value

-1

Mutex List



```
wait (mutex)
{ value = value - 1;
  if (value < 0)
  { add process to list;
    block process;
  }
}
```

```
signal (mutex)
{ value = value + 1;
  if (value <= 0)
  { remove process from list;
    wake up process;
  }
}
```

Process A

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process B

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process C

```
start
wait (mutex)
BLOCKED
critical section
signal (mutex);
do remainder
```

Context Switch, Process B resumes execution & enters its critical section.

Mutual Exclusion through Semaphores(25/34)

Mutual Exclusion Achieved Through Semaphores			
Mutex Value	Mutex List	wait (mutex)	signal (mutex)
-1	C	wait (mutex) { value = value - 1; if (value < 0) { add process to list; block process; } }	signal (mutex) { value = value + 1; if (value <= 0) { remove process from list; wake up process; } }
Process A	start wait (mutex); critical section signal (mutex); do remainder	start wait (mutex); critical section signal (mutex); do remainder	start wait (mutex); critical section signal (mutex); do remainder
Process B			BLOCKED

Process B exits its critical section and executes the signal (mutex) instruction.

Mutual Exclusion through Semaphores(26/34)

Mutual Exclusion Achieved Through Semaphores		
Mutex Value	Mutex List	
0	C	
		<pre>wait (mutex) { value = value - 1; if (value < 0) { add process to list; block process; } }</pre>
		<pre>signal (mutex) { value = value + 1; if (value <= 0) { remove process from list; wake up process; } }</pre>
Process A	Process B	Process C
start wait (mutex); critical section signal (mutex); do remainder	start wait (mutex); critical section signal (mutex); do remainder	start wait (mutex); critical section signal (mutex); BLOCKED do remainder

The Mutex Value is incremented by 1.

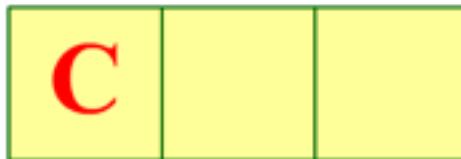
Mutual Exclusion through Semaphores(27/34)

Mutual Exclusion Achieved Through Semaphores

Mutex Value

0

Mutex List



```
wait (mutex)
{ value = value - 1;
  if (value < 0)
  { add process to list;
    block process;
  }
}
```

```
signal (mutex)
{ value = value + 1;
  if (value <= 0)
  { remove process from list;
    wake up process;
  }
}
```

Process A

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process B

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process C

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

BLOCKED

A comparison is made to determine if any processes are waiting in the Mutex List.

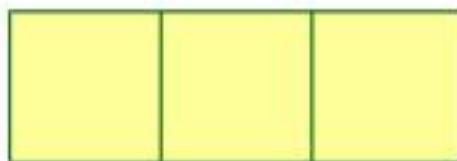
Mutual Exclusion through Semaphores(28/34)

Mutual Exclusion Achieved Through Semaphores

Mutex Value

0

Mutex List



```
wait (mutex)
{ value = value - 1;
  if (value < 0)
    { add process to list;
      block process;
    }
}
```

```
signal (mutex)
{ value = value + 1;
  if (value <= 0)
    { remove process from list;
      wake up process;
    }
}
```

Process A

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process B

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process C

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Test passes, the first process in the Mutex List is removed from the list and woken up.

Mutual Exclusion through Semaphores(29/34)

Mutual Exclusion Achieved Through Semaphores

Mutex Value	Mutex List	wait (mutex) { value = value - 1; if (value < 0) { add process to list; block process; } } }	signal (mutex) { value = value + 1; if (value <= 0) { remove process from list; wake up process; } }
0			

Process A	Process B	Process C
start wait (mutex); critical section signal (mutex); do remainder	start wait (mutex); critical section signal (mutex); do remainder	start wait (mutex); critical section signal (mutex); do remainder

Context Switch, Process C resumes execution and enters its critical section.

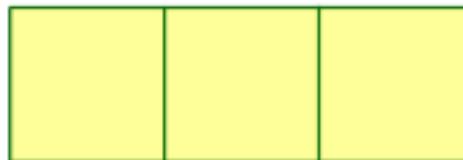
Mutual Exclusion through Semaphores(30/34)

Mutual Exclusion Achieved Through Semaphores

Mutex Value

0

Mutex List



```
wait (mutex)
{ value = value - 1;
  if (value < 0)
  { add process to list;
    block process;
  }
}
```

```
signal (mutex)
{ value = value + 1;
  if (value <= 0)
  { remove process from list;
    wake up process;
  }
}
```

Process A

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process B

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process C

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process C exits its critical section and executes the signal (mutex) instruction.

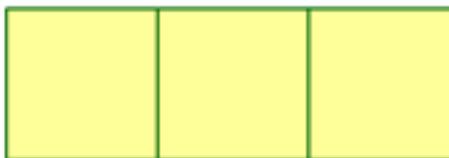
Mutual Exclusion through Semaphores(31/34)

Mutual Exclusion Achieved Through Semaphores

Mutex Value

Mutex List

1



```
wait (mutex)
{ value = value - 1;
  if (value < 0)
  { add process to list;
    block process;
  }
}
```

```
signal (mutex)
{ value = value + 1;
  if (value <= 0)
  { remove process from list;
    wake up process;
  }
}
```

Process A

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process B

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process C

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```



The Mutex Value is incremented by 1.

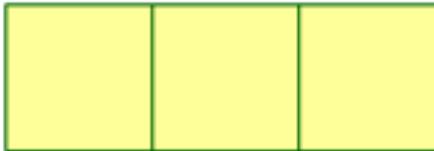
Mutual Exclusion through Semaphores(32/34)

Mutual Exclusion Achieved Through Semaphores

Mutex Value

Mutex List

1



```
wait (mutex)
{ value = value - 1;
  if (value < 0)
  { add process to list;
    block process;
  }
}
```

```
signal (mutex)
{ value = value + 1;
  if (value <= 0)
  { remove process from list;
    wake up process;
  }
}
```

Process A

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process B

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process C

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```



Test fails, there are no processes waiting in the Mutex List.

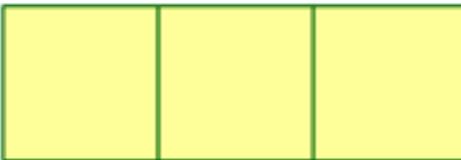
Mutual Exclusion through Semaphores(33/34)

Mutual Exclusion Achieved Through Semaphores

Mutex Value

Mutex List

0



```
wait (mutex)
{ value = value - 1;
  if (value < 0)
  { add process to list;
    block process;
  }
}
```

```
signal (mutex)
{ value = value + 1;
  if (value <= 0)
  { remove process from list;
    wake up process;
  }
}
```

Process A

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process B

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process C

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process A is allowed to enter its critical section.

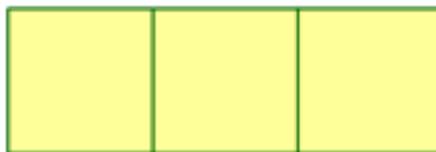
Mutual Exclusion through Semaphores(34/34)

Mutual Exclusion Achieved Through Semaphores

Mutex Value

1

Mutex List



```
wait (mutex)
{ value = value - 1;
  if (value < 0)
  { add process to list;
    block process;
  }
}
```

```
signal (mutex)
{ value = value + 1;
  if (value <= 0)
  { remove process from list;
    wake up process;
  }
}
```

Process A

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process B

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process C

```
start
wait (mutex);
critical section
signal (mutex);
do remainder
```

Process C proceeds with remainder of code.

As you can see, Semaphores can be used to achieve mutual exclusion, even when a context switch could possibly take place during execution of a critical section.

Process Synchronization with Semaphore(1/4)

- **Example 1:**
- Suppose We have 2 processes: **P1** and **P2**
 - where **P1** contains code block **B1** and **P2** contains code block **B2**
 - We want to ensure code block B1 in **P1** executes before code block B2 in **P2**?
- **Use semaphores to synchronize the order of execution of the two processes?**

Process Synchronization with Semaphore(2/4)

- **Example 1:**
- Suppose We have 2 processes: P1 and P2
 - where P1 contains code B1 and P2 contains code B2
 - We want to ensure code block B1 in P1 executes before code block B2 in P2?
- **Use semaphores to synchronize the order of execution of the two processes?**
- **Solution:**

- Define a semaphore “synch”
- Initialize synch to 0
- **Process P1**
 - B1; //executes B1
 - Signal(synch); //increment Sync
- **Process P2**
 - Wait(synch); //decrement Sync
 - B2; //executes B2

Process Synchronization with Semaphore(3/4)

- **Example 2:**
- Suppose We have 3 processes: **P1**, **P2** and **P3**
- where **P1** prints **A** and **P2** prints **B** and **P3** prints **C**
- We want to ensure to have **ABCABCABC.....** in output?

Process Synchronization with Semaphore(4/4)

- **Example 2:**

- Suppose We have 3 processes: **P1**, **P2** and **P3**
- where **P1** prints **A** and **P2** prints **B** and **P3** prints **C**
- We want to ensure to have **ABCABCABC.....** in output?

- Semaphore x=0, y=0 and z=1; //initial value of semaphore

- **Process P1**

- Wait (z)
 - Print “A”
 - Signal(x); //release(x)

- **Process P2**

- Wait(x);
 - Print B;
 - Signal(y)

- **Process P3**

- Wait(y);
 - Print “C”
 - Signal(z); //release(z)

Chapter 5: Process Synchronization

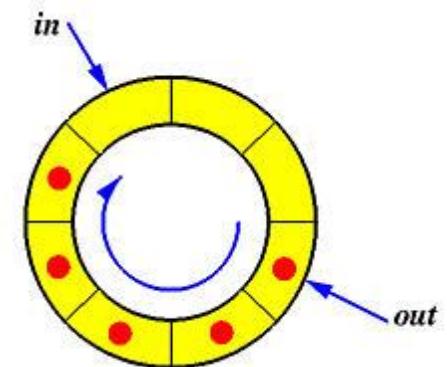
- Background
- The Critical-Section Problem
- Peterson's Solution
- Synchronization Hardware
- Mutex Locks
- Semaphores
- **Classic Problems of Synchronization**
- Monitors
- Condition Variables

Classical Problems of Synchronization

- Classical problems used to test newly-proposed synchronization schemes
 - Bounded-Buffer Problem (**also called the producer consumer problem**)
 - Readers and Writers Problem
 - Dining-Philosophers Problem

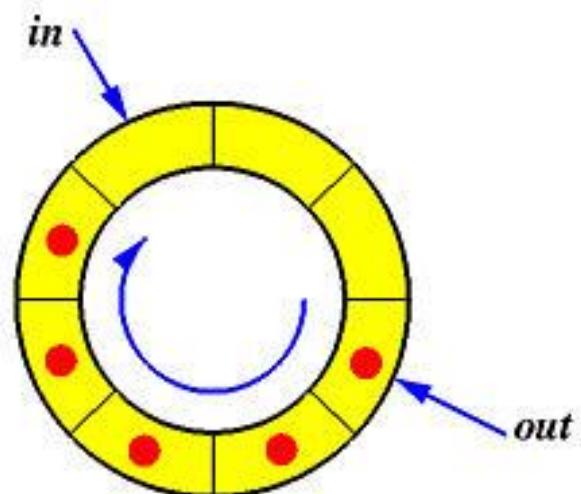
1. Bounded-Buffer Problem

- Also called producer-consumer problem
- This arises when two or more processes communicate with each other.
- Some produce data and others consume data.
- Bounded buffer: size ‘N’
 - Access entry 0... N-1, then “wrap around” to 0 again
- Producer process writes data to buffer
 - Must not write more than ‘N’ items more than consumer “consume”
- Consumer process reads data from buffer
 - Should not try to consume if there is no data
- All processes modify the same buffer
- Only one process modify buffer at the same time



1. Bounded-Buffer Problem

- Producer and Consumer share the following data:
 - n buffers, each can *hold* one item
 - Semaphore **mutex** initialized to the value 1; **for permission access to buffer pool**
 - Semaphore **full** initialized to the value 0; **number of full buffers**
 - Semaphore **empty** initialized to the value n ; **number of empty buffers**



1. Bounded Buffer Problem (Cont.)

- The structure of the *producer process*; solution using Semaphores
- Semaphore mutex=1, empty=n, full=0

```
do {  
    ...  
    /* produce an item in  
    next_produced */  
    ...  
    wait(empty);  
    wait(mutex);  
    ...  
    //produced item  
    //add to the buffer  
    ...  
    signal(mutex);  
    signal(full);  
} while (true);
```

- The structure of the *consumer process*; solution using semaphores
 - Semaphore mutex=1, empty=n, full=0
- ```
do {
 wait(full);
 wait(mutex);
 ...
 /* remove an item from buffer
 to next_consumed */
 ...
 signal(mutex);
 signal(empty);
 ...
} while (true);
```

## 2. Readers-Writers Problem

- A dataset is shared among a number of concurrent processes
  - Readers – only read the dataset; they do **not** perform any updates
    - Never a problem if only readers access the data; no waiting necessary
  - Writers – can both read and write
    - Problem if some writer accesses data; thus, exclusive access only for writing
- **Problem** – allow multiple concurrent readers to read at the same time
  - Only one single writer can access the shared data at the same time
  - And if a writer is active, readers wait for it to finish

## 2. Readers-Writers Problem

- Several variations of how readers and writers are considered – all involve some form of priorities
- Solving Reader-Writer problem using **Semaphore**
  - Semaphore **rw\_mutex** initialized to 1 //used by reader and writer
  - Semaphore **mutex** initialized to 1 //Only used by reader to inc/dec **read\_count**
  - Integer **read\_count** initialized to 0 // number of readers

## 2. Readers-Writers Problem (Cont.)

- The structure of a *writer process*

```
do {
 wait(rw_mutex);
//mutual exclusion for
writers

/* writing is performed
*/
 ...
 ...
 signal(rw_mutex),
} while (true);
```

- The structure of a *reader process*

```
do {
 wait(mutex);
read_count++; //critical section when
read_count being modified
if (read_count == 1)//if I'm first reader

wait(rw_mutex); //in case a writer is
writing
signal(mutex); //Lets another reader.

...
/* reading is performed */
 ...
 ...
 wait(mutex);
read count--; //critical section
when read_count being modified
if (read_count == 0)

signal(rw_mutex); //only last reader release
lock

signal(mutex); //release lock
} while (true);
```

### 3. Reader-writer problem

- If there is a writer
  - Only first reader blocks on wr\_mutex
  - Other readers block on mutex
- Once a reader is active, all readers get to go through
  - Which reader gets in first?
- The last reader to exit signals a writer
  - If no writer, then readers can continue
- If readers and writers are waiting on wr\_mutex, and writer exits
  - Who gets to go in first?
- Why doesn't a writer need to use mutex?
- Is the previous solution fair?
- Readers can “starve” writers!
- Building a “fair” solution is tricky!

### 3. Dining-Philosophers Problem

- Philosophers spend their lives alternating ***thinking*** and ***eating***
- Don't interact with their neighbors, occasionally try to pick up 2 chopsticks (one at a time) to eat from bowl
  - Need both chopsticks to eat, then release both when done
- In the case of ***5 philosophers***
  - Shared data
    - Bowl of rice (dataset)
    - Semaphore **chopstick [5]** initialized to **1**



### 3. Dining-Philosophers Problem Algorithm

- The structure of Philosopher *i*:

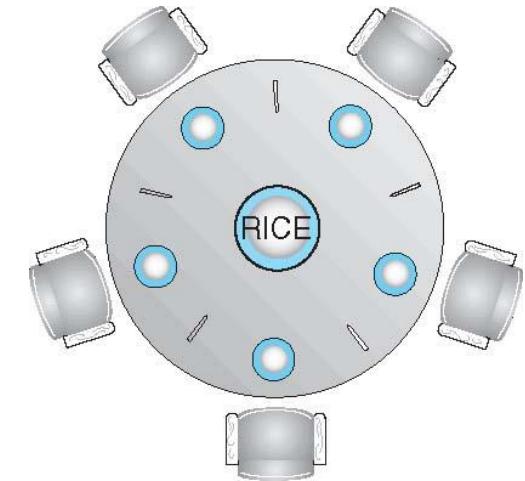
```
do {
 wait (chopstick[i]);
 wait (chopStick[(i + 1) % 5]);

 // eat

 signal (chopstick[i]);
 signal (chopstick[(i + 1) % 5]);

 // think
}

} while (TRUE);
```



- What is the problem with this algorithm?
  - Deadlock is possible if two neighbor philosophers grab only one chopstick and starts waiting for the other.
  - For example, i=2, then philosopher#2 grabs chopstick#2 first, and meanwhile philosopher#3 also becomes ready and grabs chopstick#3. Deadlock will occur

### 3. Dining-Philosophers Problem Algorithm (Cont.)

- Deadlock handling
  - Allow ***at most 4 philosophers*** to be sitting simultaneously at the table.
  - Allow a philosopher to pick up the ***forks only if both are available*** (picking must be done in a critical section).
  - Use an ***asymmetric solution*** -- an odd-numbered philosopher picks up first the left chopstick and then the right chopstick. Even-numbered philosopher picks up first the right chopstick and then the left chopstick.



# Chapter 5: Process Synchronization

- Background
- The Critical-Section Problem
- Peterson's Solution
- Synchronization Hardware
- Mutex Locks
- Semaphores
- Classic Problems of Synchronization
- **Monitors**
- Condition variables

# Problems with Semaphores

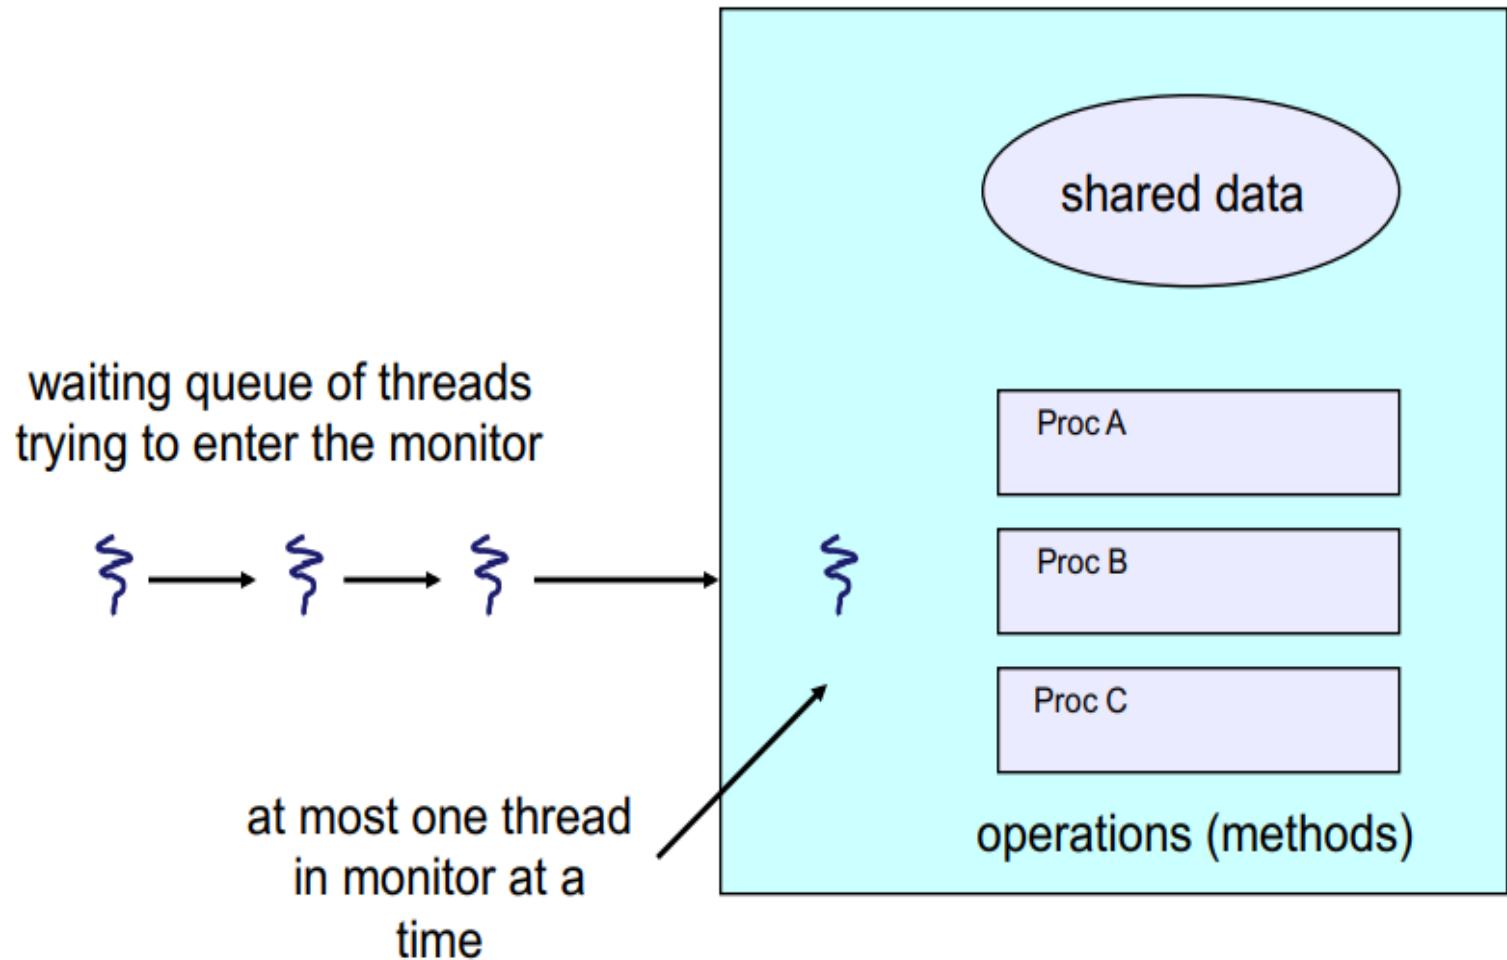
- Incorrect use of semaphore operations will result in timing errors.
- For example, If the **wait()** and **signal()** order is interchanged such as below, then there maybe several processes in the critical section
  - signal (mutex) .... wait (mutex)
- In the below cases a deadlock may occur.
  - wait (mutex) ... wait (mutex)
  - Omitting of wait (mutex) or signal (mutex) (or both)
- Wait() and Signal() are scattered among several processes. Therefore, it is difficult to understand their effects.
- Usage must be correct in all the processes.
- One bad process can kill the whole system.

# Monitors

- A **monitor** is a programming language construct that controls access to shared data
- A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- A monitor is a module that encapsulates
  - Shared **data structures**
  - **Procedures** that operate on the shared data structures
  - **Synchronization** between concurrent threads that invoke the procedures
- Only one process may be active within the monitor at a time
  - The **monitor** ensures **mutual exclusion**
  - no need to program this constraint explicitly

```
monitor monitor-name {
 // shared variable declarations
 procedure P1 (...) { }
 procedure Pn (...) {.....}
 Initialization code (...) { ... }
}
```

# Schematic view of a Monitor



# Monitor in Java (Synchronized Java methods)

- The monitor enforces mutual exclusive access to synchronized methods invoked on the associated object.
- A java class can be turned into monitor by making its one or more classes **synchronized**
- For example, a bounded buffer class can be made monitor by:

```
public void synchronized put(Object item)
```

```
public char synchronized get()
```

# Monitor for Producer-Consumer in Java

```
class Buffer {
 private char [] buffer;
 private int count = 0, in = 0, out = 0;
 Buffer(int size) {
 buffer = new char[size]; }

 public synchronized void Put(char c) {
 while(count == buffer.length);
 System.out.println("Producing " + c + " ...");
 buffer[in] = c;
 in = (in + 1) % buffer.length;
 count++; }

 public synchronized char Get() {
 while (count == 0);
 char c = buffer[out];
 out = (out + 1) % buffer.length;
 count--;
 System.out.println("Consuming " + c + " ..."); return c; }
}
```

```
public class PC
{
 public static void main(String [] args)
 {
 Buffer b = new Buffer(4);
 Producer p = new Producer(b);
 Consumer c = new Consumer(b);

 p.start(); //p puts items
 c.start(); //c gets items
 }
}
```

What are limitations of the above program?

- Execute the above program. And find out the problem? How can we handle the issue?

```
bilal@bilal-Latitude-E6540:~/Dropbox/os_prac$ java PC
Producing A ...
Consuming A ...
^Cbilal@bilal-Latitude-E6540:~/Dropbox/os_prac$ javac PCmonitor.j
```

# Condition variables

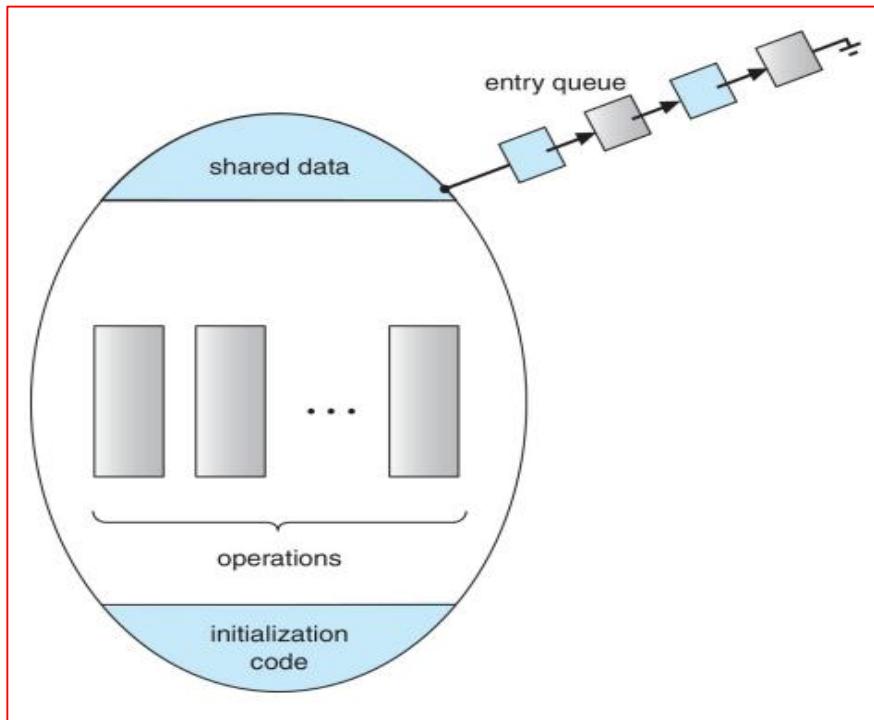
- Mutual exclusion is an easy task with monitors.
- But not powerful enough to model some synchronization schemes
- While a process is executing in a monitor, it may have to wait until an event occurs.
- Process synchronization is done using condition variables, which represent conditions (for example, **notEmpty**, **notFull**) a process may need to wait for before executing in the monitor
- A **condition variable**, or a condition, has a private waiting list, and two public methods: **signal** and **wait**.
-

# Condition Variables

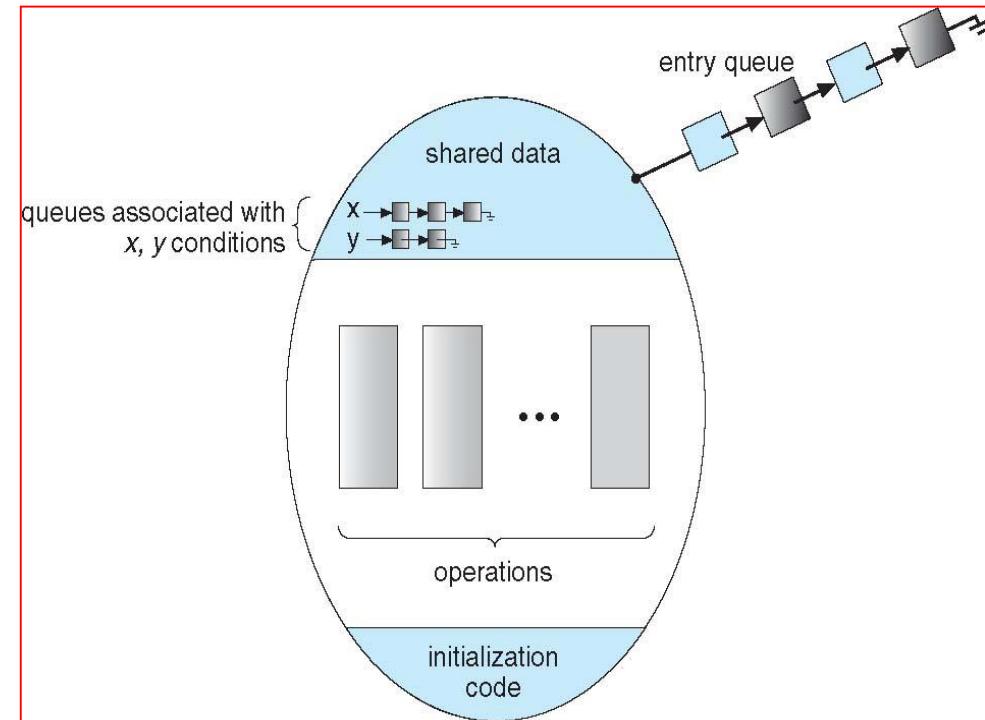
- Let `x` be a condition variable, defined as:
  - `condition x;`
- **Two** operations are allowed on a condition variable:
  - `x.wait()` – a process that invokes the operation is suspended until `x.signal()`
  - `x.signal()` – resumes one of processes (if any) that invoked `x.wait()`
    - If no one invoke `x.wait()` on the condition variable `x`, then `x.signal()` has no effect on the variable
- Condition **wait** and condition **signal** can only be used **in a monitor**.

# Monitor with Condition Variables

Monitor without condition variables



Monitor with condition variables



## Bounded Buffer (Producer Consumer) with Monitors

```
monitor BoundedBuffer {
```

```
// Private variables ...
```

```
Object buf = new Object[n];
```

```
integer out = 0, // index of first full slot
```

```
in = 0, // index of first empty slot
```

```
count = 0; // number of full slots
```

```
// Condition variables ...
```

```
condvar not_full, // signalled when count < n. Consumer sends signal.
```

```
not_empty; // signalled when count > 0. Producer sends signal on this var.
```

```
//See next slide for the remaining part.
```

# Bounded Buffer (Producer Consumer) with Monitors

```
// Monitor procedures ...
//(signal & continue signalling discipline)
```

```
procedure put(Object data) {
```

```
 while(count == n) {
```

```
 wait(not_full);
```

```
}
```

```
 buf[in] = data;
```

```
 in = (in + 1) % n;
```

```
 count++;
```

```
 signal(not_empty);
```

```
}
```

```
procedure get(Object &item) {
```

```
 while(count == 0) {
```

```
 wait(not_empty);
```

```
}
```

```
 item = buf[out];
```

```
 out = (out + 1) % n;
```

```
 count--;
```

```
 signal(not_full);
```

```
} }
```

# Synchronization using wait() and notify()

- The `wait()` method suspends the calling thread. It wakes up only when another thread calls `notify()`.

## Producer Thread

```
for(int i = 0; i < 10; i++) {
 buffer.Put((char)('A'+ i%26));
}
```

## Consumer Thread

```
for(int i = 0; i < 10; i++) {
 buffer.Get();
}
```

```
public synchronized void Put(char c) {
 if (count == buffer.length) wait();
 System.out.println("Producing " + c + " ...");
 buffer[in] = c;
 in = (in + 1) % buffer.length;
 count++;
 notify(); }
```

```
public synchronized char Get() {
 if (count == 0) wait();
 char c = buffer[out];
 out = (out + 1) % buffer.length;
 count--;
 System.out.println("Consuming " + c + " ...");
 notify(); }
```

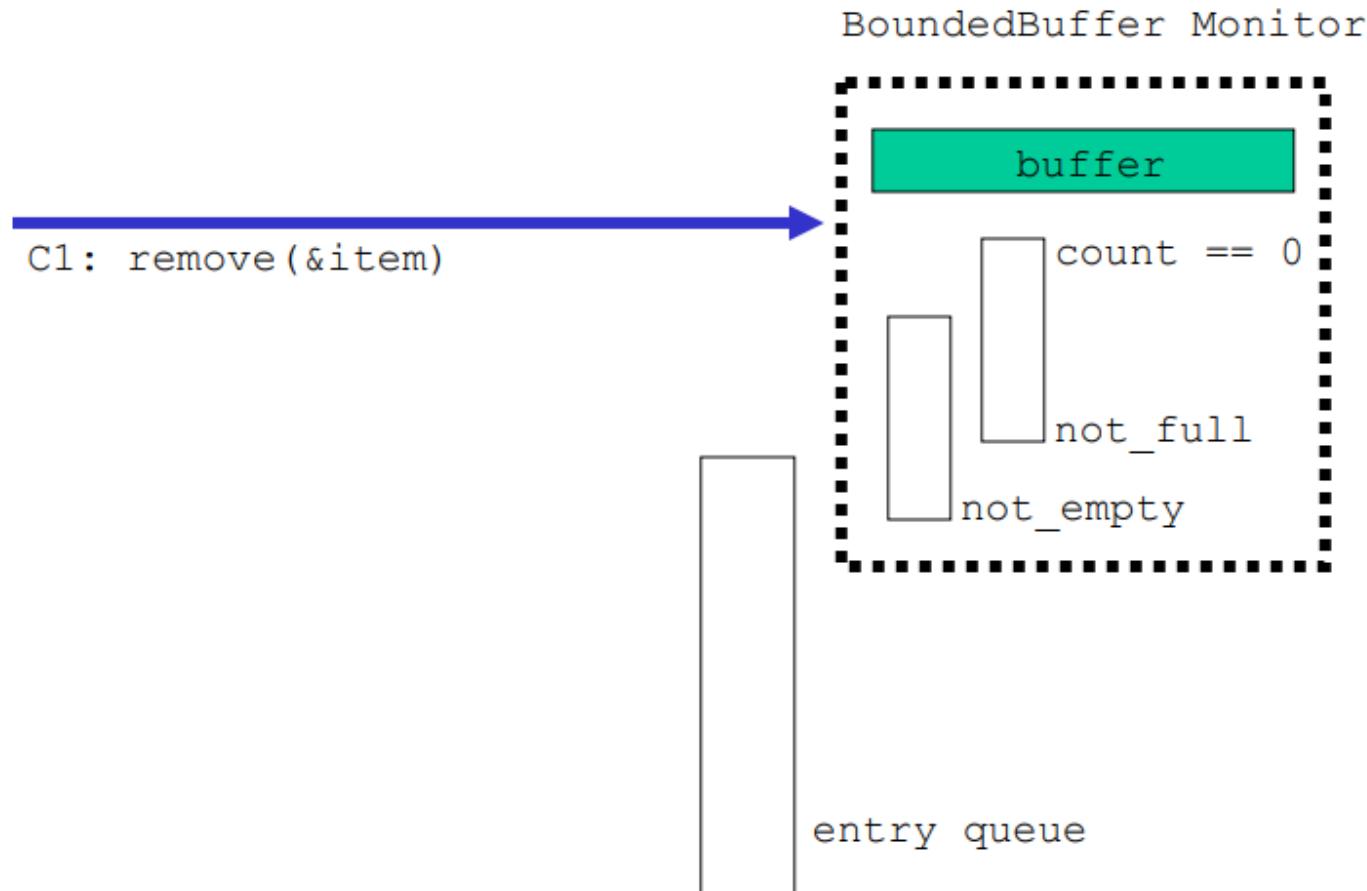
```
bilal@bilal-Latitude-E6540:~/Dropbox/os_prac$ java PCmonitor
Producing A ...
Producing B ...
Producing C ...
Producing D ...
Consuming A ...
Consuming B ...
Consuming C ...
Consuming D ...
Producing E ...
Producing F ...
Producing G ...
Producing H ...
Consuming E ...
Consuming F ...
Consuming G ...
Consuming H ...
Producing I ...
Producing J ...
Consuming I ...
Consuming J ...
bilal@bilal-Latitude-E6540:~/Dropbox/os_prac$
```

# Condition Variables Choices

- When a monitor procedure calls **signal** on a condition variable, it wakes up the first blocked process in the *delay queue* waiting on the condition.
- ***Signal and Wait***: the signaler waits until some later time and the signaled process executes now. In other words, this semantic lets the newly awakened process run, suspending the other one.
- ***Signal and Continue***: the signaler continues and the signaled process executes at some later time. In other words, the newly awakened process is moved from the signal queue back to the entry queue.
- The examples in this lecture use the *signal and continue* signaling discipline.

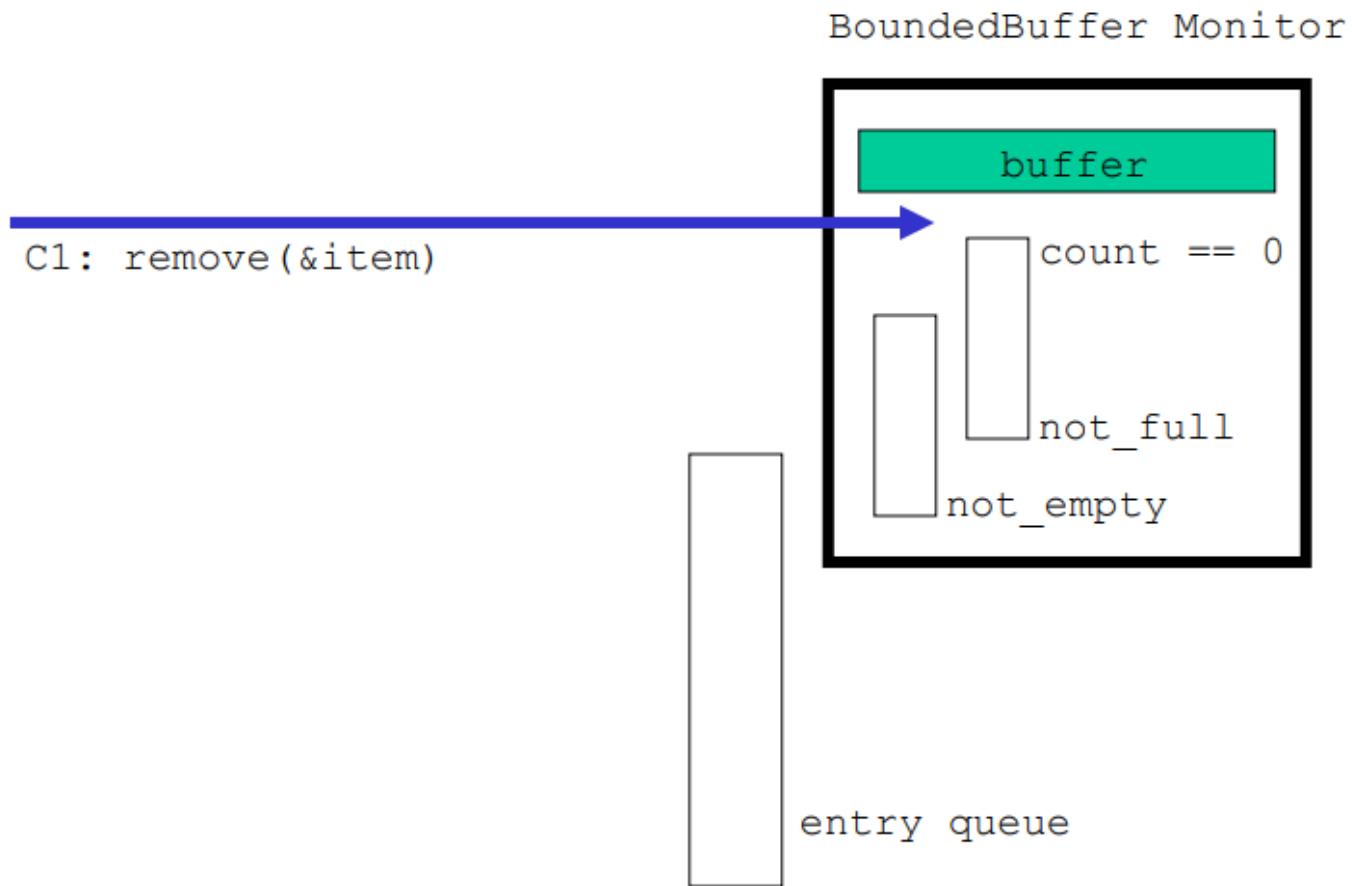
# Example Trace 1

- A consumer C1 arrives and wants to get() an item from the buffer.
- There is neither any consumer nor any producer currently in the monitor.
- So C1 will be allowed to enter monitor.



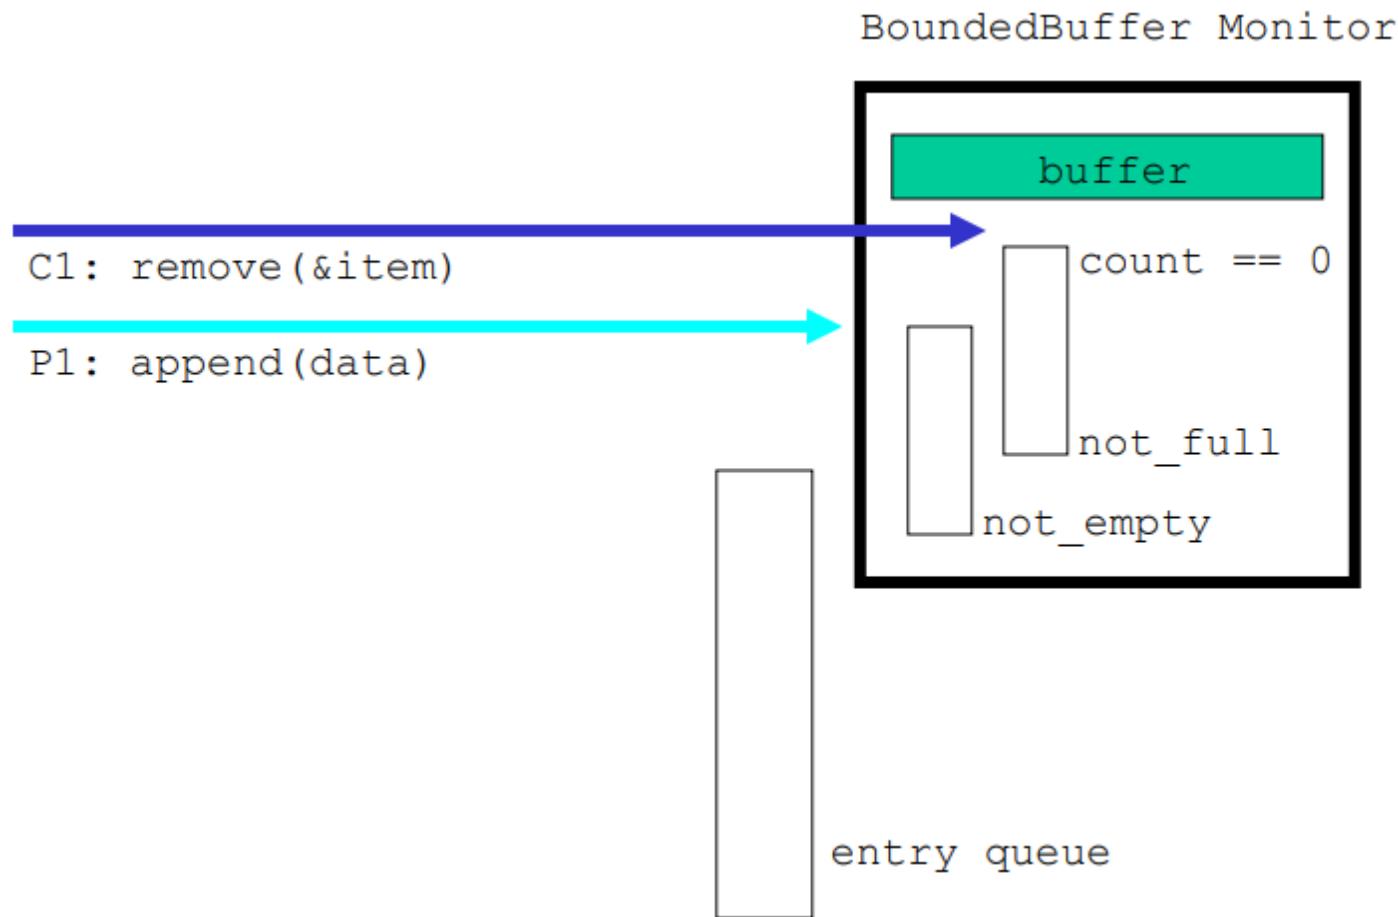
# Example Trace 2

- C1 is allowed to enter monitor.
- However, there is no item in the buffer. So C1 will wait on the condition variable **Not\_empty**.

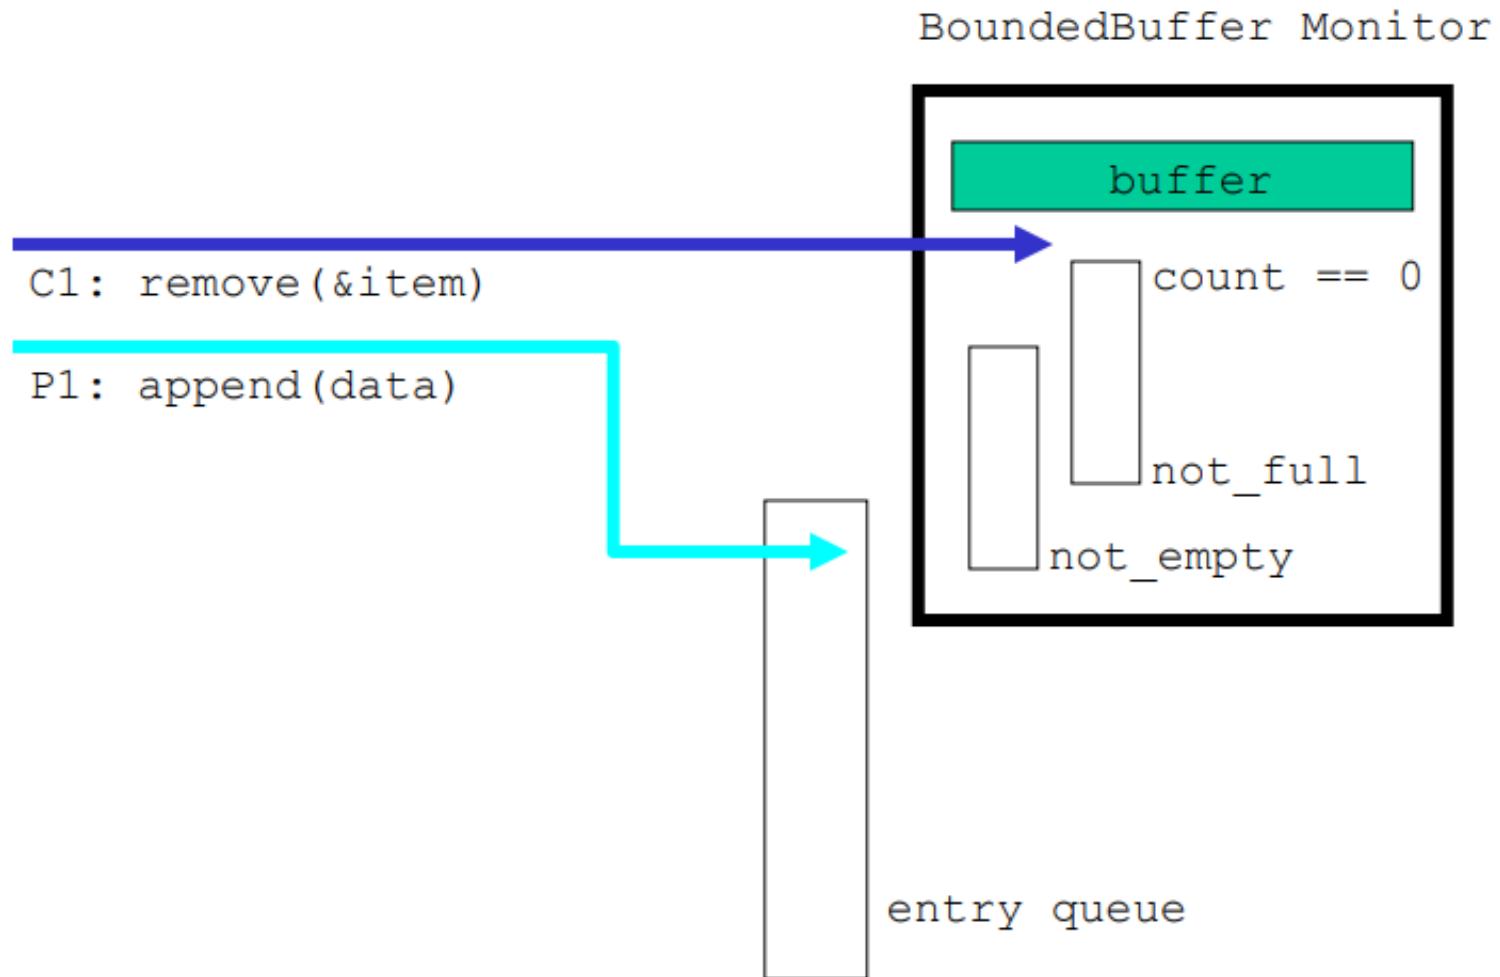


# Example Trace 3

- Next producer **P1** arrives.

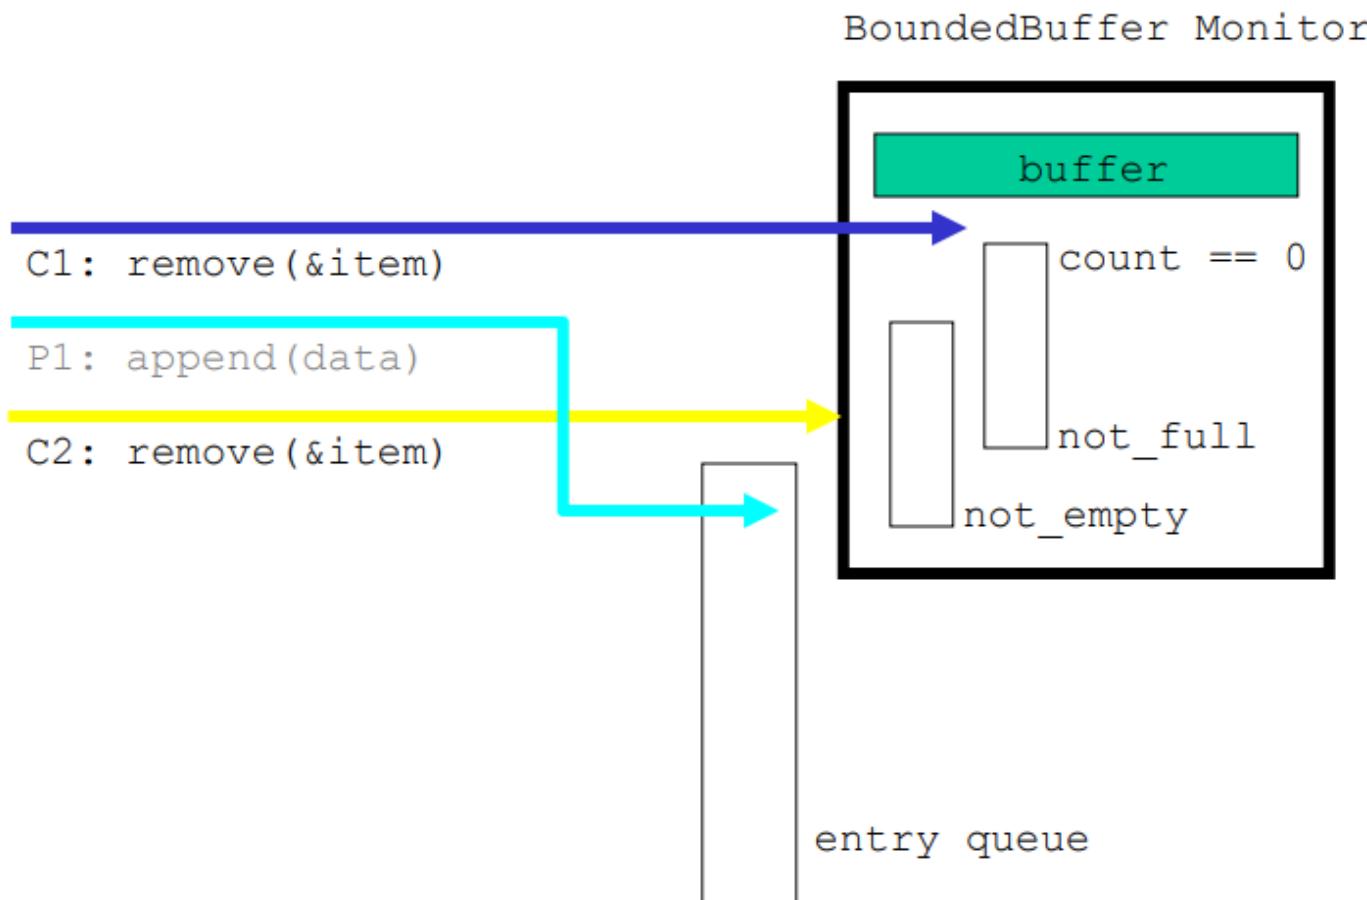


# Example Trace 4



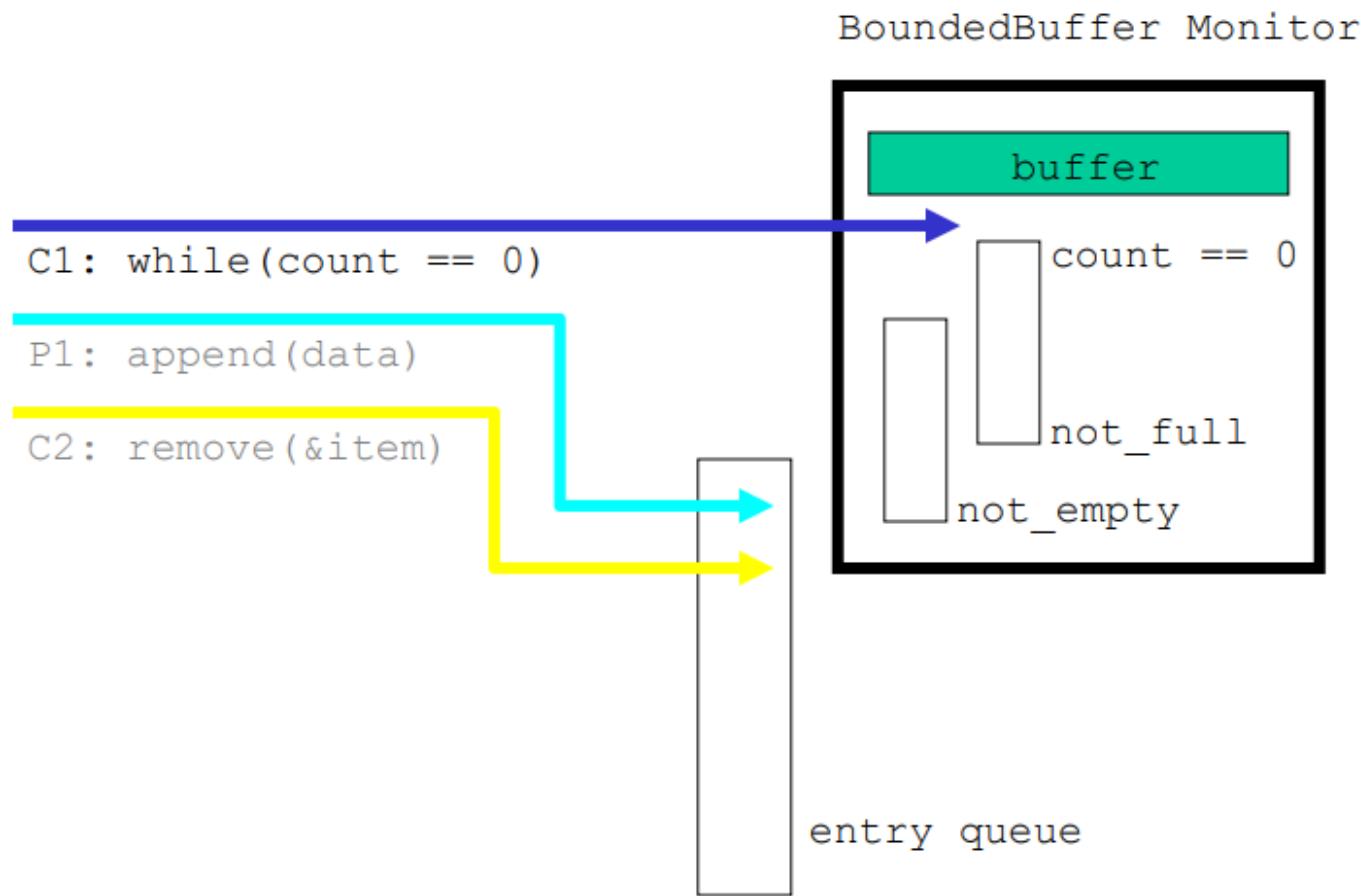
# Example Trace 5

- Meanwhile consumer **C2** arrives.

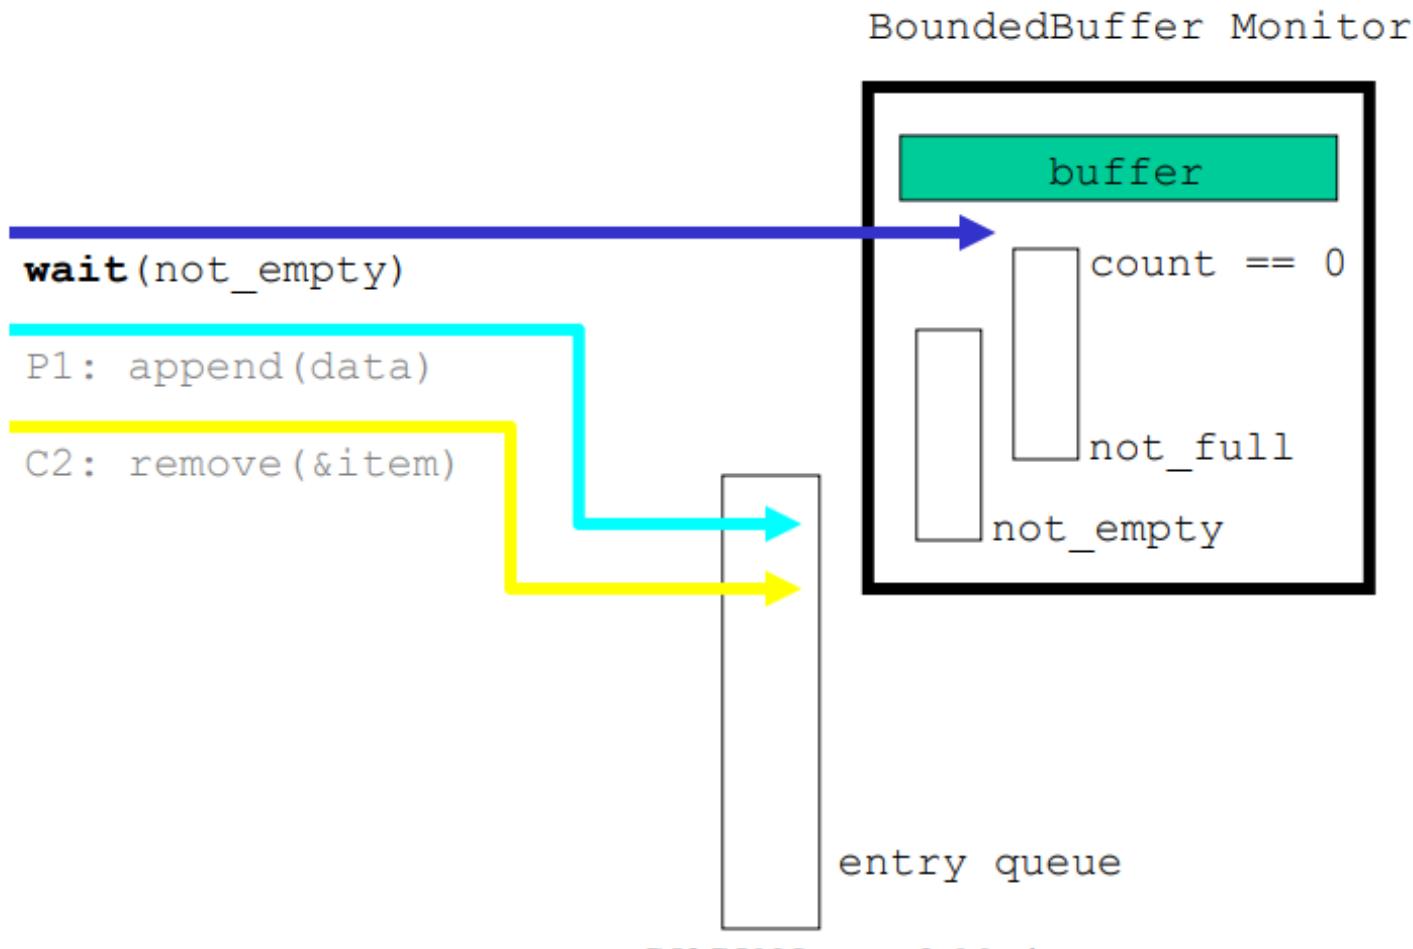


# Example Trace 6

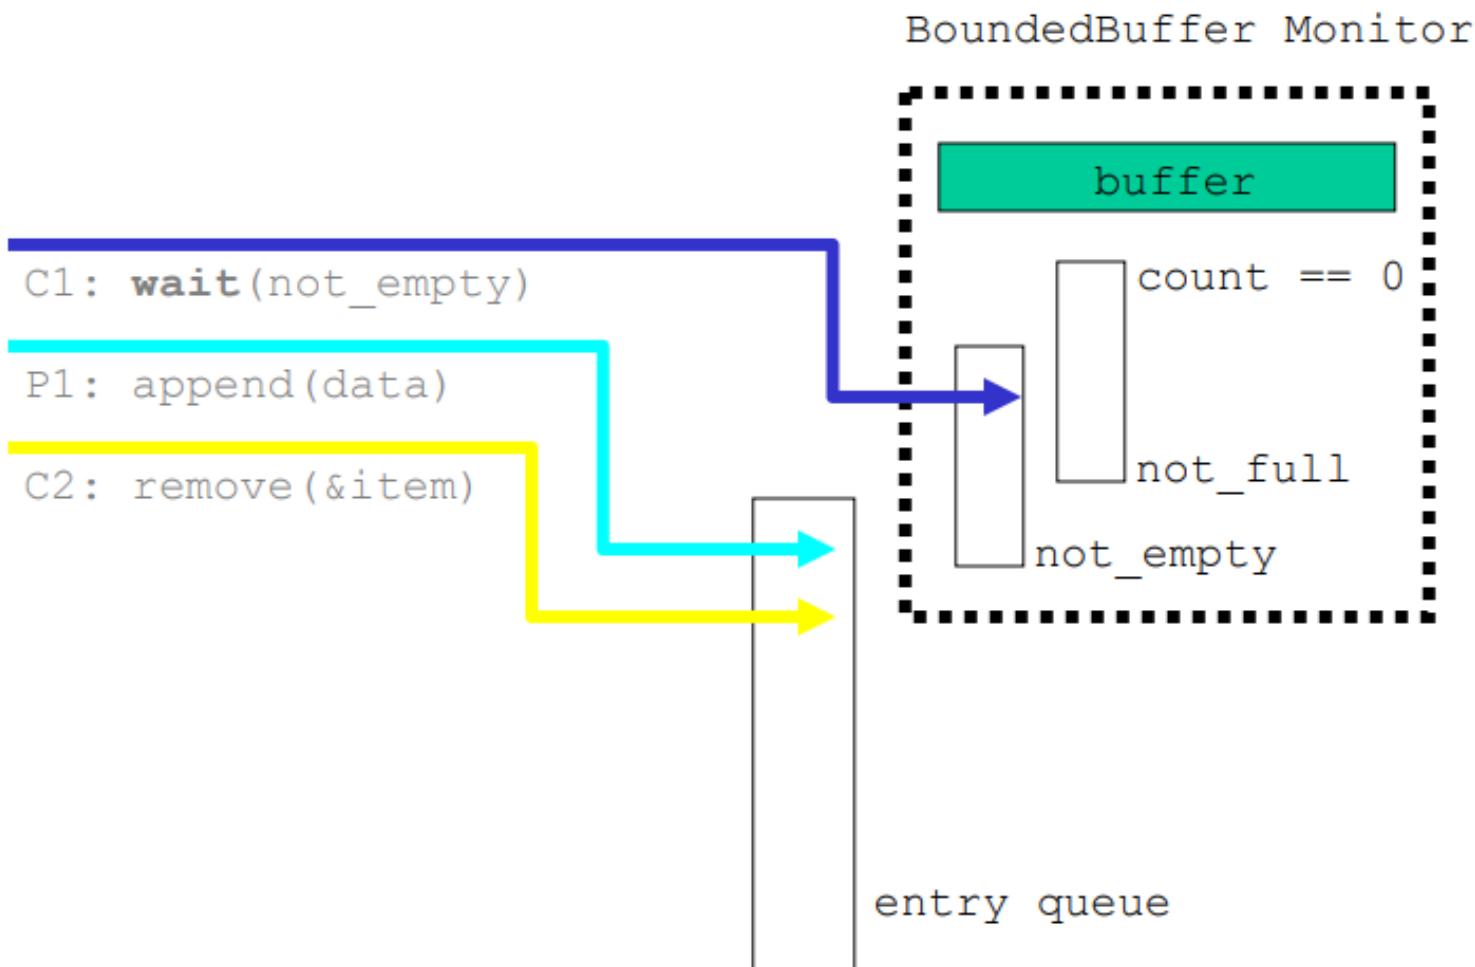
- Both P1 and C2 will wait outside the monitor to ensure Mutual exclusion.



# Example Trace 7

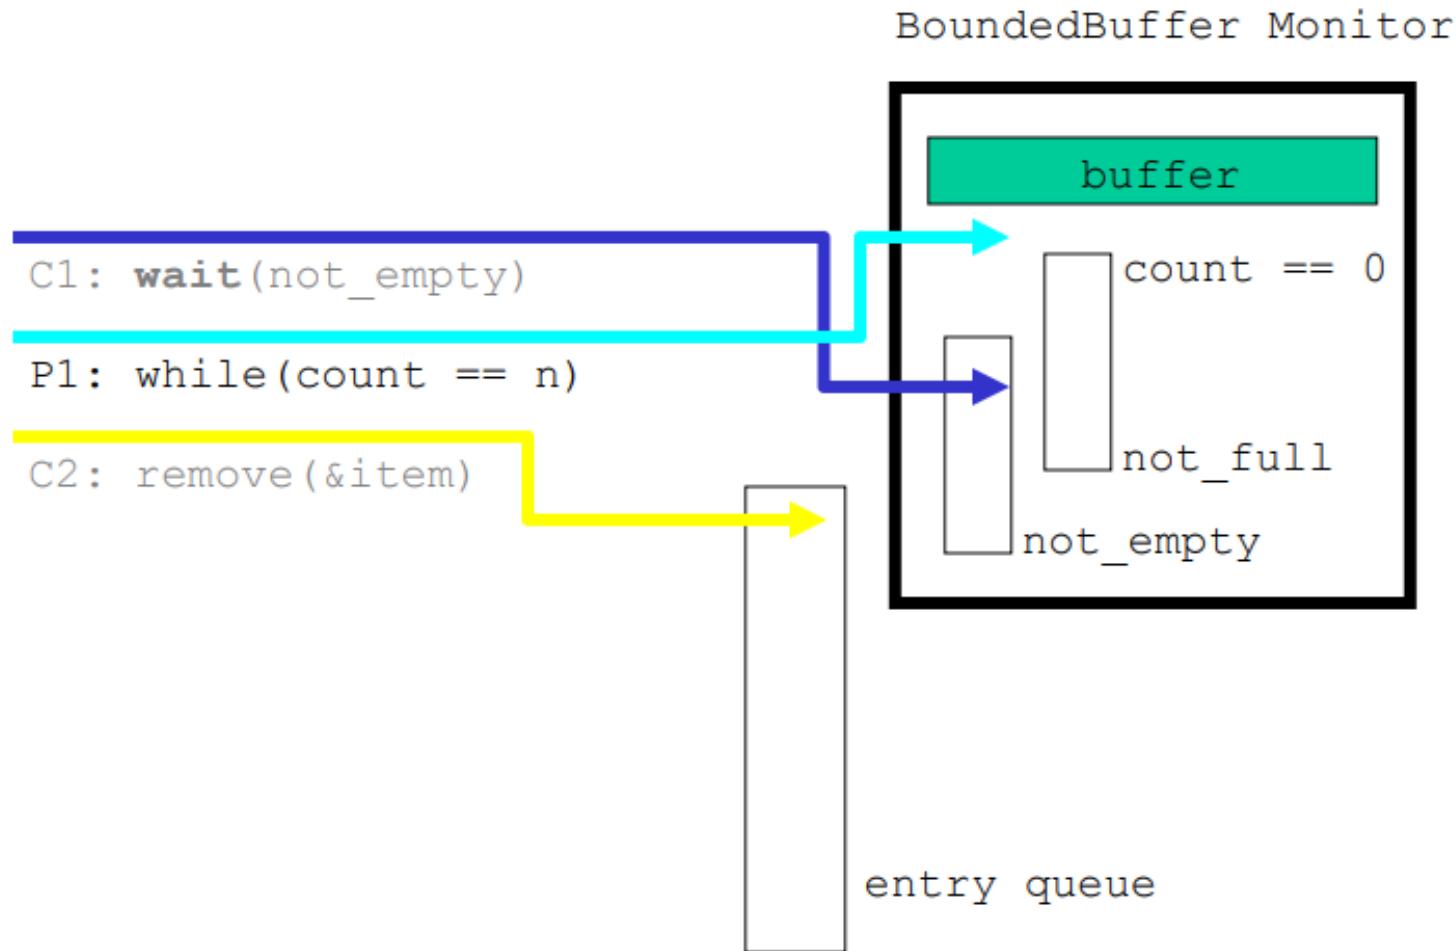


# Example Trace 8



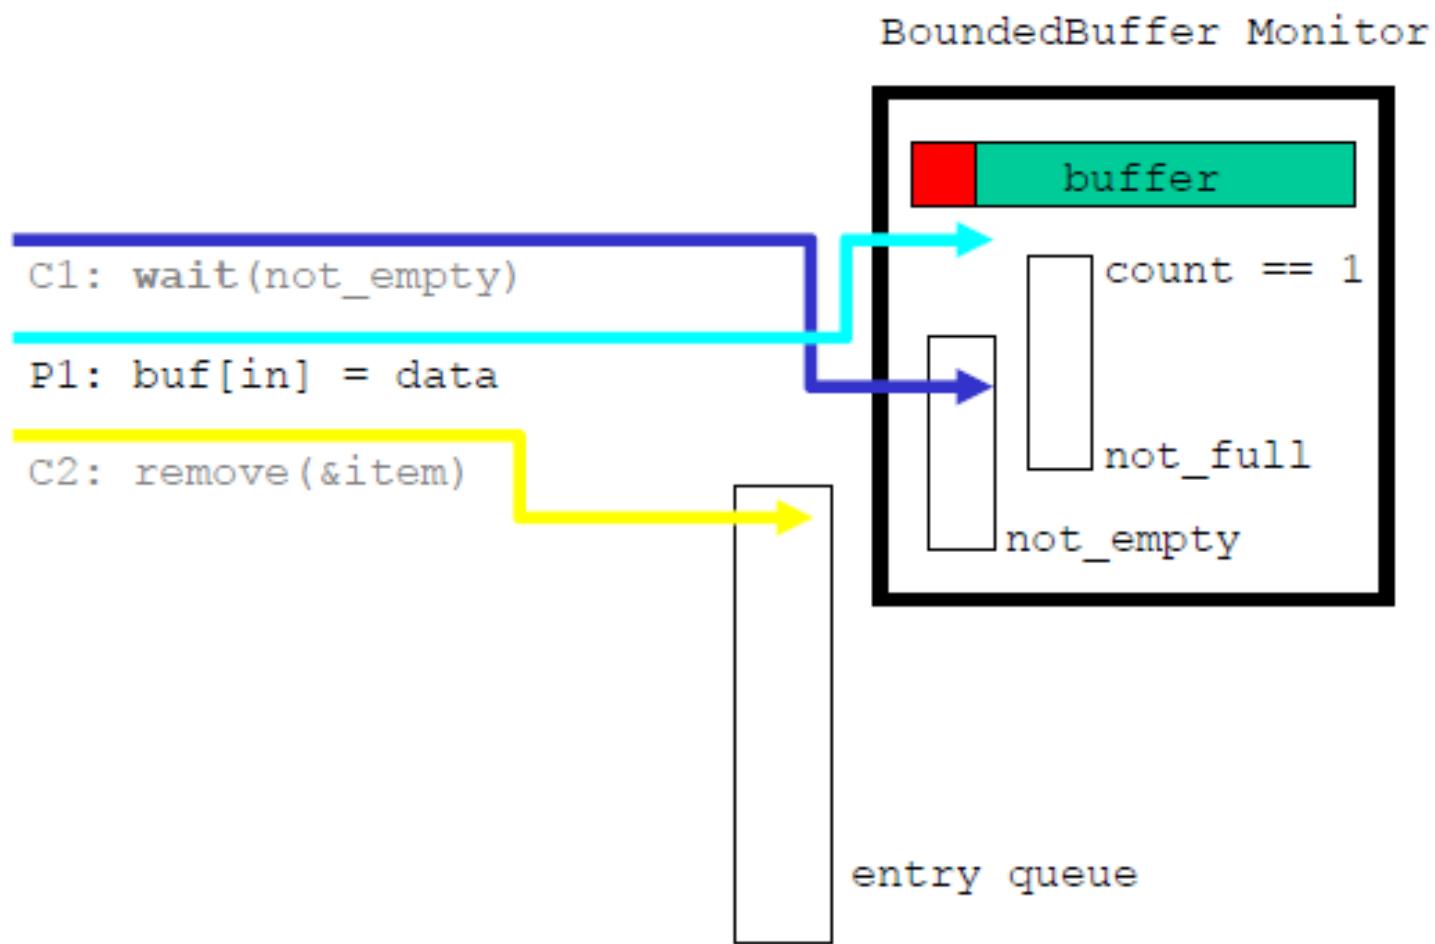
# Example Trace 9

- P1 can now enter the monitor as **C1** is waiting on the condition variable.

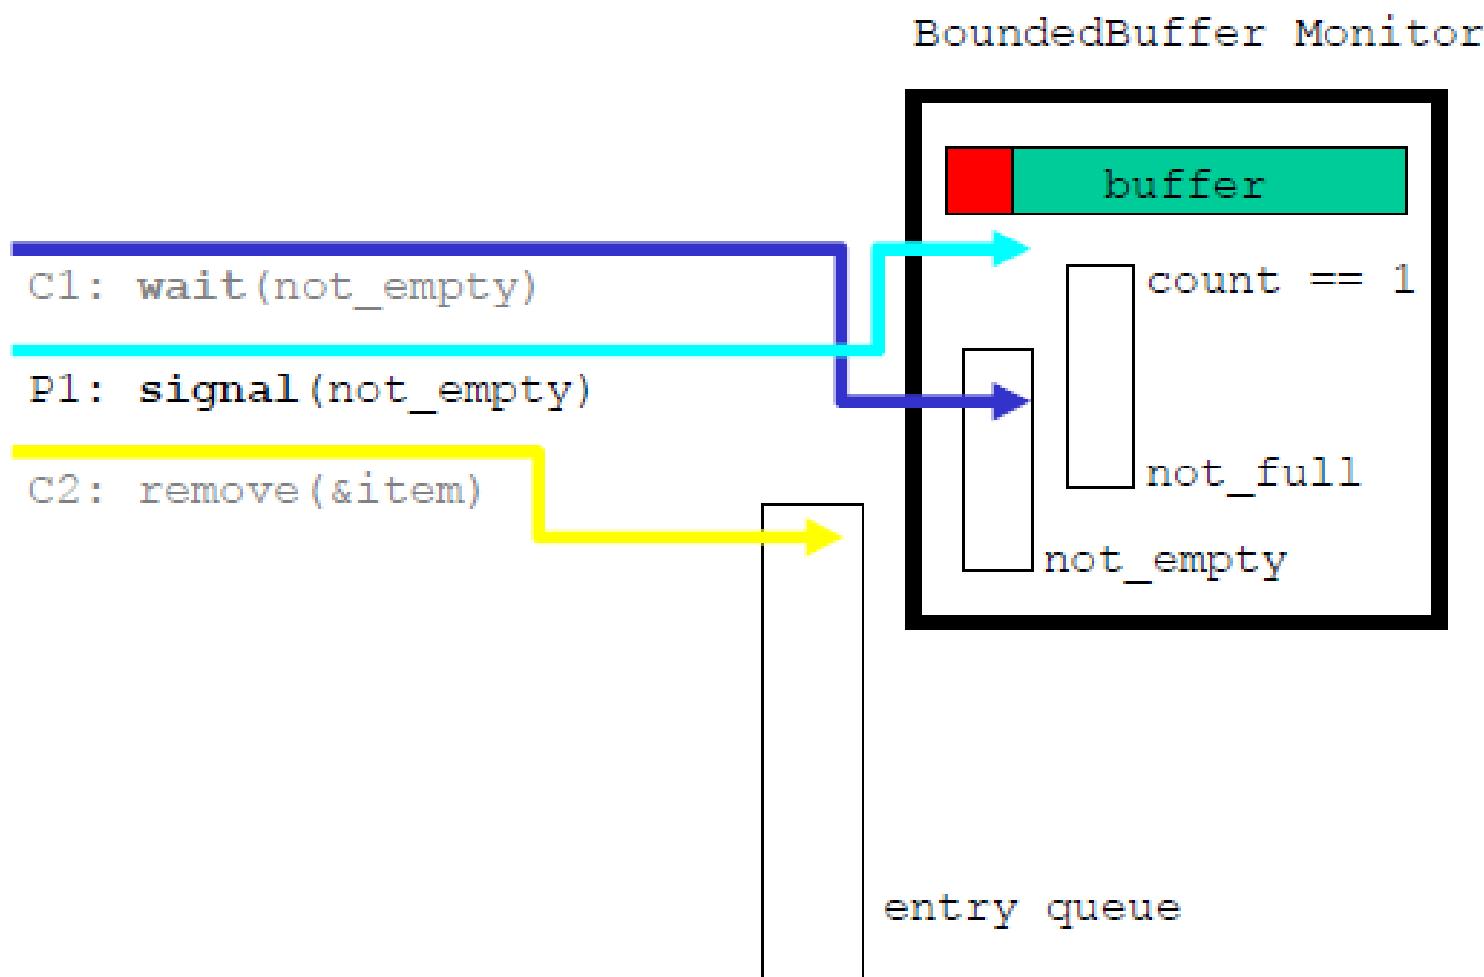


# Example Trace 10

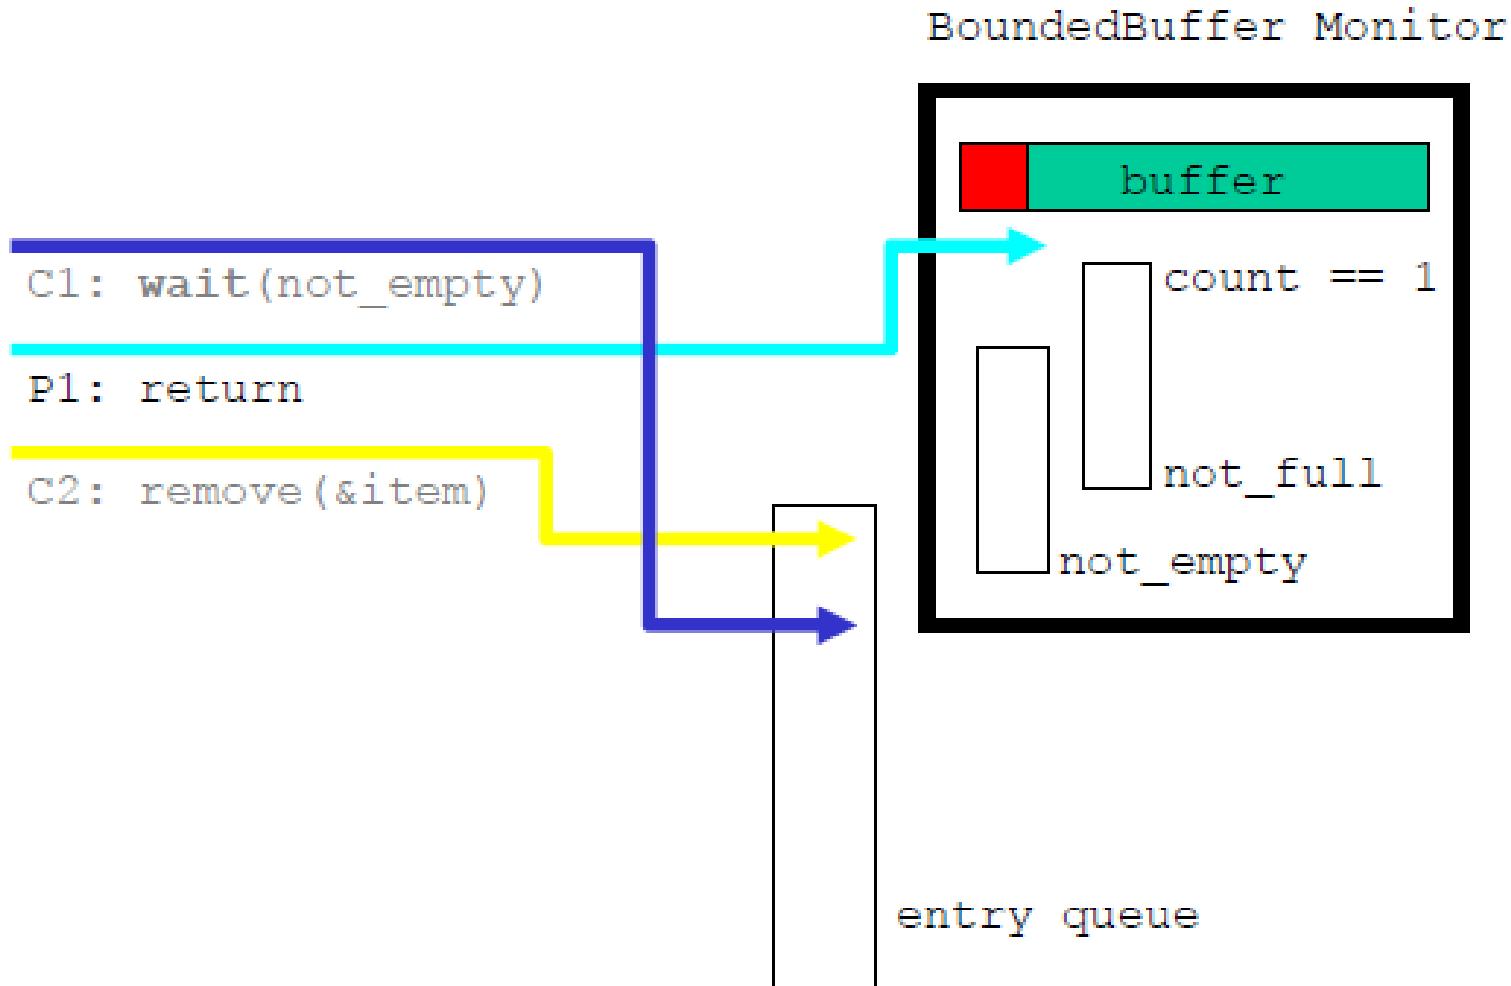
- P1 produces an item in the buffer.



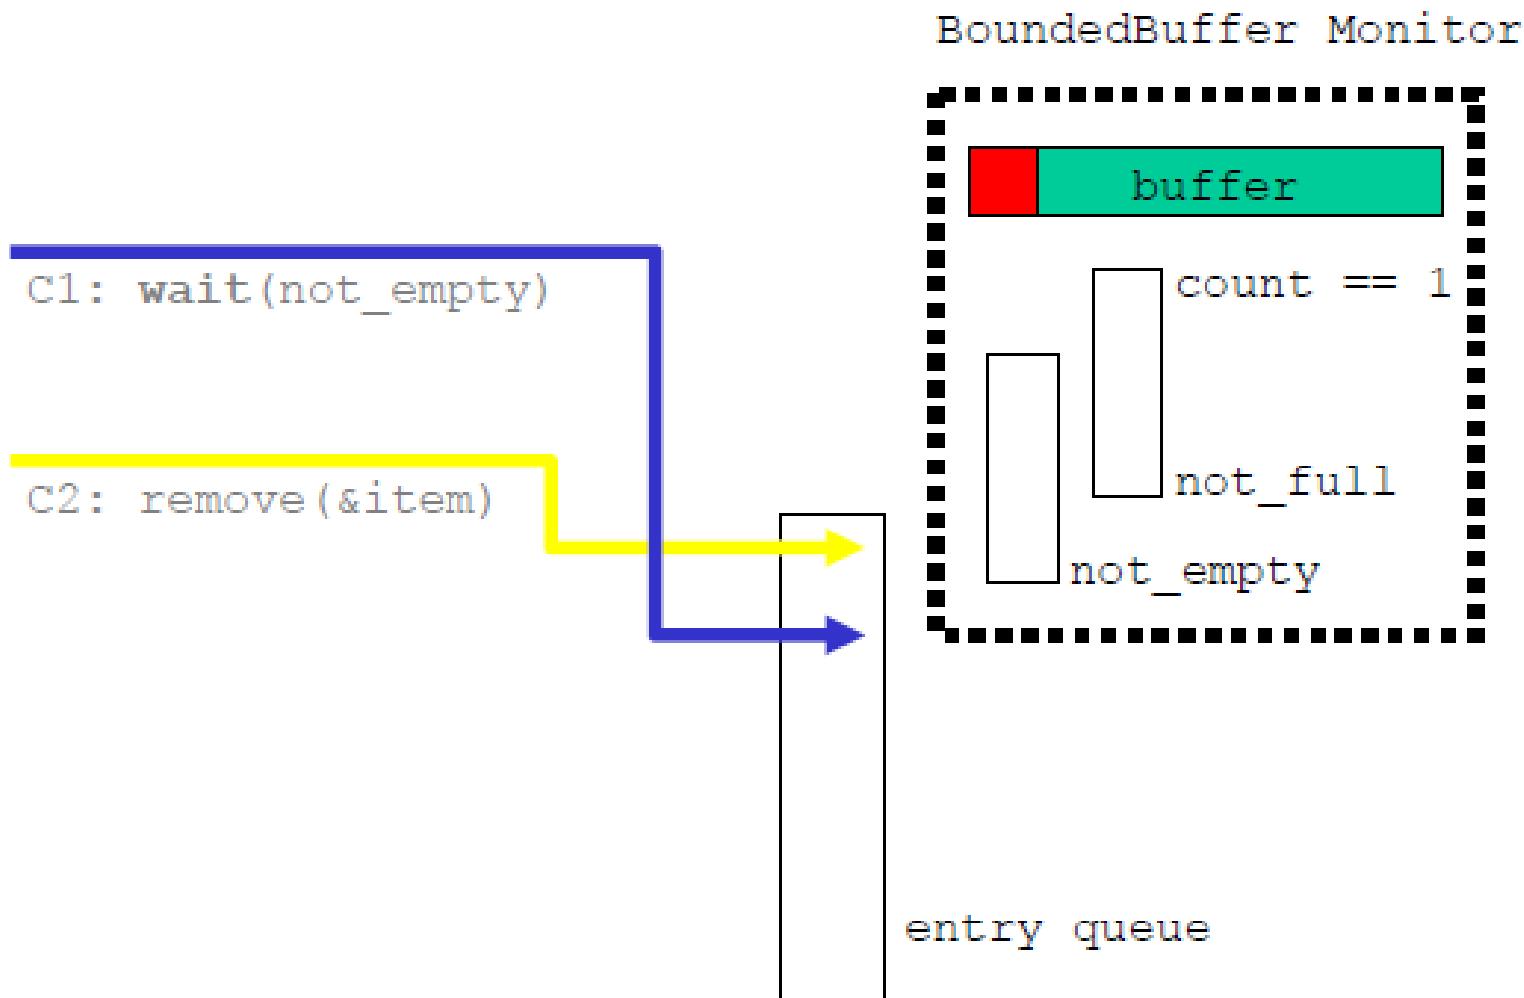
# Example Trace 11



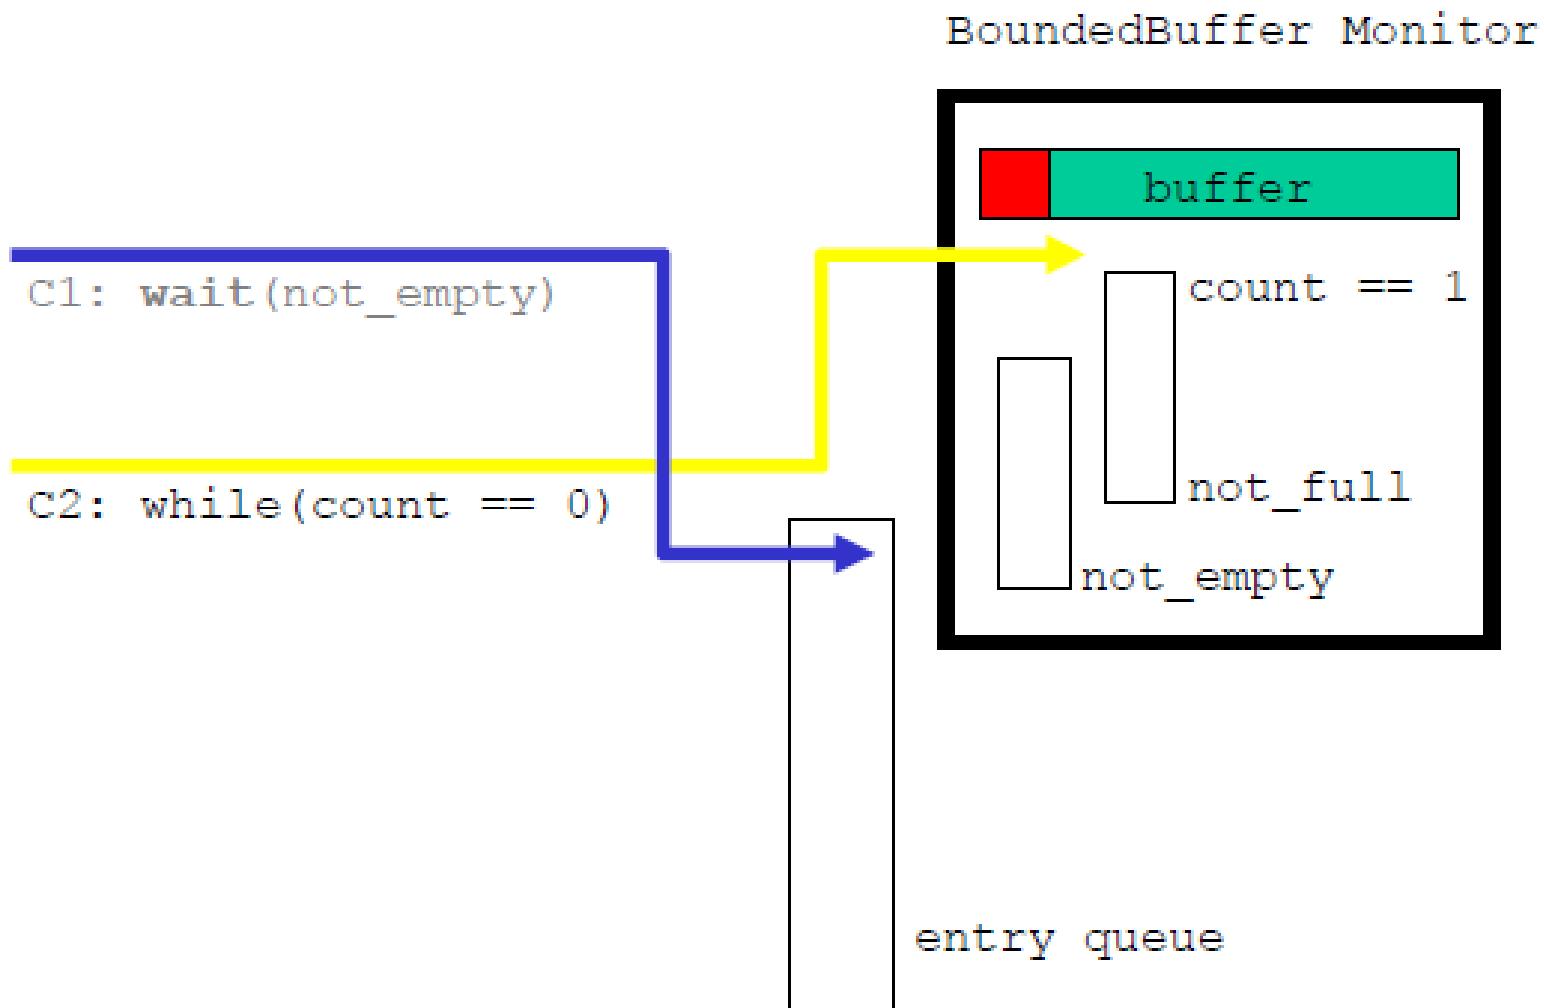
# Example Trace 12



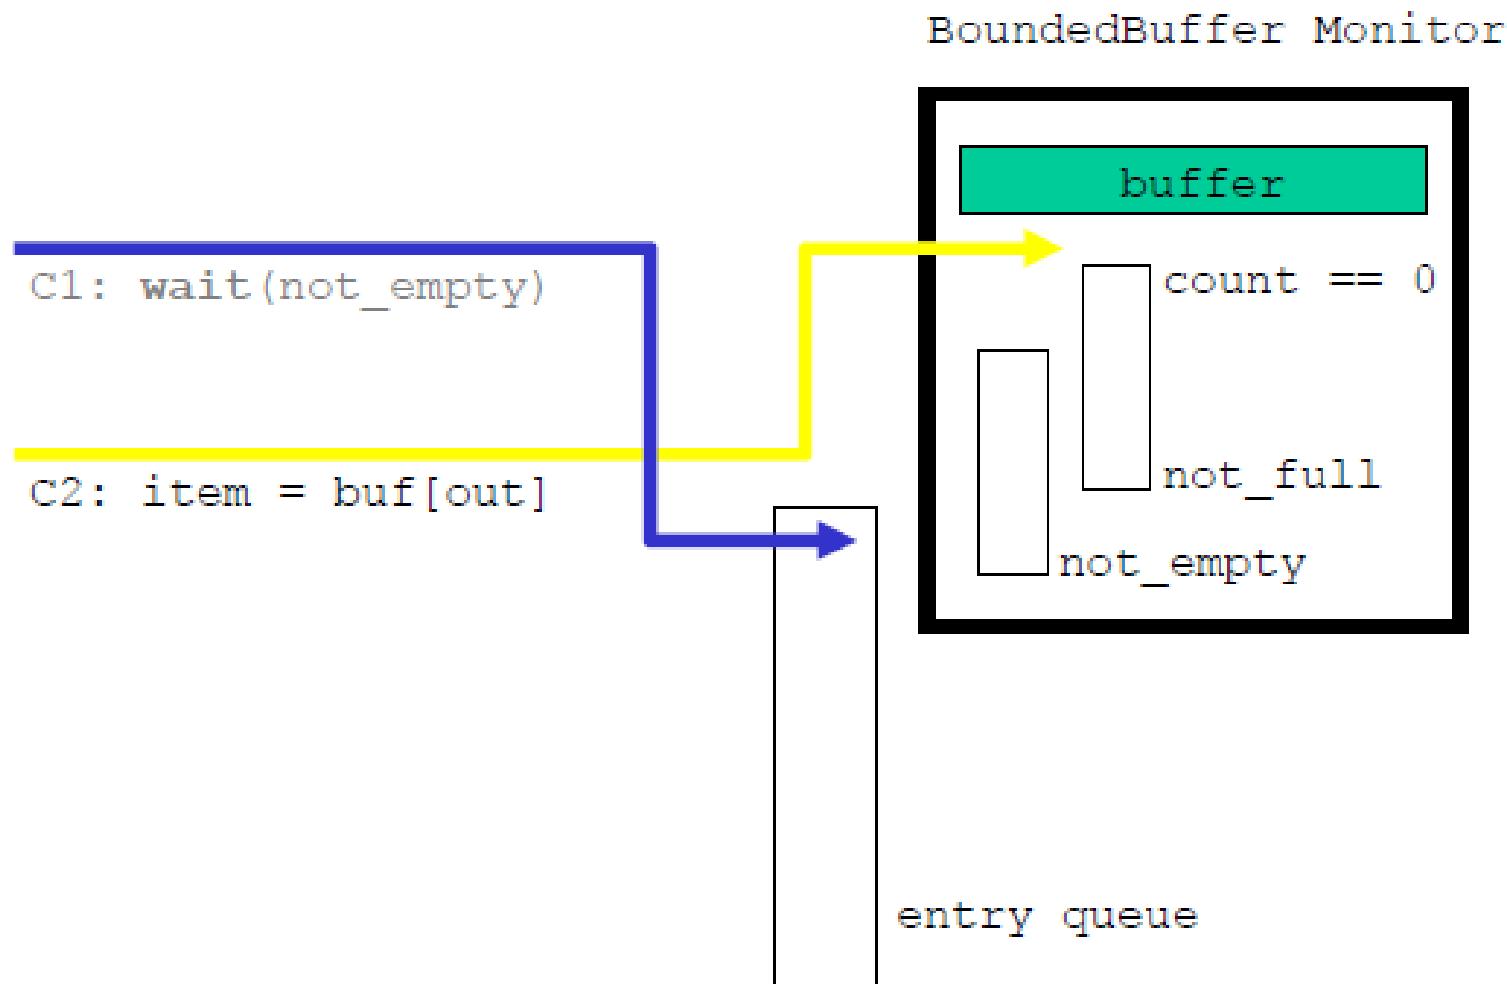
# Example Trace 13



# Example Trace 14

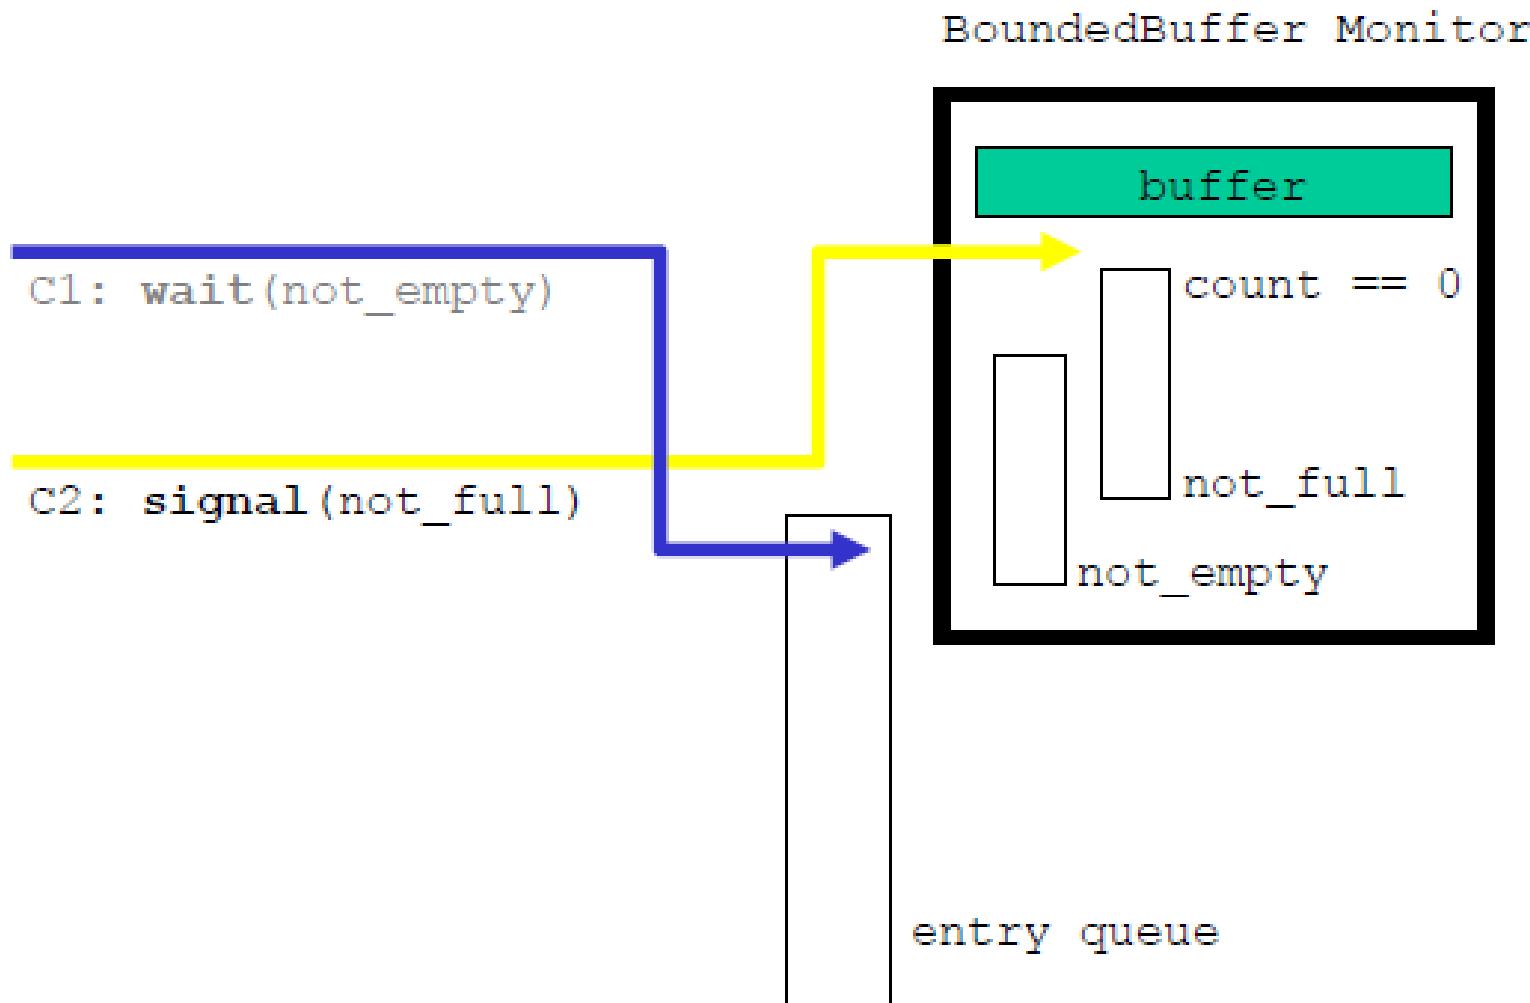


# Example Trace 15

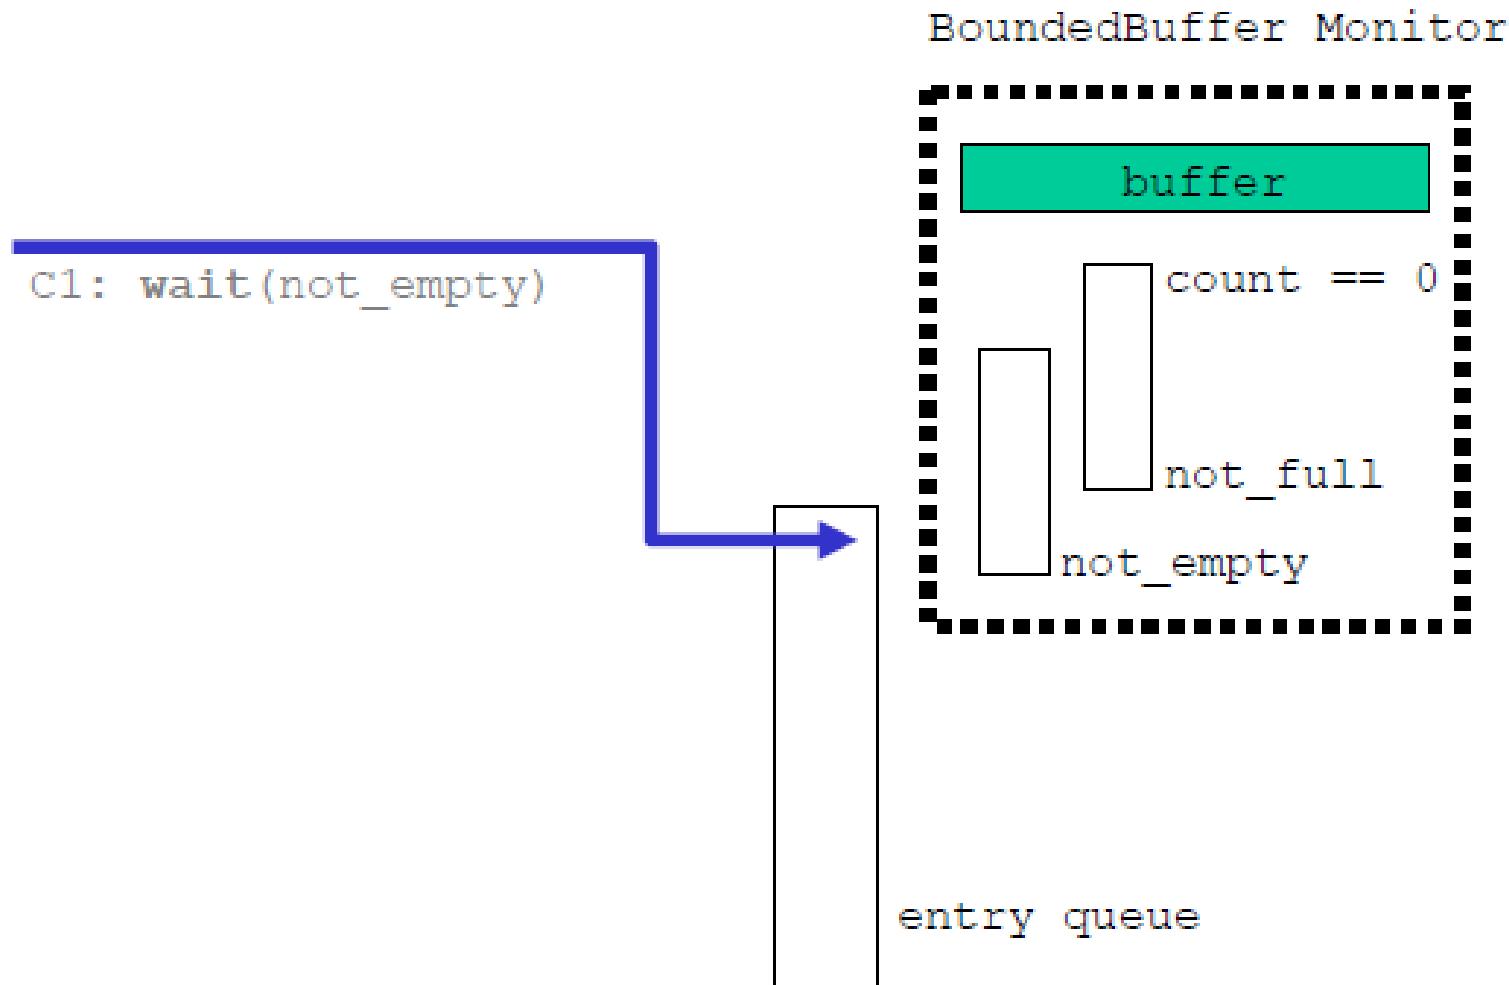


# Example Trace 16

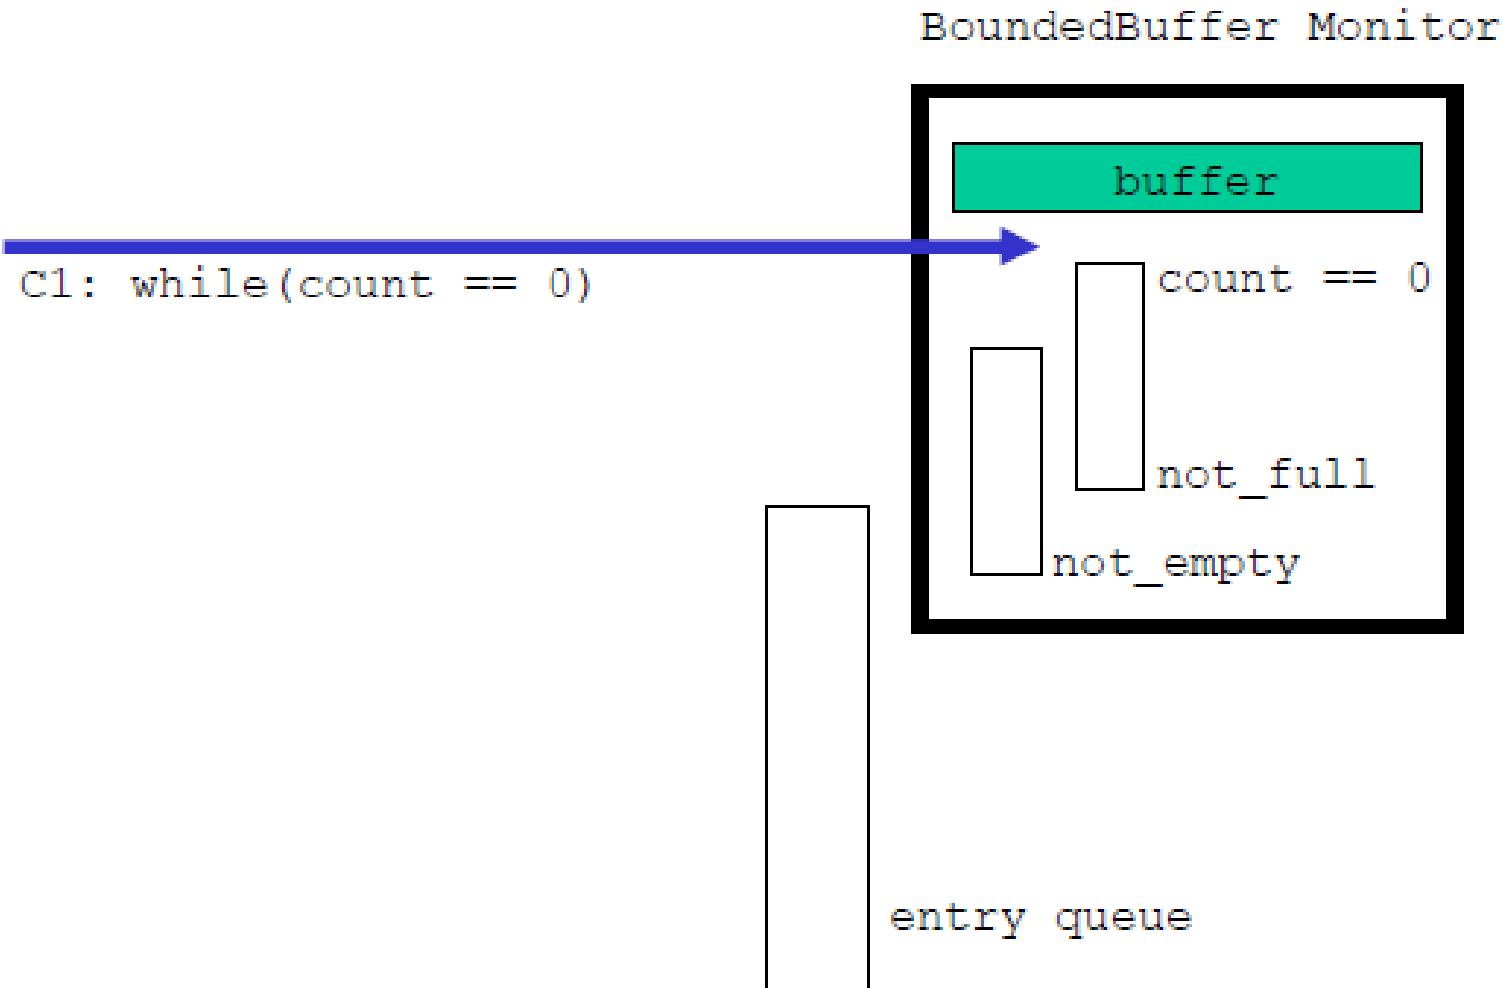
- **signal(not\_full)** does not have any effect if there is no producer waiting on the **not\_full** condition variable.



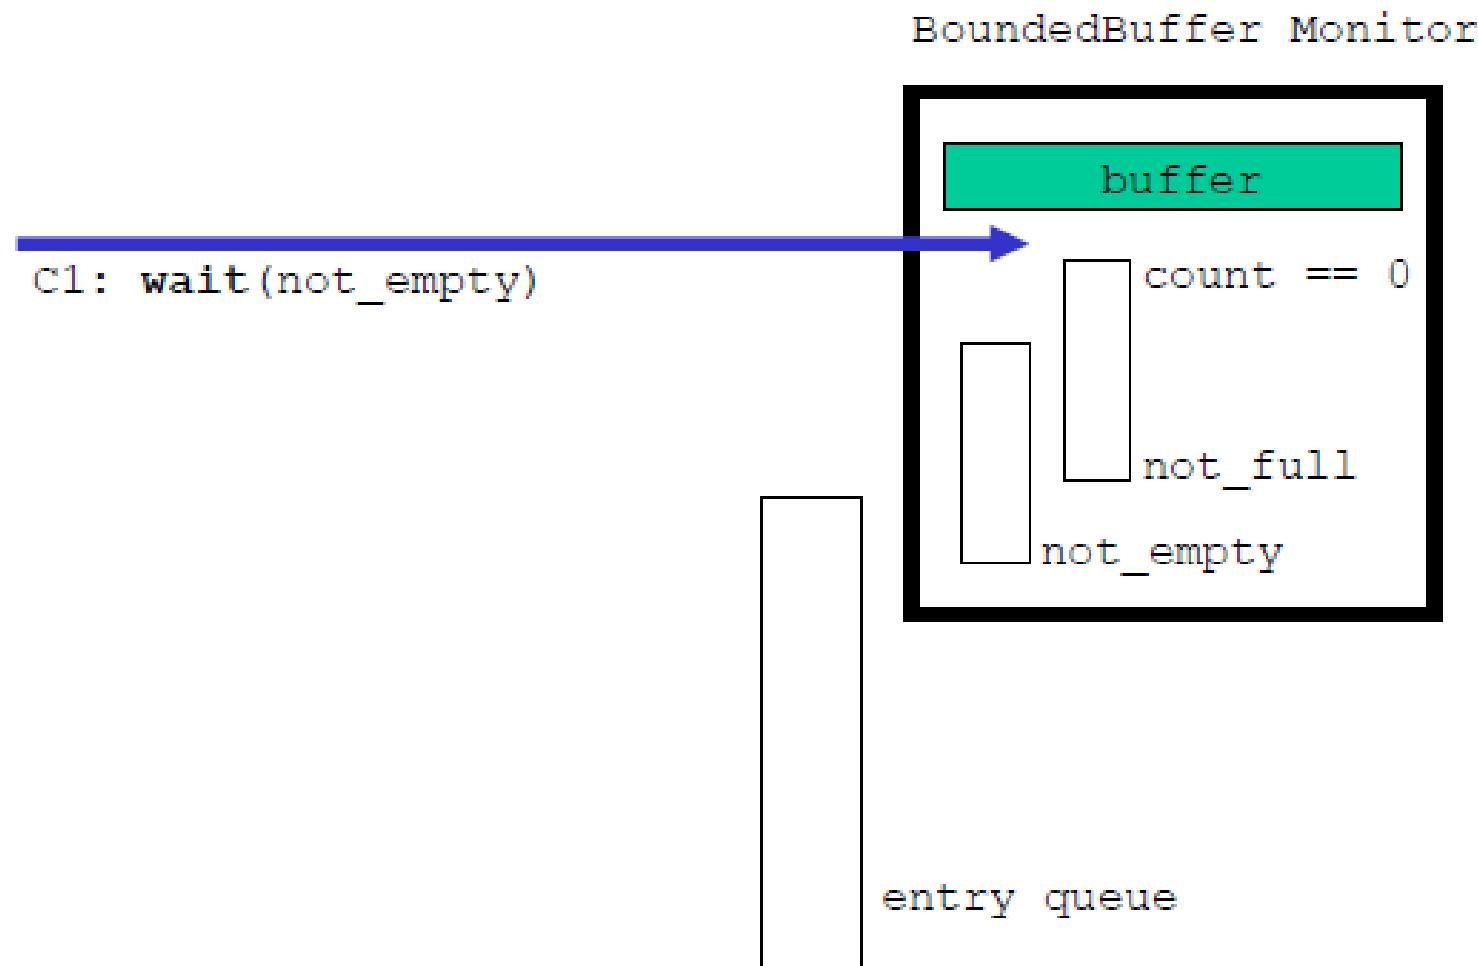
# Example Trace 17



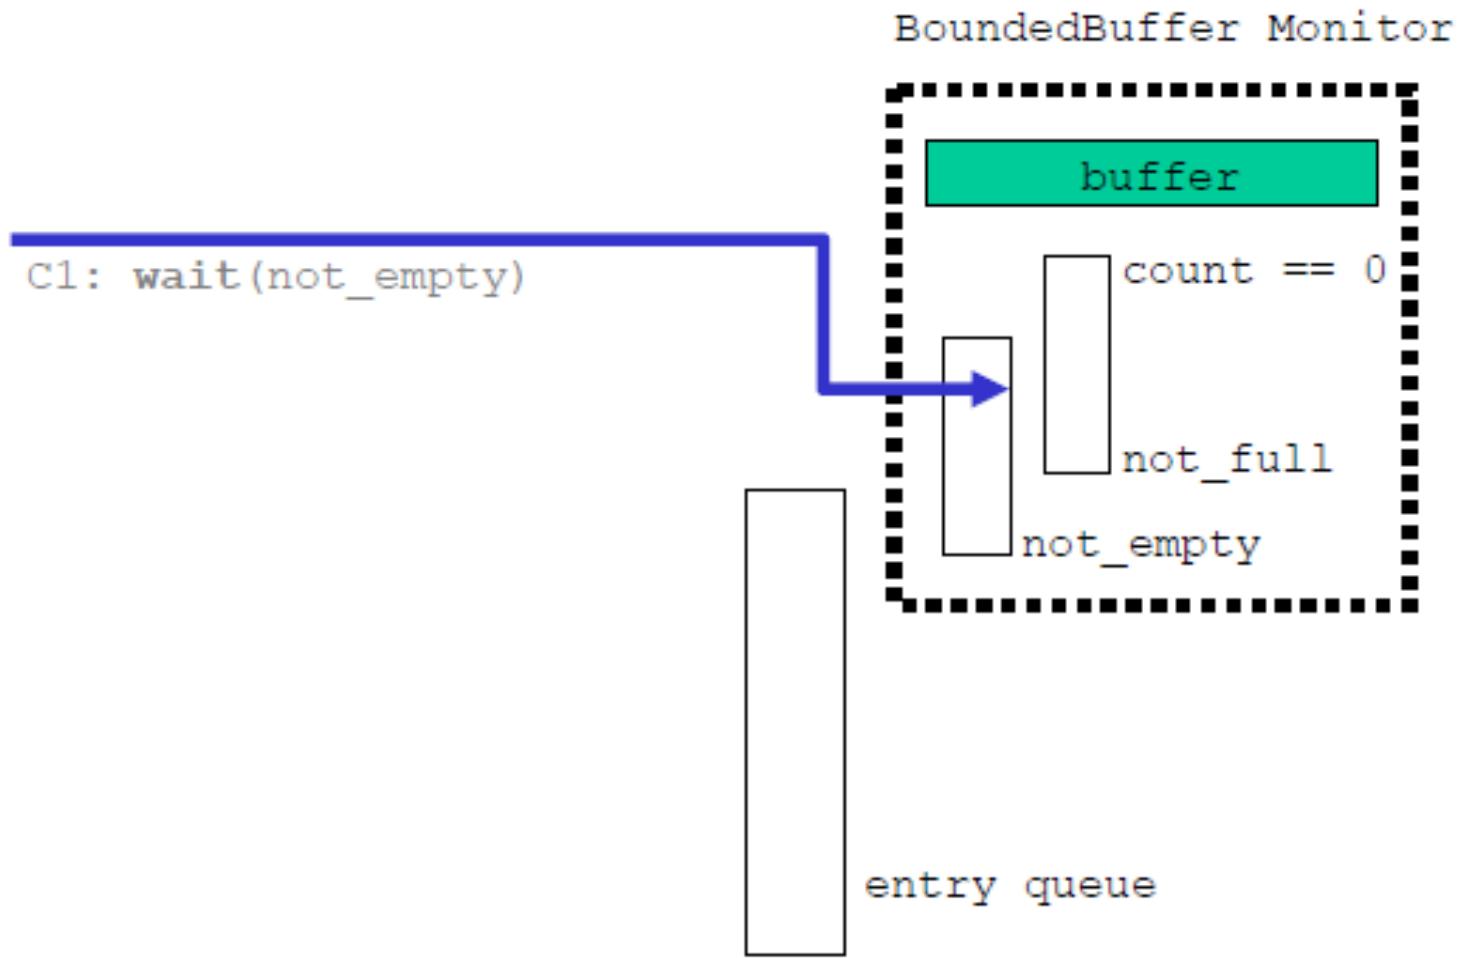
# Example Trace 18



# Example Trace 19



# Example Trace 20



# End of Chapter 5